

Contents

CHAPTER 1	WHY CMAKE?	1
1.1	The History of CMake	3
1.2	Why Not Use Autoconf?	3
1.3	Why Not Use JAM, qmake, SCons, or ANT?	4
1.4	Why Not Script It Yourself?	4
1.5	On What Platforms Does CMake Run?	5
CHAPTER 2	GETTING STARTED	7
2.1	Getting and Installing CMake on Your Computer	7
	<i>UNIX and Mac Binary Installations</i>	7
	<i>Windows Binary Installation</i>	7
2.2	Building CMake Yourself	8
2.3	Basic CMake Usage and Syntax	8
2.4	Hello World for CMake	9
2.5	How to Run CMake?	10
	<i>Running CMake's Qt Interface</i>	11
	<i>Running the ccmake Curses Interface</i>	13
	<i>Running CMake from the Command Line</i>	15
	<i>Specifying the Compiler to CMake</i>	15
	<i>Dependency Analysis</i>	16
2.6	Editing CMakeLists Files	17
2.7	Setting Initial Values for CMake	17
2.8	Building Your Project	19
CHAPTER 3	KEY CONCEPTS	21
3.1	Main Structures	21
3.2	Targets	24
3.3	Source Files	25
3.4	Directories, Generators, Tests, and Properties	26
3.5	Variables and Cache Entries	27
3.6	Build Configurations	32
CHAPTER 4	WRITING CMAKELISTS FILES	33
4.1	CMake Syntax	33
4.2	Basic Commands	34
4.3	Flow Control	35
4.4	Regular Expressions	42

4.5	Checking Versions of CMake	44
4.6	Using Modules	45
	<i>Using CMake with SWIG</i>	48
	<i>Using CMake with Qt</i>	49
	<i>Using CMake with FLTK</i>	50
4.7	Policies	50
	<i>Updating a Project For a New Version of CMake</i>	53
4.8	Linking Libraries	57
	<i>Specifying Optimized or Debug Libraries with a Target</i>	59
4.9	Shared Libraries and Loadable Modules	59
4.10	Shared Library Versioning	64
4.11	Installing Files	66
	<i>Installing Prerequisite Shared Libraries</i>	76
4.12	Advanced Commands	82
<hr/> CHAPTER 5 SYSTEM INSPECTION		85
5.1	Using Header Files and Libraries	85
5.2	System Properties	87
5.3	Finding Packages	92
5.4	Built-in Find Modules	93
5.5	How to Pass Parameters to a Compilation?	95
5.6	How to Configure a Header File	97
5.7	Creating CMake Package Configuration Files	99
<hr/> CHAPTER 6 CUSTOM COMMANDS AND TARGETS		103
6.1	Portable Custom Commands	103
6.2	Using <code>add_custom_command</code> on a Target	105
	<i>How to Copy an Executable Once it is Built?</i>	106
6.3	Using <code>add_custom_command</code> to Generate a File	107
	<i>Using an Executable to Build a Source File</i>	107
6.4	Adding a Custom Target	108
6.5	Specifying Dependencies and Outputs	111
6.6	When There Isn't One Rule For One Output	112
	<i>A Single Command Producing Multiple Outputs</i>	112
	<i>Having One Output That Can Be Generated By Different Commands</i>	112
<hr/> CHAPTER 7 CONVERTING EXISTING SYSTEMS TO CMAKE		115
7.1	Source Code Directory Structures	115
7.2	Build Directories	117
7.3	Useful CMake Commands When Converting Projects	119
7.4	Converting UNIX Makefiles	120

7.5	Converting Autoconf Based Projects	121
7.6	Converting Windows Based Workspaces	123
CHAPTER 8	CROSS COMPILING WITH CMAKE	125
8.1	Toolchain Files	126
	<i>Finding External Libraries, Programs and Other Files</i>	128
8.2	System Inspection	130
	<i>Using Compile Checks</i>	131
8.3	Running Executables Built in the Project	133
8.4	Cross Compiling Hello World	136
8.5	Cross Compiling for a Microcontroller	140
8.6	Cross Compiling an Existing Project	143
8.7	Cross Compiling a Complex Project - VTK	145
8.8	Some Tips and Tricks	147
CHAPTER 9	PACKAGING WITH CPACK	149
9.1	CPack Basics	149
	<i>Simple Example</i>	150
	<i>What Happens When CPack.cmake Is Included?</i>	151
	<i>Adding Custom CPack Options</i>	152
	<i>Options Added by CPack</i>	153
9.2	CPack Source Packages	154
9.3	CPack Installer Commands	154
9.4	CPack for Windows Installer NSIS	156
	<i>CPack Variables Used by CMake for NSIS</i>	156
	<i>Creating Windows Short Cuts in the Start Menu</i>	161
	<i>Advanced NSIS CPack Options</i>	161
	<i>Setting File Extension Associations With NSIS</i>	162
	<i>Installing Microsoft Run Time Libraries</i>	163
	<i>CPack Component Install Support</i>	163
9.5	CPack for Cygwin Setup	173
9.6	CPack for Mac OS X PackageMaker	176
9.7	CPack for Mac OS X Drag and Drop	178
9.8	CPack for Mac OS X X11 Applications	180
9.9	CPack for Debian Packages	182
9.10	CPack for RPM	183
9.11	CPack Files	183
CHAPTER 10	AUTOMATION & TESTING WITH CMAKE	185
10.1	Testing with CMake, CTest, and CDash	185
10.2	How Does CMake Facilitate Testing?	186

10.3	Additional Test Properties	187
10.4	Testing Using CTest	189
10.5	Using CTest to Drive Complex Tests	191
10.6	Handling a Large Number of Tests	192
10.7	Producing Test Dashboards	194
	<i>Adding CDash Dashboard Support to a Project</i>	196
	<i>Client Setup</i>	199
10.8	Customizing Dashboards for a Project	202
	<i>Dashboard Submissions Settings</i>	202
	<i>Filtering Errors and Warnings</i>	203
	<i>Adding Notes to a Dashboard</i>	205
10.9	Setting up Automated Dashboard Clients	206
	<i>Settings for Continuous Dashboards</i>	210
	<i>Variables Available in CTest Scripts</i>	212
10.10	Advanced CTest Scripting	212
	<i>Limitations of Traditional CTest Scripting</i>	213
	<i>Extended CTest Scripting</i>	213
10.11	Setting up a Dashboard Server	218
	<i>CDash Server</i>	218
	<i>Advanced Server Management</i>	220
	<i>Build Groups</i>	223
	<i>Email</i>	225
	<i>Sites</i>	226
	<i>Graphs</i>	227
	<i>Adding Notes to a Build</i>	228
	<i>Logging</i>	229
	<i>Test Timing</i>	229
	<i>Mobile Support</i>	230
	<i>Backing up CDash</i>	230
	<i>Upgrading CDash</i>	231
	<i>CDash Maintenance</i>	232
10.12	Subprojects	233
	<i>Using ctest_submit with PARTS and FILES</i>	236
	<i>Splitting Your Project into Multiple Subprojects</i>	237

CHAPTER 11 PORTING CMAKE TO NEW PLATFORMS AND LANGUAGES 241

11.1	The Determine System Process	241
11.2	The Enable Language Process	242
11.3	Porting to a New Platform	244
11.4	Adding a New Language	246
11.5	Rule Variable Listing	247
	<i>General Tag Variables</i>	247
	<i>Language Specific Information</i>	248

11.6	Compiler and Platform Examples <i>Como Compiler</i>	248
	<i>Borland Compiler</i>	249
11.7	Extending CMake <i>Creating a Loaded Command</i>	250
	<i>Using a Loaded Command</i>	251
CHAPTER 12 TUTORIALS		255
12.1	A Basic Starting Point (Step 1) <i>Adding a Version Number and Configured Header File</i>	255
12.2	Adding a Library (Step 2)	258
12.3	Installing and Testing (Step 3)	260
12.4	Adding System Introspection (Step 4)	262
12.5	Adding a Generated File and Generator (Step 5)	263
12.6	Building an Installer (Step 6)	267
12.7	Adding Support for a Dashboard (Step 7)	268
APPENDIX A - VARIABLES		269
	Variables That Change Behavior	269
	Variables That Describe the System	272
	Variables for Languages	274
	Variables That Control the Build	278
	Variables That Provide Information	280
APPENDIX B – COMMAND LINE REFERENCE		287
	CMake Command Line Options	287
	CMake Generators	292
	CTest Command Line Options	294
	CPack Command Line Options	298
	CPack Generators	299
APPENDIX C – LISTFILE COMMANDS		301
	Current Commands	301
	Compatibility Commands	366
APPENDIX D – SELECTED MODULES		373
	CMake Modules	373

APPENDIX E - PROPERTIES	411
Properties of Global Scope	411
Properties on Directories	414
Properties on Targets	417
Properties on Tests	431
Properties on Source Files	431
Properties on Cache Entries	434
APPENDIX F – CMAKE POLICIES	437
INDEX	447

Why CMake?

If you have ever maintained the build and installation process for a software package, you will be interested in CMake. CMake is an open source build manager for software projects that allows developers to specify build parameters in a simple portable text file format. This file is then used by CMake to generate project files for native build tools including Integrated Development Environments such as Microsoft Visual Studio or Apple's Xcode, as well as UNIX, Linux, NMake, and Borland style Makefiles. CMake handles the difficult aspects of building software such as cross platform builds, system introspection, and user customized builds, in a simple manner that allows users to easily tailor builds for complex hardware and software systems.

For any project, and especially cross platform projects, there is a need for a unified build system. Many projects today ship with both a UNIX Makefile (or Makefile.in) and a Microsoft Visual Studio workspace. This requires that developers constantly try to keep both build systems up to date and consistent with each other. To target additional build systems such as Borland or Xcode requires even more custom copies of these files, creating an even bigger problem. This problem is compounded if you try to support optional components, such as including JPEG support if libjpeg is available on the system. CMake solves this by consolidating these different operations into one simple easy to understand file format.

If you have multiple developers working on a project, or multiple target platforms, then the software will have to be built on more than one computer. Given the wide range of installed software and custom options that are involved with setting up a modern computer, the chances are that two computers running the same OS will be slightly different. CMake provides many benefits for single platform multi-machine development environments including:

- The ability to automatically search for programs, libraries, and header files that may be required by the software being built. This includes the ability to consider environment variables and Windows's registry settings when searching.
- The ability to build in a directory tree outside of the source tree. This is a useful feature found on many UNIX platforms; CMake provides this feature on Windows as well. This allows a developer to remove an entire build directory without fear of removing source files.
- The ability to create complex custom commands for automatically generated files such as Qt's moc (qt.nokia.com), The Insight Toolkit's CABLE wrappers (public.kitware.com/Cable/HTML/Index.html) and SWIG (www.swig.org) wrapper generators. These commands are used to generate new source files during the build process that are in turn compiled into the software.
- The ability to select optional components at configuration time. For example, several of VTK's libraries are optional, and CMake provides an easy way for users to select which libraries are built.
- The ability to automatically generate workspaces and projects from a simple text file. This can be very handy for systems that have many programs or test cases, each of which requires a separate project file, typically a tedious manual process to create using an IDE.
- The ability to easily switch between static and shared builds. CMake knows how to create shared libraries and modules on all platforms supported. Complicated platform-specific linker flags are handled, and advanced features like built in run time search paths for shared libraries are supported on many UNIX systems.
- Automatic generation of file dependencies and support for parallel builds on most platforms.

When developing cross platform software, CMake provides a number of additional features:

- The ability to test for machine byte order and other hardware specific characteristics.
- A single set of build configuration files that work on all platforms. This avoids the problem of developers having to maintain the same information in several different formats inside a project.
- Support for building shared libraries on all platforms that support it.
- The ability to configure files with system dependent information such as the location of data files and other information. CMake can create header files that contain information such as paths to data files and other information in the form of #define macros. System specific flags can also be placed in configured header files. This has advantages over command line -D options to the compiler because it allows other build systems to use the CMake built library without having to specify the exact same command line options used during the build.

1.1 The History of CMake

CMake development began in 1999 as part of the Insight Toolkit (ITK, www.itk.org) funded by the US National Library of Medicine. ITK is a large software project that works on many platforms and can interact with many other software packages. To support this, a powerful, yet easy to use, build tool was required. Having worked with build systems for large projects in the past, the developers designed CMake to address these needs. Since then CMake has continuously grown in popularity, with many projects and developers adopting it for its ease of use and flexibility. Since 1999 CMake has been under active development and has matured to the point where it is a proven solution for a wide range of build issues. The most telling example of this is the successful adoption of CMake as the build system of the K Desktop Environment (KDE), arguably the largest open source software project in existence.

One of the recent additions to CMake is the inclusion of software testing support in the form of CTest. Part of the process of testing software involves building the software, possibly installing it, and determining what parts of the software are appropriate for the current system. This makes CTest a logical extension of CMake as it already has most of this information. In a similar vein, a new CMake feature is CPack, which is designed to support cross platform distribution of software. It provides a cross platform approach to creating native installations for your software, making use of existing popular packages such as NSIS, RPM, Cygwin, and PackageMaker.

Other recent additions to CMake include support for Apple's Xcode IDE and support for Microsoft's Visual Studio 10. With CMake, once you write your input files you get support for new compilers and build systems for free because the support for them is built into new releases of CMake, not tied to your software distribution. Another recent addition to CMake is support for cross compiling to other operating systems or embedded devices. Many commands in CMake now properly handle the differences between the host system and the target platform when cross compiling.

1.2 Why Not Use Autoconf?

Before developing CMake its authors had experience with the existing set of available tools. Autoconf combined with automake provides some of the same functionality as CMake, but to use these tools on a Windows platform requires the installation of many additional tools not found natively on a Windows box. In addition to requiring a host of tools, autoconf can be difficult to use or extend and impossible for some tasks that are easy in CMake. Even if you do get autoconf and its required environment running on your system, it generates Makefiles that will force users to the command line. CMake on the other hand provides a choice, allowing developers to generate project files that can be used directly from the IDE to which Windows and Xcode developers are accustomed.

While autoconf supports user specified options, it does not support dependent options where one option depends on some other property or selection. For example, in CMake you could have a user option to enable multithreading be dependent on first determining if the user's system has multithreading support. CMake provides an interactive user interface, making it easy for the user to see what options are available and how to set them.

For UNIX users, CMake also provides automated dependency generation that is not done directly by autoconf. CMake's simple input format is also easier to read and maintain than a combination of `Makefile.in` and `configure.in` files. The ability of CMake to remember and chain library dependency information has no equivalent in autoconf/automake.

1.3 Why Not Use JAM, qmake, SCons, or ANT?

Other tools such as ANT, qmake, SCons, and JAM have taken different approaches to solving these problems and they have helped us to shape CMake. Of the four, qmake, is the most similar to CMake although it lacks much of the system interrogation that CMake provides. Qmake's input format is more closely related to a traditional `Makefile`. ANT, JAM and SCons are also cross-platform although they do not support generating native project files. They do break away from the traditional `Makefile` oriented input with ANT using XML, JAM using its own language, and SCons using Python. A number of these tools run the compiler directly, as opposed to letting the system's build process perform that task. Many of these tools require other tools such as Python or Java to be installed before they will work.

1.4 Why Not Script It Yourself?

Some projects use existing scripting languages such as Perl or Python to configure build processes. Although similar functionality can be achieved with systems like this, over-use of tools can make the build process more of an Easter egg hunt than a simple-to-use build system. When building your software package users are forced to find and install version 4.3.2 of this, and 3.2.4 of that, before they can even start the build process. To avoid that problem, it was decided that CMake would require no more tools than the software it was being used to build would require. At a minimum using CMake requires a C compiler, that compiler's native build tools, and a CMake executable. CMake was written in C++, requires only a C++ compiler to build and precompiled binaries are available for most systems. Scripting it yourself also typically means you will not be generating native Xcode or Visual Studio workspaces, making Mac and Windows builds limited.

1.5 On What Platforms Does CMake Run?

CMake runs on a wide variety of platforms including Microsoft Windows, Apple Mac OS X, and most UNIX or UNIX-like platforms. At the time of the writing of this book CMake was tested nightly on the following platforms: Windows 98/2000/XP/Vista/7, AIX, HPUX, IRIX, Linux, Mac OS X, Solaris, OSF, QNX, CYGWIN, MinGW, and FreeBSD. You can check www.cmake.org for a current list of tested platforms.

Likewise, CMake supports most common compilers. It supports the GNU compiler on all CMake supported platforms. Other tested compilers include Visual Studio 6 through 10, Intel C, SGI CC, Mips Pro, Borland, Sun CC and HP aCC. CMake should work for most UNIX-style compilers out of the box. If the compiler takes arguments in a strange way, then see the section *Porting CMake to New Platform* on page 241 for information on how to customize CMake for a new compiler.

Getting Started

2.1 Getting and Installing CMake on Your Computer

Before using CMake you will need to install or build the CMake binaries on your system. On many systems you may find that CMake is already installed, or is available for install with the standard package manager tool for the system. Cygwin, Debian, FreeBSD, Mac OS X Fink, and many others all have CMake distributions. If your system does not have a CMake package, you can find CMake precompiled for most common architectures at www.cmake.org. If you do not find binaries for your system precompiled, then you can build CMake from source. To build CMake you will need a modern C++ compiler.

UNIX and Mac Binary Installations

If your system provides CMake as one of its standard packages, follow your system's package installation instructions. If your system does not have CMake, or has an out of date version of CMake, you can download precompiled binaries from www.cmake.org. The binaries from www.cmake.org come in the form of a compressed tar file. The tar file contains a README file and an enclosed tar file. The README file contains a manifest of the files contained in the enclosed tar file, and some instructions. To install, simply extract the enclosed tar file into a destination directory (typically `/usr/local`). However, it can be any directory, and does not require root privileges for installation.

Windows Binary Installation

For Windows CMake has a NullSoft install file available for download from www.cmake.org. To install this file, simply run the executable on the windows machine on which you want to install CMake. You will be able to run CMake from the Start Menu after it is installed.

2.2 Building CMake Yourself

If binaries are not available for your system, or if binaries are not available for the version of CMake you wish to use, you can build CMake from the source code. You can obtain the CMake source code by following the instructions at www.cmake.org. Once you have the source code it can be built in two different ways. If you have a version of CMake on your system you can use it to build other versions of CMake. Generally the current development version of CMake can always be built from the previous release of CMake. This is how new versions of CMake are built on most Windows systems.

The second way to build CMake is by running its bootstrap build script. To do this you change directory into your CMake source directory and type

```
./bootstrap  
make  
make install
```

The make install step is optional since CMake can run directly from the build directory if desired. On UNIX, if you are not using the GNU C++ compiler, you need to tell the bootstrap script which compiler you want to use. This is done by setting the environment variable `CXX` before running bootstrap. If you need to use any special flags with your compiler, set the `CXXFLAGS` environment variable. For example, on the SGI with the 7.3X compiler, you would build CMake like this:

```
cd CMake  
(setenv CXX CC; setenv CXXFLAGS "-LANG:std"; ./bootstrap)  
make  
make install
```

2.3 Basic CMake Usage and Syntax

Using CMake is simple. The build process is controlled by creating one or more CMakeLists files (actually CMakeLists.txt but this guide will leave off the extension in most cases) in each of the directories that make up a project. The CMakeLists files should contain the project description in CMake's simple language. The language is expressed as a series of commands. Each command is evaluated in the order that it appears in the CMakeLists file. The commands have the form

```
command (args...)
```

where command is the name of the command, and args is a white-space separated list of arguments. (Arguments with embedded white-space should be double quoted.) CMake is case insensitive to command names as of version 2.2. So where you see command you could use COMMAND or Command instead. Older versions of CMake only accepted uppercase commands.

CMake supports simple variables that can be either strings or lists of strings. Variables are referenced using a `${VAR}` syntax. Multiple arguments can be grouped together into a list using the `set` command. All other commands expand the lists as if they had been passed into the command with white-space separation. For example, `set(Foo a b c)` will result in setting the variable `Foo` to `a b c`, and if `Foo` is passed into another command `command(${Foo})` it would be equivalent to `command(a b c)`. If you want to pass a list of arguments to a command as if it were a single argument simply double quote it. For example `command("${Foo}")` would be invoked passing only one argument equivalent to `command("a b c")`.

System environment variables and Windows registry values can be accessed directly in CMake. To access system environment variables the syntax `$ENV{VAR}` is used. CMake can also reference registry entries in many commands using a syntax of the form `[HKEY CURRENT USER\\Software\\path1\\path2;key]`, where the paths are built from the registry tree and key.



2.4 Hello World for CMake

For starters let us consider the simplest possible CMakeLists file. To compile an executable from one source file the CMakeLists file would contain two lines:

```
project (Hello)
add_executable (Hello Hello.c)
```

To build the Hello executable you follow the process described in Running CMake (See section 2.5) to generate the Makefiles or Microsoft project files. The `project` command indicates what the name of the resulting workspace should be and the `add_executable` command adds an executable target to the build process. That's all there is to it for this simple example. If your project requires a few files it is also quite easy, just modify the `add_executable` line as shown below.

```
add_executable (Hello Hello.c File2.c File3.c File4.c)
```

`add_executable` is just one of many commands available in CMake. Consider the more complicated example below.

```
cmake_minimum_required (2.6)
project (HELLO)

set (HELLO_SRCS Hello.c File2.c File3.c)

if (WIN32)
    set(HELLO_SRCS ${HELLO_SRCS} WinSupport.c)
else ()
    set(HELLO_SRCS ${HELLO_SRCS} UnixSupport.c)
endif ()

add_executable (Hello ${HELLO_SRCS})

# look for the Tcl library
find_library (TCL_LIBRARY
    NAMES tcl tcl84 tcl83 tcl82 tcl80
    PATHS /usr/lib /usr/local/lib
    )

if (TCL_LIBRARY)
    target_link_library (Hello ${TCL_LIBRARY})
endif ()
```

In this example the `set` command is used to group together source files into a list. The `if` command is used to add either `WinSupport.c` or `UnixSupport.c` to this list based on whether or not CMake is running on Windows. Finally, the `add_executable` command is used to build the executable with the files listed in the variable `HELLO_SRCS`. The `find_library` command looks for the Tcl library under a few different names and in a few different paths. An `if` command checks if the `TCL_LIBRARY` was found and if so adds it to the link line for the `Hello` executable target. Note the use of the `#` character to denote a comment line. All characters from the `#` to the end of the line are considered to be part of the comment.

2.5 How to Run CMake?

Once CMake has been installed on your system, using it to build a project is easy. There are two main directories CMake uses when building a project: the source directory and the binary directory. The source directory is where the source code for your project is located. This is also where the `CMakeLists` files will be found. The binary directory is where you want CMake to put the resulting object files, libraries, and executables. Typically CMake will not write any files to the source directory, only the binary directory. If you want to you can set the source and binary directories to be the same. This is known as an in-source build, in contrast to an out-of-source build where they are different.

CMake supports both in-source and out-of-source builds on all operating systems. This means that you can configure your build to be completely outside of the source code tree which makes it very easy to remove all of the files generated by a build. Having the build tree differ from the source tree also makes it easy to support having multiple builds of a single source tree. This is useful when you want to have multiple builds with different options but just one copy of the source code. Now let us consider the specifics of running CMake using its Qt based GUI and command line interfaces.

Running CMake's Qt Interface

CMake includes a Qt based user interface developed by Clinton Stimpson that can be used on most platforms, including UNIX, Mac OS X, and Windows. This interface is included in the CMake source code, but you will need an installation of Qt on your system in order to build it.

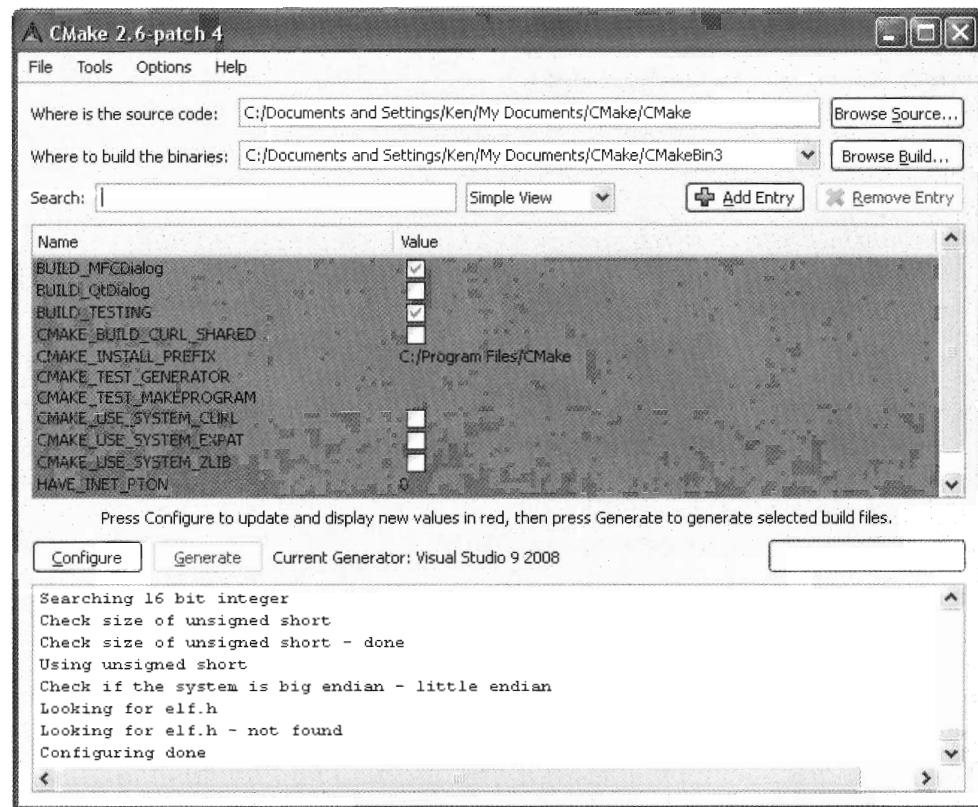


Figure 1 – Qt based CMake GUI

On Windows the executable is named `cmake-gui.exe` and should be in your Start menu under Program Files. There may also be a shortcut on your desktop, or if you built CMake from source, it will be in the build directory. For UNIX and Mac users the executable is

named cmake-gui and it can be found where you installed the CMake executables. A GUI will appear similar to what is shown in Figure 1. The top two entries are the source code and binary directories. They allow you to specify where the source code is for what you want to compile and where the resulting binaries should be placed. You should set these two values first. If the binary directory you specify does not exist, it will be created for you. If the binary directory has been configured by CMake before then it will automatically set the source tree.

The middle area is where you can specify different options for the build process. More obscure variables may be hidden, but can be seen if you select "Advanced View" from the view pulldown. You can search for values in the middle area by typing all or part of the name into the Search box. This can be handy for finding specific settings or options in a large project. The bottom area of the window includes the Configure and Generate buttons as well as a progress bar and scrollable output window.

Once you have specified the source code and binary directories you should click the Configure button. This will cause CMake to read in the CMakeLists files from the source code directory and then update the cache area to display any new options for the project. If you are running cmake-gui for the first time on this binary directory it will prompt you to determine what generator you wish to use, as shown in Figure 2. This dialog also presents options for customizing and tweaking the compilers you wish to use for this build.

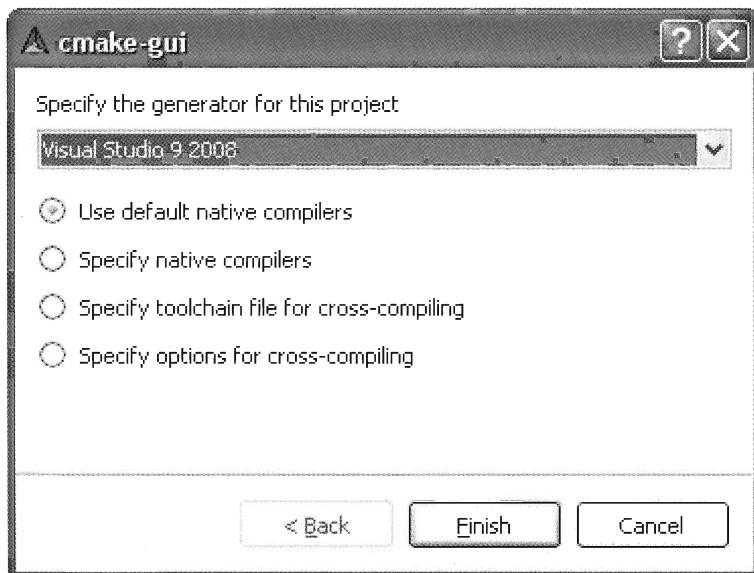


Figure 2 – Selecting a Generator

After the first configure you can adjust your cache settings if desired and click the Configure button again. New values that were created by the configure process will be colored red. To be sure you have seen all possible values you should click Configure until no values are red

and you are happy with all the settings. Once you are done configuring, click the Generate button, this will produce the appropriate files.

It is important that you make sure that your environment is suitable for running cmake-gui. If you are using an IDE such as Visual Studio then your environment will be setup correctly for you. If you are using NMake or MinGW then you need to make sure that the compiler can run from your environment. You can either directly set the required environment variables for your compiler or use a shell in which they are already set. For example, Microsoft Visual Studio has an option on the start menu for creating a Visual Studio Command Prompt. This opens up a command prompt window that has its environment already setup for Visual Studio. You should run cmake-gui from this command prompt if you want to use NMake Makefiles. The same approach applies to MinGW, you should run cmake-gui from a MinGW shell that has a working compiler in its path.

When cmake-gui finishes it will have generated the build files in the binary directory you specified. If Visual Studio was selected as the generator, a MSVC workspace (or solution) file is created. This file's name is based on the name of the project you specified in the PROJECT command at the beginning of your CMakeLists file. For many other generator types, Makefiles are generated. The next step in this process is to open the workspace with MSVC. Once open, the project can be built in the normal manner of Microsoft Visual C++. The ALL_BUILD target can be used to build all of the libraries and executables in the package. If you are using a Makefile build type, then you would build by running make or nmake on the resulting Makefiles.

Running the ccmake Curses Interface

On most UNIX platforms, if the curses library is supported, CMake provides an executable called ccmake. This interface is a terminal-based text application that is very similar to the Qt based GUI. To run ccmake, change directory (cd) to the directory where you want the binaries to be placed. This can be the same directory as the source code for what we call in-source builds or it can be a new directory you create. Then run ccmake with the path to the source directory on the command line. For in-source builds use "." for the source directory. This will start the text interface as shown in Figure 3 (in this case the cache variables are from VTK and most are set automatically).

```

andy@andoria                               Page 1 of 1
BUILD_EXAMPLES          ON
BUILD_SHARED_LIBS        ON
CMAKE_BACKWARDS_COMPATIBILITY 1.7
CMAKE_BUILD_TYPE
CMAKE_INSTALL_PREFIX
CMAKE_JAVA_ARCHIVE
CMAKE_JAVA_RUNTIME
PYTHON_INCLUDE_PATH
PYTHON_LIBRARY
TCL_INCLUDE_PATH
TCL_LIBRARY
TK_INCLUDE_PATH
TK_LIBRARY
VALGRIND_COMMAND_OPTIONS
VTK_DATA_ROOT
VTK_USE_HIERID
VTK_USE_PARALLEL
VTK_USE_PATENTED
VTK_USE_RENDERING
VTK_WRAP_JAVA
VTK_WRAP_PYTHON
VTK_WRAP_TCL

BUILD EXAMPLES: Build VTK examples.
Press [enter] to edit option
Press [c] to configure
Press [h] for help           Press [q] to quit without generating
Press [t] to toggle advanced mode (Currently off)                                     CMake Version 2.1 - development

```

Figure 3 - ccmake running on UNIX

Brief instructions are displayed in the bottom of the window. If you hit the "c" key, it will configure the project. You should always configure after changing values in the cache. To change values, use the arrow keys to select cache entries, and then the enter key to edit them. Boolean values will toggle with the enter key. Once you have set all the values as you like, you can hit the "g" key to generate the Makefiles and exit. You can also hit "h" for help, "q" to quit, and "t" to toggle the viewing of advanced cache entries. Two examples of CMake usage on the UNIX platform follow for a hello world project called Hello. In the first example, an in-source build is performed.

```

cd Hello
ccmake .
make

```

In the second example, an out-of-source build is performed.

```

mkdir Hello-Linux
cd Hello-Linux
ccmake ../Hello
make

```

Running CMake from the Command Line

From the command line, CMake can be run as an interactive question and answer session or as a non-interactive program. To run in interactive mode, just pass the "-i" option to CMake. This will cause CMake to ask you for a value for each entry in the cache file for the project. CMake will provide reasonable defaults, just like it does in the GUI and curses based interfaces. The process stops when there are no longer any more questions to ask. An example of using the interactive mode of CMake is provided below.

```
$ cmake -i -G "NMake Makefiles" ../CMake
Would you like to see advanced options? [No]:
Please wait while cmake processes CMakeLists.txt files....
```

Variable Name: BUILD_TESTING
Description: Build the testing tree.
Current Value: ON
New Value (Enter to keep current value):

Variable Name: CMAKE_INSTALL_PREFIX
Description: Install path prefix, prepended onto install directories.
Current Value: C:/Program Files/CMake
New Value (Enter to keep current value):

```
Please wait while cmake processes CMakeLists.txt files....
```

CMake complete, run make to build project.

Using CMake to build a project in non-interactive mode is a simple process if the project has few or no options. For larger projects like VTK, using `ccmake`, `cmake -i`, or `cmake-gui` is recommended. To build a project with a non-interactive CMake, first change directory to where you want the binaries to be placed. For an in-source build you then run `cmake .` and pass in any options using the `-D` flag. For out-of-source builds the process is the same except you run `cmake` and also provide the path to the source code as its argument. Then type `make` and your project should compile. Some projects will have install targets as well, you can type `make install` to install them.

Specifying the Compiler to CMake

On some systems you may have more than one compiler to choose from or your compiler may be in a non-standard place. In these cases you will need to specify to CMake where your desired compiler is located. There are three ways to specify this; the generator can specify the compiler, an environment variable can be set, or a cache entry can be set. Some generators are tied to a specific compiler, for example the Visual Studio 6 generator always uses the

Microsoft Visual Studio 6 compiler. For Makefile based generators CMake will try a list of usual compilers until it finds a working compiler. The list can be found in the files:

Modules/CMakeDetermineCCompiler.cmake and
Modules/CMakeDetermineCXXCompiler.cmake

The lists can be preempted with environment variables that can be set before CMake is run. The CC environment variable specifies the C compiler while CXX specifies the C++ compiler. You can specify the compilers directly on the command line by using -DCMAKE_CXX_COMPILER=cl for example. If those are not set, CMake will try the following list of compilers:

c++ g++ CC aCC cl bcc xlC.

Once CMake has been run and picked a compiler, you can change the selection by changing the cache entries CMAKE_CXX_COMPILER and CMAKE_C_COMPILER, although this is not recommended. The problem with doing this is that the project you are configuring may have already run some tests on the compiler to determine what it supports. Changing the compiler does not normally cause these tests to be rerun which can lead to incorrect results. If you must change the compiler, start over with an empty binary directory. The flags for the compiler and the linker can also be changed by setting environment variables. Setting LDFLAGS will initialize the cache values for link flags, while CXXFLAGS and CFLAGS will initialize CMAKE_CXX_FLAGS and CMAKE_C_FLAGS respectively.

Dependency Analysis

CMake has powerful built-in dependency analysis capabilities for C and C++ source code files. CMake also has limited support for Fortran and Java dependencies. Since Integrated Development Environments (IDEs) support and maintain dependency information, CMake skips this step for those build systems. However, Makefiles with a make program do not know how to automatically compute and keep dependency information up-to-date. For these builds, CMake automatically computes dependency information for C, C++ and Fortran files. Both the generation and maintenance of these dependencies are automatically done by CMake. Once a project is initially configured by CMake, users only need to run make, and CMake does the rest of the work. CMake's dependencies fully support parallel builds for multiprocessor systems.

Although users do not need to know how CMake does this work, it may be useful to look at the dependency information files for a project. This information for each target is stored in four files called depend.make, flags.make, build.make, and DependInfo.cmake. depend.make stores the depend information for all the object files in the directory. flags.make contains the compile flags used for the source files of this target. If they change then the files will be recompiled. DependInfo.cmake is used to keep the dependency

information up-to-date and contains information about what files are part of the project and what languages they are in. Finally, the rules for building the dependencies are stored in `build.make`. If a dependency is out of date then all of the dependencies for that target will be recomputed, keeping the dependency information current.

2.6 Editing CMakeLists Files

CMakeLists files can be edited in almost any text editor. Some editors, such as Notepad++, come with CMake syntax highlighting and indentation support built in. For editors such as Emacs or Vim CMake includes indentation and syntax highlighting modes. These can be found in the `Docs` directory of the source distribution, or downloaded from the CMake web site. The file `cmake-mode.el` is the Emacs mode, and `cmake-indent.vim` and `cmake-syntax.vim` are used by Vim. Within Visual Studio the CMakeLists files are listed as part of the project and you can edit them simply by double clicking on them. Within any of the supported generators (Makefiles, Visual Studio, etc) if you edit a CMakeLists file and rebuild, there are rules that will automatically invoke CMake to update the generated files (e.g. Makefiles or project files) as required. This helps to assure that your generated files are always in sync with your CMakeLists files.

Since CMake computes and maintains dependency information, the CMake executables must always be available (though they don't have to be in your PATH) when make or an IDE is being run on CMake generated files. This means that if a CMake input file changes on disk, your build system will automatically re-run CMake and produce up-to-date build files. For this reason you generally should not generate Makefiles or projects with CMake and move them to another machine that does not have CMake installed.

2.7 Setting Initial Values for CMake

While CMake works well in an interactive mode, sometimes you will need to setup cache entries without running a GUI. This is common when setting up nightly dashboards or if you will be creating many build trees with the same cache values. In these cases the CMake cache can be initialized in two different ways. The first way is to pass the cache values on the CMake command line using `-DCACHE_VAR:TYPE=VALUE` arguments. For example, consider the following nightly dashboard script for a UNIX machine:

```
#!/bin/tcsh

cd ${HOME}

# wipe out the old binary tree and then create it again
rm -rf Foo-Linux
mkdir Foo-Linux
```

```
cd Foo-Linux

# run cmake to setup the cache
cmake -DBUILD TESTING:BOOL=ON <etc...> .../Foo

# generate the dashboard
ctest -D Nightly
```

The same idea can be used with a batch file on Windows. The second way is to create a file to be loaded using CMake's `-C` option. In this case instead of setting up the cache with `-D` options it is done through a file that is parsed by CMake. The syntax for this file is standard CMakeLists syntax and it is typically just a series of `set` commands such as:

```
#Build the vtkHybrid kit.
set (VTK_USE_HYBRID ON CACHE BOOL "doc string")
```

In some cases there might be an existing cache and you want to force the cache values to be set a certain way. For example say you want to turn Hybrid on even if the user has previously run CMake and turned it off. Then you can do:

```
#Build the vtkHybrid kit always.
set (VTK_USE_HYBRID ON CACHE BOOL "doc" FORCE)
```

Another option is that you want to set and then hide options so the user will not be tempted to adjust them later on. This can be done using the following commands:

```
#Build the vtkHybrid kit always and don't distract
#the user by showing the option.
set (VTK_USE_HYBRID ON CACHE INTERNAL "doc" FORCE)
mark_as_advanced (VTK_USE_HYBRID)
```

You might be tempted to edit the cache file directly, or to "initialize" a project by giving it an initial cache file. This may not work and could cause additional problems in the future. First, the syntax of the CMake cache is subject to change. Second, cache files have full paths in them that make them unsuitable for moving between binary trees. So if you want to initialize a cache file use one of the two standard methods described above.

2.8 Building Your Project

After you have run CMake your project will be ready to be built. If your target generator is based on Makefiles then you can build your project by changing directory to your binary tree and typing make (or gmake or nmake as appropriate). If you generated files for an IDE such as Visual Studio, you can start your IDE, load the project files into it, and build as you normally would.

Another option is to use CMake's --build option from the command line. This option is simply a convenience that allows you to build your project from the command line, even if that requires launching an IDE. The command line options for --build include:

```
Usage: cmake --build <dir> [options] [-- [native-options]]  
Options:  
  <dir>          = Project binary directory to be built.  
  --target <tgt> = Build <tgt> instead of default targets.  
  --config <cfg> = For multi-configuration tools, choose <cfg>.  
  --clean-first   = Build target 'clean' first, then build.  
                  (To clean only, use --target 'clean'.)  
  --             = Pass remaining options to the native tool.
```

So even if you are using Visual Studio as your generator you can type the following to build your project from the command line if you wish.

```
cmake --build <your binary dir>
```

That is all there is to installing and running CMake for simple projects. In the following chapters we will consider CMake in more detail and how to use it on more complex software projects.

Key Concepts

3.1 Main Structures

This chapter provides an introduction to CMake's key concepts. As you start working with CMake you will run into a variety of concepts such as targets, generators, and commands. In CMake these concepts are implemented as C++ classes and are referenced in many of CMake's commands. Understanding these concepts will provide you with the working knowledge you need to create effective CMakeLists files.

Before going into detail about CMake's classes it is worth understanding their basic relationships. At the lowest level there are source files. These correspond to typical C or C++ source code files. Source files are combined into targets. A target is typically an executable or library. A directory represents a directory in the source tree and typically has a CMakeLists file and one or more targets associated with it. Every directory has a local generator that is responsible for generating the Makefiles or project files for that directory. All of the local generators share a common global generator that oversees the build process. Finally, the global generator is created and driven by the `cmake` class itself.

Figure 4 shows the basic class structure of CMake. We will now consider CMake's concepts in a bit more detail. CMake's execution begins by creating an instance of the `cmake` class and passing the command line arguments to it. This class manages the overall configuration process and holds information that is global to the build process such as the cache values. One of the first things the `cmake` class does is to create the correct global generator based on the user's selection of what generator to use (such as Visual Studio 10, Borland Makefiles, or UNIX Makefiles). At this point the `cmake` class passes control to the global generator it created by invoking the `configure` and `generate` methods.

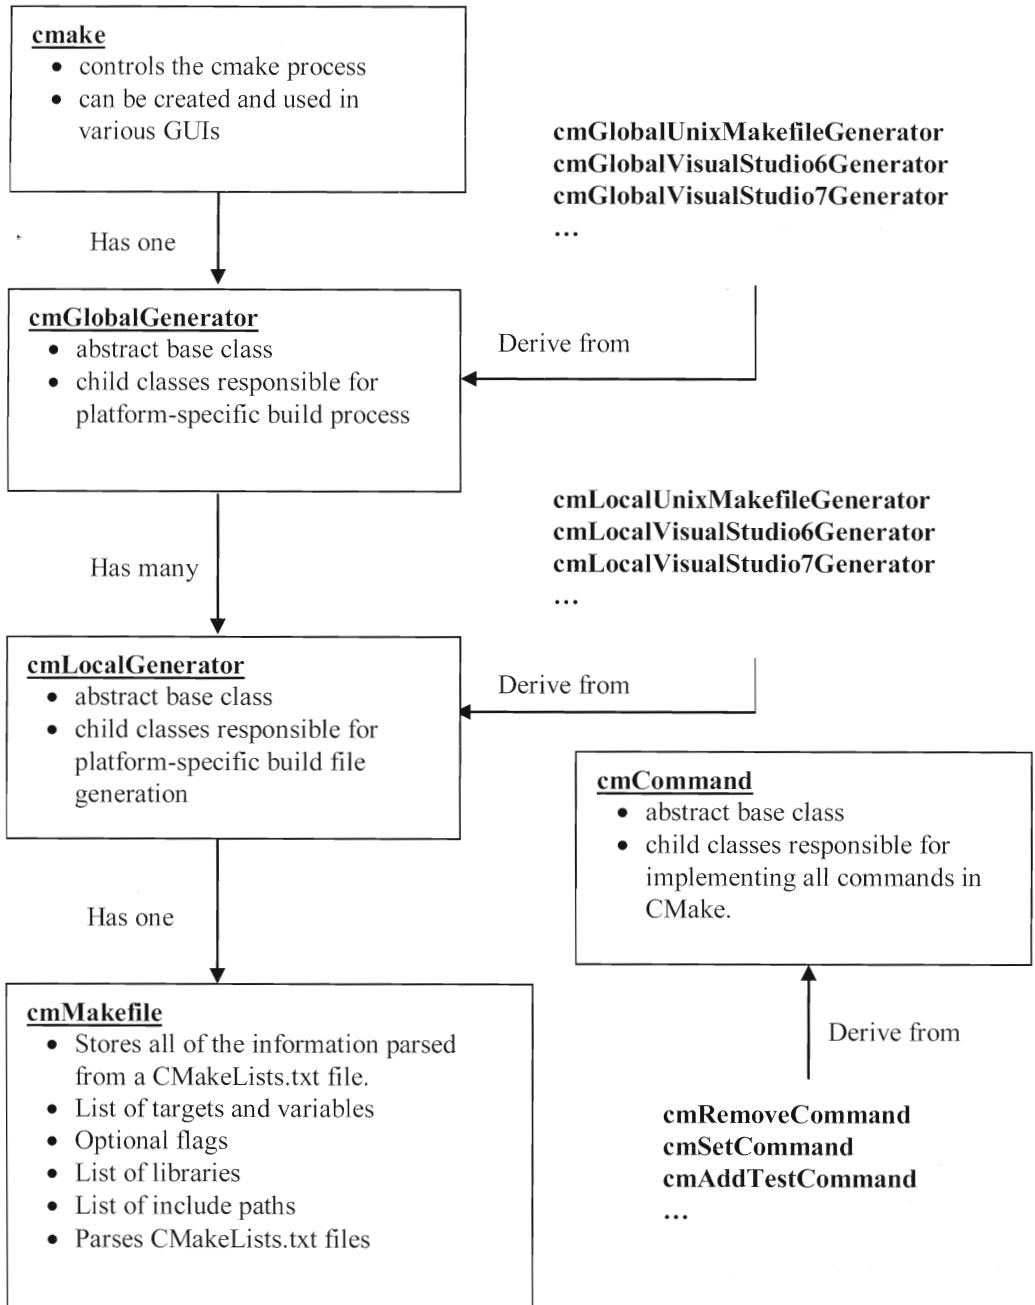


Figure 4 - CMake Internals

The global generator is responsible for managing the configuration and generation of all of the Makefiles (or project files) for a project. In practice most of the work is actually done by local generators which are created by the global generator. One local generator is created for each directory of the project that is processed. So while a project will have only one global generator it may have many local generators. For example, under Visual Studio 7 the global generator creates a solution file for the entire project while the local generators create a project file for each target in their directory.

In the case of the "Unix Makefiles" generator, the local generators create most of the Makefiles and the global generator simply orchestrates the process and creates the main top-level Makefile. Implementation details vary widely among generators. The Visual Studio 6 generators make use of .dsp and .dsw file templates and perform variable replacements on them. The generators for Visual Studio 7 and later directly generate the XML output without using any file templates. The Makefile generators including UNIX, NMake, Borland, etc use a set of rule templates and replacements to generate their Makefiles.

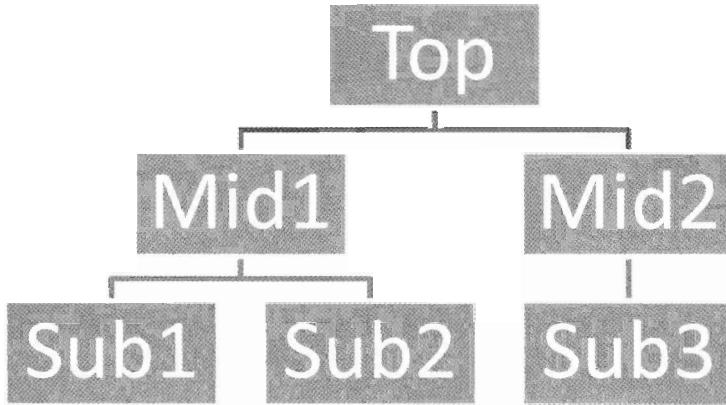


Figure 5 - Sample Directory Tree

Each local generator has an instance of the class `cmMakefile`, `cmMakefile` is where the results of parsing the CMakeLists files are stored. Specifically, for each directory in a project there will be a single `cmMakefile` instance which is why the `cmMakefile` class is often referred to as the directory. This is clearer for build systems that do not use Makefiles. That instance will hold all of the information from parsing that directory's CMakeLists file (see Figure 5). One way to think of the `cmMakefile` class is as a structure that starts out initialized with a few variables from its parent directory, and is then filled in as the CMakeLists file is processed. Reading in the CMakeLists file is simply a matter of CMake executing the commands it finds in the order it encounters them.

Each command in CMake is implemented as a separate C++ class, and has two main parts. The first part of a command is the InitialPass method. The `InitialPass` method receives the arguments and the `cmMakefile` instance for the directory currently being processed, and then

performs its operations. In the case of the `set` command, it processes its arguments and if the arguments are correct it calls a method on the `cmMakefile` to set the variable. The results of the command are always stored in the `cmMakefile` instance. Information is never stored in a command. The last part of a command is the FinalPass. The FinalPass of a command is executed after all commands (for the entire CMake project) have had their InitialPass invoked. Most commands do not have a FinalPass, but in some rare cases a command must do something with global information that may not be available during the initial pass.

Once all of the `CMakeLists` files have been processed the generators use the information collected into the `cmMakefile` instances to produce the appropriate files for the target build system (such as Makefiles).

3.2 Targets

Now that we have discussed the overall process of CMake, let us consider some of the key items stored in the `cmMakefile` instance. Probably the most important item is targets. Targets represent executables, libraries, and utilities built by CMake. Every `add_library`, `add_executable`, and `add_custom_target` command creates a target. For example, the following command will create a target named `foo` that is a static library, with `foo1.c` and `foo2.c` as source files.

```
add_library (foo STATIC foo1.c foo2.c)
```

The name `foo` is now available for use as a library name everywhere else in the project, and CMake will know how to expand the name into the library when needed. Libraries can be declared to be of a particular type such as STATIC, SHARED, MODULE, or left undeclared. STATIC indicates that the library must be built as a static library. Likewise SHARED indicates it must be built as a shared library. MODULE indicates that the library must be created so that it can be dynamically loaded into an executable. On many operating systems this is the same as SHARED, but on other systems such as Mac OS X it is different. If none of these options are specified this indicates that the library could be built as either shared or static. In that case CMake uses the setting of the variable BUILD_SHARED_LIBS to determine if the library should be SHARED or STATIC. If it is not set, then CMake defaults to building static libraries.

Likewise executables have some options. By default an executable will be a traditional console application that has a `main (int argc, const char*argv[])`. If WIN32 is specified after the executable name then the executable will be compiled as a MS Windows executable and the operating system will call `WinMain` instead of `main` at startup. WIN32 has no effect on non-Windows systems.

In addition to storing their type, targets also keep track of general properties. These properties can be set and retrieved using the set_target_properties and get_target_property

commands, or the more general `set_property` and `get_property` commands. The most commonly used property is `LINK_FLAGS`, which is used to specify link flags for a specific target. Targets store a list of libraries that they link against which are set using the `target_link_libraries` command. Names passed into this command can be libraries, full paths to libraries, or the name of a library from an `add_library` command. They also store the link directories to use when linking, the install location for the target, and custom commands to execute after linking.

For each library CMake creates, it keeps track of all the libraries on which that library depends. Since static libraries do not link to the libraries on which they depend, it is important for CMake to keep track of the libraries so they can be specified on the link line of the executable being created. For example,

```
add_library (foo foo.cxx)
target_link_libraries (foo bar)

add_executable (foobar foobar.cxx)
target_link_libraries (foobar foo)
```

This will link the libraries `foo` and `bar` into the executable `foobar` even, although only `foo` was explicitly linked into `foobar`. With shared or DLL builds this linking is not always needed, but the extra linkage is harmless. For static builds this is required. Since the `foo` library uses symbols from the `bar` library, `foobar` will most likely also need `bar` since it uses `foo`.

3.3 Source Files

The source file structure is in many ways similar to a target. It stores the filename, extension, and a number of general properties related to a source file. Like targets you can set and get properties using `set_source_files_properties` and `get_source_file_property`, or the more generic versions. The most common properties include:

COMPILE_FLAGS

Compile flags specific to this source file. These can include source specific `-D` and `-I` flags.

GENERATED

The `GENERATED` property indicates that the source file is generated as part of the build process. In this case CMake will treat it differently for computation of dependencies because the source file may not exist when CMake is first run.

OBJECT_DEPENDS

Adds additional files on which this source file should depend. CMake automatically performs dependency analysis to determine the usual C, C++ and Fortran dependencies. This parameter is used rarely in cases where there is an unconventional dependency or the source files do not exist at dependency analysis time.

ABSTRACT

WRAP_EXCLUDE

CMake doesn't directly use these properties. Some loaded commands and extensions to CMake look at these properties to determine how and when to wrap a C++ class into languages such as Tcl, Python, etc.

3.4 Directories, Generators, Tests, and Properties

In addition to targets and source files you may find yourself occasionally working with other classes such as directories, generators, and tests. Normally such interactions take the shape of setting or getting properties from these objects. All of these classes have properties associated with them, as do source files and targets. A property is a key-value pair attached to a specific object such as a target. The most generic way to access properties is through the `set_property` and `get_property` commands. These commands allow you to set or get a property from any class in CMake that has properties. Some of the properties for targets and source files have already been covered. Some useful properties for a directory include:

ADDITIONAL_MAKE_CLEAN_FILES

This property specifies a list of additional files that will be cleaned as a part of the "make clean" stage. By default CMake will clean up any generated files that it knows about, but your build process may use other tools that leave files behind. This property can be set to a list of those files so that they also will be properly cleaned up.

EXCLUDE_FROM_ALL

This property indicates if all the targets in this directory and all sub directories should be excluded from the default build target. If it is not, then with a Makefile for example typing `make` will cause these targets to be built as well. The same concept applies to the default build of other generators.

LISTFILE_STACK

This property is mainly useful when trying to debug errors in your CMake scripts. It returns a list of what list files are currently being processed, in order. So if one CMakeLists file does an `include` command then that is effectively pushing the included CMakeLists file onto the stack.

A full list of properties supported in CMake can be obtained by running `cmake` with the `-help-property-list` option. The generators and directories are automatically created for you as CMake processes your source tree.

3.5 Variables and Cache Entries

CMakeLists files use variables much like any programming language. Variables are used to store values for later use, and can be a single value such as "ON" or "OFF", or they can represent a list such as `(/usr/include /home/foo/include /usr/local/include)`. A number of useful variables are automatically defined by CMake and are discussed in Appendix A - Variables.

Variables in CMake are referenced using a `${VARIABLE}` notation, and they are defined in the order of execution of the `set` commands. Consider the following example:

```
# FOO is undefined  
  
set (FOO 1)  
# FOO is now set to 1  
  
set (FOO 0)  
# FOO is now set to 0
```

This may seem straightforward, but consider the following example:

```
set (FOO 1)  
  
if (${FOO} LESS 2)  
    set (FOO 2)  
else (${FOO} LESS 2)  
    set (FOO 3)  
endif (${FOO} LESS 2)
```

Clearly the `if` statement is true, which means that the body of the `if` statement will be executed. That will set the variable `FOO` to 2, and so when the `else` statement is encountered `FOO` will have a value of 2. Normally in CMake the new value of `FOO` would be used, but the `else` statement is a rare exception to the rule and always refers back to the value of the variable when the `if` statement was executed. So in this case the body of the `else` clause will not be executed. To further understand the scope of variables consider this example:

```

set (foo 1)

# process the dir1 subdirectory
add_subdirectory (dir1)

# include and process the commands in file1.cmake
include (file1.cmake)

set (bar 2)
# process the dir2 subdirectory
add_subdirectory (dir2)

# include and process the commands in file2.cmake
include (file2.cmake)

```

In this example because the variable `foo` is defined at the beginning, it will be defined while processing both `dir1` and `dir2`. In contrast `bar` will only be defined when processing `dir2`. Likewise `foo` will be defined when processing both `file1.cmake` and `file2.cmake`, whereas `bar` will only be defined while processing `file2.cmake`.

Variables in CMake have a scope that is a little different from most languages. When you set a variable it is visible to the current CMakeLists file or function, as well as any subdirectory's CMakeLists files, any functions or macros that are invoked, and any files that are included using the `INCLUDE` command. When a new subdirectory is processed (or a function called) a new variable scope is created and initialized with the current value of all variables in the calling scope. Any new variables created in the child scope, or changes made to existing variables, will not impact the parent scope. Consider the following example:

```

function (foo)
    message (${test}) # test is 1 here
    set (test 2)
    message (${test}) # test is 2 here, but only in this scope
endfunction()

set (test 1)
foo()
message (${test}) # test will still be 1 here

```

In some cases you might want a function or subdirectory to set a variable in its parent's scope. This is one way for CMake to return a value from a function, and it can be done by using the `PARENT_SCOPE` option with the `set` command. We can modify the prior example so that the function `foo` changes the value of `test` in its parent's scope as follows:

```
function (foo)
    message (${test}) # test is 1 here
    set (test 2 PARENT_SCOPE)
    message (${test}) # test still 1 in this scope
endfunction()

set (test 1)
foo()
message (${test}) # test will now be 2 here
```

Variables can also represent a list of values. In these cases when the variable is expanded it will be expanded into multiple values. Consider the following example:

```
# set a list of items
set (items_to_buy apple orange pear beer)

# loop over the items
foreach (item ${items_to_buy})
    message ( "Don't forget to buy one ${item}" )
endforeach ()
```

In some cases you might want to allow the user building your project to set a variable from the CMake user interface. In that case the variable must be a cache entry. Whenever CMake is run it produces a cache file in the directory where the binary files are to be written. The values of this cache file are displayed by the CMake user interface. There are a few purposes of this cache. The first is to store the user's selections and choices, so that if they should run CMake again they will not need to reenter that information. For example, the `option` command creates a Boolean variable and stores it in the cache.

```
option (USE_JPEG "Do you want to use the jpeg library")
```

The above line would create a variable called `USE_JPEG` and put it into the cache. That way the user can set that variable from the user interface and its value will remain in case the user should run CMake again in the future. To create a variable in the cache you can use commands like `option`, `find_file`, or you can use the standard `set` command with the `CACHE` option.

```
set (USE_JPEG ON CACHE BOOL "include jpeg support?")
```

When you use the cache option you must also provide the type of the variable and a documentation string. The type of the variable is used by the GUI to control how that variable

is set and displayed. Variable types include `BOOL`, `PATH`, `FILEPATH`, and `STRING`. The documentation string is used by the GUI to provide online help.

The other purpose of the cache is to store key variables that are expensive to determine. These variables may not be visible or adjustable by the user. Typically these values are system dependent variables such as `CMAKE_WORDS_BIGENDIAN`, which require CMake to compile and run a program to determine their value. Once these values have been determined, they are stored in the cache to avoid having to recompute them every time CMake is run. Generally CMake tries to limit these variables to properties that should never change (such as the byte order of the machine you are on). If you significantly change your computer, either by changing the operating system, or switching to a different compiler, you will need to delete the cache file (and probably all of your binary tree's object files, libraries, and executables).

Variables that are in the cache also have a property indicating if they are advanced or not. By default when a CMake GUI is run (such as `cmake` or `cmake-gui`) the advanced cache entries are not displayed. This is so that the user can focus on the cache entries that they should consider changing. The advanced cache entries are other options that the user can modify, but typically will not. It is not unusual for a large software project to have fifty or more options, and the advanced property lets a software project divide them into key options for most users and advanced options for advanced users. Depending on the project there may not be any non-advanced cache entries. To make a cache entry advanced the `mark as advanced` command is used with the name of the variable (a.k.a. cache entry) to make advanced.

In some cases you might want to restrict a cache entry to a limited set of predefined options. You can do this by setting the `STRINGS` property on the cache entry. The following CMakeLists code illustrates this by creating a property named `CRYPTOBACKEND` as usual, and then setting the `STRINGS` property on it to a set of three options.

```
set (CRYPTOBACKEND "OpenSSL" CACHE STRING
    "Select a cryptography backend")
set_property (CACHE CRYPTOBACKEND PROPERTY STRINGS
    "OpenSSL" "LibTomCrypt" "LibDES")
```



When `cmake-gui` is run and the user selects the `CRYPTOBACKEND` cache entry, they will be presented with a pulldown to select which option they want, as shown in Figure 6.

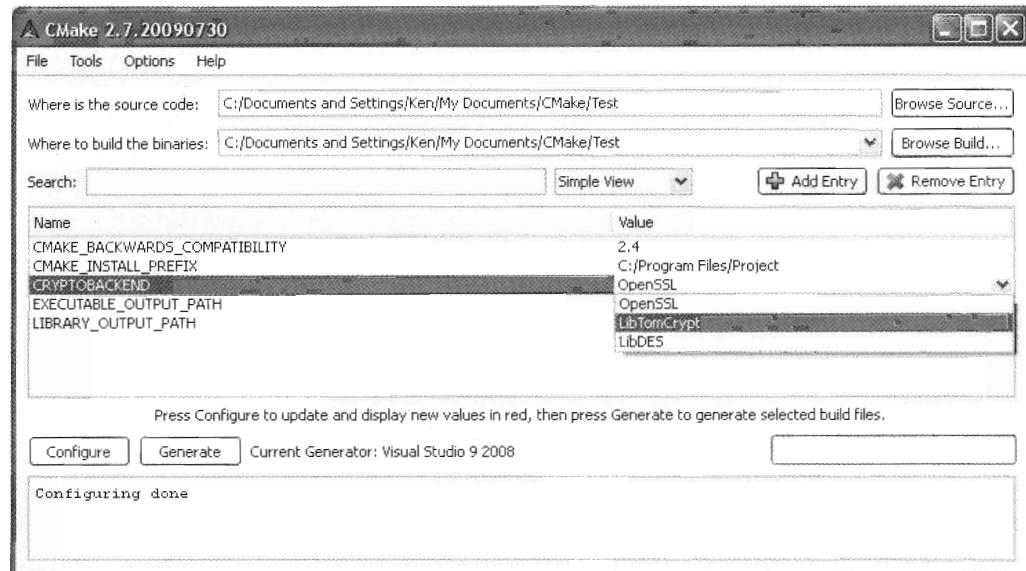


Figure 6 – Cache Value Options in cmake-gui

A few final points should be made concerning variables and their interaction with the cache. If a variable is in the cache, it can still be overridden in a CMakeLists file using the `set` command without the `CACHE` option. Cache values are checked only if the variable is not found in the current `CMakefile` instance before CMakeLists file processing begins. The `set` command will set the variable for processing the current CMakeLists file (and subdirectories as usual) without changing the value in the cache.

```
# assume that FOO is set to ON in the cache

set (FOO OFF)
# sets foo to OFF for processing this CMakeLists file
# and subdirectories; the value in the cache stays ON
```

Once a variable is in the cache, its "cache" value cannot normally be modified from a CMakeLists file. The reasoning behind this is that once CMake has put the variable into the cache with its initial value, the user may then modify that value from the GUI. If the next invocation of CMake overwrote their change back to the `set` value, the user would never be able to make a change that CMake wouldn't overwrite. So a `set (FOO ON CACHE BOOL "doc")` command will typically only do something when the cache doesn't have the variable in it. Once the variable is in the cache, that command will have no effect.

In the rare event that you really want to change a cached variable's value you can use the FORCE option in combination with the CACHE option to the `set` command. The FORCE option will cause the `set` command to override and change the cache value of a variable.

3.6 Build Configurations

Build configurations allow a project to be built in different ways for debug, optimized, or any other special set of flags. CMake supports, by default, Debug, Release, MinSizeRel, and RelWithDebInfo configurations. Debug has the basic debug flags turned on. Release has the basic optimizations turned on. MinSizeRel has the flags that produce the smallest object code, but not necessarily the fastest code. RelWithDebInfo builds an optimized build with debug information as well.

CMake handles the configurations in slightly different ways depending on what generator is being used. The conventions of the native build system are followed when possible. This means that configurations impact the build in different ways when using Makefiles versus using Visual Studio project files.

The Visual Studio IDE supports the notion of Build Configurations. A default project in Visual Studio usually has Debug and Release configurations. From the IDE you can select build Debug, and the files will be built with Debug flags. The IDE puts all of the binary files into directories with the name of the active configuration. This brings about an extra complexity for projects that build programs that need to be run as part of the build process from custom commands. See the `CMAKE_CFG_INTDIR` variable and the custom commands section for more information about how to handle this issue. The variable `CMAKE_CONFIGURATION_TYPES` is used to tell CMake which configurations to put in the workspace.

With Makefile based generators, only one configuration can be active at the time CMake is run, and it is specified by the `CMAKE_BUILD_TYPE` variable. If the variable is empty then no flags are added to the build. If the variable is set to the name of a configuration, then the appropriate variables and rules (such as `CMAKE_CXX_FLAGS_<ConfigName>`) are added to the compile lines. Makefiles do not use special configuration subdirectories for object files. To build both debug and release trees, the user is expected to create multiple build directories using the out of source build feature of CMake, and to set the `CMAKE_BUILD_TYPE` to the desired selection for each build. For example,

```
# With source code in the directory MyProject
# to build MyProject-debug create that directory, cd into it and
(ccmake .. /MyProject -DCMAKE_BUILD_TYPE:STRING=Debug)
# the same idea is used for the release tree MyProject-release
(ccmake .. /MyProject -DCMAKE_BUILD_TYPE:STRING=Release)
```

Writing CMakeLists Files

This chapter will cover the basics of writing effective CMakeLists files for your software. It will cover all of the basic commands and issues you will need to handle most projects. It will also discuss how to convert existing UNIX or Windows projects into CMakeLists files. While CMake can handle extremely complex projects, for most projects you will find this chapter's contents will tell you all you need to know. CMake is driven by the CMakeLists.txt files written for a software project. The CMakeLists files determine everything from what options to put into the cache, to what source files to compile. In addition to discussing how to write a CMakeLists file this chapter will also cover how to make them robust and maintainable. The basic syntax of a CMakeLists.txt file and key concepts of CMake have already been discussed in chapters 2 and 3. This chapter will expand on those concepts and introduce a few new ones.

4.1 CMake Syntax

CMakeLists files follow a simple syntax consisting of comments, commands, and white space. A comment is indicated using the # character and runs from that character until the end of the line. A command consists of the command name, opening parenthesis, white space separated arguments and a closing parenthesis. All white space (spaces, line feeds, tabs) are ignored except to separate arguments. Anything within a set of double quotes is treated as one argument as is typical for most languages. The backslash can be used to escape characters preventing the normal interpretation of them. The subsequent examples in this chapter will help to clear up some of these syntactic issues. You might wonder why CMake decided to have its own language instead of using an existing one such as Python, Java, or Tcl. The main reason is that we did not want to make CMake require an additional tool to run. By requiring one of these other languages all users of CMake would be required to have that language installed, and potentially a specific version of that language. This is on top of the language

extensions that would be required to do some of the CMake work, for both performance and capability reasons.

4.2 Basic Commands

While the previous chapters have already introduced many of the basic commands for CMakeLists files, this chapter will review and expand on them. The first command the top-level CMakeLists file should have is the `PROJECT` command. This command both names the project and optionally specifies what languages will be used by it. Its syntax is as follows:

```
project (projectname [CXX] [C] [Java] [NONE])
```

If no languages are specified then CMake defaults to supporting C and C++. If the `NONE` language is passed then CMake includes no language specific support. Whenever C++ language support is specified then C language support will also be loaded.

For each project command that appears in a project, CMake will create a top level IDE project file. The project will contain all targets that are in the `CMakeLists.txt` file, and any of its subdirectories as specified by the `add_subdirectory` command. If the `EXCLUDE_FROM_ALL` option is used in the `add_subdirectory` command, then the generated project will not appear in the top level Makefile or IDE project file. This is useful for generating sub projects that do not make sense as part of the main build process. Consider that a project with a number of examples could use this feature to generate the build files for each example with one run of CMake, but not have the examples built as part of the normal build process.

The `set` command is probably one of the most used commands since it is used for defining and modifying variables and lists. Complimenting the `set` command are the `remove` and `separate_arguments` commands. The `remove` command can be used to remove a value from a variable list, while the `separate_arguments` command can be used to take a single variable value (as opposed to a list) and break it into a list based on spaces.

The `add_executable` and `add_library` commands are the main commands for defining what libraries and executables to build, and what source files comprise them. For Visual Studio projects the source files will show up in the IDE as usual, but any header files the project uses will not be there. To have the header files show up as well you simply add them to the list of source files for the executable or library. This can be done for all generators. Any generators that do not use the header files directly (such as Makefile based generators) will simply ignore them.

4.3 Flow Control

In many ways writing a CMakeLists file is like writing a program in a simple language. Like most languages CMake provides flow control structures to help you along your way. CMake provides three flow control structures;

- conditional statements (e.g. `if`)
- looping constructs (e.g. `foreach` and `while`)
- procedure definitions (e.g. `macro` and `function`)

First we will consider the `if` command. In many ways the `if` command in CMake is just like the `if` command in any other language. It evaluates its expression and based on that either executes the code in its body or optionally the code in the `else` clause. For example:

```
if (FOO)
    # do something here
else (FOO)
    # do something else
endif (FOO)
```

One difference you might notice is that the conditional of the `if` statement is repeated in the `else` and `endif` clauses. This is optional and in this book you will see examples of both styles. You could just as well choose to write:

```
if (FOO)
    # do something here
else ()
    # do something else
endif ()
```

When you include conditionals in the `else` and `endif` clause they are used to provide additional error checking. As such they must exactly match the original conditional of the `if` statement. The following code would not work:

```
set (FOO 1)

if (${FOO})
    # do something
endif (1)
# ERROR, it doesn't match the original if conditional
```

Fortunately CMake provides verbose error messages in the case where an `if` statement is not properly matched with an `endif`. This should help you to track down any problems with matching conditionals. Providing the conditionals on the `else` and `endif` commands also has the added benefit of helping to document your CMakeLists file. With a long `if` statement it can be easy to lose track of what `if` statement the `endif` is closing. `if` statements can be nested to any depth, and any command can be used inside of an `if` or `else` clause.

As with many other languages, CMake supports `elseif` so that you can sequentially test for multiple conditions. For example:

```
if (MSVC80)
    # do something here
elseif (MSVC90)
    # do something else
elseif (APPLE)
    # do something else
endif ()
```

The `if` command has a limited set of operations that you can use. It does not support general purpose C style expressions such as `${FOO} && ${BAR} || ${FUBAR}`, instead it supports a limited subset of expressions that should work for most cases. Specifically `if` supports:

if (variable)

True if the variable's value is not empty, 0, FALSE, OFF, or NOTFOUND.

if (NOT variable)

True if the variable's value is empty, 0, FALSE, OFF, or NOTFOUND

if (variable1 AND variable2)

True if both variables would be considered true individually.

if (variable1 OR variable2)

True if either variable would be considered true individually.

if (COMMAND command-name)

True if the given name is a command that can be invoked.

if (DEFINED variable)

True if the given variable has been set, regardless of what value it was set to.

```
if (EXISTS file-name)
if (EXISTS directory-name)
```

True if the named file or directory exists.

```
if (IS_DIRECTORY name)
if (IS_ABSOLUTE name)
```

True if the given name is a directory, or absolute path respectively.

```
if (name1 IS_NEWER_THAN name2)
```

True if the file specified by name1 has a more recent modification time than the file specified by name2.

```
if (variable MATCHES regex)
```

```
if (string MATCHES regex)
```

True if the given string or variable's value matches the given regular expression.

Options such as EQUAL, LESS, and GREATER are available for numeric comparisons. STRLESS, STREQUAL, and STRGREATER can be used for lexicographic comparisons. VERSION_LESS, VERSION_EQUAL, and VERSION_GREATER can be used to compare versions of the form major[.minor[.patch[.tweak]]]. Similar to C and C++ these expressions can be combined to create more powerful conditionals. For example consider the following conditionals:

```
if ((1 LESS 2) AND (3 LESS 4))
    message ("sequence of numbers")
endif ()

if (1 AND 3 AND 4)
    message ("series of true values")
endif (1 AND 3 AND 4)

if (NOT 0 AND 3 AND 4)
    message ("a false value")
endif (NOT 0 AND 3 AND 4)

if (0 OR 3 AND 4)
    message ("or statements")
endif (0 OR 3 AND 4)

if (EXISTS ${PROJECT_SOURCE_DIR}/Help.txt AND COMMAND IF)
    message ("Help exists")
endif (EXISTS ${PROJECT_SOURCE_DIR}/Help.txt AND COMMAND IF)
```

```

set (fooba 0)

if (NOT DEFINED foobar)
    message ("foobar is not defined")
endif (NOT DEFINED foobar)

if (NOT DEFINED fooba)
    message ("fooba not defined")
endif (NOT DEFINED fooba)

```

In compound if statements there is an order of precedence that specifies the order that the operations will be evaluated. For example, in the statement below, the NOT will be evaluated first then the AND, not the other way around. Thus the statement will be false and the message never printed. Had the AND been evaluated first the statement would be true.

```

if (NOT 0 AND 0)
    message ("This line is never executed")
endif (NOT 0 AND 0)

```

CMake defines the order of operations such that parenthetical groups are evaluated first, then EXISTS, COMMAND, DEFINED and similar prefix operators are evaluated, then any EQUAL, LESS, GREATER, STREQUAL, STRLESS, STRGREATER, and MATCHES operators. The NOT operators are evaluated next, and finally the AND and OR expressions will be evaluated. With operations that have the same level of precedence, such as AND and OR, they will be evaluated from left to right. Once all of the expressions have been evaluated the final result will be tested to see if it is true or false. CMake considers any of the following values to be true: ON, 1, YES, TRUE, Y. The following values are all considered to be false: OFF, 0, NO, FALSE, N, NOTFOUND, -*NOTFOUND, IGNORE. This test is case insensitive so true, True, and TRUE are all treated the same.



Now let us consider the other flow control commands. The foreach, while, macro, and function commands are the best way to reduce the size of your CMakeLists files and keep them maintainable. The foreach command enables you to execute a group of CMake commands repeatedly on the members of a list. Consider the following example adapted from VTK:

```

foreach (tfile
    TestAnisotropicDiffusion2D
    TestButterworthLowPass
    TestButterworthHighPass
    TestCityBlockDistance

```

```
    TestConvolve
)
add_test(${tfile}-image ${VTK_EXECUTABLE}
${VTK_SOURCE_DIR}/Tests/rtImageTest.tcl
${VTK_SOURCE_DIR}/Tests/${tfile}.tcl
-D ${VTK_DATA_ROOT}
-V Baseline/Imaging/${tfile}.png
-A ${VTK_SOURCE_DIR}/Wrapping/Tcl
)
endforeach ( tfile )
```

The first argument of the `foreach` command is the name of the variable that will take on a different value with each iteration of the loop. The remaining arguments are the list of values over which to loop. In this example the body of the `foreach` loop is just one CMake command, `add_test`. In the body of the `foreach` loop any time the loop variable (`tfile` in this example) is referenced it will be replaced with the current value from the list. In the first iteration, occurrences of `${tfile}` will be replaced with `TestAnisotropicDiffusion2D`. In the next iteration, `${tfile}` will be replaced with `TestButterworthLowPass`. The `foreach` loop will continue to loop until all of the arguments have been processed.

It is worth mentioning that `foreach` loops can be nested and that the loop variable is replaced prior to any other variable expansion. This means that in the body of a `foreach` loop you can construct variable names using the loop variable. In the code below the loop variable `tfile` is expanded, and then concatenated with `_TEST_RESULT`. That new variable name is then expanded and tested to see if it matches `FAILED`.

```
if ( ${${tfile}}_TEST_RESULT} MATCHES FAILED)
  message ("Test ${tfile} failed.")
endif ()
```

The `while` command provides for looping based on a test condition. The format for the test expression in the `while` command is the same as that for the `if` command described earlier. Consider the following example, which is used by CTest. Note that CTest updates the value of `CTEST_ELAPSED_TIME` internally.

```
#####
# run paraview and ctest test dashboards for 6 hours
#
while (${CTEST_ELAPSED_TIME} LESS 36000)
  set (START_TIME ${CTEST_ELAPSED_TIME})
  ctest_run_script ( "dash1_ParaView_vs71continuous.cmake" )
  ctest_run_script ( "dash1_cmake_vs71continuous.cmake" )
```

```
endwhile ()
```

The `foreach` and `while` commands allow you to handle repetitive tasks that occur in sequence, whereas the `macro` and `function` commands support repetitive tasks that may be scattered throughout your CMakeLists files. Once a macro or function is defined it can be used by any CMakeLists files processed after its definition.

A function in CMake is very much like a function in C or C++. You can pass arguments into it, and the arguments passed in become variables within the function. Likewise some standard variables such as `ARGC`, `ARGV`, `ARGN`, and `ARGV0`, `ARGV1`, etc are defined. Within a function you are in a new variable scope, much like when you drop into a subdirectory using the `add_subdirectory` command you are in a new variable scope. All the variables that were defined when the function was called are still defined, but any changes to variables or new variables only exist within the function. When the function returns those variables will go away. Put more simply, when you invoke a function a new variable scope is pushed and when it returns that variable scope is popped.

The first argument is the name of the function to define. All additional arguments are formal parameters to the function.

```
function(DetermineTime _time)
    # pass the result up to whatever invoked this
    set(${_time} "1:23:45" PARENT_SCOPE)
endfunction()

# now use the function we just defined
DetermineTime( current_time )

if( DEFINED current_time )
    message(STATUS "The time is now: ${current_time}")
endif()
```

Note that in this example `_time` is used to pass the name of the return variable. The `set` command is invoked with the value of `_time`, which in this example will be `current_time`. Finally the `set` command uses the `PARENT_SCOPE` option to set that variable in the parent's scope instead of the local scope.

Macros are defined and called in the same manner as functions. The main differences are that a macro does not push and pop a new variable scope, and the arguments to a macro are not treated as variables but are string replaced prior to execution. This is very much like the differences between a macro and a function in C or C++. The first argument is the name of the macro to create. All additional arguments are formal parameters to the macro[.

```
# define a simple macro

macro (assert TEST COMMENT)
    if (NOT ${TEST})
        message ("Assertion failed: ${COMMENT}")
    endif (NOT ${TEST})
endmacro (assert)

# use the macro
find_library (FOO_LIB foo /usr/local/lib)
assert ( ${FOO_LIB} "Unable to find library foo" )
```

The simple example above creates a macro called `assert`. The macro is defined to take two arguments. The first argument is a value to test and the second argument is a comment to print out if the test fails. The body of the macro is a simple `if` command with a `message` command inside of it. The macro body ends when the `endmacro` command is found. The macro can be invoked simply by using its name as if it were a command. In the above example if `FOO_LIB` was not found a message would be displayed indicating the error condition.

The `macro` command also supports defining macros that take variable argument lists. This can be useful if you want to define a macro that has optional arguments or multiple signatures. Variable arguments can be referenced using `ARGC` and `ARGV0`, `ARGV1`, etc., instead of the formal parameters. `ARGV0` represents the first argument to the macro, `ARGV1` represents the next, and so forth. You can even use a mixture of formal arguments and variable arguments, as shown in the example below.

```
# define a macro that takes at least two arguments
# (the formal arguments) plus an optional third argument

macro (assert TEST COMMENT)
    if (NOT ${TEST})
        message ("Assertion failed: ${COMMENT}")

        # if called with three arguments then also write the
        # message to a file specified as the third argument
        if (${ARGC} MATCHES 3)
            file (APPEND ${ARGV2} "Assertion failed: ${COMMENT}")
        endif (${ARGC} MATCHES 3)

    endif (NOT ${TEST})
endmacro (ASSERT)
```

```
# use the macro
find_library (FOO_LIB foo /usr/local/lib)
assert ( ${FOO_LIB} "Unable to find library foo" )
```

In this example the two required arguments are `TEST` and `COMMENT`. These required arguments can be referenced by name, as they are in this example, or they can be referenced using `ARGV0` and `ARGV1`. If you want to process the arguments as a list you can use the `ARGV` and `ARGN` variables. `ARGV` (as opposed to `ARGV0`, `ARGV1`, etc) is a list of all the arguments to the macro, while `ARGN` is a list of all the arguments after the formal arguments. Inside your macro you can use the `foreach` command to iterate over `ARGV` or `ARGN` as desired.

CMake has two commands for interrupting the processing flow. The `break` command will break out of a `foreach` or `while` loop before it would normally end. The `return` command will return from a function or listfile before the function or listfile has reached its end.

4.4 Regular Expressions

A few CMake commands, such as `if` and `string`, make use of regular expressions, or can take a regular expression as an argument. In its simplest form, a regular-expression is a sequence of characters used to search for exact character matches. However, many times the exact sequence to be found is not known, or only a match at the beginning or end of a string is desired. Since there are several different conventions for specifying regular expressions CMake's standard is described below. The description is based on the open source regular expression class from Texas Instruments, which is used by CMake for parsing regular expressions.

Regular expressions can be specified by using combinations of standard alphanumeric characters and the following regular expression meta-characters:

- ^ Matches at beginning of a line or string.
- \$ Matches at end of a line or string.
- . Matches any single character other than a newline.
- [] Matches any character(s) inside the brackets.
- [^] Matches any character(s) not inside the brackets.
- [-] Matches any character in range on either side of a dash.
- * Matches preceding pattern zero or more times.

- + Matches preceding pattern one or more times.
- ? Matches preceding pattern zero or once only.
- () Saves a matched expression and uses it in a later replacement.
- (|) Matches either the left or right side of the bar.

Note that more than one of these meta-characters can be used in a single regular expression in order to create complex search patterns. For example, the pattern [^ab1-9] indicates to match any character sequence that does not begin with the characters "a" or "b" or numbers in the series one through nine. The following examples may help clarify regular expression usage:

- The regular expression "^hello" matches a "hello" only at the beginning of a search string. It would match "hello there", but not "hi,\nhello there".
- The regular expression "long\$" matches a "long" only at the end of a search string. It would match "so long", but not "long ago".
- The regular expression "t..t..g" will match anything that has a "t", then any two characters, another "t", any two characters, and then a "g". It would match "testing" or "test again", but would not match "toasting".
- The regular expression "[1-9ab]" matches any number one through nine, and the characters "a" and "b". It would match "hello 1" or "begin", but would not match "no-match".
- The regular expression "[^1-9ab]" matches any character that is not a number one through nine, or an "a" or "b". It would NOT match "lab2" or "b2345a", but would match "no-match".
- The regular expression "br* " matches something that begins with a "b", is followed by zero or more "r"s, and ends in a space. It would match "brrrrr " and "b ", but would not match "brrh ".
- The regular expression "br+ " matches something that begins with a "b", is followed by one or more "r"s, and ends in a space. It would match "brrrrr ", and "br ", but would not match "b " or "brrh ".
- The regular expression "br? " matches something that begins with a "b", is followed by zero or one "r"s, and ends in a space. It would match "br ", and "b ", but would not match "brrr " or "brrh ".
- The regular expression "(..p)b" matches something ending with pb and beginning with whatever the two characters before the first p encountered in the line were. It would find "repb" in "rep drepaqrepb". The regular expression "(..p)a" would find "repa qrepb" in "rep drepa qrepb"

- The regular expression "d(..p)" matches something ending with p, beginning with d, and having two characters in between that are the same as the two characters before the first p encountered in the line. It would match "drepap qrepb" in "rep drepa qrepb".

4.5 Checking Versions of CMake

CMake is an evolving program and as new versions are released, new features or commands may be introduced. As a result, there may be instances where you might want to use a command that is in a current version of CMake but not in previous versions. There are a couple of ways to handle this. One option is to use the `if` command to check whether a new command exists. For example:

```
# test if the command exists

if (COMMAND some_new_command)
    # use the command
    some_new_command ( ARGS...)
endif (COMMAND some_new_command)
```

The above approach should work in most cases, but if you need more information you can test against the actual version of CMake that is being run by evaluating the `CMAKE_VERSION` variables, as in the following example:

```
# look for newer versions of CMake

if (${CMAKE_VERSION} VERSION_GREATER 1.6.1)
    # do something special here
endif ()
```

When writing your CMakeLists files you might decide that you do not want to support old versions of CMake. To do this you can place the following command at the top of your CMakeLists file:

```
cmake_minimum_required (VERSION 2.2)
```

This indicates that the person running CMake on your project must have at least CMake version 2.2. If they are running an older version of CMake then an error message will be displayed telling them that the project requires at least the specified version of CMake.

Finally, in some cases a new release of CMake might come out that no longer supports some commands you were using (although we try to avoid this). In these cases you can use CMake policies, as discussed in section 4.7.

4.6 Using Modules

Code reuse is a valuable technique in software development and CMake has been designed to support it. Allowing CMakeLists files to make use of reusable modules enables the entire CMake community to share reusable sections of code. For CMake these sections of code are called modules and can be found in the Modules subdirectory of your CMake installation. Modules are simply sections of CMake commands put into a file. They can then be included into other CMakeLists files using the `include` command. For example, the following commands will include the `FindTCL` module from CMake and then add the Tcl library to the target FOO.

```
include (FindTCL)
target_link_libraries (FOO ${TCL_LIBRARY})
```

A module's location can be specified using the full path to the module file, or by letting CMake find the module by itself. CMake will look for modules in the directories specified by `CMAKE_MODULE_PATH` and if it cannot find it there, it will look in the Modules subdirectory of CMake. This way projects can override modules that CMake provides, to customize them for their needs. Modules can be broken into a few main categories:

Find Modules

These modules determine the location of software elements such as header files or libraries.

System Introspection Modules

These modules test the system for properties such as the size of a float, support for ANSI C++ streams, etc.

Utility Modules

These modules provide added functionality such as support for situations where one CMake project depends on another and other convenience routines.

Now let us consider these three types of modules in more detail. CMake includes a large number of Find modules. The purpose of a Find module is to locate software elements such as header or library files. If they cannot be found then they provide a cache entry so that the user can set the required properties. Consider the following module that finds the PNG library.

```
#  
# Find the native PNG includes and library  
#  
  
# This module defines  
# PNG_INCLUDE_DIR, where to find png.h, etc.  
# PNG_LIBRARIES, the libraries to link against to use PNG.  
# PNG_DEFINITIONS - You should call  
# add_definitions (${PNG_DEFINITIONS}) before compiling code  
# that includes png library files.  
# PNG_FOUND, If false, do not try to use PNG.  
  
# also defined, but not for general use are  
# PNG_LIBRARY, where to find the PNG library.  
  
# None of the above will be defined unless zlib can be found.  
  
# PNG depends on Zlib  
include (FindZLIB.cmake )  
  
if (ZLIB_FOUND)  
    find_path (PNG_PNG_INCLUDE_DIR png.h  
    /usr/local/include  
    /usr/include  
    )  
  
    find_library (PNG_LIBRARY png  
    /usr/lib  
    /usr/local/lib  
    )  
  
if (PNG_LIBRARY)  
    if (PNG_PNG_INCLUDE_DIR)  
        # png.h includes zlib.h. Sigh.  
        set (PNG_INCLUDE_DIR  
            ${PNG_PNG_INCLUDE_DIR} ${ZLIB_INCLUDE_DIR} )  
        set (PNG_LIBRARIES ${PNG_LIBRARY} ${ZLIB_LIBRARY})  
        set (PNG_FOUND "YES")  
  
        if (CYGWIN)  
            if (BUILD_SHARED_LIBS)  
                # No need to define PNG_USE_DLL here, because  
                # it's default for Cygwin.  
            else (BUILD_SHARED_LIBS)
```

```
    set (PNG_DEFINITIONS -DPNG_STATIC)
endif (BUILD_SHARED_LIBS)
endif (CYGWIN)

endif ()
endif ()

endif ()
```

The top of the module clearly documents what the module will do and what variables it will set. Next it includes another module, the `FindZLIB` module, that determines if the ZLib library is installed. Next, if ZLib is found, the `find_path` command is used to locate the PNG include files. The first argument is the name of the variable to store the result in, the second argument is the name of the header file to look for, the remaining arguments are paths to search for the header file. If it is not found in the system path then the variable is set to `PNG_PNG_INCLUDE_DIR-NOTFOUND`, allowing the user to set it.

Note that the paths to search for the PNG library can include hard coded directories, registry entries, and directories made up of other CMake variables. The next command finds the actual PNG library using the `find_library` command. This command performs additional checks to find a proper library name, such as adding "lib" in front of the name and ".so" at the end of the name on Linux systems.

After the find calls, some CMake variables are set that developers using `FindPNG` can use in their projects (such as the include paths, and library name). Finally `PNG_FOUND` is set correctly, which lets developers know that the PNG library was properly found.

This structure is fairly common to all Find modules in CMake. Usually they are fairly short, but in some cases, such as `FindOpenGL` they can be a few pages long. They are normally independent of other modules, but there is no restriction on the use of other modules.

System introspection modules provide information about the target platform or compiler. Many of these modules have names prefixed with `Test` or `Check`, such as `TestBigEndian` and `CheckTypeSize`. Many of the system introspection modules actually try to compile code in order to determine the correct result. In these cases the source code is usually named the same as the module, but with a `.c` or `.cxx` extension. System introspection modules are covered in more detail in chapter 5.

CMake includes a few Utility modules to help make using CMake a little easier. `CMakeExportBuildSettings` and `CMakeImportBuildSettings` provide tools to help verify that two C++ projects are compiled with the same compiler and key flags. The `CMakePrintSystemInformation` module prints out a number of key CMake settings to aid in debugging.

Using CMake with SWIG

One example of how modules can be used is to look at wrapping your C/C++ code in another language using SWIG. SWIG (Simplified Wrapper and Interface Generator) www.swig.org is a tool that reads annotated C/C++ header files, and creates wrapper code (glue code) in order to make the corresponding C/C++ libraries available to other programming languages such as Tcl, Python, or Java. CMake supports SWIG with the `find_package` command. Although SWIG can be used from CMake using custom commands, the SWIG package provides several macros that make building SWIG projects with CMake simpler. To use the SWIG macros, first you must call the `find_package` command with the name SWIG. Then you need to include the file referenced by the variable `SWIG_USE_FILE`. This will define several macros and set up CMake to easily build SWIG based projects.

Two very useful macros are `SWIG_ADD_MODULE` and `SWIG_LINK_LIBRARIES`. `SWIG_ADD_MODULE` works much like the `add_library` command in CMake. The command is invoked like this:

```
SWIG_ADD_MODULE (module_name language source1 source2 ... sourceN)
```

The first argument is the name of the module to create. The next argument is the target language SWIG is producing a wrapper for. The rest of the arguments consist of a list of source files used to create the shared module. The big difference is that SWIG .i interface files can be used directly as sources. The macro will create the correct custom commands to run SWIG, and generate the C or C++ wrapper code from the SWIG interface files. The sources can also be regular C or C++ files that need to be compiled in with the wrappers.

The `SWIG_LINK_LIBRARIES` macro is used to link support libraries to the module. This macro is used because depending on the language being wrapped by SWIG, the name of the module may be different. The actual name of the module is stored in a variable called `SWIG_MODULE_${name}_REAL_NAME` where `${name}` is the name passed into the `SWIG_ADD_MODULE` macro. For example, `SWIG_ADD_MODULE(foo tcl foo.i)` would create a variable called `SWIG_MODULE_foo_REAL_NAME` which would contain the name of the actual module created.

Now consider the following example that uses the SWIG example found in SWIG under Examples/python/class.

```
# Find SWIG and include the use swig file
find_package (SWIG REQUIRED)
include (${SWIG_USE_FILE})

# Find python library and add include path for python headers
find_package (PythonLibs)
include_directories (${PYTHON_INCLUDE_PATH})
```

```
# set the global swig flags to empty
set (CMAKE_SWIG_FLAGS "")

# let swig know that example.i is c++ and add the -includeall
# flag to swig
set_source_files_properties (example.i PROPERTIES CPLUSPLUS ON)
set_source_files_properties (example.i
                           PROPERTIES SWIG_FLAGS "-includeall")

# Create the swig module called example
# using the example.i source and example.cxx
# swig will be used to create wrap_example.cxx from example.i
SWIG_ADD_MODULE (example python example.i example.cxx)
SWIG_LINK_LIBRARIES (example ${PYTHON_LIBRARIES})
```

This example first uses `find_package` to locate SWIG. Next it includes the `SWIG_USE_FILE` defining the SWIG CMake macros. Then it finds the Python libraries and sets up CMake to build with the Python library. Notice that the SWIG input file `example.i` is used like any other source file in CMake, and properties are set on the file telling SWIG that the file is C++ and that the SWIG flag `-includeall` should be used when running SWIG on that source file. The module is created by telling SWIG the name of the module, the target language and the list of source files. Finally, the Python libraries are linked to the module.

Using CMake with Qt

Projects using the popular widget toolkit Qt from Nokia, qt.nokia.com, can be built with CMake. CMake supports multiple versions of Qt, including versions 3 and 4. The first step is to tell CMake what version(s) of Qt to look for. Many Qt applications are designed to work with Qt3 or Qt4, but not both. If your application is designed for Qt4 then you can use the `FindQt4` module, for Qt3 you should use the `FindQt3` module. If your project can work with either version of Qt then you can use the generic `FindQt` module. All of the modules provide helpful tools for building Qt projects. The following is a simple example of building a project that uses Qt4.

```
find_package ( Qt4 )

if (QT4_FOUND)
    include (${QT_USE_FILE})

    # what are our ui files?
    set (QTUI_SRCS qtwrapping.ui)
    QT4_WRAP_UI (QTUI_H_SRCS ${QTUI_SRCS})
    QT4_WRAP_CPP (QT_MOC_SRCS TestMoc.h)
```

```
add_library (myqtlb ${QTUI_H_SRCS} ${QT_MOC_SRCS})
target_link_libraries (myqtlb ${QT_LIBRARIES} )

add_executable (qtwrapping qtwrappingmain.cxx)
target_link_libraries (qtwrapping myqtlb)

endif (QT4_FOUND)
```

Using CMake with FLTK

CMake also supports the The Fast Light Toolkit (FLTK) with special FLTK CMake commands. The `FLTK_WRAP_UI` command is used to run the fltk fluid program on a .fl file and produce a C++ source file as part of the build. The following example shows how to use FLTK with CMake.

```
find_package (FLTK)
if (FLTK_FOUND)
    set (FLTK_SRCS
        fltk1.fl
    )
    FLTK_WRAP_UI (wraplibFLTK ${FLTK_SRCS})
    add_library (wraplibFLTK ${wraplibFLTK_UI_SRCS} )
endif (FLTK_FOUND)
```

4.7 Policies

For various reasons, sometimes a new feature or change is made to CMake that is not fully backwards compatible with older versions of CMake. This can create problems when someone tries to use an old CMakeLists file with a new version of CMake. To help both end users and developers through such issues, we have introduced policies. Policies are a mechanism in CMake to help improve backwards compatibility and track compatibility issues between different versions of CMake.

Design Goals

There were four main design goals for the CMake policy mechanism:

1. Existing projects should build with versions of CMake newer than that used by the project authors.
 - Users should not need to edit code to get the projects to build.
 - Warnings may be issued but the projects should build.

2. Correctness of new interfaces or bugs fixes in old interfaces should not be inhibited by compatibility requirements. Any reduction in correctness of the latest interface is not fair on new projects.
3. Every change made to CMake that may require changes to a project's CMakeLists files should be documented.
 - Each change should also have a unique identifier that can be referenced by warning and error messages.
 - The new behavior is enabled only when the project has somehow indicated it is supported.
4. We must be able to eventually remove code that implements compatibility with ancient CMake versions.
 - Such removal is necessary to keep the code clean and to allow for internal refactoring.
 - After such removal, attempts to build projects written for ancient versions must fail with an informative message.

All policies in CMake are assigned a name of the form CMPNNNN where NNNN is an integer value. Policies typically support both an old behavior that preserves compatibility with earlier versions of CMake, and a new behavior that is considered correct and preferred for use by new projects. Every policy has documentation detailing the motivation for the change, and the old and new behaviors

Setting Policies

Projects may configure the setting of each policy to request old or new behavior. When CMake encounters user code that may be affected by a particular policy it checks to see whether the project has set the policy. If the policy has been set (to OLD or NEW) then CMake follows the behavior specified. If the policy has not been set then the old behavior is used, but a warning is issued telling the project author to set the policy.

There are a couple ways to set the behavior of a policy. The quickest way is to set all policies to a version corresponding to the release version of CMake for which the project was written. Setting the policy version requests the new behavior for all policies introduced in the corresponding version of CMake or earlier. Policies introduced in later versions are marked as not set in order to produce proper warning messages. The policy version is set using the `cmake_policy` command's VERSION signature. For example, the code

```
cmake_policy (VERSION 2.6)
```

will request the new behavior for all policies introduced in CMake 2.6 or earlier. The `cmake_minimum_required` command will also set the policy version, which is convenient for use at the top of projects. A project should typically begin with the lines

```
cmake_minimum_required (VERSION 2.6)
project (MyProject)
# ...code using CMake 2.6 policies
```

Of course one should replace "2.6" with whatever version of CMake you are currently writing to. You can also set each policy individually if you wish. This is sometimes helpful for project authors who want to incrementally convert their projects to use the new behavior, or silence warnings about dependence on old behavior. The `cmake_policy` command's `SET` option may be used to explicitly request old or new behavior for a particular policy.

For example, CMake 2.6 introduced policy `CMP0002`, which requires all logical target names to be globally unique (duplicate target names previously worked in some cases by accident but were not diagnosed). Projects using duplicate target names and working accidentally will receive warnings referencing the policy. The warnings may be silenced by the code

```
cmake_policy (SET CMP0002 OLD)
```

which explicitly tells CMake to use the old behavior for the policy (silently accept duplicate target names). Another option is to use the code

```
cmake_policy (SET CMP0002 NEW)
```

to explicitly tell CMake to use new behavior and produce an error when a duplicate target is created. Once this is added to the project it will not build until the author removes any duplicate target names.

When a new version of CMake is released that introduces new policies it will still build old projects, because by default they do not request NEW behavior for any of the new policies. When starting a new project one should always specify the most recent release of CMake to be supported as the policy version level. This will make sure that the project is written to work using policies from that version of CMake and not using any old behavior. If no policy version is set CMake will warn and assume a policy version of 2.4. This allows existing projects that do not specify `cmake_minimum_required` to build as they would have with CMake 2.4.

The Policy Stack

Policy settings are scoped using a stack. A new level of the stack is pushed when entering a new subdirectory of the project (with `add_subdirectory`) and popped when leaving it. Therefore setting a policy in one directory of a project will not affect parent or sibling directories, but will affect subdirectories.

This is useful when a project contains subprojects maintained separately but built inside the tree. The top-level CMakeLists file in a project may write

```
cmake_policy (VERSION 2.6)
project (MyProject)
add_subdirectory (OtherProject)
# ... code requiring new behavior as of CMake 2.6 ...
```

while the OtherProject/CMakeLists.txt file contains

```
cmake_policy (VERSION 2.4)
project (OtherProject)
# ... code that builds with CMake 2.4 ...
```

This allows a project to be updated to CMake 2.6 while subprojects, modules, and included files continue to build with CMake 2.4 until their maintainers update them.

User code may use the `cmake_policy` command to push and pop its own stack levels as long as every push is paired with a pop. This is useful to temporarily request different behavior for a small section of code. For example, policy `CMP0003` removes extra link directories that used to be included when new behavior is used. While incrementally updating a project it may be difficult to build a particular target with the new behavior but all other targets are okay. The code

```
cmake_policy (PUSH)
cmake_policy (SET CMP0003 OLD) # use old-style link for now
add_executable (myexe ...)
cmake_policy (POP)
```

will silence the warning and use the old behavior for that target. You can get a list of policies and help on specific policies by running `cmake` from the command line as follows

```
cmake --help-command cmake_policy
cmake --help-policies
cmake --help-policy CMP0003
```

Updating a Project For a New Version of CMake

When a CMake release introduces new policies it may generate warnings for some existing projects. These warnings indicate that changes to the project may need to be made to deal correctly with the new policies. While old releases of the project can continue to build with

the warnings the project development tree should be updated to take the new policies into account. There are two approaches to updating a tree: one-shot and incremental. Which one is easier depends on the size of the project and what new policies produce warnings.

The One-Shot Approach

The simplest approach to updating a project for a new version of CMake is simply to change the policy version set at the top of the project, try building with the new CMake version, and fix problems. For example, to update a project to build with CMake 2.8 one might write

```
cmake_minimum_required(VERSION 2.8)
```

at the beginning of the top-level CMakeLists file. This tells CMake to use the new behavior for every policy introduced in CMake 2.8 and below. When building this project with CMake 2.8 no warnings will be produced about policies because it knows of no policies introduced in later versions. However, if the project was depending on the old behavior of a policy it may not build since CMake now uses the new behavior without warning. It is up to the project author who added the policy version line to fix these issues.

The Incremental Approach

Another approach to updating a project for a new version of CMake is to deal with each warning one-by-one. One advantage of this approach is that the project will continue to build throughout the process, so the changes can be made incrementally.

When CMake encounters a situation where it needs to know whether to use the old or new behavior for a policy, it checks whether the project has set the policy. If the policy is set CMake silently uses the corresponding behavior. If the policy is not set, CMake uses the old behavior but warns that the policy is not set.

In many cases the warning message will point at the exact line of code in the CMakeLists files that caused the warning. In some cases the situation cannot be diagnosed until CMake is generating the native build system rules for the project, so the warning will not include explicit context information. In these cases CMake will try to provide some information about where code may need to be changed. The documentation for these "generation-time" policies should indicate the point in the project code at which the policy should be set to take effect.

In order to incrementally update a project one warning should be addressed at a time. Several cases may occur as described below.

Silence a Warning When the Code is Correct

Many policy warnings may be produced simply because the project has not set the policy even though the project may work correctly with the new behavior (there is no way for CMake to know the difference). For a warning about some policy `CMP<NNNN>` one may check whether this is the case by adding

```
cmake_policy (SET CMP<NNNN> NEW)
```

to the top of the project and trying to build it. If the project builds correctly with the new behavior one may move on to the next policy warning. If the project does not build correctly one of the other cases may apply.

Silence a Warning Without Updating the Code

One may suppress all instances of a warning `CMP<NNNN>` by adding

```
cmake_policy (SET CMP<NNNN> OLD)
```

at the top of a project. However, we encourage project authors to update their code to work with the new behavior for all policies. This is especially important because versions of CMake in the (distant) future may remove support for the old behavior and produce an error for projects requesting it (which tells the user to get an older CMake to build the project).

Silence a Warning by Updating Code

When a project does not work correctly with the NEW behavior for a policy its code needs to be updated. In order to deal with a warning for some policy `CMP<NNNN>` one may add

```
cmake_policy (SET CMP<NNNN> NEW)
```

at the top of the project and then fix the code to work with the NEW behavior.

If many instances of the warning occur fixing all of them simultaneously may be too difficult. Instead a developer may fix one at a time. This may be done using the PUSH/POP signatures of the `cmake_policy` command:

```
cmake_policy (PUSH)
cmake_policy (SET CMP<NNNN> NEW)
# ... code updated for new policy behavior ...
cmake_policy (POP)
```

This will request the new behavior for a small region of code that has been fixed. Other instances of the policy warning may still appear and must be fixed separately.

Updating the Project Policy Version

After addressing all policy warnings and getting the project to build cleanly with the new CMake version one step remains. The policy version set at the top of the project should now be updated to match the new CMake version, just as in the one-shot approach above. For

example, after updating a project to build cleanly with CMake 2.8 one may update the top of the project with the line

```
cmake_minimum_required(VERSION 2.8)
```

This will set all policies introduced in CMake 2.8 or below to use the new behavior. Then one may sweep through the rest of the code and remove all the calls to the `cmake_policy` command used to request the new behavior incrementally. The end result should look the same as the one-shot approach above but could be attained step-by-step.

Supporting Multiple CMake Versions

Some projects might want to support a few releases of CMake simultaneously. The goal is to build with an older version but also work with newer versions without warnings. In order to support both CMake 2.4 and 2.6, one may write code like

```
cmake_minimum_required (VERSION 2.4)
if (COMMAND cmake_policy)
    # policy settings ...
    cmake_policy (SET CMP0003 NEW)
endif (COMMAND cmake_policy)
```

This will set the policies when building with CMake 2.6 and just ignore them for CMake 2.4. In order to support both CMake 2.6 and some policies of CMake 2.8, one may write code like

```
cmake_minimum_required (VERSION 2.6)
if (POLICY CMP1234)
    # policies not known to CMake 2.6 ...
    cmake_policy (SET CMP1234 NEW)
endif (POLICY CMP1234)
```

This will set the policies when building with CMake 2.8 and just ignore them for CMake 2.6. If it is known that the project builds with both CMake 2.6 and CMake 2.8's new policies one may write

```
cmake_minimum_required (VERSION 2.6)
if (NOT ${CMAKE_VERSION} VERSION_LESS 2.8)
    cmake_policy (VERSION 2.8)
endif ()
```

4.8 Linking Libraries

In CMake 2.6 and later a new approach to generating link lines for targets has been implemented. Consider these libraries:

```
/path/to/libfoo.a  
/path/to/libfoo.so
```

Previously if someone wrote

```
target_link_libraries (myexe /path/to/libfoo.a)
```

CMake would generate this code to link it:

```
... -L/path/to -Wl,-Bstatic -lfoo -Wl,-Bdynamic ...
```

This worked most of the time, but some platforms (such as Mac OS X) do not support the `-Bstatic` or equivalent flag. This made it impossible to link to the static version of a library without creating a symlink in another directory and using that one instead. Now CMake will generate this code:

```
... /path/to/libfoo.a ...
```

This guarantees that the correct library is chosen. However there are some caveats to keep in mind. In the past a project could write this (incorrect) code, and it would work by accident:

```
add_executable (myexe myexe.c)  
target_link_libraries (myexe /path/to/libA.so B)
```

where "B" is meant to link `"/path/to/libB.so"`. This code is incorrect because it asks CMake to link to B but does not provide the proper linker search path for it. It used to work by accident because the `-L/path/to` would get added as part of the implementation of linking to A. The correct code would be either

```
link_directories (/path/to)  
add_executable (myexe myexe.c)  
target_link_libraries (myexe /path/to/libA.so B)
```

or even better

```
add_executable (myexe myexe.c)
target_link_libraries (myexe /path/to/libA.so /path/to/libB.so)
```

Linking to System Libraries

System libraries on UNIX-like systems are typically provided in `/usr/lib` or `/lib`. These directories are considered implicit linker search paths because linkers automatically search these locations, even without a flag like `-L/usr/lib`. Consider the code

```
find_library (M_LIB m)
target_link_libraries (myexe ${M_LIB})
```

Typically the `find_library` command would find the math library `/usr/lib/libm.so`, but some platforms provide multiple versions of libraries corresponding to different architectures. For example, on an IRIX machine one might find the libraries

```
/usr/lib/libm.so          (ELF o32)
/usr/lib32/libm.so        (ELF n32)
/usr/lib64/libm.so        (ELF 64)
```

On a Solaris machine one might find

```
/usr/lib/libm.so          (sparcv8 architecture)
/usr/lib/sparcv9/libm.so   (sparcv9 architecture)
```

Unfortunately, `find_library` may not know about all of the architecture-specific system search paths used by the linker. In fact, when it finds `/usr/lib/libm.so`, it may be finding a library with the incorrect architecture. If the link computation were to produce the line

```
... /usr/lib/libm.so ...
```

the linker might complain if `/usr/lib/libm.so` does not match the architecture it wants. One solution to this problem is for the link computation to recognize that the library is in a system directory and ask the linker to search for the library. It could produce the link line

```
... -lm ...
```

and the linker would search through its architecture-specific implicit link directories to find the correct library. Unfortunately, this solution suffers from the original problem of distinguishing between static and shared versions. In order to ask the linker to find a static system library with the correct architecture it must produce the link line

```
... -Wl,-Bstatic -lm ... -Wl,-Bshared ...
```

Since not all platforms support such flags CMake compromises. Libraries that are not in implicit system locations are linked by passing the full library path to the linker. Libraries that are in implicit system locations (such as `/usr/lib`) are linked by passing the `-l` option if a flag like `-Bstatic` is available, and by passing the full library path to the linker otherwise.

Specifying Optimized or Debug Libraries with a Target

On Windows platforms it is often required to link debug libraries with debug libraries, and optimized libraries with optimized libraries. CMake helps satisfy this requirement with the `target_link_libraries` command, which accepts an optional flag that is debug or optimized. So, if a library is preceded with either debug or optimized, then that library will only be linked in with the like configuration type. For example:

```
add_executable (foo foo.c)
target_link_libraries (foo debug libdebug optimized libopt)
```

In this case `foo` will be linked against `libdebug` if a debug build was selected, or against `libopt` if an optimized build was selected.

4.9 Shared Libraries and Loadable Modules

Shared libraries and loadable modules are very powerful tools for software developers. They can be used to create extension modules or plugins for off-the-shelf software, and can be used to decrease the compile/link/run cycles for C and C++ programs. However, despite years of use, the cross platform creation of shared libraries and modules remains a black art understood by only a few developers. CMake has the ability to aid developers in the creation of shared libraries and modules. CMake knows the correct tools and flags to use in order to produce the shared libraries for most modern operating systems that support them. Unfortunately, CMake cannot do all the work, and developers must sometimes alter source code and understand the basic concepts and common pitfalls associated with shared libraries before they can be used effectively. This section will describe many of the issues required to take advantage of shared libraries and loadable modules.

A shared library should be thought of more like an executable than a static library, and on most systems actually requires executable permissions to be set on the shared library file. This

means that shared libraries can link to other shared libraries when they are created in the same way as an executable. Unlike a static library where the atomic unit is the object file, for shared libraries, the entire library is the atomic unit. This can cause some unexpected linker errors when converting from static to shared libraries. If an object file is part of a static library, but the executable linking to the library does not use any of the symbols in that object file, then the file is simply excluded from the final linked executable. With shared libraries, all the object files that make up the library and all of the dependencies that they require come as one unit. For example, suppose you had a library with an object file defining the function `DisplayOnXWindow()` which required the X11 library. If you linked an executable to that library, but did not call the `DisplayOnXWindow()` function, the static library version would not require X11, but the shared library version would require the X11 library. This is because a shared library has to be taken as one unit, and a static library is only an archive of object files from which linkers can choose which objects are needed. This means that static linked executables can be smaller, as they only contain the object code actually used.

Another difference between shared and static libraries is library order. With static libraries the order on the link line can make a difference. This is because most linkers only use the symbols that are needed in a single pass over all the given libraries. So, the library order should go from the library that uses the most other libraries to the library that uses no other libraries. CMake will preserve and remember the order of libraries and library dependencies of a project. This means that each library in a project should use the `target_link_libraries` command to specify all of the libraries that it directly depends on. The libraries will be linked with each other for shared builds, but not static builds. However, the link information is used in static builds when executables are linked. An executable that only links library libA will get libA plus libB and libC as long as libA's dependency on libB and libC was properly specified using `target_link_libraries` (`libA libB libC`).

At this point, one might wonder why shared libraries would be preferred over static libraries. There are several reasons. First, shared libraries can decrease the compile/link/run cycle time. This is because the linker does not have to do as much work when linking to shared libraries because there are fewer decisions to be made about which object files to keep. Also, often times, the executable does not even need to be re-linked after the shared library is rebuilt. So, developers can work on a library compiling and linking only the small part of the program that is currently being developed, and then re-run the executable after each build of the shared library. Also, if a library is used by many different executables on a system, then there only needs to be one copy of the library on disk, and often in memory too.

In addition to the concept of a software library, shared libraries can also be used on many systems as run time loadable modules. This means that a program can at run time, load and execute object code that was not part of the original software. This allows developers to create software that is both open and closed. (For more information see Object Oriented Software Construction by Bertrand Meyer.) Closed software is software that cannot be modified. It has been through a testing cycle and can be certified to perform specific tasks with regression

tests. However, a seemingly opposite goal is sought after by developers of object oriented software. This is the concept of Open software that can be extended by future developers. This can be done via inheritance and polymorphism with object systems. Shared libraries that can be loaded at run time, allow for these seemingly opposing goals to be achieved in the same software package. Many common applications support the idea of plugins. The most common of these applications is the web browser. Internet Explorer uses plugins to support video over the web and 3D visualization. In addition to plugins, loadable factories can be used to replace C++ objects at run time, as is done in VTK.

Once it is decided that shared libraries or loadable modules are the right choice for a particular project, there are a few issues that developers need to be aware of. The first question that must be answered is which symbols are exported by the shared library? This may sound like a simple question, but the answer is different from platform to platform. On many, but not all UNIX systems, the default behavior is to export all the symbols much like a static library. However, on Windows systems, developers must explicitly tell the linker and compiler which symbols are to be exported and imported from shared libraries. This is often a big problem for UNIX developers moving to Windows. There are two ways to tell the compiler/linker which symbols to export/import on Windows. The most common approach is to decorate the code with a Microsoft™ C/C++ language extension. An alternative is to create an extra file called a .def file. This file is a simple ASCII file containing the names of all the symbols to be exported from a library.

The Microsoft™ extension uses the `__declspec` directive. If a symbol has `__declspec(dllexport)` in front of it, it will be exported, and if it has `__declspec(dllimport)` it will be imported. Since the same file may be shared during the creation and use of a library, it must be both exported and imported in the same source file. This can only be done with the preprocessor. The developer can create a macro called `LIBRARY_EXPORT` that is defined to `dllexport` when building the library and `dllimport` when using the library. CMake helps this process by automatically defining `${LIBNAME}_EXPORTS` when building a DLL (dynamic link library, a.k.a. a shared library) on Windows.

The following code snippet is from the VTK library `vtkCommon`, and is included by all files in the `vtkCommon` library:

```
#if defined(WIN32)

#if defined(vtkCommon_EXPORTS)
#define VTK_COMMON_EXPORT __declspec( dllexport )
#else
#define VTK_COMMON_EXPORT __declspec( dllimport )
#endif
#else
#define VTK_COMMON_EXPORT
#endif
```

The example checks for Windows and checks the `vtkCommon_EXPORTS` macro provided by CMake. So, on UNIX `VTK_COMMON_EXPORT` is defined to nothing, and on Windows during the building of `vtkCommon.dll` it is defined as `__declspec(dllexport)`, and when the file is being used by another file, it is defined to `__declspec(dllimport)`.

The second approach requires a .def file to specify the symbols to be exported. This file could be created by hand, but for a large and changing C++ library that could be time consuming and error prone. CMake's custom commands can be used to run a pre-link program that will create a .def file from the compiled object files automatically. In the following example, a Perl script called `makedef.pl` is used, the script runs the `DUMPBIN` program on the .obj files and extracts all of the exportable symbols and writes a .def file with the correct exports for all the symbols in the library `mylib`.

```
-----CMakeLists.txt-----

cmake_minimum_required (VERSION 2.6)
project (myexe)

set (SOURCES mylib.cxx mylib2.cxx)

# create a list of all the object files
string (REGEX REPLACE "\\.cxx" ".obj" OBJECTS "${SOURCES}")

# create a shared library with the .def file
add_library (mylib SHARED ${SOURCES}
    ${CMAKE_CURRENT_BINARY_DIR}/mylib.def
)
# set the .def file as generated
set_source_files_properties (
    ${CMAKE_CURRENT_BINARY_DIR}/mylib.def
    PROPERTIES GENERATED 1
)

# create an executable
add_executable (myexe myexe.cxx)

# link the executable to the dll
target_link_libraries(myexe mylib)

#convert to windows slashes
set (OUTDIR
    ${CMAKE_CURRENT_BINARY_DIR}/${CMAKE_CFG_INTDIR}
)
```

```
string (REGEN REPLACE "/" "\\" OUTDIR ${OUTDIR})\n\n# create a custom pre link command that runs\n# a perl script to create a .def file using dumpbin\nadd_custom_command (\n    TARGET mylib PRE_LINK\n    COMMAND perl\n    ARGS ${CMAKE_CURRENT_SOURCE_DIR}/makedef.pl\n    ${CMAKE_CURRENT_BINARY_DIR}\\mylib.def mylib\n    ${OUTDIR} ${OBJECTS}\n    COMMENT "Create .def file"\n)\n
```

```
---myexe.cxx---\n#include <iostream>\n#include "mylib.h\"\nint main()\n{\n    std::cout << myTen() << "\\n";\n    std::cout << myEight() << "\\n";\n}
```

```
---mylib.cxx--\nint myTen()\n{\n    return 10;\n}
```

```
---mylib2.cxx---\nint myEight()\n{\n    return 8;\n}
```

There is a significant difference between Windows and most UNIX systems with respect to the requirements of symbols. DLLs on Windows are required to be fully resolved, this means that they must link every symbol at creation. UNIX systems allow shared libraries to get symbols from the executable or other shared libraries at run time. On UNIX systems that support this feature, CMake will compile with the flags that allow executable symbols to be used by shared libraries. This small difference can cause large problems. A common, but hard to track down bug with DLLs happens with C++ template classes and static members. Two DLLs can end up with separate copies of what is supposed to be a single global static member of a class. There are also problems with the approach taken on most UNIX systems. The start up time for large applications with many symbols can be long since much of the linking is deferred to run time.

Another common pitfall occurs with C++ global objects. These objects require that constructors must be called before they can be used. The `main` that links or loads C++ shared libraries MUST be linked with the C++ compiler, or globals like `cout` may not be initialized before they are used, causing strange crashes at start up time.

Since executables that link to shared libraries must be able to find the libraries at run time, special environment variables and linker flags must be used. There are tools that can be used to show which libraries an executable is actually using. On many UNIX systems there is a tool called `ldd` (`otool -L` on Mac OS X) that shows which libraries are used by an executable. On Windows, a program called `depends` can be used to find the same type of information. On many UNIX systems there are also environment variables like `LD_LIBRARY_PATH` that tell the program where to find the libraries at run time. Where supported CMake will add run time library path information into the linked executables, so that `LD_LIBRARY_PATH` is not required. This feature can be turned off by setting the cache entry `CMAKE_SKIP_RPATH` to false. This may be desirable for installed software that should not be looking in the build tree for shared libraries. On Windows there is only one `PATH` environment variable that is used for both DLLs and finding executables.

4.10 Shared Library Versioning

When an executable is linked to a shared library, it is important that the copy of the shared library loaded at runtime matches that expected by the executable. On some UNIX systems, a shared library has an associated "soname" intended to solve this problem. When an executable links against the library, its soname is copied into the executable. At runtime, the dynamic linker uses this name from the executable to search for the library.

Consider a hypothetical shared library "foo" providing a few C functions that implement some functionality. The interface to foo is called an Application Programming Interface (API). If the implementation of these C functions changes in a new version of foo, but the API remains the same, then executables linked against foo will still run correctly. When the API changes,

old executables will no longer run with a new copy of foo, so a new API version number must be associated with foo.

This can be implemented by creating the original version of foo with a soname and file name such as libfoo.so.1. A symbolic link such as libfoo.so -> libfoo.so.1 will allow standard linkers to work with the library and create executables. The new version of foo can be called libfoo.so.2 and the symbolic link updated so that new executables use the new library. When an old executable runs, the dynamic linker will look for libfoo.so.1, find the old copy of the library, and run correctly. When a new executable runs, the dynamic linker will look for libfoo.so.2 and correctly load the new version.

This scheme can be expanded to handle the case of changes to foo that do not modify the API. We introduce a second set of version numbers that is totally independent of the first. This new set corresponds to the software version providing foo. For example, some larger project may have introduced the existence of library foo starting in version 3.4. In this case, the file name for foo might be libfoo.so.3.4, but the soname would still be libfoo.so.1 because the API for foo is still on its first version. A symbolic link from libfoo.so.1 -> libfoo.so.3.4 will allow executables linked against the library to run. When a bug is fixed in the software without changing the API to foo, then the new library file name might be libfoo.so.3.5, and the symbolic link can be updated to allow existing executables to run.

CMake supports this soname-based version number encoding on platforms supporting soname natively. A target property for the shared library named "VERSION" specifies the version number used to create the file name for the library. This version should correspond to that of the software package providing foo. On Windows the VERSION property is used to set the binary image number, using major.minor format. Another target property named "SOVERSION" specifies the version number used to create the soname for the library. This version should correspond to the API version number for foo. These target properties are ignored on platforms where CMake does not support this scheme.

The following CMake code configures the version numbers of the shared library foo:

```
set_target_properties (foo PROPERTIES VERSION 1.2 SOVERSION 4)
```

This results in the following library and symbolic links:

```
libfoo.so.1.2
libfoo.so.4 -> libfoo.so.1.2
libfoo.so -> libfoo.so.4
```

If only one of the two properties is specified, the other defaults to its value automatically. For example, the code

```
set_target_properties (foo PROPERTIES VERSION 1.2)
```

results in the following shared library and symbolic link:

```
libfoo.so.1.2  
libfoo.so -> libfoo.so.1.2
```

CMake makes no attempt to enforce sensible version numbers. It is up to the programmer to utilize this feature in a productive manner.

4.11 Installing Files

Software is typically installed into a directory separate from the source and build trees. This allows it to be distributed in a clean form and isolates users from the details of the build process. CMake provides the `install` command to specify how a project is to be installed. This command is invoked by a project in the CMakeLists file and tells CMake how to generate installation scripts. The scripts are executed at install time to perform the actual installation of files. For Makefile generators (UNIX, NMake, Borland, MinGW, etc.), the user simply runs "`make install`" (or "`nmake install`") and the make tool will invoke CMake's installation module. With GUI based systems (Visual Studio, Xcode, etc.) the user simply builds the target called `INSTALL`.

Each call to the `install` command defines some installation rules. Within one CMakeLists file (source directory) these rules will be evaluated in the order in which the corresponding commands are invoked. The order across multiple directories is not specified.

The `install` command has several signatures designed for common installation use cases. A particular invocation of the command specifies the signature as the first argument. The signatures are `TARGETS`, `FILES`, `PROGRAMS`, `DIRECTORY`, `SCRIPT`, and `CODE`.

install (TARGETS ...)

Install the binary files corresponding to targets built inside the project.

install (FILES ...)

General-purpose file installation. It is typically used for installation of header files, documentation, and data files required by your software.

install (PROGRAMS ...)

Installs executable files not built by the project, such as shell scripts. It is identical to `install (FILES)` except that the default permissions of the installed file include the executable bit.

install (DIRECTORY ...)

Install an entire directory tree. This may be used for installing directories with resources such as icons and images.

install (SCRIPT ...)

Specify a user-provided CMake script file to be executed during installation. Typically this is used to define pre-install or post-install actions for other rules.

install (CODE ...)

Specify user-provided CMake code to be executed during the installation. This is similar to `install (SCRIPT)` but the code is provided inline in the call as a string.

The TARGETS, FILES, PROGRAMS, DIRECTORY signatures are all meant to create install rules for files. The targets, files, or directories to be installed are listed immediately after the signature name argument. Additional details can be specified using keyword arguments followed by corresponding values. Keyword arguments provided by most of the signatures are as follows.

DESTINATION

Specifies the location in which the installation rule will place files. This argument must be followed by a directory path indicating the location. If the directory is specified as a full path it will be evaluated at install time as an absolute path. If the directory is specified as a relative path it will be evaluated at install time relative to the installation prefix. The prefix may be set by the user through the cache variable `CMAKE_INSTALL_PREFIX`. A platform-specific default is provided by CMake: “`/usr/local`” on UNIX and “`<SystemDrive>/Program Files/<ProjectName>`” on Windows, where SystemDrive is something like “`C:`” and ProjectName is the name given to the top-most `PROJECT` command.

PERMISSIONS

Specifies file permissions to be set on the installed files. This option is needed only to override the default permissions selected by a particular `INSTALL` command signature. Valid permissions are `OWNER_READ`, `OWNER_WRITE`, `OWNER_EXECUTE`, `GROUP_READ`, `GROUP_WRITE`, `GROUP_EXECUTE`, `WORLD_READ`, `WORLD_WRITE`, `WORLD_EXECUTE`, `SETUID`, and `SETGID`. Some platforms do not support all of these permissions, on such platforms those permission names are ignored.

CONFIGURATIONS

Specifies a list of build configurations for which an installation rule applies (Debug, Release, etc.). For Makefile generators the build configuration is specified by the `CMAKE_BUILD_TYPE` cache variable. For Visual Studio and Xcode generators the configuration is selected when the `INSTALL` target is built. An installation rule will

be evaluated only if the current install configuration matches an entry in the list provided to this argument. Configuration name comparison is case-insensitive.

COMPONENT

Specifies the installation component for which the installation rule applies. Some projects divide their installations into multiple components for separate packaging. For example, a project may define a “Runtime” component that contains the files needed to run a tool, a “Development” component containing the files needed to build extensions to the tool, and a “Documentation” component containing the manual pages and other help files. The project may then package each component separately for distribution by installing only one component at a time. By default all components are installed. Component-specific installation is an advanced feature intended for use by package maintainers. It requires manual invocation of the installation scripts with an argument defining the `COMPONENT` variable to name the desired component. Note that component names are not defined by CMake. Each project may define its own set of components.

OPTIONAL

Specifies that it is not an error if the input file to be installed does not exist. If the input file exists it will be installed as requested. If it does not exist it will be silently not installed.

Projects typically install some of the library and executable files created during their build process. The `install` command provides the `TARGETS` signature for this purpose:

```
install (TARGETS targets...
    [[ARCHIVE|LIBRARY|RUNTIME|FRAMEWORK|BUNDLE|
      PRIVATE_HEADER|PUBLIC_HEADER|RESOURCE]
    [DESTINATION <dir>]
    [PERMISSIONS permissions...]
    [CONFIGURATIONS [Debug|Release|...]]
    [COMPONENT <component>]
    [OPTIONAL]
    [EXPORT <export name>]
    [NAMELINK_ONLY|NAMELINK_SKIP]
  ] [...])
```

The `TARGETS` keyword is immediately followed by a list of the targets created using `add_executable` or `add_library` to be installed. One or more files corresponding to each target will be installed.

Files installed with this signature may be divided into three categories: `ARCHIVE`, `LIBRARY`, and `RUNTIME`. These categories are designed to group target files by typical installation

destination. The corresponding keyword arguments are optional, but if present specify that other arguments following them apply only to target files of that type. Target files are categorized as follows:

executables - RUNTIME

Created by `add_executable` (.exe on Windows, no extension on UNIX)

loadable modules - LIBRARY

Created by `add_library` with the `MODULE` option (.dll on Windows, .so on UNIX)

shared libraries - LIBRARY

Created by `add_library` with the `SHARED` option on UNIX-like platforms (.so on most UNIX, .dylib on Mac)

dynamic-link libraries - RUNTIME

Created by `add_library` with the `SHARED` option on Windows platforms (.dll)

import libraries - ARCHIVE

Linkable file created by a dynamic-link library that exports symbols (.lib on most Windows, .dll.a on Cygwin and MinGW).

static libraries - ARCHIVE

Created by `add_library` with the `STATIC` option (.lib on Windows, .a on UNIX, Cygwin, and MinGW)

Consider a project that defines an executable `myExecutable` that links to a shared library `mySharedLib`. It also provides a static library `myStaticLib` and a plugin module to the executable called `myPlugin` that also links to the shared library. The executable, static library, and plugin file may be installed individually using the commands

```
install (TARGETS myExecutable DESTINATION bin)
install (TARGETS myStaticLib DESTINATION lib/myproject)
install (TARGETS myPlugin DESTINATION lib)
```

The executable will not be able to run from the installed location until the shared library to which it links is also installed. Installation of the library requires a bit more care in order to support all platforms. It must be installed to a location searched by the dynamic linker on each platform. On UNIX-like platforms the library is typically installed to `lib`, while on Windows it should be placed next to the executable in `bin`. An additional challenge is that the import library associated with the shared library on Windows should be treated like the static library and installed to `lib/myproject`. In other words we have three different kinds of files created with a single target name that must be installed to three different destinations!

Fortunately this problem can be solved using the category keyword arguments. The shared library may be installed using the command

```
install (TARGETS mySharedLib
         RUNTIME DESTINATION bin
         LIBRARY DESTINATION lib
         ARCHIVE DESTINATION lib/myproject)
```

This tells CMake that the `RUNTIME` file (.dll) should be installed to `bin`, the `LIBRARY` file (.so) should be installed to `lib`, and the `ARCHIVE` (.lib) file should be installed to `lib/myproject`. On UNIX the `LIBRARY` file will be installed and on Windows the `RUNTIME` and `ARCHIVE` files will be installed.

If the above sample project is to be packaged into separate runtime and development components we must assign the appropriate component to each target file installed. The executable, shared library, and plugin are required in order to run the application, so they belong in a `Runtime` component. Meanwhile the import library (corresponding to the shared library on Windows) and the static library are only required to develop extensions to the application, and therefore belong in a `Development` component.

Component assignments may be specified by adding the `COMPONENT` argument to each of the commands above. We may also combine all of the installation rules into a single command invocation. This single command is equivalent to all of the above commands with components added. The files generated by each target are installed using the rule for their category.

```
install (TARGETS myExecutable mySharedLib myStaticLib myPlugin
         RUNTIME DESTINATION bin           COMPONENT Runtime
         LIBRARY DESTINATION lib          COMPONENT Runtime
         ARCHIVE DESTINATION lib/myproject COMPONENT Development)
```

Either `NAMELINK_ONLY` or `NAMELINK_SKIP` may be specified as a `LIBRARY` option. On some platforms a versioned shared library has a symbolic link such as

```
lib<name>.so -> lib<name>.so.1
```

where `lib<name>.so.1` is the soname of the library and `lib<name>.so` is a "namelink" that helps linkers to find the library when given `-l<name>`. The `NAMELINK_ONLY` option causes installation of only the namelink when a library target is installed. The `NAMELINK_SKIP` option causes installation of library files other than the namelink when a library target is installed. When neither option is given both portions are installed. On

platforms where versioned shared libraries do not have namelinks, or when a library is not versioned, the `NAMELINK_SKIP` option installs the library and the `NAMELINK_ONLY` option installs nothing. See the `VERSION` and `SOVERSION` target properties for details on creating versioned shared libraries.

Projects may install files other than those that are created with `add_executable` or `add_library`, such as header files or documentation. General-purpose installation of files is specified using the `FILES` signature:

```
install (FILES files... DESTINATION <dir>
          [PERMISSIONS permissions...]
          [CONFIGURATIONS [Debug|Release|...]]
          [COMPONENT <component>]
          [RENAME <name>] [OPTIONAL])
```

The `FILES` keyword is immediately followed by a list of files to be installed. Relative paths are evaluated with respect to the current source directory. Files will be installed to the given `DESTINATION` directory. For example, the command

```
install (FILES my-api.h ${CMAKE_CURRENT_BINARY_DIR}/my-config.h
          DESTINATION include)
```

Installs the file `my-api.h` from the source tree and the file `my-config.h` from the build tree into the `include` directory under the installation prefix. By default installed files are given permissions `OWNER_WRITE`, `OWNER_READ`, `GROUP_READ`, and `WORLD_READ`, but this may be overridden by specifying the `PERMISSIONS` option. Consider the case in which we want to install a global configuration file on a UNIX system that is readable only by its owner (such as root). We may accomplish this with the command

```
install (FILES my-rc DESTINATION /etc
          PERMISSIONS OWNER_WRITE OWNER_READ)
```

which installs the file `my-rc` with owner read/write permission into the absolute path `/etc`.

The `RENAME` argument specifies a name for an installed file that may be different from the original file. Renaming is allowed only when a single file is installed by the command. For example, the command

```
install(FILES version.h DESTINATION include RENAME my-version.h)
```

will install the file `version.h` from the source directory to `include/my-version.h` under the installation prefix.

Projects may also install helper programs such as shell scripts or python scripts that are not actually compiled as targets. These may be installed with the `FILES` signature using the `PERMISSIONS` option to add execute permission. However this case is common enough to justify a simpler interface. CMake provides the `PROGRAMS` signature for this purpose:

```
install (PROGRAMS files... DESTINATION <dir>
          [PERMISSIONS permissions...]
          [CONFIGURATIONS [Debug|Release|...]]
          [COMPONENT <component>]
          [RENAME <name>] [OPTIONAL])
```

The `PROGRAMS` keyword is immediately followed by a list of scripts to be installed. This command is identical to the `FILES` signature except that the default permissions additionally include `OWNER_EXECUTE`, `GROUP_EXECUTE`, and `WORLD_EXECUTE`. For example, we may install a python utility script with the command

```
install (PROGRAMS my-util.py DESTINATION bin)
```

which installs `my-util.py` to the `bin` directory under the installation prefix and gives it owner, group, and world read and execute permission plus owner write.

Projects may also provide a whole directory full of resource files such as icons or html documentation. An entire directory may be installed using the `DIRECTORY` signature:

```
install (DIRECTORY dirs... DESTINATION <dir>
          [FILE_PERMISSIONS permissions...]
          [DIRECTORY_PERMISSIONS permissions...]
          [USE_SOURCE_PERMISSIONS]
          [CONFIGURATIONS [Debug|Release|...]]
          [COMPONENT <component>]
          [[PATTERN <pattern> | REGEX <regex>]
          [EXCLUDE] [PERMISSIONS permissions...]] [...])
```

The `DIRECTORY` keyword is immediately followed by a list of directories to be installed. Relative paths are evaluated with respect to the current source directory. Each named directory is installed to the destination directory. The last component of each input directory

name is appended to the destination directory as that directory is copied. For example, the command

```
install (DIRECTORY data/icons DESTINATION share/myproject)
```

will install the `data/icons` directory from the source tree into `share/myproject/icons` under the installation prefix. A trailing slash will leave the last component empty and install the contents of the input directory to the destination. The command

```
install (DIRECTORY doc/html/ DESTINATION doc/myproject)
```

installs the contents of `doc/html` from the source directory into `doc/myproject` under the installation prefix. If no input directory names are given, as in

```
install (DIRECTORY DESTINATION share/myproject/user)
```

the destination directory will be created but nothing will be installed into it.

Files installed by the `DIRECTORY` signature are given the same default permissions as the `FILE` signature. Directories installed by the `DIRECTORY` signature are given the same default permissions as the `PROGRAMS` signature. The `FILE_PERMISSIONS` and `DIRECTORY_PERMISSIONS` options may be used to override these defaults. Consider the case in which a directory full of example shell scripts is to be installed into a directory that is both owner and group writable. We may use the command

```
install (DIRECTORY data/scripts DESTINATION share/myproject  
FILE_PERMISSIONS  
    OWNER_READ OWNER_EXECUTE OWNER_WRITE  
    GROUP_READ GROUP_EXECUTE  
    WORLD_READ WORLD_EXECUTE  
DIRECTORY_PERMISSIONS  
    OWNER_READ OWNER_EXECUTE OWNER_WRITE  
    GROUP_READ GROUP_EXECUTE GROUP_WRITE  
    WORLD_READ WORLD_EXECUTE)
```

which installs the directory `data/scripts` into `share/myproject/scripts` and sets the desired permissions. In some cases a fully prepared input directory created by the project may have the desired permissions already set. The `USE_SOURCE_PERMISSIONS` option tells CMake to use the file and directory permissions from the input directory during installation. If

in the previous example the input directory were to have already been prepared with correct permissions the following command may have been used instead.

```
install (DIRECTORY data/scripts DESTINATION share/myproject  
        USE_SOURCE_PERMISSIONS)
```

If the input directory to be installed is under source management, such as CVS, there may be extra subdirectories in the input that we do not wish to install. There may also be specific files which should not be installed, or be installed with different permissions, while most files get the defaults. The `PATTERN` and `REGEX` options may be used for this purpose. A `PATTERN` option is followed first by a globbing pattern and then by an `EXCLUDE` or `PERMISSIONS` option. A `REGEX` option is followed first by a regular expression and then by `EXCLUDE` or `PERMISSIONS`. The `EXCLUDE` option skips installation of those files or directories matching the preceding pattern or expression, while the `PERMISSIONS` option assigns specific permissions to them.

Each input file and directory is tested against the pattern or regular expression as a full path with forward slashes. A pattern will match only complete file or directory names occurring at the end of the full path while a regular expression may match any portion. For example, the pattern “`foo`” will match “`.../foo.txt`” but not “`.../myfoo.txt`” or “`.../foo/bar.txt`” but the regular expression “`foo`” will match all of them.

Returning to the above example of installing an icons directory, consider the case in which the input directory is managed by CVS and also contains some extra text files that we do not want to install. The command

```
install (DIRECTORY data/icons DESTINATION share/myproject  
        PATTERN "CVS" EXCLUDE  
        PATTERN "*.txt" EXCLUDE)
```

installs the icons directory while ignoring any CVS directory or text file contained. The equivalent command using the `REGEX` option is

```
install (DIRECTORY data/icons DESTINATION share/myproject  
        REGEX "/CVS$" EXCLUDE  
        REGEX "/[^/]*.txt$" EXCLUDE)
```

which uses ‘/’ and ‘\$’ to constrain the match in the same way as the patterns. Consider a similar case in which the input directory contains shell scripts and text files that we wish to install with different permissions than the other files. The command

```
install (DIRECTORY data/other/ DESTINATION share/myproject  
        PATTERN "CVS" EXCLUDE  
        PATTERN "*.txt"  
        PERMISSIONS OWNER_READ OWNER_WRITE  
        PATTERN "*.sh"  
        PERMISSIONS OWNER_READ OWNER_WRITE OWNER_EXECUTE)
```

will install the contents of `data/other` from the source directory to `share/myproject` while ignoring CVS directories and giving specific permissions to `.txt` and `.sh` files.

Project installations may need to perform tasks other than just placing files in the installation tree. Third-party packages may provide their own mechanisms to register new plugins which must be invoked during project installation. The `SCRIPT` signature is provided for this purpose:

```
install (SCRIPT <file>)
```

The `SCRIPT` keyword is immediately followed by the name of a CMake script. CMake will execute the script during installation. If the file name given is a relative path it will be evaluated with respect to the current source directory. A simple use case is printing a message during installation. We first write a `message.cmake` file containing the code

```
message ("Installing My Project")
```

and then reference this script using the command

```
install (SCRIPT message.cmake)
```

Custom installation scripts are not executed during the main CMakeLists file processing. They are executed during the installation process itself. Variables and macros defined in the code containing the `install (SCRIPT)` call will not be accessible from the script. However there are a few variables defined during the script execution which may be used to get information about the installation. The variable `CMAKE_INSTALL_PREFIX` is set to the actual installation prefix. This may be different from the corresponding cache variable value because the installation scripts may be executed by a packaging tool that uses a different prefix. An environment variable `ENV{DESTDIR}` may be set by the user or packaging tool. Its value is prepended to the installation prefix and to absolute installation paths to determine the location to which files are installed. In order to reference an install location on disk the custom script may use `$ENV{DESTDIR}${CMAKE_INSTALL_PREFIX}` as the top portion of the path. The variable `CMAKE_INSTALL_CONFIG_NAME` is set to the name of the build configuration

currently being installed (Debug, Release, etc.). During component-specific installation the variable `CMAKE_INSTALL_COMPONENT` is set to the name of the current component.

Custom installation scripts, as simple as the message above, may be more easily created with the script code placed inline in the call to the `INSTALL` command. The `CODE` signature is provided for this purpose:

```
install (CODE "<code>")
```

The `CODE` keyword is immediately followed by a string containing the code to place in the installation script. An install-time message may be created using the command

```
install (CODE "MESSAGE(\"Installing My Project\")")
```

which has the same effect as the `message.cmake` script but contains the code inline.

Installing Prerequisite Shared Libraries

Executables are frequently built using shared libraries as building blocks. When you install such an executable, you must also install its prerequisite shared libraries, called “prerequisites” because the executable requires their presence in order to load and run properly. The three main sources for shared libraries are the operating system itself, the build products of your own project and third party libraries belonging to an external project. The ones from the operating system may be relied upon to be present without installing anything: they are on the base platform on which your executable runs. The build products in your own project presumably have `add_library` build rules in the CMakeLists files, and so it should be straightforward to create CMake install rules for them. It is the third party libraries that frequently become a high maintenance item when there are more than a handful of them or when the set of them fluctuates from version to version of the third party project. Libraries may be added, code may be reorganized, and the third party shared libraries themselves may actually have additional prerequisites that are not obvious at first glance.

CMake provides two modules to make it easier to deal with required shared libraries. The first module, `GetPrerequisites.cmake`, provides the `get_prerequisites` function to analyze and classify the prerequisite shared libraries upon which an executable depends. Given an executable file as input, it will produce a list of the shared libraries required to run that executable, including any prerequisites of the discovered shared libraries themselves. It uses native tools on the various underlying platforms to perform this analysis: `dumpbin` (Windows), `otool` (Mac) and `ldd` (Linux). The second module, `BundleUtilities.cmake`, provides the `fixup_bundle` function to copy and fixup prerequisite shared libraries using well-defined locations relative to the executable. For Mac bundle applications, it embeds the libraries inside the bundle, fixing them up with `install_name_tool` to make a self-

contained unit. On Windows, it copies the libraries into the same directory with the executable since executables will search in their own directories for their required DLLs.

The `fixup_bundle` function helps you create relocatable install trees. Mac users appreciate self-contained bundle applications: you can drag them anywhere, double click them and they still work. They do not rely on anything being installed in a certain location other than the operating system itself. Similarly Windows users without administrative privileges appreciate a relocatable install tree where an executable and all of its required DLLs are installed in the same directory and it works no matter where you install it. You can even move things around after installing them and it will still work.

To use `fixup_bundle`, first install one of your executable targets. Then, configure a CMake script that can be called at install time. Inside the configured CMake script, simply include `BundleUtilities` and call the `fixup_bundle` function with appropriate arguments.

In `CMakeLists.txt`:

```
install (TARGETS myExecutable DESTINATION bin)

# To install, for example, MSVC runtime libraries:
include (InstallRequiredSystemLibraries)

# To install other/non-system 3rd party required libraries:
configure_file (
    ${CMAKE_CURRENT_SOURCE_DIR}/FixBundle.cmake.in
    ${CMAKE_CURRENT_BINARY_DIR}/FixBundle.cmake
    @ONLY
)

install (SCRIPT ${CMAKE_CURRENT_BINARY_DIR}/FixBundle.cmake)
```

In `FixBundle.cmake.in`:

```
include (BundleUtilities)

# Set bundle to the full path name of the executable already
# existing in the install tree:
set (bundle
    "${CMAKE_INSTALL_PREFIX}/myExecutable@CMAKE_EXECUTABLE_SUFFIX@")

# Set other_libs to a list of full path names to additional
# libraries that cannot be reached by dependency analysis.
# (Dynamically loaded PlugIns, for example.)
```

```

set (other_libs "")

# Set dirs to a list of directories where prerequisite libraries
# may be found:
set (dirs "@LIBRARY_OUTPUT_PATH@")

fixup_bundle ("${bundle}" "${other_libs}" "${dirs}")

```

You are responsible for verifying that you have permission to copy and distribute the prerequisite shared libraries for your executable. Some libraries may have restrictive software licenses that prohibit making copies a la `fixup_bundle`.

Exporting and Importing Targets

CMake 2.6 introduced support for exporting targets from one CMake-based project and importing them into another. The main feature allowing this functionality is the notion of an `IMPORTED` target. Here we present imported targets and then show how CMake files may be generated by a project to export its targets for use by other projects.

Importing Targets

Imported targets are used to convert files outside of the project on disk into logical targets inside a CMake project. They are created using the `IMPORTED` option to the `add_executable` and `add_library` commands. No build files are generated for imported targets. They are used simply for convenient, flexible reference to outside executables and libraries. Consider the following example which creates and uses an `IMPORTED` executable target:

```

add_executable (generator IMPORTED)                                     # 1
set_property (TARGET generator PROPERTY
              IMPORTED_LOCATION "/path/to/some_generator") # 2

add_custom_command (OUTPUT generated.c
                   COMMAND generator generated.c)           # 3

add_executable (myexe src1.c src2.c generated.c)

```

Line #1 creates a new CMake target called `generator`. Line #2 tells CMake the location of the target on disk to import. Line #3 references the target in a custom command. Once CMake is run the generated build system will contain a command line such as

```
/path/to/some_generator /project/binary/dir/generated.c
```

in the rule to generate the source file. In a similar manner libraries from other projects may be used through `IMPORTED` targets:

```
add_library (foo IMPORTED)
set_property (TARGET foo PROPERTY
             IMPORTED_LOCATION "/path/to/libfoo.a")
add_executable (myexe src1.c src2.c)
target_link_libraries (myexe foo)
```

On Windows a .dll and its .lib import library may be imported together:

```
add_library (bar IMPORTED)
set_property (TARGET bar PROPERTY
             IMPORTED_LOCATION "c:/path/to/bar.dll")
set_property (TARGET bar PROPERTY
             IMPORTED_IMPLIB "c:/path/to/bar.lib")
add_executable (myexe src1.c src2.c)
target_link_libraries (myexe bar)
```

A library with multiple configurations may be imported with a single target:

```
add_library (foo IMPORTED)
set_property (TARGET foo PROPERTY
             IMPORTED_LOCATION_RELEASE "c:/path/to/foo.lib")
set_property (TARGET foo PROPERTY
             IMPORTED_LOCATION_DEBUG "c:/path/to/foo_d.lib")
add_executable (myexe src1.c src2.c)
target_link_libraries (myexe foo)
```

The generated build system will link `myexe` to `foo.lib` when it is built in the release configuration and `foo_d.lib` when built in the debug configuration.

Exporting Targets

Imported targets on their own are useful, but they still require the project that imports them to know the locations of the target files on disk. The real power of imported targets is when the project providing the target files also provides a file to help import them.

The `install(TARGETS)` and `install(EXPORT)` commands work together to install both a target and a CMake file to help import it. For example, the code

```
add_executable (generator generator.c)
```

```
install (TARGETS generator DESTINATION lib/myproj/generators
        EXPORT myproj-targets)
install (EXPORT myproj-targets DESTINATION lib/myproj)
```

will install the two files

- <prefix>/lib/myproj/generators/generator
- <prefix>/lib/myproj/myproj-targets.cmake

The first is the regular executable named generator. The second file, myproj-targets.cmake, is a CMake file designed to make it easy to import generator. This file contains code such as

```
get_filename_component (_self "${CMAKE_CURRENT_LIST_FILE}" PATH)
get_filename_component (PREFIX "${_self}/../../" ABSOLUTE)
add_executable (generator IMPORTED)
set_property (TARGET generator PROPERTY
    IMPORTED_LOCATION "${PREFIX}/lib/myproj/generators/generator")
```

(note that \${PREFIX} is computed relative to the file location). An outside project may now use generator as follows:

```
include (${PREFIX}/lib/myproj/myproj-targets.cmake) # 1
add_custom_command (OUTPUT generated.c
    COMMAND generator generated.c) # 2
add_executable (myexe src1.c src2.c generated.c)
```

Line #1 loads the target import script (see section 5.7 to make this automatic). The script may import any number of targets. Their locations are computed relative to the script location so the install tree may be easily moved. Line #2 references the generator executable in a custom command. The resulting build system will run the executable from its installed location. Libraries may also be exported and imported:

```
add_library (foo STATIC fool.c)
install (TARGETS foo DESTINATION lib EXPORTS myproj-targets)
install (EXPORT myproj-targets DESTINATION lib/myproj)
```

This installs the library and an import file referencing it. Outside projects may simply write

```
include (${PREFIX}/lib/myproj/myproj-targets.cmake)
add_executable (myexe src1.c)
target_link_libraries (myexe foo)
```

and the executable will be linked to the library `foo` exported and installed by the original project.

Any number of target installations may be associated with the same export name. The export names are considered global so any directory may contribute a target installation. Only one call to the `install (EXPORT)` command is needed to install an import file that references all targets. Both of the examples above may be combined into a single export file, even if they are in different subdirectories of the project as shown in the code below.

```
# A/CMakeLists.txt
add_executable (generator generator.c)
install (TARGETS generator DESTINATION lib/myproj/generators
        EXPORT myproj-targets)

# B/CMakeLists.txt
add_library (foo STATIC fool.c)
install (TARGETS foo DESTINATION lib EXPORTS myproj-targets)

# Top CMakeLists.txt
add_subdirectory (A)
add_subdirectory (B)
install (EXPORT myproj-targets DESTINATION lib/myproj)
```

Typically projects are built and installed before being used by an outside project. However in some cases it is desirable to export targets directly from a build tree. The targets may then be used by an outside project that references the build tree with no installation involved. The `export` command is used to generate a file exporting targets from a project build tree. For example, the code

```
add_executable (generator generator.c)
export (TARGETS generator FILE myproj-exports.cmake)
```

will create a file in the project build tree called `myproj-exports.cmake` that contains the required code to import the target. This file may be loaded by an outside project that is aware of the project build tree in order to use the executable to generate a source file. An example application of this feature is for building a generator executable on a host platform when cross compiling. The project containing the generator executable may be built on the host platform and then the project that is being cross-compiled for another platform may load it.

4.12 Advanced Commands

There are a few commands that can be very useful but are not typically used in writing CMakeLists files. This section will discuss a few of these commands and when they are useful. First consider the `add_dependencies` command which creates a dependency between two targets. CMake automatically creates dependencies between targets when it can determine them. For example, CMake will automatically create a dependency for an executable target that depends on a library target. The `add_dependencies` command is typically used to specify inter target dependencies between targets where at least one of the targets is a custom target (see section 6.4 for more information on custom targets).

The `include_regular_expression` command also relates to dependencies. This command controls the regular expression that is used for tracing source code dependencies. By default CMake will trace all the dependencies for a source file including system include files such as `stdio.h`. If you specify a regular expression with the `include_regular_expression` command that regular expression will be used to limit what include files are processed. For example; if your software project's include files all started with the prefix `foo` (e.g. `fooMain.c` `fooStruct.h` etc) then you could specify a regular expression of `^foo.*$` to limit the dependency checking to just the files of your project.

Occasionally you might want to get a listing of all the source files that another source file depends on. This is useful when you have a program that uses pieces of a large library but you are not sure what pieces it is using. The `output_required_files` command will take a source file and produce a list of all the other source files it depends on. You could then use this list to produce a reduced version of the library that only contains the necessary files for your program.

Some tools such as Rational Purify on the Sun platform are run by inserting an extra command before the final link step. So, instead of

```
CC foo.o -o foo
```

The link step would be

```
purify CC foo.o -o foo
```

It is possible to do this with CMake. To run an extra program in front of the link line change the rule variables `CMAKE_CXX_LINK_EXECUTABLE`, and `CMAKE_C_LINK_EXECUTABLE`. Rule variables are described in chapter 11. The values for these variables are contained in the file `Modules/CMakeDefaultMakeRuleVariables.cmake`, and they are sometimes redefined

in Modules/Platform/*.cmake. Make sure it is set after the PROJECT command in the CMakeLists file. Here is a small example of using purify to link a program called foo:

```
project (foo)

set (CMAKE_CXX_LINK_EXECUTABLE
    "purify ${CMAKE_CXX_LINK_EXECUTABLE}")
add_executable (foo foo.cxx)
```

Of course, for a generic CMakeLists file you should have some if checks for the correct platform. This will only work for the Makefile generators because the rule variables are not used by the IDE generators. Another option would be to use \$(PURIFY) instead of plain purify. This would pass through CMake into the Makefile and be a make variable. The variable could be defined on the command line like this: make PURIFY=purify. If not specified then it would just use the regular rule for linking a C++ executable as PURIFY would be expanded by make to nothing.

System Inspection

This chapter will describe how you can use CMake to inspect the environment of the system on which the software is being built. This is a critical factor in creating cross-platform applications or libraries. It covers how to find and use system and user installed header files and libraries. It also covers some of the more advanced features of CMake including the `try_compile` and `try_run` commands. These commands are extremely powerful tools for determining the capabilities of the system and compiler that is hosting your software. This chapter also describes how to generate configured files and how to cross compile with CMake. Finally, the steps required to enable a project for the `find_package` command are covered, explaining how to create a `<Package>Config.cmake` file and other required files.

5.1 Using Header Files and Libraries

Many C and C++ programs depend on external libraries. However, when it comes to the practical aspects of compiling and linking a project, taking advantage of existing libraries can be difficult for both developers and users. Problems usually show up as soon as the software is built on a system other than that on which it was developed. Assumptions regarding where libraries and header files are located become obvious when they are not installed in the same place on the new computer and the build system is unable to find them. CMake has many features to aid developers in the integration of external software libraries into a project.

The CMake commands that are most relevant to this type of integration are the `find_file`, `find_library`, `find_path`, `find_program`, and `find_package` commands. For most C and C++ libraries, a combination of `find_library` and `find_path` will be enough to compile and link with an installed library. `find_library` can be used to locate, or allow a

user to locate a library, and `find_path` can be used to find the path to a representative include file from the project. For example, if you wanted to link to the tiff library, you could use the following commands in your `CMakeLists.txt` file:

```
# find libtiff, looking in some standard places
find_library (TIFF_LIBRARY
    NAMES tiff tiff2
    PATHS /usr/local/lib /usr/lib
)

# find tiff.h looking in some standard places
find_path (TIFF_INCLUDES tiff.h
    /usr/local/include
    /usr/include
)

include_directories (${TIFF_INCLUDES})

add_executable (mytiff mytiff.c )

target_link_libraries (myprogram ${TIFF_LIBRARY})
```

The first command used is `find_library` which in this case will look for a library with the name `tiff` or `tiff2`. The `find_library` command only requires the library's base name without any platform specific prefixes or suffixes, such as `lib` and `.dll`. The appropriate prefixes and suffixes for the system running CMake will be added to the library name automatically when CMake attempts to find it. All the `FIND_*` commands will look in the `PATH` environment variable. In addition, the commands allow the specification of additional search paths as arguments listed after the `PATHS` marker argument. As well as supporting standard paths, windows registry entries and environment variables can be used to construct search paths. The syntax for registry entries is the following:

```
[HKEY_CURRENT_USER\\Software\\Kitware\\Path;Build1]
```

Because software can be installed in many different places, it is impossible for CMake to find the library every time, but most standard installations should be covered. The `find_*` commands automatically create a cache variable so that users can override or specify the location from the CMake GUI. This way if CMake is unable to locate the files it is looking for users will still have an opportunity to specify them. If CMake does not find a file, the value is set to `VAR-NOTFOUND`. This value tells CMake that it should continue looking each time CMake's configure step is run. Note that in `if` statements, values of `VAR-NOTFOUND` will evaluate as false.

The next command used is `find_path`. This is a general purpose command that, in this example, is used to locate a header file from the library. Header files and libraries are often installed in different locations, and both locations are required to compile and link programs that use them. `find_path` is similar to `find_library`, although it only supports one name. It supports a list of search paths.

The next part of the CMakeLists file uses the variables created by the `find_*` commands. The variables can be used without checking for valid values as CMake will print an error message notifying the user if any of the required variables have not been set. The user can then set the cache values and reconfigure until the message goes away. Optionally, a CMakeLists file could use the `if` command to use alternative libraries or options to build the project without the library if it cannot be found.

From the above example you should be able to see how using the `find_*` commands can help your software to compile on a wide variety of systems. It is worth noting that the `find_*` commands search for a match starting with the first argument and first path. So when listing paths and library names you should list your preferred paths and names first. If there are multiple versions of a library, and you would prefer `tiff` over `tiff2`, make sure you list them in that order.

5.2 System Properties

Although it is a common practice in C and C++ code to add platform-specific code inside preprocessor `ifdef` directives, for maximum portability this should be avoided. Software should not be tuned to specific platforms with `ifdefs`, but rather to a canonical system consisting of a set of features. Coding to specific systems makes the software less portable, because systems and the features they support change with time, and even from system to system. A feature that may not have worked on a platform in the past may be a required feature for the platform in the future. The following code fragments illustrate the difference between coding to a canonical system and a specific system:

```
// coding to a feature
#ifndef HAS_FOOBAR_CALL
    foobar();
#else
    myfoobar();
#endif

// coding to specific platforms
#if defined(SUN) && defined(HPUX) && !defined(GNUC)
    foobar();
#else
    myfoobar();

```

```
#endif
```

The problem with the second approach is that the code will have to be modified for each new platform on which the software is compiled. For example, a future version of SUN may no longer have the `foobar` call. Using the `HAS_FOOBAR_CALL` approach, the software will work as long as `HAS_FOOBAR_CALL` is defined correctly, and this is where CMake can help. CMake can be used to define `HAS_FOOBAR_CALL` correctly and automatically by making use of the `try_compile` and `try_run` commands. These commands can be used to compile and run small test programs during the CMake configure step. The test programs will be sent to the compiler that will be used to build the project, and if errors occur the feature can be disabled. These commands require that you write a small C or C++ program to test the feature. For example, to test if the `foobar` call is provided on the system, try compiling a simple program that uses `foobar`. First write the simple test program (`testNeedFoobar.c` in this example) and then add the CMake calls to the CMakeLists file to try compiling that code. If the compilation works then `HAS_FOOBAR_CALL` will be set to true.

```
--- testNeedFoobar.c ----

#include <foobar.h>
main()
{
    foobar();
}
```

```
--- testNeedFoobar.cmake ---

try_compile (HAS_FOOBAR_CALL
${CMAKE_BINARY_DIR}
${PROJECT_SOURCE_DIR}/testNeedFoobar.c
)
```

Now that `HAS_FOOBAR_CALL` is set correctly in CMake you can use it in your source code through either the `add_definitions` command or by configuring a header file. We recommend configuring a header file as that file can be used by other projects that depend on your library. This is discussed further in section 5.6.

Sometimes, just compiling a test program is not enough. In some cases, you may actually want to compile and run a program to get its output. A good example of this is testing the byte order of a machine. The following example shows how you can write a small program that CMake will compile and then run to determine the byte order of a machine.

```
---- TestByteOrder.c -----

int main () {
    /* Are we most significant byte first or last */
    union
    {
        long l;
        char c[sizeof (long)];
    } u;
    u.l = 1;
    exit (u.c[sizeof (long) - 1] == 1);
}
```

```
---- TestByteOrder.cmake-----

try_run (RUN_RESULT_VAR
         COMPILE_RESULT_VAR
         ${CMAKE_BINARY_DIR}
         ${PROJECT_SOURCE_DIR}/Modules/TestByteOrder.c
         OUTPUT_VARIABLE OUTPUT
     )
```

The return result of the run will go into `RUN_RESULT_VAR` and the result of the compile will go into `COMPILE_RESULT_VAR`, and any output from the run will go into `OUTPUT`. You can use these variables to report debug information to the users of your project.

For small test programs the `FILE` command with the `WRITE` option can be used to create the source file from the CMakeLists file. The following example tests the C compiler to verify that it can be run.

```
file (WRITE
      ${CMAKE_BINARY_DIR}/CMakeTmp/testCCompiler.c
      "int main(){return 0;}"
      )

try_compile (CMAKE_C_COMPILER_WORKS
            ${CMAKE_BINARY_DIR}
            ${CMAKE_BINARY_DIR}/CMakeTmp/testCCompiler.c
            OUTPUT_VARIABLE OUTPUT
        )
```

There are several predefined try-run and try-compile macros in the `CMake/Modules` directory, some of which are listed below. These macros allow some common checks to be performed without having to create a source file for each test. For detailed documentation or to see how these macros work look at the implementation files for them in the `CMake/Modules` directory of your CMake installation. Many of these macros will look at the current value of the `CMAKE_REQUIRED_FLAGS` and `CMAKE_REQUIRED_LIBRARIES` variables to add additional compile flags or link libraries to the test.

CheckFunctionExists.cmake

Checks to see if a C function is on a system. This macro takes two arguments, the first is the name of the function to check for. The second is the variable to store the result into. This macro does use `CMAKE_REQUIRED_FLAGS` and `CMAKE_REQUIRED_LIBRARIES` if they are set.

CheckIncludeFile.cmake:

Checks for an include file on a system. This macro takes two arguments. The first is the include file to look for and the second is the variable to store the result into. Additional CFlags can be passed in as a third argument or by setting `CMAKE_REQUIRED_FLAGS`.

CheckIncludeFileCXX.cmake

Check for an include file in a C++ program. This macro takes two arguments. The first is the include file to look for and the second is the variable to store the result into. Additional CFlags can be passed in as a third argument.

CheckIncludeFiles.cmake

Check for a group of include files. This macro takes two arguments. The first is the include files to look for and the second is the variable to store the result into. This macro does use `CMAKE_REQUIRED_FLAGS` if it is set. This macro is useful when a header file you are interested in checking for is dependent on including another header file first.

CheckLibraryExists.cmake

Check to see if a library exists. This macro takes four arguments; the first is the name of the library to check for. The second is the name of a function that should be in that library. The third argument is the location of where the library should be found. The fourth argument is a variable to store the result into. This macro uses `CMAKE_REQUIRED_FLAGS` and `CMAKE_REQUIRED_LIBRARIES` if they are set.

CheckSymbolExists.cmake

Check to see if a symbol is defined in a header file. This macro takes three arguments. The first argument is the symbol to look for. The second argument is a list of header files to try including. The third argument is where the result is stored.

This macro uses `CMAKE_REQUIRED_FLAGS` and `CMAKE_REQUIRED_LIBRARIES` if they are set.

CheckTypeSize.cmake

Determines the size in bytes of a variable type. This macro takes two arguments. The first argument is the type to evaluate. The second argument is where the result is stored. Both `CMAKE_REQUIRED_FLAGS` and `CMAKE_REQUIRED_LIBRARIES` are used if they are set.

CheckVariableExists.cmake

Checks to see if a global variable exists. This macro takes two arguments. The first argument is the variable to look for. The second argument is the variable to store the result in. This macro will prototype the named variable and then try to use it. If the test program compiles then the variable exists. This will only work for C variables. This macro uses `CMAKE_REQUIRED_FLAGS` and `CMAKE_REQUIRED_LIBRARIES` if they are set.

Consider the following example that shows a variety of these modules being used to compute properties of the platform. At the beginning of the example four modules are loaded from CMake. The remainder of the example uses the macros defined in those modules to test for header files, libraries, symbols, and type sizes respectively.

```
# Include all the necessary files for macros
include (CheckIncludeFiles)
include (CheckLibraryExists)
include (CheckSymbolExists)
include (CheckTypeSize)

# Check for header files
set (INCLUDES "")
CHECK_INCLUDE_FILES ("${INCLUDES};winsock.h" HAVE_WINSOCK_H)

if (HAVE_WINSOCK_H)
    set (INCLUDES ${INCLUDES} winsock.h)
endif (HAVE_WINSOCK_H)

CHECK_INCLUDE_FILES ("${INCLUDES};io.h" HAVE_IO_H)
if (HAVE_IO_H)
    set (INCLUDES ${INCLUDES} io.h)
endif (HAVE_IO_H)

# Check for all needed libraries
set (LIBS "")
CHECK_LIBRARY_EXISTS ("dl;${LIBS}" dlopen "" HAVE_LIBDL)
```

```
if (HAVE_LIBDL)
    set (LIBS ${LIBS} dl)
endif (HAVE_LIBDL)

CHECK_LIBRARY_EXISTS ("ucb;${LIBS}" gethostname "" HAVE_LIBUCB)
if (HAVE_LIBUCB)
    set (LIBS ${LIBS} ucb)
endif (HAVE_LIBUCB)

# Add the libraries we found to the libraries to use when
# looking for symbols with the CHECK_SYMBOL_EXISTS macro
set (CMAKE_REQUIRED_LIBRARIES ${LIBS})

# Check for some functions that are used
CHECK_SYMBOL_EXISTS (socket      "${INCLUDES}" HAVE_SOCKET)
CHECK_SYMBOL_EXISTS (poll        "${INCLUDES}" HAVE_POLL)

# Various type sizes
CHECK_TYPE_SIZE (int      SIZEOF_INT)
CHECK_TYPE_SIZE (size_t   SIZEOF_SIZE_T)
```

For more advanced `try_compile` and `try_run` operations, it may be desirable to pass flags to the compiler, or to CMake. Both commands support the optional arguments `CMAKE_FLAGS` and `COMPILER_DEFINITIONS`. `CMAKE_FLAGS` can be used to pass `-DVAR:TYPE=VALUE` flags to CMake. The value of `COMPILER_DEFINITIONS` is passed directly to the compiler command line.

5.3 Finding Packages

Many software projects provide tools and libraries meant as building blocks for other projects and applications. CMake projects that depend on outside packages locate their dependencies using the `find_package` command. A typical invocation is of the form

```
find_package(<Package> [version])
```

where “`<Package>`” is the name of the package to be found, and “[version]” is an optional version request (of the form `major[.minor.[patch]]`). See Appendix C – Listfile Commands for the full command documentation. The command’s notion of a package is distinct from that of CPack, which is meant for creating source and binary distributions and installers.

The command operates in two modes: `Module` mode and `Config` mode. In `Module` mode the command searches for a `find-module`: a file named "`Find<Package>.cmake`". It looks first in the `CMAKE_MODULE_PATH` and then in the CMake installation. If a `find-module` is found, it is loaded to search for individual components of the package. `Find-modules` contain package-specific knowledge of the libraries and other files they expect to find, and internally use commands like `find_library` to locate them. CMake provides `find-modules` for many common packages; see Appendix D – Selected Modules. `Find-modules` are tedious and difficult to write and maintain because they need very specific knowledge of every version of the package to be found.

The `Config` mode of `find_package` provides a powerful alternative through cooperation with the package to be found. It enters this mode after failing to locate a `find-module` or when explicitly requested by the caller. In `Config` mode the command searches for a `package configuration file`: a file named "`<Package>Config.cmake`" or "`<package>-config.cmake`" that is provided by the package to be found. Given the name of a package, the `find_package` command knows how to search deep inside installation prefixes for locations like

```
<prefix>/lib/<package>/<package>-config.cmake
```

(see documentation of `find_package` in Appendix C – Listfile Commands for a complete list of locations). CMake creates a cache entry called "`<Package>_DIR`" to store the location found or allow the user to set it. Since a `package configuration file` comes with an installation of its package, it knows exactly where to find everything provided by the installation. Once the `find_package` command locates the file it provides the locations of package components without any additional searching.

The "[`version`]" option asks `find_package` to locate a particular version of the package. In `Module` mode, the command passes the request on to the `find-module`. In `Config` mode the command looks next to each candidate `package configuration file` for a `package version file`: a file named "`<Package>ConfigVersion.cmake`" or "`<package>-config-<version>.cmake`". The `version file` is loaded to test whether the package version is an acceptable match for the version requested (see documentation of `find_package` for the `version file API specification`). If the `version file` claims compatibility the `configuration file` is accepted, otherwise it is ignored. This approach allows each project to define its own rules for version compatibility.

5.4 Built-in Find Modules

CMake has many predefined modules that can be found in the `Modules` subdirectory of CMake. The modules can find many common software packages. See Appendix D – Selected Modules for a detailed list.

Each `Find<XX>.cmake` module defines a set of variables that will allow a project to use the software package once it is found. Those variables all start with the name of the software being found `<XX>`. With CMake we have tried to establish a convention for naming these variables, but you should read the comments at the top of the module for a more definitive answer. The following variables are used by convention when needed:

`<XX>_INCLUDE_DIRS`

Where to find the package's header files, typically `<XX>.h`, etc.

`<XX>_LIBRARIES`

The libraries to link against to use `<XX>`. These include full paths.

`<XX>_DEFINITIONS`

Preprocessor definitions to use when compiling code that uses `<XX>`.

`<XX>_EXECUTABLE`

Where to find the `<XX>` tool that is part of the package.

`<XX>-<YY>_EXECUTABLE`

Where to find the `<YY>` tool that comes with `<XX>`.

`<XX>_ROOT_DIR`

Where to find the base directory of the installation of `<XX>`. This is useful for large packages where you want to reference many files relative to a common base (or root) directory.

`<XX>_VERSION_-<YY>`

Version `<YY>` of the package was found if true. Authors of find modules should make sure at most one of these is ever true. For example `TCL_VERSION_84`

`<XX>-<YY>_FOUND`

If false, then the optional `<YY>` part of `<XX>` package is not available.

`<XX>_FOUND`

Set to false, or undefined, if we haven't found or don't want to use `<XX>`.

Not all of the variables are present in each of the `FindXX.cmake` files. However, the `<XX>_FOUND` should exist under most circumstances. If `<XX>` is a library, then `<XX>_LIBRARIES` should also be defined, and `<XX>_INCLUDE_DIR` should usually be defined.

Modules can be included in a project either with the `include` command or the `find_package` command.

```
find_package(OpenGL)
```

is equivalent to

```
include(${CMAKE_ROOT}/Modules/FindOpenGL.cmake)
```

and

```
include(FindOpenGL)
```

If the project converts over to CMake for its build system, then the `find_package` will still work if the package provides a `<XX>Config.cmake` file. How to create a CMake package is described in section 5.7.

5.5 How to Pass Parameters to a Compilation?

Once you have determined all features of the system in which you are interested, it is time to configure the software based on what has been found. There are two common ways to pass this information to the compiler: on the compile line, and using a preconfigured header. The first way is to pass definitions on the compile line. A preprocessor definition can be passed to the compiler from a CMakeLists file with the `add_definitions` command. For example, a common practice in C code is to have the ability to selectively compile in/out debug statements.

```
#ifdef DEBUG_BUILD
    printf("the value of v is %d", v);
#endif
```

A CMake variable could be used to turn on or off debug builds using the `OPTION` command:

```
option (DEBUG_BUILD
        "Build with extra debug print messages.")

if (DEBUG_BUILD)
    add_definitions (-DDEBUG_BUILD)
endif (DEBUG_BUILD)
```

Another example would be to tell the compiler the result of the previous HAS_FOOBAR_CALL test that was discussed earlier in this chapter. You could do this with the following:

```
if (HAS_FOOBAR_CALL)
    add_definitions (-DHAS_FOOBAR_CALL)
endif (HAS_FOOBAR_CALL)
```

If you want to pass preprocessor definitions at a finer level of granularity, you can use the COMPILE_DEFINITIONS property that is defined for directories, targets, and source files. For example, the code

```
add_library (mylib src1.c src2.c)
add_executable (myexe main1.c)
set_property (
    DIRECTORY
    PROPERTY COMPILE_DEFINITIONS A AV=1
)
set_property (
    TARGET mylib
    PROPERTY COMPILE_DEFINITIONS B BV=2
)
set_property (
    SOURCE src1.c
    PROPERTY COMPILE_DEFINITIONS C CV=3
)
```

will build the source files with these definitions:

```
src1.c: -DA -DAV=1 -DB -DBV=2 -DC -DCV=3
src2.c: -DA -DAV=1 -DB -DBV=2
main2.c: -DA -DAV=1
```

When the `add_definitions` command is called with flags like "`-DX`" the definitions are extracted and added to the current directory's `COMPILE_DEFINITIONS` property. When a new subdirectory is created with `add_subdirectory` the current state of the directory-level property is used to initialize the same property in the subdirectory.

Note in the above example that the `set_property` command will actually set the property and replace any existing value. The command provides the `APPEND` option to add more definitions without removing existing ones. For example, the code

```
set_property (
    SOURCE src1.c
    APPEND PROPERTY COMPILE_DEFINITIONS D DV=4
)
```

will add the definitions `"-DD -DDV=4"` when building `src1.c`. Definitions may also be added on a per-configuration basis using the `COMPILE_DEFINITIONS_<CONFIG>` property. For example, the code

```
set_property (
    TARGET mylib
    PROPERTY COMPILE_DEFINITIONS_DEBUG MYLIB_DEBUG_MODE
)
```

will build sources in `mylib` with `-DMYLIB_DEBUG_MODE` only when compiling in a Debug configuration.

The second approach for passing definitions to source code is to configure a header file. For maximum portability of a toolkit, it is recommended that `-D` options are not required for the compiler command line. Instead of command line options, CMake can be used to configure a header file that applications can include. The header file will include all of the `#define` macros needed to build the project. The problem with using compile line definitions can be seen when building an application that in turn uses a library. If building the library correctly relies on compile line definitions, then chances are that an application that uses the library will also require the exact same set of compile line definitions. This puts a large burden on the application writer to make sure they add the correct flags to match the library. If instead the library's build process configures a header file with all of the required definitions, then any application that uses the library will automatically get the correct definitions when that header file is included. A definition can often change the size of a structure or class, and if the macros are not exactly the same during the build process of the library and the application linking to the library, the application may reference the “wrong part” of a class or struct and crash unexpectedly.

5.6 How to Configure a Header File

Hopefully we have convinced you that configured header files are the right choice for most software projects. To configure a file with CMake the `configure_file` command is used. This command requires an input file that is parsed by CMake which then produces an output file with all variables expanded or replaced. There are three ways to specify a variable in an input file for `configure_file`.

```
#cmakedefine VARIABLE
```

If VARIABLE is true, then the result will be:

```
#define VARIABLE
```

If VARIABLE is false, then the result will be:

```
/* #undef VARIABLE */
```

`\${VARIABLE}`

This is simply replaced by the value of VARIABLE.

@VARIABLE@

This is simply replaced by the value of VARIABLE.

Since the `{}$` syntax is commonly used by other languages, there is a way to tell the `configure_file` command to only expand variables using the `@var@` syntax. This is done by passing the `@ONLY` option to the command. This is useful if you are configuring a script that may contain `{}${var}` strings that you want to preserve. This is important because CMake will replace all occurrences of `{}${var}` with the empty string if `var` is not defined in CMake.

The following example configures a .h file for a project that contains preprocessor variables. The first definition indicates if the `FOOBAR` call exists in the library, and the next one contains the path to the build tree.

```
----- CMakeLists.txt file -----
```

```
# Configure a file from the source tree
# called projectConfigure.h.in and put
# the resulting configured file in the build
# tree and call it projectConfigure.h
configure_file (
    ${PROJECT_SOURCE_DIR}/projectConfigure.h.in
    ${PROJECT_BINARY_DIR}/projectConfigure.h)
```

```
----projectConfigure.h.in file-----
/* define a variable to tell the code if the */
/* foobar call is available on this system */
#definecmakedefine HAS_FOOBAR_CALL

/* define a variable with the path to the */
/* build directory */
#define PROJECT_BINARY_DIR "${PROJECT_BINARY_DIR}"
```

It is important to configure files into the binary tree, not the source tree. A single source tree may be shared by multiple build trees or platforms. By configuring files into the binary tree the differences between builds or platforms will be kept isolated in the build tree, where it will not corrupt other builds. This means that you will need to include the directory of the build tree where you configured the header file into the project's list of include directories using the `include_directories` command.

5.7 Creating CMake Package Configuration Files

Projects must provide package configuration files so that outside applications can find them. Consider a simple project “Gromit” providing an executable to generate source code and a library against which the generated code must link. The `CMakeLists.txt` file might start with:

```
cmake_minimum_required (VERSION 2.6.3)
project (Gromit C)
set (version 1.0)

# Create library and executable.
add_library (gromit STATIC gromit.c gromit.h)
add_executable (gromit-gen gromit-gen.c)
```

In order to install Gromit and export its targets for use by outside projects, one adds the code

```
# Install and export the targets.
install (FILES gromit.h DESTINATION include/gromit-${version})
install (TARGETS gromit gromit-gen
        DESTINATION lib/gromit-${version}
        EXPORT gromit-targets)
install (EXPORT gromit-targets
        DESTINATION lib/gromit-${version})
```

as described in Section 0. Finally, Gromit must provide a package configuration file in its installation tree so that outside projects can locate it with `find_package`:

```
# Create and install package configuration and version files.
configure_file (
  ${Gromit_SOURCE_DIR}/pkg/gromit-config.cmake.in
  ${Gromit_BINARY_DIR}/pkg/gromit-config.cmake @ONLY)

configure_file (
  ${Gromit_SOURCE_DIR}/gromit-config-version.cmake.in
  ${Gromit_BINARY_DIR}/gromit-config-version.cmake @ONLY)

install (FILES ${Gromit_BINARY_DIR}/pkg/gromit-config.cmake
         ${Gromit_BINARY_DIR}/gromit-config-version.cmake
         DESTINATION lib/gromit-$version))
```

This code configures and installs the package configuration file and a corresponding package version file. The package configuration input file `gromit-config.cmake.in` has the code,

```
# Compute installation prefix relative to this file.
get_filename_component (_dir "${CMAKE_CURRENT_LIST_FILE}" PATH)
get_filename_component (_prefix "${_dir}/../../" ABSOLUTE)

# Import the targets.
include ("${_prefix}/lib/gromit-@version@gromit-targets.cmake")

# Report other information.
set (gromit_INCLUDE_DIRS "${_prefix}/include/gromit-@version@")
```

After installation the configured package configuration file `gromit-config.cmake` knows the locations of other installed files relative to itself. The corresponding package version file is configured from its input file `gromit-config-version.cmake.in` which contains code such as,

```
set (PACKAGE_VERSION "@version@")
if (NOT "${PACKAGE_FIND_VERSION}" VERSION_GREATER "@version@")
  set (PACKAGE_VERSION_COMPATIBLE 1) # compatible with older
  if ("${PACKAGE_FIND_VERSION}" VERSION_EQUAL "@version@")
    set (PACKAGE_VERSION_EXACT 1) # exact match for this version
  endif ()
endif ()
```

An application that uses the Gromit package might create a CMake file that looks like this:

```
cmake_minimum_required (VERSION 2.6.3)
project (MyProject C)

find_package (gromit 1.0 REQUIRED)
include_directories (${gromit_INCLUDE_DIRS})
# run imported executable
add_custom_command (OUTPUT generated.c
                     COMMAND gromit-gen generated.c)
add_executable (myexe generated.c)
target_link_libraries (myexe gromit) # link to imported library
```

The call to `find_package` locates an installation of Gromit or terminates with an error message if none can be found (due to `REQUIRED`). After the command succeeds, the Gromit package configuration file `gromit-config.cmake` has been loaded, so Gromit targets have been imported and variables like `gromit_INCLUDE_DIRS` have been defined.

The above example creates a package configuration file and places it in the `install` tree. One may also create a package configuration file in the `build` tree to allow applications to use the project without installation. In order to do this, one extends Gromit's CMake file with the code:

```
# Make project useable from build tree.
export (TARGETS gromit gromit-gen FILE gromit-targets.cmake)
configure_file (${Gromit_SOURCE_DIR}/gromit-config.cmake.in
               ${Gromit_BINARY_DIR}/gromit-config.cmake @ONLY)
```

This `configure_file` call uses a different input file, `gromit-config.cmake.in`, containing

```
# Import the targets.
include("@Gromit_BINARY_DIR@/gromit-targets.cmake")

# Report other information.
set(gromit_INCLUDE_DIRS "@Gromit_SOURCE_DIR@")
```

The package configuration file `gromit-config.cmake` placed in the build tree provides the same information to an outside project as that in the install tree, but refers to files in the source and build trees. It shares an identical package version file `gromit-config-version.cmake` with that placed in the install tree.

Custom Commands and Targets

Frequently, the build process for a software project goes beyond simply compiling libraries and executables. In many cases additional tasks may be required during or after the build process. Common examples include: compiling documentation using a documentation package, generating source files by running another executable, generating files using tools for which CMake doesn't have rules (such as lex, and yacc), moving the resulting executables, post processing the executable, etc. CMake supports these additional tasks using custom commands and custom targets. This chapter will describe how to use custom commands and targets to perform complex tasks that CMake does not inherently support.

6.1 Portable Custom Commands

Before going into detail on how to use custom commands, we will discuss how to deal with some of their portability issues. Custom commands typically involve running programs with files as inputs or outputs. Even a simple command such as copying a file can be tricky to do in a cross-platform way. For example, copying a file on UNIX is done with the `cp` command, while on windows it is done with the `copy` command. To make matters worse, frequently the names of the files will change on different platforms. Executables on Windows end with `.exe` while on UNIX they do not. Even between UNIX implementations there are differences such as what extensions are used for shared libraries; `.so`, `.sl`, `.dylib`, etc.

CMake provides two main tools for handling these differences. The first is the `-E` option (short for execute) to `cmake`. When the `cmake` executable is passed the `-E` option it acts as a general purpose cross-platform utility command. The arguments following the `-E` option indicate what `cmake` should do. Some of the options include:

chdir dir command args

Changes the current directory to `dir` and then execute the command with the provided arguments.

copy file destination

Copies a file from one directory or filename to another.

copy_if_different in-file out-file

`copy_if_different` first checks to see if the files are different before copying them. `copy_if_different` is critical in many rules since the build process is based on file modification times. If the copied file is used as the input to another build rule then `copy_if_different` can eliminate unnecessary recompilations.

copy_directory source destination

This option copies the source directory including any subdirectories to the destination directory.

remove file1 file2 ...

Removes the listed files from the disk.

echo string

Echos a string to the console. This is useful for providing output during the build process.

time command args

Runs the command and times its execution.

These options provide a platform independent way to perform a few common tasks. The `cmake` executable can be referenced by using the `CMAKE_COMMAND` variable in your CMakeLists files as later examples will show.

The second tool CMake provides to address portability issues is a number of variables describing the characteristics of the platform. While most of these are covered in Appendix A - Variables, some are particularly useful for custom commands.

EXE_EXTENSION

This is the file extension for executables. Typically nothing on UNIX and .exe on Windows

CMAKE_CFG_INDIR

Development environments such as Visual Studio and Xcode use subdirectories based on the build type selected, such as Release or Debug. When performing a

custom command on a library, executable or object file you will typically need the full path to the file. `CMAKE_CFG_INDIR` on UNIX will typically be `". /"` while for Visual Studio it will be set to `"$(INTDIR) /"`, which at build time will be replaced by the selected configuration.

CMAKE_CURRENT_BINARY_DIR

This is the full path to the output directory associated with the current CMakeLists file. This may be different from `PROJECT_BINARY_DIR` which is the full path to the top of the current project's binary tree.

CMAKE_CURRENT_SOURCE_DIR

This is the full path to the source directory associated with the current CMakeLists file. This may be different from `PROJECT_SOURCE_DIR` which is the full path to the top of the current project's source tree.

EXECUTABLE_OUTPUT_PATH

Some projects specify a directory into which all the executables should be built. This variable, if defined, holds the full path to that directory.

LIBRARY_OUTPUT_PATH

Some projects specify a directory into which all the libraries should be built. This variable, if defined, holds the full path to that directory.

There are also a series of variables such as `CMAKE_SHARED_MODULE_PREFIX` and `...SUFFIX` that describe the current platform's prefix and suffix for that type of file. These variables are defined for `SHARED_MODULE`, `SHARED_LIBRARY`, and `LIBRARY`. Using these variables, you can typically construct the full path to any CMake generated file that you need to. For libraries and executable targets you can also use `get_target_property` with the `LOCATION` argument to get the full path to the target.

CMake doesn't limit you to using `cmake -E` in all your commands. You can use any command that you like, although you should consider portability issues when doing it. A common practice is to use `find_program` to find an executable (perl for example), and then use that executable in your custom commands.

6.2 Using add_custom_command on a Target

Now we will consider the signature for `add_custom_command`. In Makefile terminology `add_custom_command` adds a rule to a Makefile. For those more familiar with Visual Studio, it adds a custom build step to a file. `add_custom_command` has two main signatures, one for adding a custom command to a target and one for adding a custom command to build a file. When adding a custom command to a target the signature is as follows:

```
add_custom_command (
    TARGET target
    PRE_BUILD | PRE_LINK | POST_BUILD
    COMMAND command [ARGS arg1 arg2 arg3 ...]
    [COMMAND command [ARGS arg1 arg2 arg3 ...] ...]
    [COMMENT comment]
)
```

The target is the name of a CMake target (executable, library, or custom) to which you want to add the custom command. There is a choice of when the custom command should be executed. `PRE_BUILD` indicates that the command should be executed before any other dependencies for the target are built. `PRE_LINK` indicates that the command should be executed after the dependencies are all built, but before the actual link command. `POST_BUILD` indicates that the custom command should be executed after the target has been built. The `COMMAND` argument is the command (executable) to run and `ARGS` provides an optional list of arguments to the command. Finally the `COMMENT` argument can be used to provide a quoted string to be used as output when this custom command is run. This is useful if you want to provide some feedback or documentation on what is happening during the build. You can specify as many commands as you want for a custom command. They will be executed in the order specified.

How to Copy an Executable Once it is Built?

Now let us consider a simple custom command for copying an executable once it has been built.

```
# first define the executable target as usual
add_executable (Foo bar.c)

# get where the executable will be located
get_target_property (EXE_LOC Foo LOCATION)

# then add the custom command to copy it
add_custom_command (
    TARGET Foo
    POST_BUILD
    COMMAND ${CMAKE_COMMAND}
    ARGS -E copy ${EXE_LOC} /testing_department/files
)
```

The first command in this example is the standard command for creating an executable from a list of source files. In this case an executable named `Foo` is created from the source file `bar.c`. The next command is `get_target_property` which will set the variable called

EXE_LOC to where the executable will be built. Next is the add_custom_command invocation. In this case the target is simply Foo and we are adding a post build command. The command to execute is cmake which has its full path specified in the CMAKE_COMMAND variable. Its arguments are "-E copy" and the source and destination locations. In this case it will copy the Foo executable from where it was built into the /testing_department/files directory. Note that the TARGET parameter accepts a CMake target (Foo in this example), but most commands such as cmake -E copy will require the full path to the executable which can be retrieved using GET_TARGET_PROPERTY.

6.3 Using add_custom_command to Generate a File

The second use for add_custom_command is to add a rule for how to build an output file. In this case the rule provided will replace any current rules for building that file. The signature is as follows:

```
add_custom_command (OUTPUT output1 [output2 ...]
COMMAND command [ARGS [args...]]
[COMMAND command [ARGS arg1 arg2 arg3 ...] ...]
[MAIN_DEPENDENCY depend]
[DEPENDS [depends...]]
[COMMENT comment]
)
```

The OUTPUT is the file (or files) that will result from running this custom command, the COMMAND and ARGS parameters are the command to execute and the arguments to pass to it. As with the prior signature you can have as many commands as you wish. The DEPENDS are files or executables on which this custom command depends. If any of these dependencies change this custom command will re-execute. The MAIN_DEPENDENCY is an optional argument that acts as a regular dependency and under Visual Studio it provides a suggestion for what file to hang this custom command onto. If the MAIN_DEPENDENCY is not specified then one will be created automatically by CMake. The MAIN_DEPENDENCY should not be a regular .c or .cxx file since the custom command will override the default build rule for the file. Finally the optional COMMENT is a comment that may be used by some generators to provide additional information during the build process.

Using an Executable to Build a Source File

Sometimes a software project builds an executable, which is then used for generating source files that are then used to build other executables or libraries. This may sound like an odd case, but it occurs quite frequently. One example is the build process for the TIFF library, which creates an executable that is then run to generate a source file that has system specific information in it. This file is then used as a source file in building the main TIFF library. Another example is the Visualization Toolkit that builds an executable called vtkWrapTcl

that wraps C++ classes into Tcl. The executable is built and then used to create more source files for the build process.

```
#####
# Test using a compiled program to create a file
#####

# add the executable that will create the file
# build creator executable from creator.cxx
add_executable (creator creator.cxx)

get_target_property (creator EXE_LOC LOCATION)

# add the custom command to produce created.c
add_custom_command (
    OUTPUT ${PROJECT_BINARY_DIR}/created.c
    DEPENDS creator
    COMMAND ${EXE_LOC}
    ARGS ${PROJECT_BINARY_DIR}/created.c
)

# add an executable that uses created.c
add_executable (Foo ${PROJECT_BINARY_DIR}/created.c)
```

The first part of this example produces the `creator` executable from the source file `creator.cxx`. The custom command then sets up a rule for producing the source file `created.c` by running the executable `creator`. The custom command depends on the `creator` target and writes its result into the output tree (`PROJECT_BINARY_DIR`). Finally, an executable target called `Foo` is added that is built using the `created.c` source file. CMake will create all the required rules in the Makefile (or Visual Studio workspace) such that when you build the project, first the `creator` executable will be built, then it will be run to create `created.c`, which will then be used to build the `Foo` executable.

6.4 Adding a Custom Target

In the discussion so far CMake targets have generally referred to executables and libraries. CMake supports a more general notion of targets, called custom targets, that can be used whenever you want the notion of a target but the end product will not be a library or executable. Examples of custom targets include targets to build documentation, run tests, or update web pages. To add a custom target you use the `add_custom_target` command with the following signature:

```
add_custom_target ( name [ALL]
    [command arg arg arg ... ]
    [DEPENDS depend depend depend ... ]
)
```

The name specified will be the name given to the target. You can use that name to specifically build that target with Makefiles (make name) or Visual Studio (right click on the target and then select Build). If the optional ALL argument is specified then this target will be included in the ALL_BUILD target and will automatically be built whenever the Makefile or Project is built. The command and arguments are optional, and if specified will be added to the target as a post build command. For custom targets that will only execute a command this is all you will need. More complex custom targets may depend on other files, in these cases the DEPENDS arguments are used to list what files this target depends on. We will consider examples of both cases. First let us look at a custom target that has no dependencies:

```
add_custom_target ( FooJAR ALL
    ${JAR} -cvf "\"${PROJECT_BINARY_DIR}/Foo.jar\""
        "\"${PROJECT_SOURCE_DIR}/Java\""
)
```

With the above definition, whenever the FooJAR target is built it will run Java's Archiver (jar) to create the Foo.jar file from java classes in the \${PROJECT_SOURCE_DIR}/Java directory. In essence this type of custom target allows the developer to tie a command to a target so that it can be conveniently invoked during the build process. Now let us consider a more complex example that roughly models the generation of PDF files from LaTeX. In this case the custom target depends on other generated files (mainly the end product .pdf files):

```
# Add the rule to build the .dvi file from the .tex
# file. This relies on LATEX being set correctly
#
add_custom_command(
    OUTPUT  ${PROJECT_BINARY_DIR}/doc1.dvi
    DEPENDS ${PROJECT_SOURCE_DIR}/doc1.tex
    COMMAND ${LATEX}
    ARGS    ${PROJECT_SOURCE_DIR}/doc1.tex
)

# Add the rule to produce the .pdf file from the .dvi
# file. This relies on DVIPDF being set correctly
#
add_custom_command(
    OUTPUT  ${PROJECT_BINARY_DIR}/doc1.pdf
```

```
DEPENDS ${PROJECT_BINARY_DIR}/doc1.dvi
COMMAND ${DVIPDF}
ARGS    ${PROJECT_BINARY_DIR}/doc1.dvi
)

# finally add the custom target that when invoked
# will cause the generation of the pdf file
#
add_custom_target ( TDocument ALL
    DEPENDS ${PROJECT_BINARY_DIR}/doc1.pdf
)
```

This example makes use of both `add_custom_command` and `add_custom_target`. The two `add_custom_command` invocations are used to specify the rules for producing a .pdf file from a .tex file. In this case there are two steps and two custom commands. First a .dvi file is produced from the .tex file by running LaTeX, then the .dvi file is processed to produce the desired .pdf file. Finally a custom target is added called TDocument. Its command simply echoes out what it is doing, while the real work is done by the two custom commands. The `DEPENDS` argument sets up a dependency between the custom target and the custom commands. When TDocument is built it will first look to see if all of its dependencies are built. If any are not built it will invoke the appropriate custom commands to build them. This example can be shortened by combining the two custom commands into one custom command, as shown in the following example:

```
# Add the rule to build the .pdf file from the .tex
# file. This relies on LATEX and DVIPDF being set correctly
#
add_custom_command(
    OUTPUT  ${PROJECT_BINARY_DIR}/doc1.pdf
    DEPENDS ${PROJECT_SOURCE_DIR}/doc1.tex
    COMMAND ${LATEX}
    ARGS    ${PROJECT_SOURCE_DIR}/doc1.tex
    COMMAND ${DVIPDF}
    ARGS    ${PROJECT_BINARY_DIR}/doc1.dvi
)

# finally add the custom target that when invoked
# will cause the generation of the pdf file
#
add_custom_target ( TDocument ALL
    DEPENDS ${PROJECT_BINARY_DIR}/doc1.pdf
)
```

Now consider the case in which the documentation consists of multiple files. The above example can be modified to handle many files using a list of inputs and a foreach loop. For example:

```
# set the list of documents to process
set (DOCS doc1 doc2 doc3)

# add the custom commands for each document
foreach (DOC ${DOCS})
    add_custom_command (
        OUTPUT  ${PROJECT_BINARY_DIR}/${DOC}.pdf
        DEPENDS ${PROJECT_SOURCE_DIR}/${DOC}.tex
        COMMAND ${LATEX}
        ARGS    ${PROJECT_SOURCE_DIR}/${DOC}.tex
        COMMAND ${DVIPDF}
        ARGS    ${PROJECT_BINARY_DIR}/${DOC}.dvi
    )

    # build a list of all the results
    set (DOC_RESULTS ${DOC_RESULTS}
        ${PROJECT_BINARY_DIR}/${DOC}.pdf
    )

endforeach (DOC)

# finally add the custom target that when invoked
# will cause the generation of the pdf file
#
add_custom_target ( TDocument ALL
    DEPENDS ${DOC_RESULTS}
)
```

In this example building the custom target `TDocument` will cause all of the specified `.pdf` files to be generated. Adding a new document to the list is simply a matter of adding its filename to the `DOCS` variable at the top of the example.

6.5 Specifying Dependencies and Outputs

When using custom commands and custom targets you will often be specifying dependencies. When you specify a dependency or the output of a custom command you should always specify the full path. For example, if the command produces `foo.h` in the binary tree then its

output should be something like `${PROJECT_BINARY_DIR}/foo.h`. CMake will try to determine the correct path for the file if it is not specified, complex projects frequently end up using files in both the source and build trees, this can eventually lead to errors if the full paths are not specified.

When specifying a target as a dependency you can leave off the full path and executable extension, referencing it simply by its name. Consider the specification of the generator target as an `add_custom_command` dependency in the example on page 108. CMake recognizes `creator` as matching an existing target and properly handles the dependencies.

6.6 When There Isn't One Rule For One Output

There are a couple of unusual cases that can arise when using custom commands that warrant further explanation. The first is a case where one command (or executable) can create multiple outputs, and the second is the case where multiple commands can be used to create a single output.

A Single Command Producing Multiple Outputs

In CMake a custom command can produce multiple outputs simply by listing multiple outputs after the `OUTPUT` keyword. CMake will create the correct rules for your build system so that no matter which output is required for a target the right rules will be run. If the executable happens to produce a few outputs, but the build process is only using one of them, then you can simply ignore the other outputs when creating your custom command. Say that the executable produces a source file that is used in the build process and also an execution log that is not used. The custom command should specify the source file as the output and ignore the fact that a log file is also generated.

Another case of having one command with multiple outputs is the case where the command is the same but the arguments to it change. This is effectively the same as having a different command and each case should have its own custom command. An example of this was the documentation example on page 111, where a custom command was added for each `.tex` file. The command is the same but the arguments passed to it change each time.

Having One Output That Can Be Generated By Different Commands

In rare cases you may find that you have more than one command that you can use to generate an output. Most build systems such as `make` and Visual Studio do not support this and likewise CMake does not. There are two common approaches that can be used to resolve this. If you truly have two different commands that produce the same output, and no other significant outputs, then you can simply pick one of them and create a custom command for that one.

In the more complex case there are multiple commands with multiple outputs. For example:

```
Command1 produces foo.h and bar.h  
Command2 produces widget.h and bar.h
```

There are a few approaches that can be used in this case. In some cases you might just combine both commands and all three outputs into a single custom command, so that whenever one output is required all three are built at the same time. You could also create three custom commands, one for each unique output. The custom command for `foo.h` would invoke Command1 while the one for `widget.h` would invoke Command2. When specifying the custom command for `bar.h` you could choose either Command1 or Command2.

Converting Existing Systems to CMake

For many people the first thing they will do with CMake is to convert an existing project from using an older build system to use CMake. In many cases this can be a fairly easy process, but there are a few issues to consider. This section will address those issues and provide some suggestions for effectively converting a project over to CMake. The first issue to consider when converting to CMake is the project's directory structure.

7.1 Source Code Directory Structures

Most small projects will have their source code in either the top level directory or in a directory named `src` or `source`. Even if all of the source code is in a subdirectory we highly recommend creating a `CMakeLists` file for the top level directory. There are two reasons for this. First it can be confusing to some people that they must run CMake on the subdirectory of the project instead of the main directory. Second, you may want to install documentation or other support files from the other directories. By having a `CMakeLists` file at the top of the project you can use the `add_subdirectory` command to step down into the documentation directory where its `CMakeLists` file can install the documentation (you can have a `CMakeLists` file for a documentation directory with no targets or source code).

For projects that have source code in multiple directories there are a few options. One option that many Makefile based projects use is to have a single Makefile at the top level directory that lists all the source files to compile in their subdirectories. For example:

```
SOURCES=\
  subdir1/foo.cxx \
  subdir1/foo2.cxx \
  subdir2/gah.cxx \
  subdir2/bar.cxx
```

This approach works just as well with CMake using a similar syntax:

```
set (SOURCES
  subdir1/foo.cxx
  subdir1/foo2.cxx
  subdir1/gah.cxx
  subdir2/bar.cxx
)
```

Another option is to have each subdirectory build a library or libraries that can then be linked into the executables. In that case each subdirectory would define its own list of source files and add the appropriate targets. A third option is a mixture of the first two. Each subdirectory can have a CMakeLists file that lists its sources, but the top level CMakeLists file will not use the `add_subdirectory` command to step into the subdirectories. Instead the top level CMakeLists file will use the `include` command to include each of the subdirectory's CMakeLists files. For example, a top level CMakeLists file might include the following code:

```
# collect the files for subdir1
include (subdir1/CMakeLists.txt)
foreach (FILE ${FILES})
  set (subdir1Files ${subdir1Files} subdir1/${FILE})
endforeach (FILE)

# collect the files for subdir2
include (subdir2/CMakeLists.txt)
foreach (FILE ${FILES})
  set (subdir2Files ${subdir2Files} subdir2/${FILE})
endforeach (FILE)

# add the source files to the executable
add_executable (foo ${subdir1Files} ${subdir2Files})
```

While the CMakeLists files in the subdirectories might look like the following:

```
# list the source files for this directory
set (FILES
    foo1.cxx
    foo2.cxx
)
```

The approach you use is entirely up to you. For large projects, having multiple shared libraries can certainly improve build times when changes are made. For smaller projects the other two approaches have their advantages. The main suggestion here is to choose a strategy and stick with it.

7.2 Build Directories

The next issue to consider is where to put the resulting object files, libraries, and executables. There are a few different approaches commonly used, and some work better with CMake than others. Probably the most common approach is to produce the binary files in the same directory as the source files. For some Windows generators such as Visual Studio they are actually kept in a subdirectory matching the selected configuration, e.g. debug, release, etc. CMake supports this approach by default. A closely related approach is to put the binary files into a separate tree that has the same structure as the source tree. For example if the source tree looked like the following:

```
foo/
  subdir1
  subdir2
```

the binary tree might look like:

```
foobin/
  subdir1
  subdir2
```

CMake also supports this structure by default. Switching between in-source builds and out-of-source builds is simply a matter of changing the binary directory in CMake (see [How to Run CMake?](#) on page 10). Note that if you have already done an in-source build and want to switch to an out of source build, you should start with a fresh copy of the source tree. If you need to support multiple architectures from one source tree we highly recommend a directory structure like the following:

```
project foo/
  foo/
    subdir1
    subdir2
  foo-linux/
    subdir1
    subdir2
  foo-solaris/
    subdir1
    subdir2
  foo-hpux/
    subdir1
    subdir2
```

That way each architecture has its own build directory which will not interfere with any other architecture. Remember that not only are the object files kept in the binary directories but also any configured files are typically written to the binary tree. Another tree structure found primarily on UNIX projects is one where the binary files for different architectures are kept in subdirectories of the source tree. (see below) CMake does not work well with this structure, so we recommend switching to the separate build tree structure shown above.

```
foo/
  subdir1/
    linux
    solaris
    hpux
  subdir2/
    linux
    solaris
    hpux
```

CMake provides two variables for controlling where binary targets are written. They are the `EXECUTABLE_OUTPUT_PATH` and `LIBRARY_OUTPUT_PATH` variables. These variables control where the resulting libraries and executables will be written respectively. Setting these enables a project to place all the libraries and executables into a single directory. For projects with many subdirectories this can be a real time saver. A typical implementation is listed below:

```
# Setup output directories.  
set (LIBRARY_OUTPUT_PATH  
    ${PROJECT_BINARY_DIR}/bin  
    CACHE PATH  
    "Single directory for all libraries."  
)  
  
set (EXECUTABLE_OUTPUT_PATH  
    ${PROJECT_BINARY_DIR}/bin  
    CACHE PATH  
    "Single directory for all executables."  
)  
  
mark_as_advanced (  
    LIBRARY_OUTPUT_PATH  
    EXECUTABLE_OUTPUT_PATH  
)
```

The two variables are set to the `bin` subdirectory of the project's binary tree (in this example the `bin` subdirectory could just as easily be named "binaries"). These variables are cached so that they will impact the entire project. Finally the variables are marked as advanced for two reasons; first to reduce the number of choices the average user will see when running CMake, and second to reduce the chances of someone changing them to another value. Setting these variables is very useful for projects that make use of shared libraries (DLLs) since it collects all of the shared libraries into one directory. If the executables are placed in the same directory then they can find the required shared libraries more easily when run on Windows.

One final note on directory structures: with CMake it is perfectly acceptable to have a project within a project. For example, within the Visualization Toolkit's source tree is a directory that contains a complete copy of the zlib compression library. In writing the CMakeLists file for that library we use the `PROJECT` command to create a project named `VTKZLIB` even though it is within the VTK source tree and project. This has no real impact on VTK, but it does allow us to build zlib independent of VTK without having to modify its CMakeLists file.

7.3 Useful CMake Commands When Converting Projects

There are a few CMake commands that can make the job of converting an existing project easier and faster. The `file` command with the `GLOB` argument allows you to quickly set a variable containing a list of all the files that match the glob expression. For example:

```
# collect up the source files
file (GLOB SRC_FILES *.cxx)

# create the executable
add_executable (foo ${SRC_FILES})
```

will set the `SRC_FILES` variable to a list of all the `".cxx"` files in the current source directory. Then it will create an executable using those source files. Windows developers should be aware that glob matches are case sensitive.

Two other useful commands are `make_directory` and `exec_program`. By default CMake will create all the output directories it needs for the object files, libraries, and executables. With existing projects there may be some part of the build process that creates directories that CMake would not normally create. In these cases the `make_directory` command can be used. As soon as CMake executes that command it will create the directory specified if it does not already exist. The `exec_program` command will execute a program when it is encountered by CMake. This is useful if you want to quickly convert a UNIX autoconf configured header file to CMake. Instead of doing the full conversion to CMake you could run `configure` from an `exec_program` command to generate the configured header file (on UNIX only of course).

7.4 Converting UNIX Makefiles

If your project is currently based on standard UNIX Makefiles (not autoconf and `Makefile.in` or `imake`) then their conversion to CMake should be fairly straightforward. Essentially for every directory in your project that has a `Makefile` you will create a matching `CMakeLists` file. How you handle multiple `Makefile`s in a directory really depends on their function. If the additional `Makefile`s (or `Makefile` type files) are simply included in the main `Makefile` then you can create matching CMake syntax files and include them into your main `CMakeLists` file in a similar manner. If the different `Makefile`s are meant to be invoked on the command line for different situations then you should consider creating a main `CMakeLists` file that uses some logic to choose which one to `include` based on a CMake option.

Converting the `Makefile` syntax to CMake is fairly easy. Frequently `Makefile`s have a list of object files to compile. These can be converted to CMake variables as follows:

```
OBJS= \
  foo1.o \
  foo2.o \
  foo3.o
```

becomes

```
set (SOURCES  
    foo1.c  
    foo2.c  
    foo3.c  
)
```

While the object files are typically listed in a Makefile, in CMake the focus is on the source files. If you used conditional statements in your Makefiles they can be converted over to CMake if commands. Since CMake handles generating dependencies, most dependencies or rules to generate dependencies can be eliminated. Where you have rules to build libraries or executables replace them with add_library or add_executable commands. Some UNIX build systems (and source code) make heavy use of the system architecture to determine what files to compile or what flags to use. Typically this information is stored in a Makefile variable called ARCH or UNAME.

The first choice in these cases is to replace the architecture specific code with a more generic test. For example, instead of switching your handling of byte order based on operating system, make the decision based on the results of a byte order test such as CheckBigEndian.cmake. With some software packages, there is too much architecture specific code for such a change to be reasonable, or you may want to make decisions based on architecture for other reasons. In those cases you can use the variables defined in the CMakeDetermineSystem module. They provide fairly detailed information on the operating system and version of the host computer.

7.5 Converting Autoconf Based Projects

Autoconf based projects primarily consist of three key pieces. The first is the configure.in file that drives the process. The second is Makefile.in which will become the resulting Makefile and the third piece is the remaining configured files that result from running configure. In converting an autoconf based project to CMake you will start with the configure.in and Makefile.in files.

The Makefile.in file can be converted to CMake syntax as explained in the preceding section on converting UNIX Makefiles. Once this has been done you need to convert the configure.in file into CMake syntax. Most functions (macros) in autoconf have corresponding commands in CMake. A short table of some of the basic conversions is listed below:

AC_ARG_WITH

use the option command

AC_CHECK_HEADER

use the `CHECK_INCLUDE_FILE` macro from the `CheckIncludeFile` module

AC_MSG_CHECKING

use the `message` command with the `STATUS` argument

AC_SUBST

done automatically when using the `configure_file` command

AC_CHECK_LIB

use the `CHECK_LIBRARY_EXISTS` macro from the `CheckLibraryExists` module

AC_CONFIG_SUBDIRS

use the `add_subdirectory` command

AC_OUTPUT

use the `configure_file` command

AC_TRY_COMPILE

use the `try_compile` command

If your `configure` script performs test compilations using `AC_TRY_COMPILE` you can use the same code for CMake. You can either put it directly into your `CMakeLists` file if it is short, or preferably put it into a source code file for your project. We typically put such files into a CMake subdirectory for large projects that require such testing.

Where you are relying on autoconf to configure files you can use CMake's `configure_file` command. The basic approach is the same and we typically name input files to be configured with a `.in` extension just as autoconf does. This command replaces any variables in the input file referenced as `$(VAR)` or `@VAR@` with their values as determined by CMake. If a variable is not defined, it will be replaced with nothing. Optionally, only variables of the form `@VAR@` will be replaced and `$(VAR)` will be ignored. This is useful for configuring files for languages that use `$(VAR)` as a syntax for evaluating variables. You can also conditionally define variables using the C pre processor by using `#cmakedefine VAR`. If the variable is defined then `configure_file` will convert the `#cmakedefine` into a `#define`, if it is not defined it will become a commented out `#undef`. For example:

```
/* what byte order is this system */
#define CMAKE_WORDS_BIGENDIAN

/* what size is an INT */
#define SIZEOF_INT @SIZEOF_INT@
```

7.6 Converting Windows Based Workspaces

To convert a Visual Studio workspace (or solution for Visual Studio .Net) to CMake involves a few steps. First you will need to create a `CMakeLists` file at the top of your source code directory. This file should start with a `project` command that defines the name of the project. This will become the name of the resulting workspace (or solution for Visual Studio .Net). Next you need to add all of your source files into CMake variables. For large projects that have multiple directories you can create a `CMakeLists` file in each directory as described previously in the section on source directory structures (page 115). You will then add your libraries and executables using `add_library` and `add_executable`. By default `add_executable` assumes that your executable is a console application. Adding the `WIN32` argument to `add_executable` indicates that it is a Windows application (using `WinMain` instead of `main`).

There are a few nice features that Visual Studio supports that CMake can take advantage of. One is support for class browsing. Typically in CMake only source files are added to a target, not header files. If you add header files to a target, they will show up in the workspace and then you will be able to browse them as usual. Visual Studio also supports the notion of groups of files. By default CMake creates groups for source files and header files. Using the `source_group` command you can create your own groups and assign files to them. If you have any custom build steps in your workspace, these can be added to your `CMakeLists` files using the `add_custom_command` command. Custom targets (Utility Targets) in Visual Studio can be added with the `add_custom_target` command.

Cross Compiling with CMake

Cross compiling a piece of software means that the software is built on one system but intended to run on a different system. The system which is used to build the software will be called the build host, the system for which the software is built will be called the target system or target platform. The target system usually runs a different operating system (or none at all) and/or runs on different hardware. A typical use case is software development for embedded devices like network switches, mobile phones or engine control units. In these cases the target platform doesn't have or is not able to run the required software development environment.

Starting with CMake 2.6.0 cross compiling is fully supported by CMake, ranging from cross compiling from Linux to Windows, cross compiling for supercomputers through to cross compiling for small embedded devices without an operating system (OS).

Cross compiling has several consequences for CMake;

- CMake cannot automatically detect the target platform
- CMake cannot find libraries and headers in the default system directories
- executables built during cross compiling cannot be executed

Cross compiling support doesn't mean that all CMake-based projects can be magically cross compiled out-of-the-box (some are), but that CMake separates between information about the build platform and target platform and gives the user mechanisms to solve cross compiling issues without additional requirements such as running virtual machines, etc.

To support cross compiling for a specific software project, CMake must be told about the target platform via a so called toolchain file. The CMakeLists.txt may have to be adjusted so they are aware that the build platform may have different properties to the target platform, and it has to deal with the cases where a compiled executable tries to execute on the build host.

8.1 Toolchain Files

In order to use CMake for cross compiling, a CMake file that describes the target platform has to be created, called the "toolchain file". This file tells CMake everything it needs to know about the target platform. Here is an example that uses the MinGW cross compiler for Windows under Linux, the contents will be explained line by line afterwards:

```
# the name of the target operating system
set (CMAKE_SYSTEM_NAME Windows)

# which compilers to use for C and C++
set (CMAKE_C_COMPILER i586-mingw32msvc-gcc )
set (CMAKE_CXX_COMPILER i586-mingw32msvc-g++ )

# where is the target environment located
set (CMAKE_FIND_ROOT_PATH /usr/i586-mingw32msvc
    /home/alex/mingw-install )

# adjust the default behavior of the FIND_XXX() commands:
# search programs in the host environment
set (CMAKE_FIND_ROOT_PATH_MODE_PROGRAM NEVER)

# search headers and libraries in the target environment
set (CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set (CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)
```

Assuming that this file is saved with the name TC-mingw.cmake in your home directory, you instruct CMake to use this file by setting the CMAKE_TOOLCHAIN_FILE variable:

```
~/src$ cd build
~/src/build$ cmake -DCMAKE_TOOLCHAIN_FILE=~/TC-mingw.cmake ..
...
```

CMAKE_TOOLCHAIN_FILE has to be specified only on the initial CMake run, after that the results are reused from the CMake cache. You don't need to write a separate toolchain file for every piece of software you want to build. The toolchain files are per target platform, i.e. if

you are building several software packages for the same target platform, you only have to write one toolchain file that can be used for all packages. What do the settings in the toolchain file mean? We will examine them one by one. Since CMake cannot guess the target operating system or hardware, you have to set the following CMake variables:

CMAKE_SYSTEM_NAME

This variable is mandatory; it sets the name of the target system, i.e. to the same value as `CMAKE_SYSTEM_NAME` would have if CMake were run on the target system. Typical examples are "Linux" and "Windows". It is used for constructing the file names of the platform files like `Linux.cmake` or `Windows-gcc.cmake`. If your target is an embedded system without an OS then set `CMAKE_SYSTEM_NAME` to "Generic". Presetting `CMAKE_SYSTEM_NAME` this way instead of being detected automatically causes CMake to consider the build a cross compiling build and the CMake variable `CMAKE_CROSSCOMPILING` will be set to `TRUE`. `CMAKE_CROSSCOMPILING` is the variable which should be tested in CMake files to determine whether the current build is a cross compiled build or not.

CMAKE_SYSTEM_VERSION

This variable is optional, it sets the version of your target system. CMake does not currently use `CMAKE_SYSTEM_VERSION`.

CMAKE_SYSTEM_PROCESSOR

This variable is optional, it sets the processor or hardware name of the target system. It is used in CMake for one purpose, load the

```
${CMAKE_SYSTEM_NAME}-COMPILER_ID-${CMAKE_SYSTEM_PROCESSOR}.cmake
```

file. This file can be used to modify settings like compiler flags for the target. You should only have to set this variable if you are using a cross compiler where each target needs special build settings. The value can be chosen freely, so it could be e.g. i386, or IntelPXA255, or MyControlBoardRev42.

In CMake code the `CMAKE_SYSTEM_XXX` variables always describe the target platform. The same is true for the short `WIN32`, `UNIX`, `APPLE` variables. These variables can be used to test the properties of the target. If it is necessary to test the build host system, there is a corresponding set of variables: `CMAKE_HOST_SYSTEM`, `CMAKE_HOST_SYSTEM_NAME`, `CMAKE_HOST_SYSTEM_VERSION`, `CMAKE_HOST_SYSTEM_PROCESSOR` and also the short forms `CMAKE_HOST_WIN32`, `CMAKE_HOST_UNIX` and `CMAKE_HOST_APPLE`.

Since CMake cannot guess the target system, it cannot guess which compiler it should use. Setting the following variables defines what compilers to use for the target system.

CMAKE_C_COMPILER

This specifies the C compiler executable as either a full path or just the filename. If it is specified with full path, then this path will be preferred when searching for the C++ compiler and the other tools (binutils, linker, etc.). If the compiler is a GNU cross compiler with a prefixed name (e.g. "arm-elf-gcc") CMake will detect this and automatically find the corresponding C++ compiler (i.e. "arm-elf-c++"). The compiler can also be set via the CC environment variable. Setting CMAKE_C_COMPILER directly in a toolchain file has the advantage that the information about the target system is completely contained in this file, and it does not depend on environment variables.

CMAKE_CXX_COMPILER

This specifies the C++ compiler executable as either a full path or just the filename. It is handled the same way as CMAKE_C_COMPILER. If the toolchain is a GNU toolchain, it should suffice to set only CMAKE_C_COMPILER, CMake should find the corresponding C++ compiler automatically. As for CMAKE_C_COMPILER, also for C++ the compiler can be set via the CXX environment variable.

Once the system and the compiler are determined by CMake, it will load the corresponding files in the order described in section 11.2, The Enable Language Process.

Finding External Libraries, Programs and Other Files

Most non-trivial projects make use of external libraries or tools. CMake offers the `find_program`, `find_library`, `find_file`, `find_path`, and `find_package` commands for this purpose. They search the file system in common places for these files and return the results. `find_package` is a bit different in that it does not actually search itself, but executes `Find<*>.cmake` modules, which in turn usually call the `find_program`, `find_library`, `find_file` and `find_path` commands.

When cross compiling these commands become more complicated. For example, when cross compiling to Windows on a Linux system, getting `/usr/lib/libjpeg.so` as the result of the command `find_package(JPEG)` would be useless, since this would be the JPEG library for the host system and not the target system. In some cases you want to find files that are meant for the target platform, in other cases you will want to find files for the build host. The following variables are designed to give you the flexibility to change how the typical find commands in CMake work, so that you can find both build host and target files as necessary.

The toolchain will come with its own set of libraries and headers for the target platform, which are usually installed under a common prefix. It is also a good idea to set up a directory where all the software that is built for the target platform will be installed, so that the software packages don't get mixed up with the libraries that come with the toolchain.

The `find_program` command is usually used to find a program which will be executed during the build, so this should still search in the host file system, not in the environment of the target platform. `find_library` is normally used to find a library which is then used for linking purposes, so this command should only search in the target environment. For `find_path` and `find_file` it is not so obvious, in many cases they are used to search for headers, so by default they should only search in the target environment. The following CMake variables can be set to adjust the behavior of the find commands for cross compiling.

CMAKE_FIND_ROOT_PATH

This is a list of the directories that contain the target environment. Each of the directories listed here will be prepended to each of the search directories of every `find` command. Assuming your target environment is installed under `/opt/eldk/ppc_74xx` and your installation for that target platform goes to `~/install-eldk-ppc74xx`, set `CMAKE_FIND_ROOT_PATH` to these two directories. Then `find_library (JPEG_LIB jpeg)` will search in `/opt/eldk/ppc_74xx/lib`, `/opt/eldk/ppc_74xx/usr/lib`, `~/install-eldk-ppc74xx/lib`, `~/install-eldk-ppc74xx/usr/lib`, and should result in `/opt/eldk/ppc_74xx/usr/lib/libjpeg.so`.

By default `CMAKE_FIND_ROOT_PATH` is empty. If set, first the directories prefixed with the path given in `CMAKE_FIND_ROOT_PATH` will be searched, and after that the unprefixed versions of the same directories will be searched.

By setting this variable you are basically adding a new set of search prefixes to all of the find commands in CMake, but for some find commands you may not want to search the target or host directories. You can control how each find command invocation works by passing in one of the three following options `NO_CMAKE_FIND_ROOT_PATH`, `ONLY_CMAKE_FIND_ROOT_PATH` or `CMAKE_FIND_ROOT_PATH_BOTH` when you call it. You can also control how the find commands work using the following three variables.

CMAKE_FIND_ROOT_PATH_MODE_PROGRAM

This sets the default behavior for the `find_program` command. It can be set to `NEVER`, `ONLY` or `BOTH`. The default setting is `BOTH`. When set to `NEVER`, `CMAKE_FIND_ROOT_PATH` will not be used for `find_program` calls except where it is enabled explicitly. If set to `ONLY`, only the search directories with the prefixes coming from `CMAKE_FIND_ROOT_PATH` will be used by `find_program`. The default is `BOTH`, which means that first the prefixed directories, and then the unprefixed directories, will be searched.

In most cases `find_program` is used to search for an executable which will then be executed, e.g. using `execute_process` or `add_custom_command`. So in most cases an executable from the build host is required, so setting `CMAKE_FIND_ROOT_PATH_MODE_PROGRAM` to `NEVER` is normally preferred.

CMAKE_FIND_ROOT_PATH_MODE_LIBRARY

This is the same as above, but for the `find_library` command. In most cases this is used to find a library which will then be used for linking, so a library for the target is required. So in the common case it should be set to `ONLY`.

CMAKE_FIND_ROOT_PATH_MODE_INCLUDE

This is the same as above and used for both `find_path` and `find_file`. In most cases this is used for finding include directories, so the target environment should be searched. In the common case it should be set to `ONLY`. If you also need to find files in the file system of the build host, e.g. some data files that will be processed during the build. You may need to adjust the behavior for those `find_path` or `find_file` calls using the `NO_CMAKE_FIND_ROOT_PATH`, `ONLY_CMAKE_FIND_ROOT_PATH` and `CMAKE_FIND_ROOT_PATH_BOTH` options.

With a toolchain file set up as described, CMake now knows how to handle the target platform and the cross compiler. We should now be able to build software for the target platform. For complex projects there are more issues that must be taken care of.

8.2 System Inspection

Most portable software projects have a set of system inspection tests for determining the properties of the (target) system. The simplest way to check for a system feature with CMake is by testing variables. For this purpose CMake provides the variables `UNIX`, `WIN32` and `APPLE`. When cross compiling, these variables apply to the target platform, for testing the build host platform there are corresponding variables `CMAKE_HOST_UNIX`, `CMAKE_HOST_WINDOWS` and `CMAKE_HOST_APPLE`.

If this granularity is too coarse, the variables `CMAKE_SYSTEM_NAME`, `CMAKE_SYSTEM`, `CMAKE_SYSTEM_VERSION` and `CMAKE_SYSTEM_PROCESSOR` can be tested, along with their counterparts `CMAKE_HOST_SYSTEM_NAME`, `CMAKE_HOST_SYSTEM`, `CMAKE_HOST_SYSTEM_VERSION` and `CMAKE_HOST_SYSTEM_PROCESSOR`, which contain the same information, but for the build host and not for the target system.

```
if (CMAKE_SYSTEM MATCHES Windows)
    message (STATUS "Target system is Windows")
endif ()

if (CMAKE_HOST_SYSTEM MATCHES Linux)
    message (STATUS "Build host runs Linux")
endif ()
```

Using Compile Checks

In CMake there are macros such as `CHECK_INCLUDE_FILES` and `CHECK_C_SOURCE_RUNS` that are used to test the properties of the platform. Most of these macros internally use either the `try_compile` or the `try_run` commands. The `try_compile` command works as expected when cross compiling, it tries to compile the piece of code with the cross compiling toolchain, which will give the expected result.

All tests using `try_run` will not work since the created executables cannot normally run on the build host. In some cases this might be possible, e.g. using virtual machines, emulation layers like Wine or interfaces to the actual target, CMake does not depend on such mechanisms. Depending on emulators during the build process would introduce a new set of potential problems, they may have a different view on the file system, use other line endings, require special hardware or software, etc.

If `try_run` is invoked when cross compiling, it will first try to compile the software, which will work the same way as when not cross compiling. If this succeeds, it will check the variable `CMAKE_CROSSCOMPILING` to determine whether the resulting executable can be executed or not. If it cannot, it will create two cache variables, which then have to be set by the user or via the CMake cache. Assume the command looks like this:

```
try_run (SHARED_LIBRARY_PATH_TYPE
          SHARED_LIBRARY_PATH_INFO_COMPILED
          ${PROJECT_BINARY_DIR}/CMakeTmp
          ${PROJECT_SOURCE_DIR}/CMake/SharedLibraryPathInfo.cxx
          OUTPUT_VARIABLE OUTPUT
          ARGS "LDPATH"
      )
```

In this example the source file `SharedLibraryPathInfo.cxx` will be compiled and if that succeeds, the resulting executable should be executed. The variable `SHARED_LIBRARY_PATH_INFO_COMPILED` will be set to the result of the build, i.e. `TRUE` or `FALSE`. CMake will create a cache variable `SHARED_LIBRARY_PATH_TYPE` and preset it to `PLEASE_FILL_OUT-FAILED_TO_RUN`. This variable must be set to what the exit code of the executable would have been if it had been executed on the target. Additionally, CMake will create a cache variable `SHARED_LIBRARY_PATH_TYPE__TRYRUN_OUTPUT` and preset it to `PLEASE_FILL_OUT-NOTFOUND`. This variable should be set to the output that the executable prints to `stdout` and `stderr` if it were executed on the target. This variable is only created if the `try_run` command was used with the `RUN_OUTPUT_VARIABLE` or the `OUTPUT_VARIABLE` argument. You have to fill in the appropriate values for these variables. To help you with this CMake tries its best to give you useful information. To accomplish this CMake creates a file `$(CMAKE_BINARY_DIR)/TryRunResults.cmake`, that you can see an example of here:

```
# SHARED_LIBRARY_PATH_TYPE
#   indicates whether the executable would have been able to run
#   on its target platform. If so, set SHARED_LIBRARY_PATH_TYPE
#   to the exit code (in many cases 0 for success), otherwise
#   enter "FAILED_TO_RUN".
# SHARED_LIBRARY_PATH_TYPE__TRYRUN_OUTPUT
#   contains the text the executable would have printed on
#   stdout and stderr. If the executable would not have been
#   able to run, set SHARED_LIBRARY_PATH_TYPE__TRYRUN_OUTPUT
#   empty. Otherwise check if the output is evaluated by the
#   calling CMake code. If so, check what the source file would
#   have printed when called with the given arguments.
# The SHARED_LIBRARY_PATH_INFO_COMPILED variable holds the build
# result for this TRY_RUN().
#
# Source file: ~/src/SharedLibraryPathInfo.cxx
# Executable : ~/build/cmTryCompileExec-SHARED_LIBRARY_PATH_TYPE
# Run arguments: LDPATH
#     Called from: [1]    ~/src/CMakeLists.cmake

set (SHARED_LIBRARY_PATH_TYPE
    "0"
    CACHE STRING "Result from TRY_RUN" FORCE)

set (SHARED_LIBRARY_PATH_TYPE__TRYRUN_OUTPUT
    ""
    CACHE STRING "Output from TRY_RUN" FORCE)
```

You can find all of the variables that CMake could not determine, from which CMake file they were called, the source file, the arguments for the executable and the path to the executable. CMake will also copy the executables to the build directory, they have the names cmTryCompileExec-<name of the variable>, e.g. in this case cmTryCompileExec-SHARED_LIBRARY_PATH_TYPE. You can then try to run this executable manually on the actual target platform and check the results.

Once you have these results, they have to be put into the CMake cache. This can be done by using ccmake/cmake-gui/"make edit_cache" and editing the variables directly in the cache. It is not possible to reuse these changes in another build directory or if CMakeCache.txt is removed.

The recommended approach is to use the TryRunResults.cmake file created by CMake. You should copy it to a safe location (i.e. where it will not be removed if the build directory is deleted), and give it a useful name, e.g. TryRunResults-MyProject-eldk-ppc.cmake. The contents of this file have to be edited so that the set commands set the required variables

to the appropriate values for the target system. This file can then be used to preload the CMake cache by using the -C option of cmake:

```
src/build/ $ cmake -C ~/TryRunResults-MyProject-eldk-ppc.cmake .
```

You do not have to use the other CMake options again, they are now in the cache. This way you can use `MyProjectTryRunResults-eldk-ppc.cmake` in multiple build trees, and it could be distributed with your project so that it is easier for other users to cross compile it.

8.3 Running Executables Built in the Project

In some cases it is necessary that during a build an executable is invoked that was built earlier in the same build, this is usually the case for code generators and similar tools. This does not work when cross compiling, as the executables are built for the target platform and cannot run on the build host (without the use of virtual machines, compatibility layers, emulators, etc.). With CMake these programs are created using `add_executable`, and executed with `add_custom_command` or `add_custom_target`. The following three options can be used to support these executables with CMake. The old version of the CMake code could look something like this:

```
add_executable (mygen gen.c)
get_target_property (mygenLocation mygen LOCATION)
add_custom_command (
    OUTPUT "${CMAKE_CURRENT_BINARY_DIR}/generated.h"
    COMMAND ${mygenLocation}
        -o "${CMAKE_CURRENT_BINARY_DIR}/generated.h" )
```

Now we will show how this file can be modified so that it works when cross compiling. The basic idea is that the executable is built only when doing a native build for the build host and then exported as an executable target to a CMake script file. This file is then included when cross compiling, and the executable target for the executable `mygen` will be loaded. An imported target with the same name as the original target will be created. Since CMake 2.6 `add_custom_command` recognizes target names as executables, so for the command in `add_custom_command` simply the target name can be used, it is not necessary to use the `LOCATION` property to obtain the path of the executable:

```
if (CMAKE_CROSSCOMPILING)
    find_package (MyGen)
endif ()

if (NOT CMAKE_CROSSCOMPILING)
```

```

add_executable (mygen gen.c)
export (TARGETS mygen FILE
        "${CMAKE_BINARY_DIR}/MyGenConfig.cmake")
endif ()

add_custom_command (
    OUTPUT "${CMAKE_CURRENT_BINARY_DIR}/generated.h"
    COMMAND mygen -o "${CMAKE_CURRENT_BINARY_DIR}/generated.h" )

```

With the CMakeLists.txt modified like this the project can be cross compiled. First, a native build has to be done in order to create the necessary mygen executable. After that the cross compiling build can begin. The build directory of the native build has to be given to the cross compiling build as the location of the MyGen package, so that `find_package(MyGen)` can find it:

```

mkdir build-native; cd build-native
cmake ..
make
cd ..
mkdir build-cross; cd build-cross
cmake -DCMAKE_TOOLCHAIN_FILE=MyToolchain.cmake \
      -DMyGen_DIR=~/src/build-native/ ..
make

```

This code works, but CMake versions prior to 2.6 will not be able to process it, as they do not know the `export` command and they do not recognize the target name `mygen` in `add_custom_command`. A compatible version that works with CMake 2.4 looks like this:

```

if (CMAKE_CROSSCOMPILING)
    find_package (MyGen)
endif (CMAKE_CROSSCOMPILING)

if (NOT CMAKE_CROSSCOMPILING)
    add_executable (mygen gen.c)
    if (COMMAND EXPORT)
        export (TARGETS mygen FILE
                "${CMAKE_BINARY_DIR}/MyGenConfig.cmake")
    endif (COMMAND EXPORT)
endif (NOT CMAKE_CROSSCOMPILING)

get_target_property (mygenLocation mygen LOCATION)

```

```
add_custom_command (
    OUTPUT "${CMAKE_CURRENT_BINARY_DIR}/generated.h"
    COMMAND ${mygenLocation}
    -o "${CMAKE_CURRENT_BINARY_DIR}/generated.h" )
```

In this case the target is only exported if the `export` command exists and the location of the executable is retrieved using the `LOCATION` target property.

```
mkdir build-native; cd build-native
cmake ..
make
cd ..
mkdir build-cross; cd build-cross
cmake -DCMAKE_TOOLCHAIN_FILE=MyToolchain.cmake \
      -DMyGen_DIR=~/src/build-native/ ..
make
```

The “old” CMake code could also be using the `utility_source` command:

```
subdirs (mygen)
utility_source (MYGEN_LOCATION mygen mygen gen.c)
add_custom_command (
    OUTPUT "${CMAKE_CURRENT_BINARY_DIR}/generated.h"
    COMMAND ${MYGEN_LOCATION}
    -o "${CMAKE_CURRENT_BINARY_DIR}/generated.h" )
```

In this case the CMake script doesn't have to be changed, but the invocation of CMake is more complicated, since each executable location has to be specified manually:

```
mkdir build-native; cd build-native
cmake ..
make
cd ..
mkdir build-cross; cd build-cross
cmake -DCMAKE_TOOLCHAIN_FILE=MyToolchain.cmake
      -DMYGEN_LOCATION=~/src/build-native/bin/mygen ..
make
```

8.4 Cross Compiling Hello World

Now let's actually start with the cross compiling. The first step is to install a cross compiling toolchain. If this is already installed, you can skip the next paragraph.

There are many different approaches and projects that deal with cross compiling for Linux, ranging from free software projects working on Linux based PDAs to commercial embedded Linux vendors. Most of these projects come with their own way to build and use the respective toolchain. Any of these toolchains can be used with CMake, the only requirement is that it works in the normal file system and does not expect a "sandboxed" environment, like for example the Scratchbox project.

An easy to use toolchain with a relatively complete target environment is the Embedded Linux Development Toolkit (<http://www.denx.de/wiki/DULG/ELDK>). It supports ARM, PowerPC and MIPS as target platforms. ELDK can be downloaded from <ftp://ftp.sunet.se/pub/Linux/distributions/eldk/>. The easiest way is to download the ISOs, mount them and then install them:

```
mkdir mount-iso/
sudo mount -t iso9660 mips-2007-01-21.iso mount-iso/ -o loop
cd mount-iso/
./install -d /home/alex/eldk-mips/
...
Preparing...
#####
1:appWeb-mips_4KCle
#####
Done
ls /opt/eldk-mips/
bin eldk_init etc mips_4KC mips_4KCle usr var version
```

ELDK (and other toolchains) can be installed anywhere, either in the home directory or system wide if there are more users working with them. In this example the toolchain will now be located in `/home/alex/eldk-mips/usr/bin/` and the target environment is in `/home/alex/eldk-mips/mips_4KC/`.

Now that a cross compiling toolchain is installed, CMake has to be set up to use it. As already described, this is done by creating a toolchain file for CMake. In this example the toolchain file looks like this:

```
# the name of the target operating system
set (CMAKE_SYSTEM_NAME Linux)
```

```
# which C and C++ compiler to use
set (CMAKE_C_COMPILER /home/alex/eldk-mips/usr/bin/mips_4KC-gcc)
set (CMAKE_CXX_COMPILER
     /home/alex/eldk-mips/usr/bin/mips_4KC-g++)

# location of the target environment
set (CMAKE_FIND_ROOT_PATH /home/alex/eldk-mips/mips_4KC
                           /home/alex/eldk-mips-extra-install )

# adjust the default behavior of the FIND_XXX() commands:
# search for headers and libraries in the target environment,
# search for programs in the host environment
set (CMAKE_FIND_ROOT_PATH_MODE_PROGRAM NEVER)
set (CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set (CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)
```

The toolchain files can be located anywhere, but it is a good idea to put them in a central place so that they can be reused in multiple projects. We will save this file as `~/Toolchains/Toolchain-eldk-mips4K.cmake`. The variables mentioned above are set here: `CMAKE_SYSTEM_NAME`, the C/C++ compilers, `CMAKE_FIND_ROOT_PATH` to specify where libraries and headers for the target environment are located. The find modes are also set up so that libraries and headers are searched for in the target environment only, whereas programs are searched for in the host environment only. Now we will cross compile the hello world project from Chapter 2:

```
project (Hello)
add_executable (Hello Hello.c)
```

Run CMake, this time telling it to use the toolchain file from above:

```
mkdir Hello-eldk-mips
cd Hello-eldk-mips
cmake -DCMAKE_TOOLCHAIN_FILE=~/Toolchains/Toolchain-eldk-
mips4K.cmake ..
make VERBOSE=1
```

This should give you an executable that can run on the target platform. Thanks to the `VERBOSE=1` option you should see that the cross compiler is used. Now we will make the example a bit more sophisticated by adding system inspection and install rules. We will build and install a shared library named Tools, and then build the Hello application which links to the Tools library.

```
include (CheckIncludeFiles)
check_include_files (stdio.h HAVE_STDIO_H)

set (VERSION_MAJOR 2)
set (VERSION_MINOR 6)
set (VERSION_PATCH 0)

configure_file (config.h.in ${CMAKE_BINARY_DIR}/config.h)

add_library (Tools SHARED tools.cxx)
set_target_properties (Tools PROPERTIES
    VERSION ${VERSION_MAJOR}.${VERSION_MINOR}.${VERSION_PATCH}
    SOVERSION ${VERSION_MAJOR})

install (FILES tools.h DESTINATION include)
install (TARGETS Tools DESTINATION lib)
```

There is no difference to a normal CMakeLists.txt, no special prerequisites are required for cross compiling. The CMakeLists.txt checks that the header stdio.h is available and sets the version number for the Tools library. These are configured into config.h, which is then used in tools.cxx. The version number is also used to set the version number of the Tools library. The library and headers are installed to \${CMAKE_INSTALL_PREFIX}/lib and \${CMAKE_INSTALL_PREFIX}/include respectively. Running CMake gives this result:

```
mkdir build-eldk-mips
cd build-eldk-mips
cmake -DCMAKE_TOOLCHAIN_FILE=~/Toolchains/Toolchain-eldk-
mips4K.cmake -DCMAKE_INSTALL_PREFIX=~/eldk-mips-extra-install ..
-- The C compiler identification is GNU
-- The CXX compiler identification is GNU
-- Check for working C compiler: /home/alex/eldk-
mips/usr/bin/mips_4KC-gcc
-- Check for working C compiler: /home/alex/eldk-
mips/usr/bin/mips_4KC-gcc -- works
-- Check size of void*
-- Check size of void* - done
-- Check for working CXX compiler: /home/alex/eldk-
mips/usr/bin/mips_4KC-g++
-- Check for working CXX compiler: /home/alex/eldk-
mips/usr/bin/mips_4KC-g++ -- works
-- Looking for include files HAVE_STDIO_H
-- Looking for include files HAVE_STDIO_H - found
-- Configuring done
```

```
-- Generating done
-- Build files have been written to:
/home/alex/src/tests/Tools/build-mips
make install
Scanning dependencies of target Tools
[100%] Building CXX object CMakeFiles/Tools.dir/tools.o
Linking CXX shared library libTools.so
[100%] Built target Tools
Install the project...
-- Install configuration: ""
-- Installing /home/alex/eldk-mips-extra-install/include/tools.h
-- Installing /home/alex/eldk-mips-extra-install/lib/libTools.so
```

As can be seen in the output above, CMake detected the correct compiler, found the `stdio.h` header for the target platform and successfully generated the Makefiles. The `make` command was invoked, which then successfully built and installed the library in the specified installation directory. Now we can build an executable that uses the `Tools` library and does some system inspection:

```
project (HelloTools)

find_package (ZLIB REQUIRED)

find_library (TOOLS_LIBRARY Tools)
find_path (TOOLS_INCLUDE_DIR tools.h)

if (NOT TOOLS_LIBRARY OR NOT TOOLS_INCLUDE_DIR)
    message (FATAL_ERROR "Tools library not found")
endif (NOT TOOLS_LIBRARY OR NOT TOOLS_INCLUDE_DIR)

set (CMAKE_INCLUDE_CURRENT_DIR TRUE)
set (CMAKE_INCLUDE_DIRECTORIES_PROJECT_BEFORE TRUE)
include_directories ("${TOOLS_INCLUDE_DIR}"
                     "${ZLIB_INCLUDE_DIR}")

add_executable (HelloTools main.cpp)
target_link_libraries (HelloTools ${TOOLS_LIBRARY}
                      ${ZLIB_LIBRARIES})
set_target_properties (HelloTools PROPERTIES
                      INSTALL_RPATH_USE_LINK_PATH TRUE)

install (TARGETS HelloTools DESTINATION bin)
```

Building works in the same way as with the library, the toolchain file has to be used, and then it should just work:

```
cmake -DCMAKE_TOOLCHAIN_FILE=~/Toolchains/Toolchain-eldk-mips4K.cmake -DCMAKE_INSTALL_PREFIX=~/eldk-mips-extra-install ..  
-- The C compiler identification is GNU  
-- The CXX compiler identification is GNU  
-- Check for working C compiler: /home/alex/denx-mips/usr/bin/mips_4KC-gcc  
-- Check for working C compiler: /home/alex/denx-mips/usr/bin/mips_4KC-gcc -- works  
-- Check size of void*  
-- Check size of void* - done  
-- Check for working CXX compiler: /home/alex/denx-mips/usr/bin/mips_4KC-g++  
-- Check for working CXX compiler: /home/alex/denx-mips/usr/bin/mips_4KC-g++ -- works  
-- Found ZLIB: /home/alex/denx-mips/mips_4KC/usr/lib/libz.so  
-- Found Tools library: /home/alex/denx-mips-extra-install/lib/libTools.so  
-- Configuring done  
-- Generating done  
-- Build files have been written to:  
/home/alex/src/tests>HelloTools/build-eldk-mips  
make  
[100%] Building CXX object CMakeFiles>HelloTools.dir/main.o  
Linking CXX executable HelloTools  
[100%] Built target HelloTools
```

Obviously CMake found the correct zlib and also libTools.so, that had been installed in the previous step.

8.5 Cross Compiling for a Microcontroller

CMake can be used for more than cross compiling to targets with operating systems, it is also possible to use it in development for deeply embedded devices with small microcontrollers and no operating system at all. As an example we will use the Small Devices C Compiler (<http://sdcc.sourceforge.net>), which runs under Windows, Linux and Mac OS X, that supports 8 and 16 Bit microcontrollers. For driving the build we will use MS NMake under Windows. As before, the first step is to write a toolchain file so that CMake knows about the target platform. For sdcc it should look something like this:

```
set (CMAKE_SYSTEM_NAME Generic)
set (CMAKE_C_COMPILER "c:/Program Files/SDCC/bin/sdcc.exe")
```

The system name for targets that do not have an operating system, "Generic", should be used as the `CMAKE_SYSTEM_NAME`. The CMake platform file for "Generic" doesn't set up any specific features. All that it assumes is that the target platform does not support shared libraries, and so all properties will depend on the compiler and `CMAKE_SYSTEM_PROCESSOR`. The toolchain file above does not set the FIND-related variables. As long as none of the find commands is used in the CMake commands, this is fine. In many projects for small microcontrollers this will be the case. The `CMakeLists.txt` should look like the following:

```
project (Blink C)

add_library (blink blink.c)

add_executable (hello main.c)
target_link_libraries (hello blink)
```

There are no major differences to other `CMakeLists.txt` files. One important point is that the language "C" is enabled explicitly using the `PROJECT` command. If this is not done, CMake will also try to enable support for C++, which will fail as `sdcc` only has support for C. Running CMake and building the project should work as usual:

```
cmake -G"NMake Makefiles"
-DCMAKE_TOOLCHAIN_FILE=c:/Toolchains/Toolchain-sdcc.cmake ..
-- The C compiler identification is SDCC
-- Check for working C compiler: c:/program
files/sdcc/bin/sdcc.exe
-- Check for working C compiler: c:/program
files/sdcc/bin/sdcc.exe -- works
-- Check size of void*
-- Check size of void* - done
-- Configuring done
-- Generating done
-- Build files have been written to: C:/src/tests/blink/build

nmake
Microsoft (R) Program Maintenance Utility Version 7.10.3077
Copyright (C) Microsoft Corporation. All rights reserved.

Scanning dependencies of target blink
[ 50%] Building C object CMakeFiles/blink.dir/blink.rel
```

```

Linking C static library blink.lib
[ 50%] Built target blink
Scanning dependencies of target hello
[100%] Building C object CMakeFiles/hello.dir/main.rel
Linking C executable hello.ihx
[100%] Built target hello

```

This was a simple example using NMake with sdcc with the default settings of sdcc. Of course more sophisticated project layouts are possible. For this kind of project it is also a good idea to setup an install directory where reusable libraries can be installed, so it is easier to use them in multiple projects. It is normally necessary to choose the correct target platform for sdcc, not everybody uses i8051, which is the default for sdcc. The recommended way to do this is via setting `CMAKE_SYSTEM_PROCESSOR`.

This will cause CMake to search for and load the platform file `Platform/Generic-SDCC-C-$CMAKE_SYSTEM_PROCESSOR.cmake`. As this happens right before loading `Platform/Generic-SDCC-C.cmake`, it can be used to setup the compiler and linker flags for the specific target hardware and project. Therefore, a slightly more complex toolchain file is required:

```

get_filename_component (_ownDir
                      "${CMAKE_CURRENT_LIST_FILE}" PATH)
set (CMAKE_MODULE_PATH "${_ownDir}/Modules"
${CMAKE_MODULE_PATH})

set (CMAKE_SYSTEM_NAME Generic)
set (CMAKE_C_COMPILER "c:/Program Files/SDCC/bin/sdcc.exe")
set (CMAKE_SYSTEM_PROCESSOR "Test_DS80C400_Rev_1")

# here is the target environment located
set (CMAKE_FIND_ROOT_PATH "c:/Program Files/SDCC"
                           "c:/ds80c400-install" )

# adjust the default behavior of the FIND_XXX() commands:
# search for headers and libraries in the target environment
# search for programs in the host environment
set (CMAKE_FIND_ROOT_PATH_MODE_PROGRAM NEVER)
set (CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set (CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)

```

This toolchain file contains a few new settings, it is also about the most complicated toolchain file you should ever need. `CMAKE_SYSTEM_PROCESSOR` is set to `Test_DS80C400_Rev_1`, which is just an identifier for the specific target hardware. This has the effect that CMake will

try to load Platform/Generic-SDCC-C-Test_DS80C400_Rev_1.cmake. As this file does not exist in the CMake system module directory, the CMake variable CMAKE_MODULE_PATH has to be adjusted so that this file can be found. If this toolchain file is saved to c:/Toolchains/sdcc-ds400.cmake, the hardware specific file should be saved in c:/Toolchains/Modules/Platform/. An example of this is shown below:

```
set (CMAKE_C_FLAGS_INIT "-mds390 --use-accelerator")
set (CMAKE_EXE_LINKER_FLAGS_INIT "")
```

This will select the DS80C390 as the target platform and add the --use-accelerator argument to the default compile flags. In this example the "NMake Makefiles" generator was used. In the same way e.g. the "MinGW Makefiles" generator could be used if GNU make from MinGW, or another Windows version of GNU make, are available. At least version 3.78 is required, or the "Unix Makefiles" generator under UNIX. Also any Makefile-based IDE-project generators could be used, e.g. the Eclipse, CodeBlocks, or the KDevelop3 generator.

8.6 Cross Compiling an Existing Project

Existing CMake based projects may need some work so that they can be cross compiled, other projects may work without any modifications. One such project is FLTK, the Fast Lightweight Toolkit. We will compile FLTK on a Linux machine using the MinGW cross compiler for Windows.

The first step is to install the MinGW cross compiler. For some Linux distributions there are ready-to-use binary packages, for Debian the package name is mingw32. Once this is installed you need to setup a toolchain file for this as described above. It should look something like this:

```
# the name of the target operating system
set (CMAKE_SYSTEM_NAME Windows)

# which compiler to use
set (CMAKE_C_COMPILER i586-mingw32msvc-gcc)
set (CMAKE_CXX_COMPILER i586-mingw32msvc-g++)

# where are the target libraries and headers installed ?
set(CMAKE_FIND_ROOT_PATH /usr/i586-mingw32msvc
                           /home/alex/mingw-install )

# find_program() should by default NEVER search the target tree
# adjust the default behavior of the FIND_XX() commands:
# search for headers and libraries in the target environment
```

```
# search for programs in the host environment
set (CMAKE_FIND_ROOT_MODE_PROGRAM NEVER)
set (CMAKE_FIND_ROOT_MODE_LIBRARY ONLY)
set (CMAKE_FIND_ROOT_MODE_INCLUDE ONLY)
```

Once this is working, run CMake with the appropriate options on FLTK:

```
mkdir build-mingw
cd build-mingw
cmake -DCMAKE_TOOLCHAIN_FILE=~/Toolchains/Toolchain-
mingw32.cmake -DCMAKE_INSTALL_PREFIX=~/mingw-install ..
-- The C compiler identification is GNU
-- The CXX compiler identification is GNU
-- Check for working C compiler: /usr/bin/i586-mingw32msvc-gcc
-- Check for working C compiler: /usr/bin/i586-mingw32msvc-gcc -
- works
...
```

In FLTK the `utility_source` command is used to build the executable fluid, whose location is put into the CMake variable `FLUID_COMMAND`. If you intend to run this executable, you need to preload the cache with the full path to a version of that program that can be run on the build host.

```
-- Configuring done
-- Generating done
-- Build files have been written to: /home/alex/src/fltk-1.1.x-
r5940/build-mingw
```

Below you can see a warning from CMake about the use of the `utility_source` command. To find out more CMake offers the `--debug-output` argument:

```
rm -rf *
cmake -DCMAKE_TOOLCHAIN_FILE=~/Toolchains/Toolchain-
mingw32.cmake -DCMAKE_INSTALL_PREFIX=~/mingw-install .. --debug-
output
...
UTILITY_SOURCE is used in cross compiling mode for
FLUID_COMMAND. If your intention is to run this executable, you
need to preload the cache with the full path to a version of
that program, which runs on this build machine.
Called from: [1] /home/alex/src/fltk-1.1.x-r5940/CMakeLists.txt
```

This tells us that `utility_source` has been called from `/home/alex/src/fltk-1.1.x-r5940/CMakeLists.txt`, then CMake processed some more directories, finally it created Makefiles in each subdirectory. Examining the top level CMakeLists.txt shows the following:

```
# Set the fluid executable path
utility_source (FLUID_COMMAND fluid fluid fluid.cxx)
set (FLUID_COMMAND "${FLUID_COMMAND}" CACHE INTERNAL "" FORCE)
```

Apparently `FLUID_COMMAND` is used to hold the path for the executable fluid, which is built by the project. Fluid is used during the build to generate code, so the cross compiled executable will not work, instead a native fluid has to be used. In the following example the variable `FLUID_COMMAND` is set to the location of a fluid executable for the build host, which is then used in the cross compiling build to generate code that will be compiled for the target system:

```
cmake . -DFLUID_COMMAND=/home/alex/src/download/fltk-1.1.x-r5940/build-native/bin/fluid
...
-- Configuring done
-- Generating done
make
Scanning dependencies of target fltk_zlib
[ 0%] Building C object
zlib/CMakeFiles/fltk_zlib.dir/adler32.obj
[ 0%] Building C object
zlib/CMakeFiles/fltk_zlib.dir/compress.obj
...
Scanning dependencies of target valuators
[100%] Building CXX object
test/CMakeFiles/valuators.dir/valuators.obj
Linking CXX executable ../bin/valuators.exe
[100%] Built target valuators
```

That's it, the executables are now in `mingw-bin/`, and can be run via wine or by copying them to a Windows system.

8.7 Cross Compiling a Complex Project - VTK

Building a complex project is a multi-step process. Complex in this case means that the project uses tests that run executables, and that it builds executables which are used later in the build to generate code (or something similar). One such project is VTK, the Visualization

Toolkit. It uses several `try_run` tests and creates several code generators. When running CMake on the project, every `try_run` command will produce an error message and at the end there will be a `TryRunResults.cmake` file in the build directory. You need to go through all of the entries of this file and fill in the appropriate values. If you are uncertain about the correct result, you can also try to execute the test binaries on the real target platform, they are saved in the binary directory.

- VTK contains several code generators, one of which is `ProcessShader`. These code generators are added using `add_executable` and `get_target_property(LOCATION)` is used to get the locations of the resulting binaries, which are then used in `add_custom_command` or `add_custom_target` commands. Since the cross compiled executables cannot be executed during the build, the `add_executable` calls are surrounded by `if (NOT CMAKE_CROSSCOMPILING)` commands and the executable targets are imported into the project using the `add_executable` command with the `IMPORTED` option. These import statements are in the file `VTKCompileToolsConfig.cmake`, which does not have to be created manually, but it is created by a native build of VTK.

So in order to cross compile VTK you need

- install a toolchain and create a toolchain file for CMake
- build VTK natively for the build host
- run CMake for the target platform
- complete `TryRunResults.cmake`
- use the `VTKCompileToolsConfig.cmake` file from the native build
- finally build

So first, build a native VTK for the build host using the standard procedure.

```
cvs -d :pserver:anonymous@public.kitware.com:/cvsroot/VTK co VTK
cd VTK
mkdir build-native; cd build-native
cmake ..
make
```

Ensure that all required options are enabled using `ccmake`, e.g. if you need Python wrapping for the target platform you must enable Python wrapping in `build-native/`. Once this build has finished, there will be a `VTKCompileToolsConfig.cmake` file in `build-native/`. If this succeeded, we can continue to cross compiling the project, in this example for an IBM BlueGene supercomputer.

```
cd VTK
mkdir build-bgl-gcc
cd build-bgl-gcc
cmake -DCMAKE_TOOLCHAIN_FILE=~/Toolchains/Toolchain-BlueGeneL-
gcc.cmake -DVTKCompileTools_DIR=~/VTK/build-native/ ..
```

This will finish with an error message for each `try_run` and a `TryRunResults.cmake` file, that you have to complete as described above. You should save the file to a safe location, otherwise it will be overwritten on the next CMake run.

```
cp TryRunResults.cmake ../TryRunResults-VTK-BlueGeneL-gcc.cmake
ccmake -C ../TryRunResults-VTK-BlueGeneL-gcc.cmake .
...
make
```

On the second run of `ccmake` all the other arguments can be skipped as they are now in the cache. It is possible to point CMake to the build directory that contains a `CMakeCache.txt`, so CMake will figure out that this is the build directory.

8.8 Some Tips and Tricks

Dealing with `try_run` tests

In order to make cross compiling your project easier, try to avoid `try_run` tests and use other methods to test something instead. For examples of how this can be done consider the tests for endianess in `CMake/Modules/TestBigEndian.cmake`, and the test for the compiler id using the source file `CMake/Modules/CMakeCCompilerId.c`. In both `try_compile` is used to compile the source file into an executable, where the desired information is encoded into a text string. Using the `COPY_FILE` option of `try_compile` this executable is copied to a temporary location and then all strings are extracted from this file using `file (STRINGS)`. The test result is obtained using regular expressions to get the information from the string.

If you cannot avoid `try_run` tests, try to use just the exit code from the run, not the output of the process. That way it will not be necessary to set both the exit code and the `stdout` and `stderr` variables for the `try_run` test when cross compiling. This allows the `OUTPUT_VARIABLE` or the `RUN_OUTPUT_VARIABLE` options for `try_run` to be omitted.

If you have done that, created and completed a correct `TryRunResults.cmake` file for the target platform, you might consider adding this file to the sources of the project, so that it can be reused by others. These files are per-target per-toolchain.

Target platform and toolchain issues

If your toolchain is not able to build a simple program without special arguments, like e.g. a linker script file or a memory layout file, the tests CMake does initially will fail. To make it work anyway, there is a CMake module, `CMakeForceCompiler`, that offers the following macros:

```
CMAKE_FORCE_SYSTEM (name version processor),  
CMAKE_FORCE_C_COMPILER (compiler compiler_id sizeof_void_p)  
CMAKE_FORCE_CXX_COMPILER (compiler compiler_id).
```

These macros can be used in a toolchain file so that the required variables will be preset and the CMake tests avoided.

RPATH handling under UNIX

For native builds CMake builds executables and libraries by default with RPATH. In the build tree the RPATH is set so that the executables can be run from the build tree, i.e. the RPATH points into the build tree. When installing the project, CMake links the executables again, this time with the RPATH for the install tree, which is empty by default.

When cross compiling you probably want to set up RPATH handling differently, as the executable cannot run on the build host it makes more sense to build it with the install RPATH right from the start. There are several CMake variables and target properties for adjusting RPATH handling.

```
set (CMAKE_BUILD_WITH_INSTALL_RPATH TRUE)  
set (CMAKE_INSTALL_RPATH "<whatever you need>")
```

With these two settings the targets will be built with the install RPATH instead of the build RPATH, this avoids the need to link them again when installing. If you don't need RPATH support in your project, you don't need to set `CMAKE_INSTALL_RPATH`, it is empty by default.

Setting `CMAKE_INSTALL_RPATH_USE_LINK_PATH` to `TRUE` is useful for native builds, since it automatically collects the RPATH from all libraries against which a targets links. For cross compiling it should be left at the default setting, which is `FALSE`, because on the target the automatically generated RPATH will be wrong in most cases, it will probably have a different file system layout to the build host.

Packaging with CPack

CPack is a powerful, easy to use, cross-platform software packaging tool distributed with CMake since version 2.4.2. It uses the generators concept from CMake to abstract package generation on specific platforms. It can be used with or without CMake, but it may depend on some software being installed on the system. Using a simple configuration file, or using a CMake module, the author of a project can package a complex project into a simple installer. This chapter will describe how to apply CPack to a CMake project.

9.1 CPack Basics

Users of your software may not always want to, or be able to, build the software in order to install it. The software maybe closed source, or it may take a long time to compile, or in the case of an end user application the users may have no interest in building the application. For these cases, what is needed is a way to build the software on one machine, and then move the install tree to a different machine. The most basic way to do this is to use the DESTDIR environment variable to install the software into a temporary location, then to tar or zip up that directory and move it to the another machine. Another more powerful approach is to use the CPack tool included in CMake.

CPack is a tool included with CMake, it can be used to create installers and packages for projects. CPack can create two basic types of packages, source and binary. CPack works in much the same way as CMake does for building software. It does not aim to replace native packaging tools, rather it provides a single interface to a variety of tools. Currently CPack supports the creation of Windows installers using NullSoft installer NSIS, Mac OS X Package Maker tool, OS X Drag and Drop, OS X X11 Drag and Drop, Cygwin Setup packages,

Debian packages, RPMs, .tar.gz, .sh (self extracting .tar.gz files), .zip compressed files. The implementation of CPack works in a similar way to CMake. For each type of packaging tool supported, there is a CPack generator written in C++ that is used to run the native tool and create the package. For simple tar based packages, CPack includes a library version of tar and does not require tar to be installed on the system. For many of the other installers, native tools must be present for CPack to function.

With source packages, CPack makes a copy of the source tree and creates a zip or tar file. For binary packages, the use of CPack is tied to the install commands working correctly for a project. When setting up install commands, the first step is to make sure the files go into the correct directory structure with the correct permissions. The next step is to make sure the software is relocatable and can run in an installed tree. This may require changing the software itself, and there are many techniques to do that for different environments that go beyond the scope of this book. Basically, executables should be able to find data or other files using relative paths to the location of where it is installed. CPack installs the software into a temporary directory, and copies the install tree into the format of the native packaging tool. Once the install commands have been added to a project, enabling CPack in the simplest case is done by including the CPack.cmake file into the project.

Simple Example

The most basic CPack project would look like this:

```
project(CoolStuff)
add_executable(coolstuff coolstuff.cxx)
install(TARGETS coolstuff RUNTIME DESTINATION bin)
include(CPack)
```

In the CoolStuff project, an executable is created and installed into the directory bin. Then the CPack file is included by the project. At this point project CoolStuff will have CPack enabled. To run CPack for a CoolStuff, you would first build the project as you would any other CMake project. CPack adds several targets to the generated project. These targets in Makefiles are package and package_source, and PACKAGE in Visual Studio and Xcode. For example, to build a source and binary package for CoolStuff using a Makefile generator you would run the following commands:

```
mkdir build
cd build
cmake ..../CoolStuff
make
make package
make package_source
```

This would create a source zip file called `CoolStuff-0.1.1-Source.zip`, a NSIS installer called `CoolStuff-0.1.1-win32.exe`, and a binary zip file `CoolStuff-0.1.1-win32.zip`. The same thing could be done using the CPack command line.

```
cd build  
cpack -C CPackConfig.cmake  
cpack -C CPackSourceConfig.cmake
```

What Happens When CPack.cmake Is Included?

When the `include(CPack)` command is executed, the `CPack.cmake` file is included into the project. By default this will use the `configure_file` command to create `CPackConfig.cmake` and `CPackSourceConfig.cmake` in the binary tree of the project. These files contain a series of set commands that set variables for use when CPack is run during the packaging step. The names of the files that are configured by the `CPack.cmake` file can be customized with these two variables; `CPACK_OUTPUT_CONFIG_FILE` which defaults to `CPackConfig.cmake` and `CPACK_SOURCE_OUTPUT_CONFIG_FILE` which defaults to `CPackSourceConfig.cmake`.

The source for these files can be found in the `Templates/CPackConfig.cmake.in`. This file contains some comments, and a single variable that is set by `CPack.cmake`. The file contains this line of CMake code:

```
@_CPACK_OTHER_VARIABLES_@
```

If the project contains the file `CPackConfig.cmake.in` in the top level of the source tree, that file will be used instead of the file in the `Templates` directory. If the project contains the file `CPackSourceConfig.cmake.in`, then that file will be used for the creation of `CPackSourceConfig.cmake`.

The configuration files created by `CPack.cmake` will contain all variables that begin with “`CPACK_`” in the current project. This is done using the command:

```
get_cmake_property(res VARIABLES)
```

The above command gets all variables defined for the current CMake project. Some CMake code then looks for all variables starting with “`CPACK_`”, and each variable found is configured into the two configuration files as CMake code. For example, if you had a variable set like this in your CMake project:

```
set(CPACK_PACKAGE_NAME "CoolStuff")
```

`CPackConfig.cmake` and `CPackSourceConfig.cmake` would have the same thing in them:

```
set(CPACK_PACKAGE_NAME "CoolStuff")
```

It is important to keep in mind that CPack is run after CMake on the project. CPack uses the same parser as CMake, but will not have the same variable values as the CMake project. It will only have the variables that start with `CPACK_`, and these variables will be configured into a configuration file by CMake. This can cause some errors and confusion if the values of the variables use escape characters. Since they are getting parsed twice by the CMake language, they will need double the level of escaping. For example, if you had the following in your CMake project:

```
set(CPACK_PACKAGE_NAME "Cool \"Stuff\"")
```

The resulting CPack files would have this:

```
set(CPACK_FOOBAR "Cool \"Stuff\"")
```

That would not be exactly what you would want or might expect. To get around this problem, there are two solutions. The first is to add an additional level of escapes to the original set command like this:

```
(CPACK_PACKAGE_NAME "Cool \\\\"Stuff\\\\\"")
```

This would result in the correct `set` command which would look like this:

```
set(CPACK_FOOBAR "Cool \"Stuff\"")
```

The second solution to the escaping problem is explained in the next section.

Adding Custom CPack Options

To avoid the escaping problem a project specific CPack configure file can be specified. This file will be loaded by CPack after `CPackConfig.cmake` or `CPackSourceConfig.cmake` is loaded, and `CPACK_GENERATOR` will be set to the CPack generator being run. Variables set in this file only require one level of CMake escapes. This file can be configured or not, and contains regular CMake code. In the example above, you could move `CPACK_FOOBAR` into a file `MyCPackOptions.cmake.in` and configure that file into the build tree of the project. Then set the project configuration file path like this:

```
configure_file ("${PROJECT_SOURCE_DIR}/MyCPackOptions.cmake.in"
                "${PROJECT_BINARY_DIR}/MyCPackOptions.cmake"
                @ONLY)
set (CPACK_PROJECT_CONFIG_FILE
    "${PROJECT_BINARY_DIR}/CMakeCPackOptions.cmake")
```

Where `MyCPackOptions.cmake.in` contained:

```
(CPACK_FOOBAR "Cool \"Stuff\\\"")
```

The `CPACK_PROJECT_CONFIG_FILE` variable should contain the full path to the CPack config file for the project, as seen in the above example. This has the added advantage that the CMake code can contain if statements based on the `CPACK_GENERATOR` value, so that packager specific values can be set for a project. For example, the CMake project sets the icon for the installer in this file:

```
set (CPACK_NSIS_MUI_ICON
    "@${CMAKE_SOURCE_DIR@}/Utilities/Release\\CMakeLogo.ico")
```

Note that the path has forward slashes except for the last part which has an escaped \ as the path separator. As of the writing of this book, NSIS needed the last part of the path to have a Windows style slash. If you do not do this, you may get the following error:

```
File: ".../Release/CMakeLogo.ico" -> no files found.
Usage: File [/nonfatal] [/a] ([/r] [/x filespec [...]]]
      filespec [...] | /oname=outfile one_file_only)
```

Options Added by CPack

In addition to creating the two configuration files, `CPack.cmake` will add some advanced options to your project. The options added depend on the environment and OS that CMake is running on, and control the default packages that are created by CPack. These options are of the form `CPACK_<CPack Generator Name>`, where generator names available on each platform can be found in the following table:

Windows	Cygwin	Linux/UNIX	Mac OS X
NSIS	CYGWIN_BINARY	DEB	PACKAGEMAKER
ZIP	SOURCE_CYGWIN	RPM	OSXX11
SOURCE_ZIP		STGZ	Drag and Drop
		TBZ2	
		TZ	
		SOURCE_TGZ	
		SOURCE_TZ	

Turning these options on or off affects the packages that are created when running CPack with no options. If the option is off in the CMakeCache.txt file for the project, you can still build that package type by specifying the -G option to the CPack command line.

9.2 CPack Source Packages

Source packages in CPack simply copy the entire source tree for a project into a package file, and no install rules are used as they are in the case of binary packages. Out of source builds should be used to avoid having extra binary stuff polluting the source package. If you have files or directories in your source tree that are not wanted in the source package, you can use the variable `CPACK_SOURCE_IGNORE_FILES` to exclude things from the package. This variable contains a list of regular expressions. Any file or directory that matches a regular expression in that list will be excluded from the sources. The default setting is as follows:

```
"/CVS/;\\\\\\\\\\\\\\\\.svn/;\\\\\\\\\\\\\\\\.swp$;\\\\\\\\\\\\\\\\.#;/#"
```

There are many levels of escapes used in the default value as this variable is parsed by CMake once and CPack again. It is important to realize that the source tree will not use any install commands, it will simply copy the entire source tree minus the files it is told to ignore into the package. To avoid the multiple levels of escape, the file referenced by `CPACK_PROJECT_CONFIG_FILE` should be used to set this variable. The expression is a regular expression and not a wild card statement, see section 4.4 for more information about CMake regular expressions.

9.3 CPack Installer Commands

Since binary packages require CPack to interact with the install rules of the project being packaged, this section will cover some of the options CPack provides to interact with the

install rules of a project. CPack can work with CMake's install scripts or with external install commands.

CPack and CMake install commands

In most CMake projects, using the `cmake install` rules will be sufficient to create the desired package. By default CPack will run the install rule for the current project. However, if you have a more complicated project, you can specify sub-projects and install directories with the variable `CPACK_INSTALL_CMAKE_PROJECTS`. This variable should hold quadruplets of install directory, install project name, install component, and install subdirectory. For example, if you had a project with a sub project called `MySub` that was compiled into a directory called `SubProject`, and you wanted to install all of its components, you would have this:

```
SET(CPACK_INSTALL_CMAKE_PROJECTS "SubProject;MySub;ALL;/")
```

CPack and DESTDIR

By default CPack does not use the `DESTDIR` option during the installation phase. Instead it sets the `CMAKE_INSTALL_PREFIX` to the full path of the temporary directory being used by CPack to stage the install package. This can be changed by setting `CPACK_SET_DESTDIR` to on. If the `DESTDIR` option is on, CPack will use the project's cache value for `CPACK_INSTALL_PREFIX`, and set `DESTDIR` to the temporary staging area. This allows absolute paths to be installed under the temporary directory. Relative paths are installed into `DESTDIR/${project's CMAKE_INSTALL_PREFIX}` where `DESTDIR` is set to the temporary staging area.

When doing a non-`DESTDIR` install for packaging, which is the default, any absolute paths are installed into absolute directories, and not into the package. Therefore, projects that do not use the `DESTDIR` option, must not use any absolute paths in install rules. Conversely, projects that use absolute paths, must use the `DESDIR` option.

One other variable can be used to control the root path into which projects are installed into, the `CPACK_PACKAGING_INSTALL_PREFIX`. By default many of the generators install into the directory `/usr`. That variable can be used to change that to any directory, including just `/`.

CPack and other installed directories

It is possible to run other install rules if the project is not CMake based. This can be done by using the variables `CPACK_INSTALL_COMMANDS`, and `CPACK_INSTALLED_DIRECTORIES`. `CPACK_INSTALL_COMMANDS` are commands that will be run during the installation phase of the packaging. `CPACK_INSTALLED_DIRECTORIES` should contain pairs of directory and subdirectory. The subdirectory can be `'.'` to be installed in the top-level directory of the installation. The files in each directory will be copied to the corresponding subdirectory of the CPack staging directory and packaged with the rest of the files.

9.4 CPack for Windows Installer NSIS

To create Windows style wizard based self extracting executables, CPack uses NSIS (NullSoft Scriptable Install System). More information about NSIS can be found at the NSIS home page: <http://nsis.sourceforge.net/> NSIS is a powerful tool with a scripting language used to create professional Windows installers. To create Windows installers with CPack, you will need NSIS installed on your machine.

CPack uses configured template files to control NSIS. There are two files configured by CPack during the creation of a NSIS installer. Both files are found in the CMake Modules directory. `Modules/NSIS.template.in` is the template for the NSIS script, and `Modules/NSIS.InstallOptions.ini.in` is the template for the modern user interface or MUI used by NSIS. The install options file contains the information about the pages used in the install wizard. This section will describe how to configure CPack to create an NSIS install wizard.

CPack Variables Used by CMake for NSIS

This section contains screen captures from the CMake NSIS install wizard. For each part of the installer that can be changed or controlled from CPack, the variables and values used are given.

The first thing that a user will see of the installer in Windows is the icon for the installer executable itself. By default the installer will have the Null Soft Installer icon, as seen in Figure 7 for the 20071023 cmake installer. This icon can be changed by setting the variable `CPACK_NSIS_MUI_ICON`. The installer for 20071025 in the same figure shows the CMake icon being used for the installer.

	cmake-2.5.20071023-win32-x86.exe	6,290 KB	Application
	cmake-2.5.20071025-win32-x86.exe	6,324 KB	Application

Figure 7 Icon for installer in Windows Explorer

The last thing a users will see of the installer in Windows is the icon for the uninstall executable, as seen in Figure 8. This option can be set with the `CPACK_NSIS_MUI_UNIICON` variable. Both the install and uninstall icons must be the same size and format, a valid windows .ico file usable by Windows Explorer. The icons are set like this:

```
# set the install/uninstall icon used for the installer itself
set (CPACK_NSIS_MUI_ICON
    "${CMake_SOURCE_DIR}/Utilities/Release\\CMakeLogo.ico")
set (CPACK_NSIS_MUI_UNIICON
    "${CMake_SOURCE_DIR}/Utilities/Release\\CMakeLogo.ico")
```

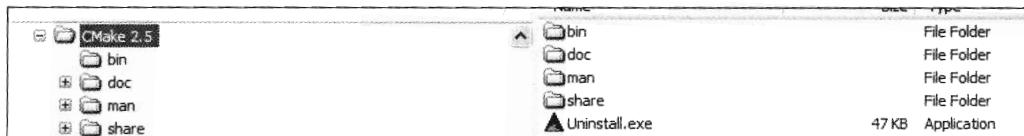


Figure 8 Uninstall Icon for NSIS installer

On Windows, programs can also be removed using the Add or Remove Programs tool from the control panel as seen in Figure 9. The icon for this should be embedded in one of the installed executables. This can be set like this:

```
# set the add/remove programs icon using an installed executable
SET(CPACK_NSIS_INSTALLED_ICON_NAME "bin\\cmake-gui.exe")
```

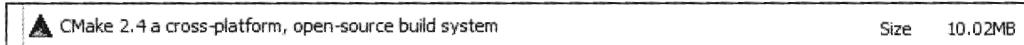


Figure 9 Add or Remove Programs Entry



Figure 10 First Screen of Install Wizard

When running the installer, the first screen of the wizard will look like Figure 10. In this screen you can control the name of the project that shows up in two places on the screen. The name used for the project is controlled by the variable

CPACK_PACKAGE_INSTALL_DIRECTORY or CPACK_NSIS_PACKAGE_NAME. In this example, it was set to “CMake 2.5” like this:

```
set (CPACK_PACKAGE_INSTALL_DIRECTORY "CMake
${CMake_VERSION_MAJOR}.${CMake_VERSION_MINOR}")

# or this
set(CPACK_NSIS_PACKAGE_NAME "CMake
${CMake_VERSION_MAJOR}.${CMake_VERSION_MINOR}")
```

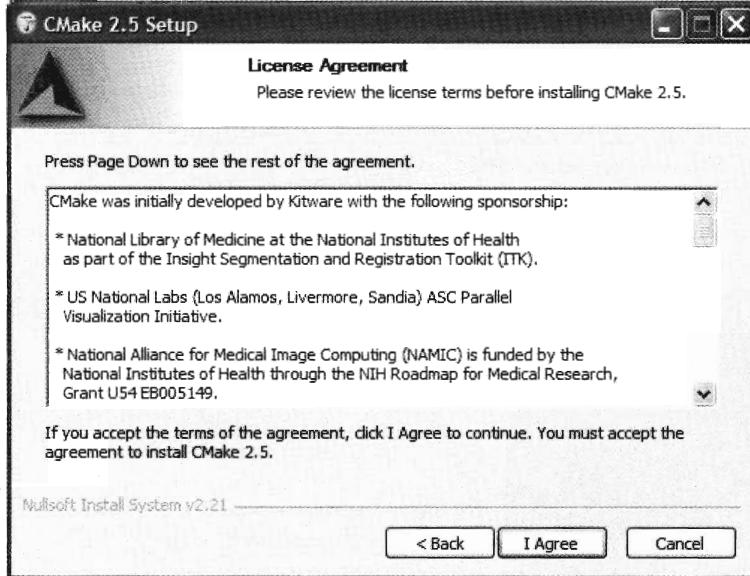


Figure 11 Second Screen of Install Wizard

The second page of the install wizard can be seen in Figure 11. This screen contains the license agreement and there are several things that can be configured on this page. The banner bitmap to the left of the “License Agreement” label is controlled by the variable CPACK_PACKAGE_ICON like this:

```
set (CPACK_PACKAGE_ICON
"${CMake_SOURCE_DIR}/Utilities/Release\\CMakeInstall.bmp")
```

CPACK_PACKAGE_INSTALL_DIRECTORY is used again on this page everywhere you see the text “CMake 2.5”. The text of the license agreement is set to the contents of the file specified in the CPACK_RESOURCE_FILE_LICENSE variable. CMake does the following:

```
set (CPACK_RESOURCE_FILE_LICENSE  
    "${CMAKE_CURRENT_SOURCE_DIR}/Copyright.txt")
```

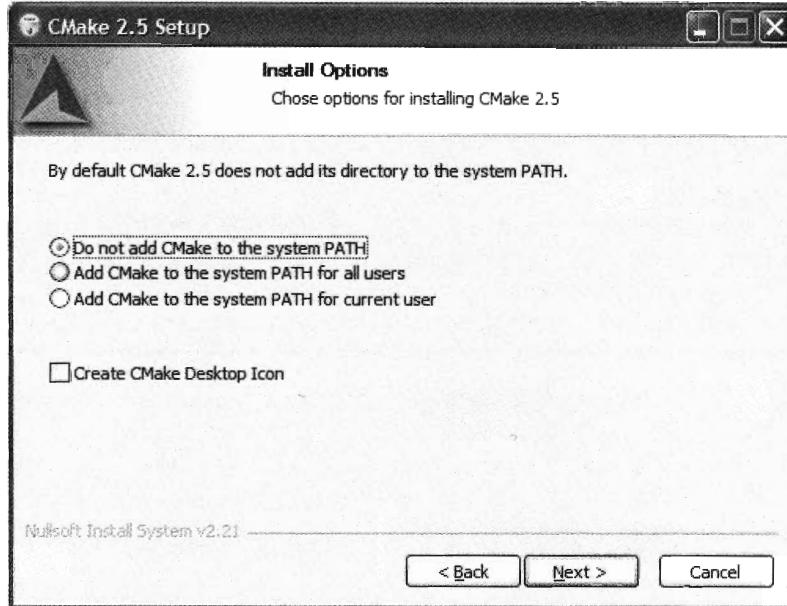


Figure 12 Third page of installer wizard

The third page of the installer can be seen in Figure 12. This page will only show up if CPACK_NSIS MODIFY_PATH is set to on. If you check the Create “name” Desktop Icon button, and you put executable names in the variable CPACK_CREATE_DESKTOP_LINKS, then a desktop icon for those executables will be created. For example, to create a desktop icon for the cmake-gui program of CMake, the following is done:

```
set (CPACK_CREATE_DESKTOP_LINKS cmake-gui)
```

Multiple desktop links can be created if your application contains more than one executable. The link will be created to the Start Menu entry, so CPACK_PACKAGE_EXECUTABLES, which is described later in this section, must also contain the application in order for a desktop link to be created.

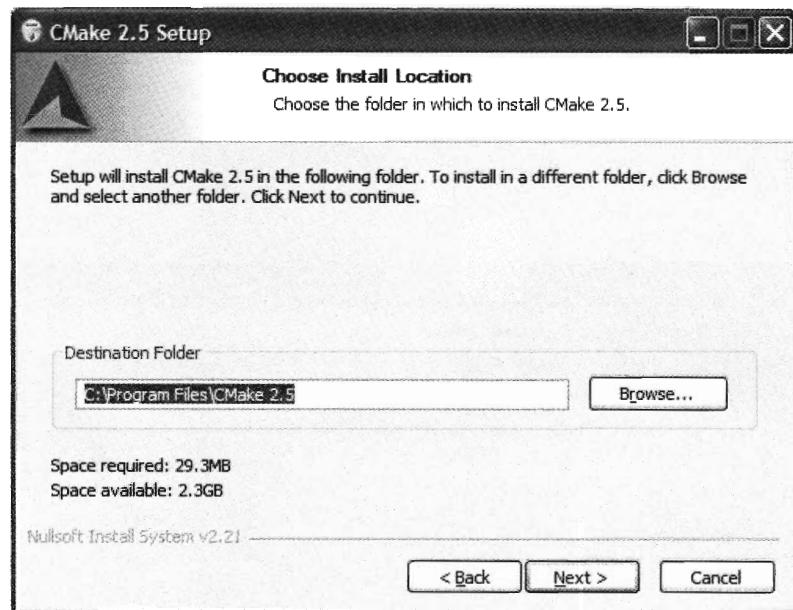


Figure 13 Fourth page of installer wizard

The fourth page of the installer seen in Figure 13 uses the variable `CPACK_PACKAGE_INSTALL_DIRECTORY` to specify the default destination folder in Program Files. The following cmake code was used to set that default:

```
set (CPACK_PACKAGE_INSTALL_DIRECTORY "CMake
${CMake_VERSION_MAJOR}.${CMake_VERSION_MINOR}")
```

The remaining pages of the installer wizard do not use any additional CPack variables, and are not included in this section. Another important option that can be set by the NSIS CPack generator is the registry key used. There are several CPack variables that control the default key used. The key is defined in the `NSIS.template.in` file as follows:

```
Software\@CPACK_PACKAGE_VENDOR\@CPACK_PACKAGE_INSTALL_REGISTRY_KEY@
```

Where `CPACK_PACKAGE_VENDOR` defaults to `Humanity`, and `CPACK_PACKAGE_INSTALL_REGISTRY_KEY` defaults to `"${CPACK_PACKAGE_NAME}${CPACK_PACKAGE_VERSION}"`

So for CMake 2.5.20071025 the registry key would look like this:

```
HKEY_LOCAL_MACHINE\SOFTWARE\Kitware\CMake 2.5.20071025
```

Creating Windows Short Cuts in the Start Menu

There are two variables that control the short cuts that are created in the Windows Start menu by NSIS. The variables contain lists of pairs, and must have an even number of elements to work correctly. The first is `CPACK_PACKAGE_EXECUTABLES`, it should contain the name of the executable followed by the name of the short cut. For example in the case of CMake, the executable is called `cmake-gui`, but the shortcut is named “CMake”. CMake does the following to create that short cut:

```
set (CPACK_PACKAGE_EXECUTABLES "cmake-gui" "CMake" )
```

The second is `CPACK_NSIS_MENU_LINKS`. This variable contains arbitrary links into the install tree, or to external web pages. The first of the pair is always the existing source file or location, and the second is the name that will show up in the Start menu. To add a link to the help file for `cmake-gui` and a link to the CMake web page the following add the following:

```
set (CPACK_NSIS_MENU_LINKS  
"doc/cmake-${VERSION_MAJOR}.${VERSION_MINOR}/cmake-gui.html"  
"cmake-gui Help" "http://www.cmake.org" "CMake Web Site")
```

Advanced NSIS CPack Options

In addition to the variables already discussed, CPack provides a few variables that are directly configured into the NSIS script file. These can be used to add NSIS script fragments to the final NSIS script used to create the installer. They are as follows:

CPACK_NSIS_EXTRA_INSTALL_COMMANDS

Extra commands used during install.

CPACK_NSIS_EXTRA_UNINSTALL_COMMANDS

Extra commands used during uninstall.

CPACK_NSIS_CREATE_ICONS_EXTRA

Extra NSIS commands in the icon section of the script.

CPACK_NSIS_DELETE_ICONS_EXTRA

Extra NSIS commands in the delete icons section of the script.

When using these variables the NSIS documentation should be referenced, and the author should look at the NSIS.template.in file for the exact placement of the variables.

Setting File Extension Associations With NSIS

One example of a useful thing to do with the extra install commands is to create associations from file extensions to the installed application. For example, if you had an application CoolStuff that could open files with the extension .cool, you would set the following extra install and uninstall commands:

```
set (CPACK_NSIS_EXTRA_INSTALL_COMMANDS "
    WriteRegStr HKCR '.cool' '\"CoolFile'
    WriteRegStr HKCR 'CoolFile' '\"Cool Stuff File'
    WriteRegStr HKCR 'CoolFile\\shell' '\"open'
    WriteRegStr HKCR 'CoolFile\\DefaultIcon' \\
        '\"$INSTDIR\\bin\\coolstuff.exe,0'
    WriteRegStr HKCR 'CoolFile\\shell\\open\\command' \\
        '\"$INSTDIR\\bin\\coolstuff.exe \"%1\"'
    WriteRegStr HKCR '\"CoolFile\\shell\\edit' \\
        '\"Edit Cool File'
    WriteRegStr HKCR 'CoolFile\\shell\\edit\\command' \\
        '\"$INSTDIR\\bin\\coolstuff.exe \"%1\"'
    System::Call \\
        'Shell32::SHChangeNotify(i 0x8000000, i 0, i 0, i 0)'
    ")

set (CPACK_NSIS_EXTRA_UNINSTALL_COMMANDS "
    DeleteRegKey HKCR '.cool'
    DeleteRegKey HKCR 'CoolFile'
")
```

This creates a Windows file association to all files ending in .cool, so that when a user double clicks on a .cool file coolstuff.exe is run with the full path to the file as an argument. This also sets up an association for editing the file from the windows right click menu to the same coolstuff.exe program. The Windows explorer icon for the file is set to the icon found in the coolstuff.exe executable. When it is uninstalled, the registry keys are removed. Since the double quotes and Windows path separators must be escaped, it is best put this code into the CPACK_PROJECT_CONFIG_FILE for the project.

```
configure_file(
    ${CoolStuff_SOURCE_DIR}/CoolStuffCPackOptions.cmake.in
    ${CoolStuff_BINARY_DIR}/CoolStuffCPackOptions.cmake @ONLY)
```

```
set (CPACK_PROJECT_CONFIG_FILE  
    ${CoolStuff_BINARY_DIR}/CoolStuffCPackOptions.cmake)  
include (CPack)
```

Installing Microsoft Run Time Libraries

Although not strictly an NSIS CPack command, if you are creating applications on Windows with the Microsoft compiler, you will most likely want to distribute the run time libraries from Microsoft alongside your project. In CMake, all you need to do is the following:

```
include (InstallRequiredSystemLibraries)
```

This will add the compiler run time libraries as install files that will go into the bin directory of your application. If you do not want the libraries to go into the bin directory, you would do this:

```
set CMAKE_INSTALL_SYSTEM_RUNTIME_LIBS_SKIP TRUE)  
include InstallRequiredSystemLibraries)  
install PROGRAMS ${CMAKE_INSTALL_SYSTEM_RUNTIME_LIBS}  
DESTINATION mydir)
```

It is important to note that the run time libraries must be right next to the executables of your package in order for Windows to find them. As of Visual Studio 2005, side by side manifest files are also required to be installed with your application when distributing the run time libraries. If you want to package a debug version of your software you will need to set CMAKE_INSTALL_DEBUG_LIBRARIES to ON prior to the include.

CPack Component Install Support

By default, CPack's installers consider all of the files installed by a project as a single, monolithic unit: either the whole set of files is installed, or none of the files are installed. However, with many projects it makes sense for the installation to be subdivided into distinct, user-selectable components: some users may want to install only the command-line tools for a project, while other users might want the GUI or the header files.

This section describes how to configure CPack to generate component-based installers that allow users to select the set of project components that they wish to install. As an example a simple installer will be created for a library that has three components: a library binary, a sample application, and a C++ header file. When finished the resulting installers for Windows and Mac OS X look like the ones in Figure 14.

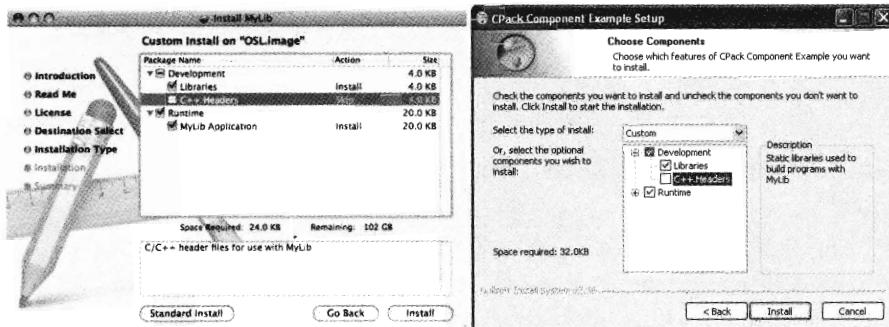


Figure 14 Mac and Windows Component Installers

The simple example we will be working with is as follows, it has a library and an executable. CPack commands that have already been covered are used.

```

cmake_minimum_required(VERSION 2.6.0 FATAL_ERROR)
project(MyLib)

add_library(mylib mylib.cpp)

add_executable(mylibapp mylibapp.cpp)
target_link_libraries(mylibapp mylib)

install(TARGETS mylib ARCHIVE DESTINATION lib)
install(TARGETS mylibapp RUNTIME DESTINATION bin)
install(FILES mylib.h DESTINATION include)
# add CPack to project
set(CPACK_PACKAGE_NAME "MyLib")
set(CPACK_PACKAGE_VENDOR "CMake.org")
set(CPACK_PACKAGE_DESCRIPTION_SUMMARY "MyLib - CPack Component Installation Example")
set(CPACK_PACKAGE_VERSION "1.0.0")
set(CPACK_PACKAGE_VERSION_MAJOR "1")
set(CPACK_PACKAGE_VERSION_MINOR "0")
set(CPACK_PACKAGE_VERSION_PATCH "0")
set(CPACK_PACKAGE_INSTALL_DIRECTORY "CPack Component Example")

# This must always be last!
include(CPack)

```

Specifying Components

The first step in building a component-based installation is to identify the set of installable components. In this example, three components will be created: the library binary, the application, and the header file. This decision is arbitrary and project-specific, but be sure to identify the components that correspond to units of functionality important to your user, rather than basing the components on the internal structure of your program.

For each of these components, we need to identify which component each of the installed files belong in. For each `INSTALL` command in `CMakeLists.txt`, add an appropriate `COMPONENT` argument stating which component the installed files will be associated with:

```
install(TARGETS mylib
        ARCHIVE
        DESTINATION lib
        COMPONENT libraries)
install(TARGETS mylibapp
        RUNTIME
        DESTINATION bin
        COMPONENT applications)
install(FILES mylib.h
        DESTINATION include
        COMPONENT headers)
```

Note that the `COMPONENT` argument to the `INSTALL` command is not new; it has been a part of CMake's `INSTALL` command to allow installation of only part of a project. If you are using any of the older installation commands (`INSTALL_TARGETS`, `INSTALL_FILES`, etc.), you will need to convert them to `INSTALL` commands in order to use components.

The next step is to notify CPack of the names of all of the components in your project by calling the `cpack_add_component` function for each component of the package:

```
cpack_add_component(applications)
cpack_add_component(libraries)
cpack_add_component(headers)
```

At this point, you can build a component-based installer with CPack that will allow one to independently install the applications, libraries, and headers of MyLib. The Windows and Mac OS X installers will look like the ones shown in Figure 15.

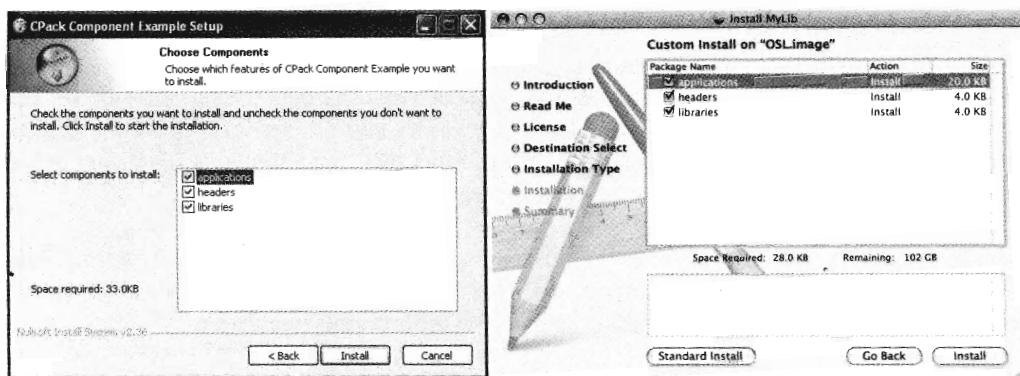


Figure 15 Windows and Mac OS X Component Installer First Page

Naming Components

At this point, you may have noted that the names of the actual components in the installer are not very descriptive: they just say "applications", "libraries", or "headers", as specified in the component names. These names can be improved by using the `DISPLAY_NAME` option in the `cpack_add_component` function:

```
cpack_add_component(applications DISPLAY_NAME
"MyLib Application")
cpack_add_component(libraries DISPLAY_NAME "Libraries")
cpack_add_component(headers DISPLAY_NAME "C++ Headers")
```

Any macro prefixed with `CPACK_COMPONENT_${COMPNAME}`, where `_${COMPNAME}` is the uppercase name of a component, is used to set a particular property of that component in the installer. Here, we set the `DISPLAY_NAME` property of each of our components, so that we get human-readable names. These names will be listed in the selection box rather than the internal component names "applications", "libraries", "headers".

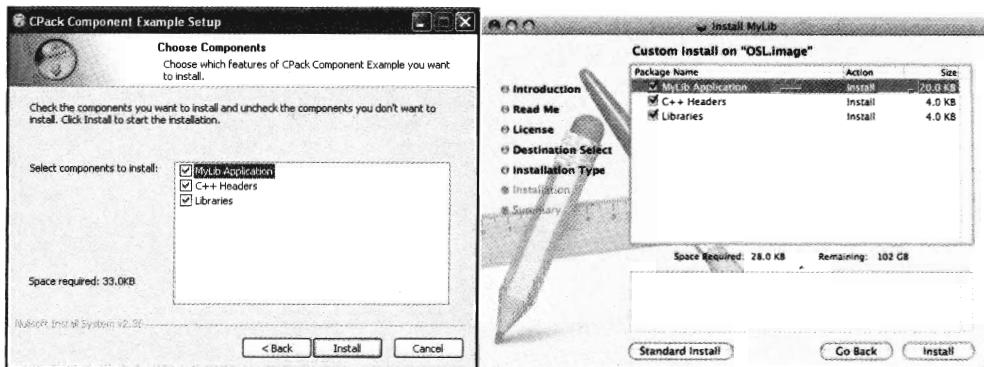


Figure 16 Windows and Mac OS X Installers with named components

Adding Component Descriptions

There are several other properties associated with components, including the ability to make a component hidden, required, or disabled by default, to provide additional descriptive information. Of particular note is the `DESCRIPTION` property, which provides some descriptive text for the component. This descriptive text will show up in a separate "description" box in the installer, and will be updated either when the user's mouse hovers over the name of the corresponding component (Windows) or when the user clicks on a component (Mac OS X). We will add a description for each of our components below:

```
cpack_add_component(applications DISPLAY_NAME "MyLib"
Application"
DESCRIPTION
"An extremely useful application that makes use of MyLib"
)
cpack_add_component(libraries DISPLAY_NAME "Libraries"
DESCRIPTION
"Static libraries used to build programs with MyLib"
)
cpack_add_component(headers DISPLAY_NAME "C++ Headers"
DESCRIPTION "C/C++ header files for use with MyLib"
)
```

Generally, descriptions should provide enough information for the user to make a decision on whether to install the component, but should not themselves be more than a few lines long (the "Description" box in the installers tends to be small). Figure 17 shows the description display for both the Windows and Mac OS X installers.

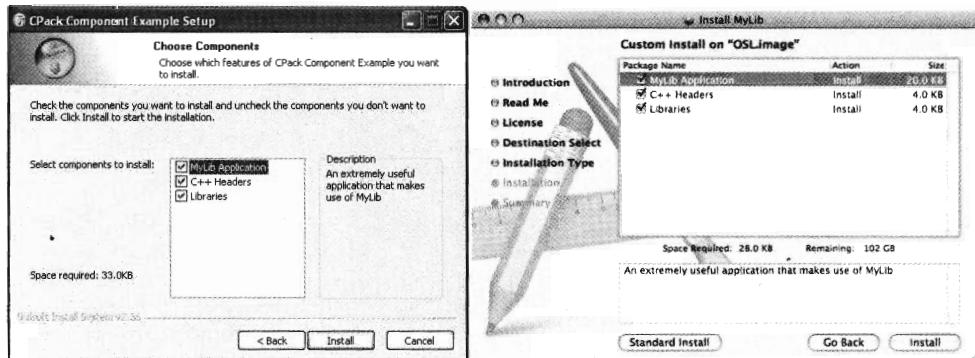


Figure 17 Component Installers with descriptions

Component Interdependencies

With most projects, the various components are not completely independent. For example, an application component may depend on the shared libraries in another component to execute properly, such that installing the application component without the corresponding shared libraries would result in an unusable installation. CPack allows you to express the dependencies between components, so that a component will only be installed if all of the other components it depends on are also installed.

To illustrate component dependencies, we will place a simple restriction on our component-based installer. Since we do not provide source code in our installer, the C++ header files we distribute can only actually be used if the user also installs the library binary to link their program against. Thus, the "headers" component depends on the availability of the "libraries" component. We can express this notion by setting the `DEPENDS` property for the `HEADERS` component as such:

```
cpack_add_component(headers DISPLAY_NAME "C++ Headers"
DESCRIPTION
"C/C++ header files for use with MyLib"
DEPENDS libraries
)
```

The `DEPENDS` property for a component is actually a list, as such a component can depend on several other components. By expressing all of the component dependencies in this manner, you can ensure that users will not be able to select an incomplete set of components at installation time.

Grouping Components

When the number of components in your project grows large, you may need to provide additional organization for the list of components. To help with this organization, CPack

includes the notion of component groups. A component group is, simply, a way to provide a name for a group of related components. Within the user interface, a component group has its own name, and underneath that group are the names of all of the components in that group. Users will have the option to (de-)select the installation of all components in the group with a single click, or expand the group to select individual components.

We will expand our example by categorizing its three components, "applications", "libraries", and "headers", into "Runtime" and "Development" groups. We can place a component into a group by using the GROUP option to the `cpack_add_component` function as follows:

```
cpack_add_component(applications
    DISPLAY_NAME "MyLib Application"
    DESCRIPTION
        "An extremely useful application that makes use of MyLib"
    GROUP Runtime)
cpack_add_component(libraries
    DISPLAY_NAME "Libraries"
    DESCRIPTION
        "Static libraries used to build programs with MyLib"
    GROUP Development)
cpack_add_component(headers
    DISPLAY_NAME "C++ Headers"
    DESCRIPTION "C/C++ header files for use with MyLib"
    GROUP Development
    DEPENDS libraries
)
```

Like components, component groups have various properties that can be customized, including the `DISPLAY_NAME` and `DESCRIPTION`. For example, the following code adds an expanded description to the "Development" group:

```
cpack_add_component_group(Development
    EXPANDED
    DESCRIPTION
        "All of the tools you'll ever need to develop software")
```

Once you have customized the component groups to your liking, rebuild the binary installer to see the new organization: the MyLib application will show up under the new "Runtime" group, while the MyLib library and C++ header will show up under the new "Development" group. One can easily turn on/off all of the components within a group using the installer's GUI. This can be seen in Figure 18.

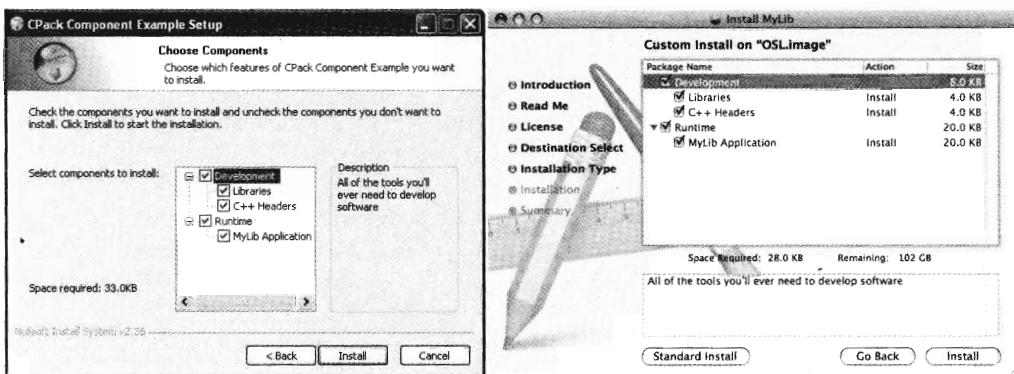


Figure 18 Component Grouping

Installation Types (NSIS Only)

When a project contains a large number of components, it is common for a Windows installer to provide pre-selected sets of components based on specific user needs. For example, a user wanting to develop software against a library will want one set of components, while an end user might use an entirely different set. CPack supports this notion of pre-selected component sets via installation types. An installation type is, simply, a set of components. When the user selects an installation type, exactly that set of components is selected---then the user is permitted to further customize their installation as desired. Currently this is only supported by the Windows NSIS generator.

For our simple example, we will create two installation types: a "Full" installation type that contains all of the components, and a "Developer" installation type that includes only the libraries and headers. To do this, we use the function `cpack_add_install_type` to add the types.

```
cpack_add_install_type(Full    DISPLAY_NAME "Everything")
cpack_add_install_type(Developer)
```

Next, we set the `INSTALL_TYPES` property of each component to state which installation types will include that component. This is done with the `INSTALL_TYPES` option to the `cpack_add_component` function.

```
cpack_add_component(libraries DISPLAY_NAME "Libraries"
DESCRIPTION
"Static libraries used to build programs with MyLib"
GROUP Development
INSTALL_TYPES Developer Full)
cpack_add_component(applications
```

```

DISPLAY_NAME "MyLib Application"
DESCRIPTION
  "An extremely useful application that makes use of MyLib"
GROUP Runtime
INSTALL_TYPES Full)
cpack_add_component(headers
  DISPLAY_NAME "C++ Headers"
  DESCRIPTION "C/C++ header files for use with MyLib"
GROUP Development
DEPENDS libraries
INSTALL_TYPES Developer Full)

```

Components can be in any number of installation types. If you rebuild the Windows installer, the components page will contain a combo box that allows you to select the installation type, and therefore its corresponding set of components as shown in Figure 19.

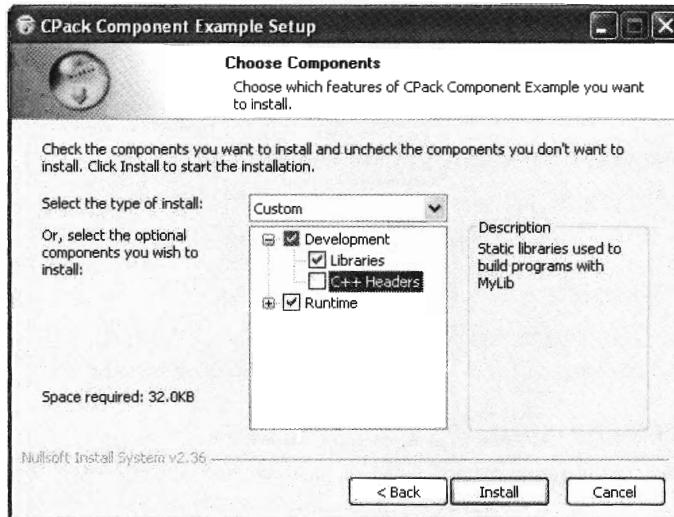


Figure 19 NSIS Installation Types

Variables that control CPack components

The functions `cpack_add_install_type`, `cpack_add_component_group`, and `cpack_add_component` just set CPACK_ variables. Those variables are described in the following list:

CPACK_COMPONENTS_ALL

This is a list containing the names of all components that should be installed by CPack. The presence of this macro indicates that CPack should build a component-

based installer. Files associated with any components not listed here or any installation commands not associated with any component will not be installed.

CPACK_COMPONENT_\${COMPNAME}_DISPLAY_NAME

The displayed name of the component \${COMPNAME}, used in graphical installers to display the component name. This value can be any string.

CPACK_COMPONENT_\${COMPNAME}_DESCRIPTION

An extended description of the component \${COMPNAME}, used in graphical installers to give the user additional information about the component. Descriptions can span multiple lines using "\n" as the line separator.

CPACK_COMPONENT_\${COMPNAME}_HIDDEN

Flag that indicates that this component will be hidden in the graphical installer, and therefore cannot be selected or installed. Only available with NSIS.

CPACK_COMPONENT_\${COMPNAME}_REQUIRED

Flag that indicates that this component is required, and therefore will always be installed. It will be visible in the graphical installer, but it cannot be unselected.

CPACK_COMPONENT_\${COMPNAME}_DISABLED

Flag that indicates that this component should be disabled (unselected) by default. The user is free to select this component for installation.

CPACK_COMPONENT_\${COMPNAME}_DEPENDS

Lists the components on which this component depends. If this component is selected, then each of the components listed must also be selected.

CPACK_COMPONENT_\${COMPNAME}_GROUP

Names the component group of which this component is a part. If not provided, the component will be a standalone component, not part of any component group.

CPACK_COMPONENT_\${COMPNAME}_INSTALL_TYPES

Lists the installation types of which this component is a part. When one of these installations types is selected, this component will automatically be selected. Only available with NSIS.

CPACK_COMPONENT_GROUP_\${GROUPNAME}_DISPLAY_NAME

The displayed name of the component group \${GROUPNAME}, used in graphical installers to display the component group name. This value can be any string.

CPACK_COMPONENT_GROUP_\${GROUPNAME}_DESCRIPTION

An extended description of the component group \${GROUPNAME}, used in graphical installers to give the user additional information about the components contained within this group. Descriptions can span multiple lines using "\n" as the line separator.

CPACK_COMPONENT_GROUP_\${GROUPNAME}_BOLD_TITLE

Flag indicating whether the group title should be in bold. Only available with NSIS.

CPACK_COMPONENT_GROUP_\${GROUPNAME}_EXPANDED

Flag indicating whether the group should start out "expanded", showing its components. Otherwise only the group name itself will be shown until the user clicks on the group. Only available with NSIS.

CPACK_INSTALL_TYPE_\${INSTNAME}_DISPLAY_NAME

The displayed name of the installation type. This value can be any string.

9.5 CPack for Cygwin Setup

Cygwin (<http://www.cygwin.com/>) is a Linux-like environment for Windows that consists of a run time DLL and a collection of tools. To add tools to the official cygwin the cygwin setup program is used. The setup tool has very specific layouts for the source and binary trees that are to be included. CPack can create the source and binary tar files correctly bzip'ed that can be uploaded to the cygwin mirror sites. You must of course have your package accepted by the cygwin community before that is done. Since the layout of the package is more restrictive than other packaging tools, you may have to change some of the install options for your project.

The cygwin setup program requires that all files be installed into /usr/bin, /usr/share/project-version, /usr/share/man and /usr/share/doc/package-version. The cygwin CPack generator will automatically add the /usr to the install directory for the project. The project must install things into /share and /bin, and CPack will add the /usr prefix automatically.

Cygwin also requires that you provide a shell script that can be used to create the package from the sources. Any cygwin specific patches that are required for the package must also be provided in a diff file. CMake's configure_file command can be used to create both of these files for a project. Since CMake is a cygwin package, the CMake code used to configure CMake for the cygwin CPack generators is as follows:

```

set (CPACK_PACKAGE_NAME cmake)

# setup the name of the package for cygwin cmake-2.4.3
set (CPACK_PACKAGE_FILE_NAME
    "${CPACK_PACKAGE_NAME}-2.2.${CPACK_PACKAGE_VERSION_PATCH}")

#. the source has the same name as the binary
set (CPACK_SOURCE_PACKAGE_FILE_NAME ${CPACK_PACKAGE_FILE_NAME})

# Create a cygwin version number in case there are changes
# for cygwin that are not reflected upstream in CMake
set (CPACK_CYGWIN_PATCH_NUMBER 1)

# if we are on cygwin and have cpack, then force the
# doc, data and man dirs to conform to cygwin style directories
set (CMAKE_DOC_DIR "/share/doc/${CPACK_PACKAGE_FILE_NAME}")
set (CMAKE_DATA_DIR "/share/${CPACK_PACKAGE_FILE_NAME}")
set (CMAKE_MAN_DIR "/share/man")

# These files are required by the cmCPackCygwinSourceGenerator and
# the files put into the release tar files.
set (CPACK_CYGWIN_BUILD_SCRIPT
    "${CMAKE_BINARY_DIR}/@CPACK_PACKAGE_FILE_NAME@-
     @CPACK_CYGWIN_PATCH_NUMBER@.sh")
set (CPACK_CYGWIN_PATCH_FILE
    "${CMAKE_BINARY_DIR}/@CPACK_PACKAGE_FILE_NAME@-
     @CPACK_CYGWIN_PATCH_NUMBER@.patch")

# include the sub directory for cygwin releases
include (Utilities/Release/Cygwin/CMakeLists.txt)

# when packaging source make sure to exclude the .build directory
set (CPACK_SOURCE_IGNORE_FILES
    "/CVS/" "/*\\.build/" "/*\\.svn/" "/*\\.swp$" "/*\\.#" "/*~$")


```

Utilities/Release/Cygwin/CMakeLists.txt:

```

# create the setup.hint file for cygwin
configure_file (
    "${CMAKE_SOURCE_DIR}/Utilities/Release/Cygwin/cygwin-setup.hint.in"
    "${CMAKE_BINARY_DIR}/setup.hint")

configure_file (

```

```
"${CMAKE_SOURCE_DIR}/Utilities/Release/Cygwin/README.cygwin.in"
"${CMAKE_BINARY_DIR}/Docs/@CPACK_PACKAGE_FILE_NAME@-
@CPACK_CYGWIN_PATCH_NUMBER@.README"

install_files (/share/doc/Cygwin FILES
${CMAKE_BINARY_DIR}/Docs/@CPACK_PACKAGE_FILE_NAME@-
@CPACK_CYGWIN_PATCH_NUMBER@.README)

# create the shell script that can build the project
configure_file (
"${CMAKE_SOURCE_DIR}/Utilities/Release/Cygwin/cygwin-package.sh.in"
${CPACK_CYGWIN_BUILD_SCRIPT})

# Create the patch required for cygwin for the project
configure_file (
"${CMAKE_SOURCE_DIR}/Utilities/Release/Cygwin/cygwin-patch.diff.in"
${CPACK_CYGWIN_PATCH_FILE})
```

The file Utilities/Release/Cygwin/cygwin-package.sh.in, can be found in the CMake source tree. It is a shell script that can be used to re-create the cygwin package from source. For other projects, there is a template install script that can be found in Templates/cygwin-package.sh.in. This script should be able to configure and package any cygwin based CPack project, and it is required for all official cygwin packages.

Another important file for cygwin binaries is share/doc/Cygwin/project-version.README. This file should contain the information required by cygwin about the project. In the case of CMake, the file is configured so that it can contain the correct version information. For example part of that file for CMake looks like this:

Build instructions:

```
unpack cmake-2.5.20071029-1-src.tar.bz2
if you use setup to install this src package, it will be
    unpacked under /usr/src automatically
cd /usr/src
./cmake-2.5.20071029-1.sh all
```

This will create:

```
/usr/src/cmake-2.5.20071029.tar.bz2
/usr/src/cmake-2.5.20071029-1-src.tar.bz2
```

9.6 CPack for Mac OS X PackageMaker

On the Apple Mac OS X operating system, CPack provides the ability to use the system Package Maker tool. This section will show the CMake application install screens users will see when installing the CMake package on OS X. The CPack variables set to change the text in the installer will be given for each screen of the installer.

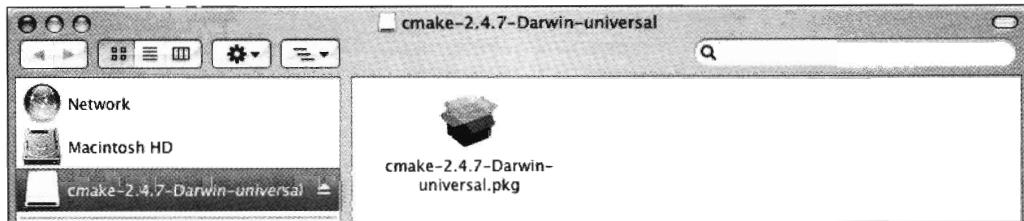


Figure 20 Mac Package inside .dmg

In Figure 20, the .pkg file found inside the .dmg disk image created by the CPack package maker for Mac OS X is seen. The name of this file is controlled by the CPACK_PACKAGE_FILE_NAME variable. If this is not set CPack will use a default name based on the package name and version settings.

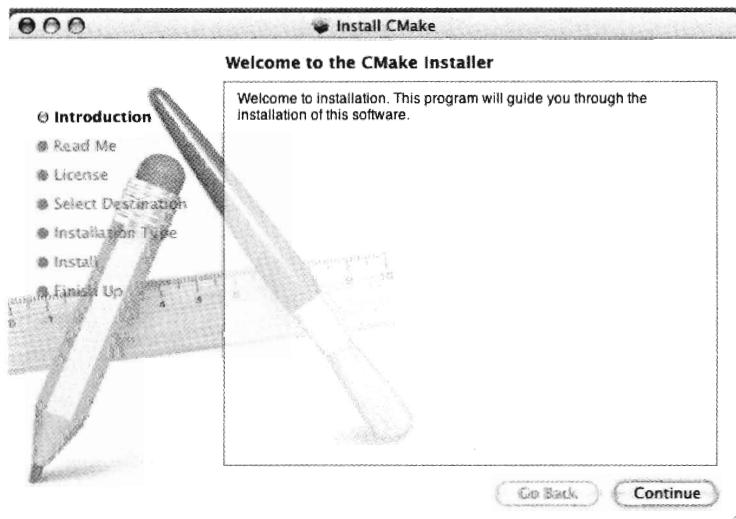


Figure 21 Introduction Screen Mac PackageMaker

When the .pkg file is run, the package wizard starts with the screen seen in Figure 21. The text in this window is controlled by the file pointed to by the CPACK_RESOURCE_FILE_WELCOME variable.

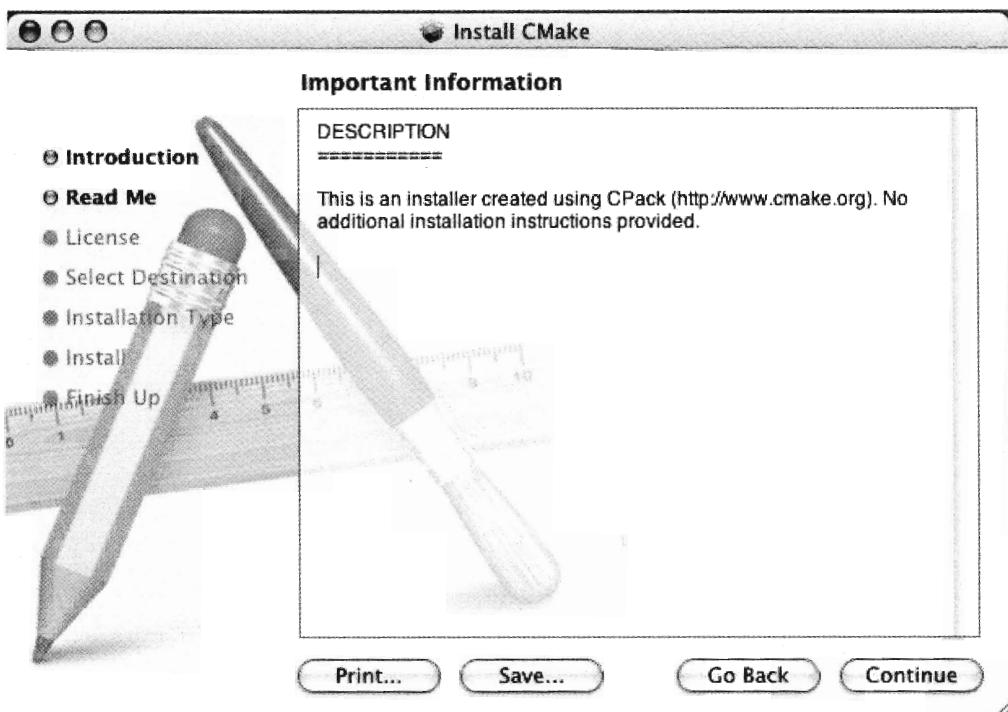


Figure 22 Readme section of Mac package wizard

Figure 22 shows the read me section of the package wizard. The text for this window is customized by using the CPACK_RESOURCE_FILE_README variable. It should contain a path to the file containing the text that should be displayed on this screen.

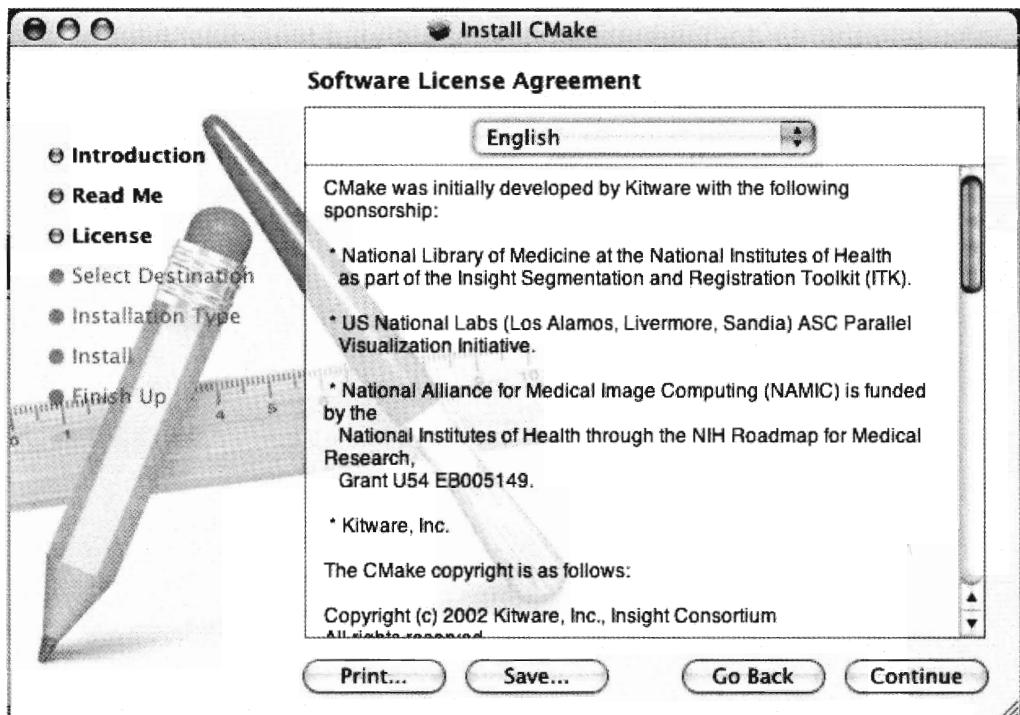


Figure 23 License screen Mac packager

Figure 23 contains the license text for the package. Users must accept the license for the installation process to continue. The text for the license comes from the file pointed to by the `CPACK_RESOURCE_FILE_LICENSE` variable.

The other screens in the installation process are not customizable from CPack. To change more advanced features of this installer there are two template CPack template files that you can modify, `Modules/CPack.Info.plist.in` and `Modules/CPack.Description.plist.in`. These files can be replaced by using the `CMAKE_MODULE_PATH` variable to point to a directory in your project containing a modified copy of `CPackInfo.plist.in`.

9.7 CPack for Mac OS X Drag and Drop

CPack also support the creation of a Drag and Drop installer for the Mac. In this case a .dmg disk image is created. Inside the image, it contains a symbolic link to the /Applications directory and a copy of the install tree. In this case it is best to use a Mac application bundle or a single folder containing your relocatable installation as the only install target for the

project. The variable `CPACK_PACKAGE_EXECUTABLES` is used to point to the application bundle for the project.

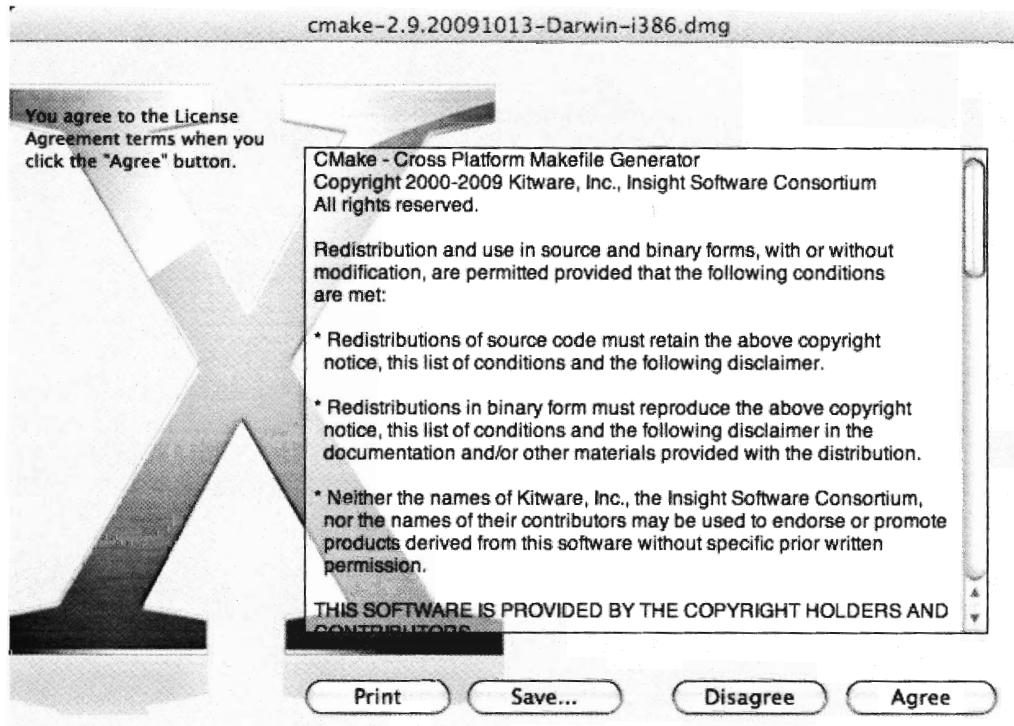


Figure 24 Drag and Drop License dialog

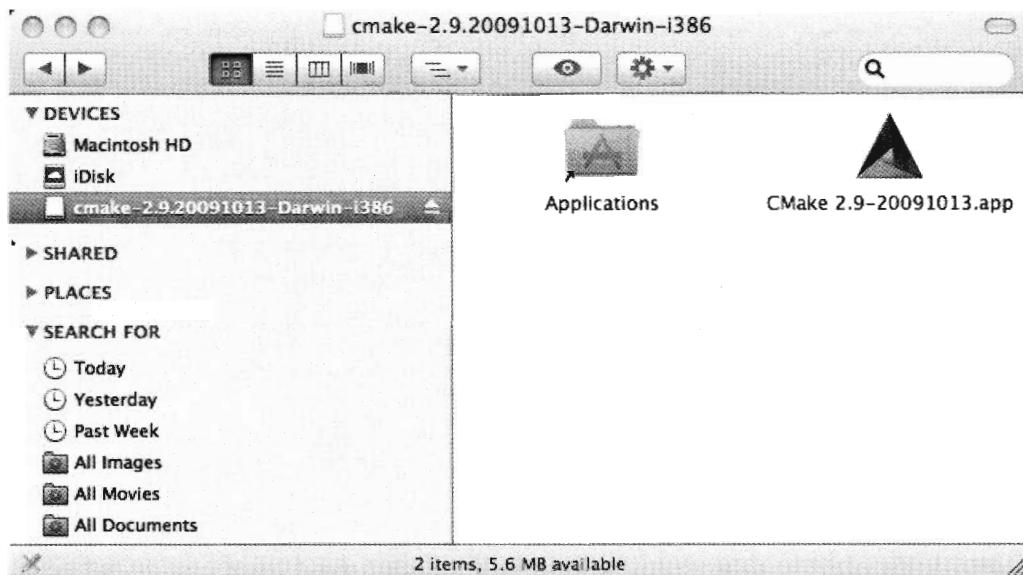


Figure 25 Resulting Drag and Drop folders

9.8 CPack for Mac OS X X11 Applications

CPack also includes an OS X X11 package maker generator. This can be used to package X11 based applications, and make them act more like native OS X applications. CPack not only packages them, but wraps them with a script that will allow users to run them as they would any native OS X application. Much like the OS X PackageMaker generator, the OS X X11 generator creates a disk image .dmg file. In this example, an X11 application called KWPolygonalObjectViewerExample is packaged with the OS X X11 CPack generator.

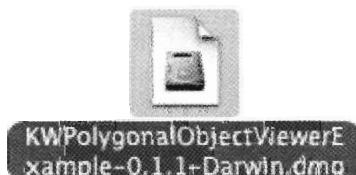


Figure 26 Mac OS X X11 package disk image

Figure 26 shows the disk image created. In this case the `CPACK_PACKAGE_NAME` was set to `KWPolygonalObjectViewerExample`, and the version information was left with the CPack default of 0.1.1. The variable `CPACK_PACKAGE_EXECUTABLES` was set to the pair

KWPolygonalObjectViewerExample and KWPolygonalObjectViewerExample, the installed X11 application is called KWPolygonalObjectViewerExample.

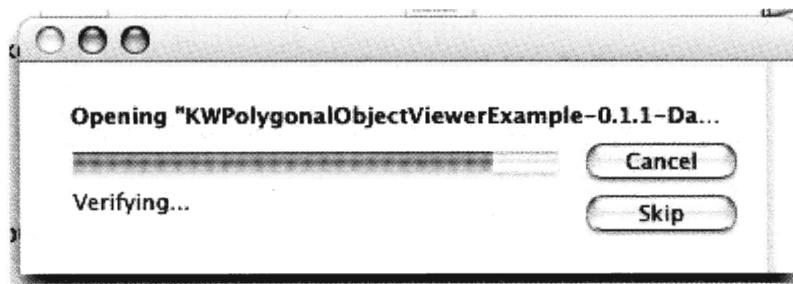


Figure 27 Opening OS X X11 disk image

Figure 27 shows what a user would see after clicking on the .dmg file created by CPack. Mac OS X is mounting this disk image as a disk

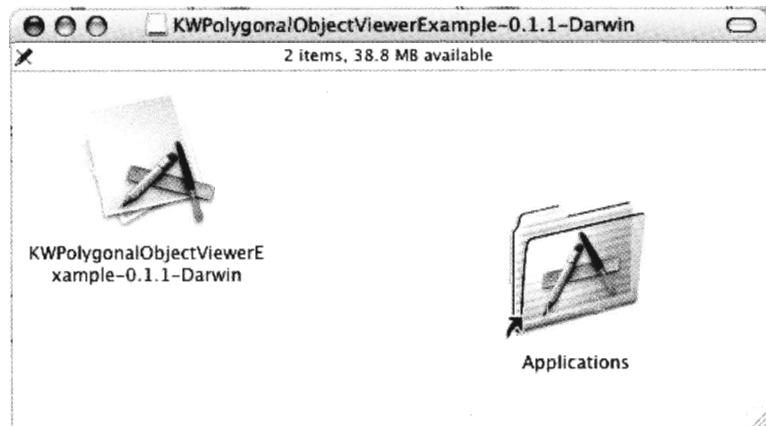


Figure 28 Mounted .dmg disk image

Figure 28 shows the mounted disk image. It will contain a symbolic link to the /Applications directory for the system, and it will contain an application bundle for each executable found in CPACK_PACKAGE_EXECUTABLES. The users can then drag and drop the applications into the Applications folder as seen in Figure 29.

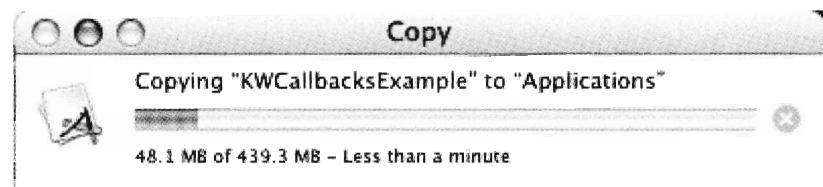


Figure 29 Drag and drop application to Applications

CPack actually provides a C++ based executable that can run an X11 application via the Apple scripting language. The application bundle installed will run that forwarding application when the user double clicks on KWPolygonalObjectViewerExample. This script will make sure that the X11 server is started. The script that is run can be found in CMake/Modules/CPack.RuntimeScript.in. The source for the script launcher C++ program can be found in Source/CPack/OSXScriptLauncher.cxx.

9.9 CPack for Debian Packages

A Debian package .deb is simply an “ar” archive. CPack includes the code for the BSD style ar that is required by Debian packages. The Debian packager used the standard set of CPack variables to initialize a set of Debian specific variables. These can be overridden in the CPACK_PROJECT_CONFIG_FILE, the name of the generator is “DEB”. The variables used by the DEB generator are as follows:

CPACK_DEBIAN_PACKAGE_NAME

defaults to lower case of CPACK_PACKAGE_NAME.

CPACK_DEBIAN_PACKAGE_ARCHITECTURE

defaults to i386.

CPACK_DEBIAN_PACKAGE_DEPENDS

Must be set to other packages that this package depends on, and if empty a warning is emitted.

CPACK_DEBIAN_PACKAGE_MAINTAINER

defaults to CPACK_PACKAGE_CONTACT

CPACK_DEBIAN_PACKAGE_DESCRIPTION

defaults to CPACK_PACKAGE_DESCRIPTION_SUMMARY

CPACK_DEBIAN_PACKAGE_SECTION

defaults to devl

CPACK_DEBIAN_PACKAGE_PRIORITY

defaults to optional

9.10 CPack for RPM

CPack has support for creating Linux RPM files. The name of the generator as set in `CPACK_GENERATOR` is “RPM”. The RPM package capability requires that `rpmbuild` is installed on the machine and is in PATH. The RPM packager uses the standard set of CPack variables to initialize RPM specific variables. The RPM specific variables are as follows:

CPACK_RPM_PACKAGE_SUMMARY

defaults to `${CPACK_PACKAGE_DESCRIPTION_SUMMARY}`

CPACK_RPM_PACKAGE_NAME

defaults to lower case of `${CPACK_PACKAGE_NAME}`

CPACK_RPM_PACKAGE_VERSION

defaults to `${CPACK_PACKAGE_VERSION}`.

CPACK_RPM_PACKAGE_ARCHITECTURE

defaults to i386

CPACK_RPM_PACKAGE_RELEASE

defaults to 1. This is the version of the RPM file, and not the version of the software being packaged.

CPACK_RPM_PACKAGE_GROUP

defaults to none.

CPACK_RPM_PACKAGE_VENDOR

defaults to `${CPACK_PACKAGE_VENDOR}`

9.11 CPack Files

There are a number of files that are used by CPack that can be useful for learning more about how CPack works and what options you can set. These files can also be used as the starting point for other generators for CPack. These files can mostly be found in the Modules and Templates directories of CMake and typically start with the prefix CPack.

Automation & Testing with CMake

10.1 Testing with CMake, CTest, and CDash

Testing is a key tool for producing and maintaining robust, valid software. This chapter will examine the tools that are part of CMake to support software testing. We will begin with a brief discussion of testing approaches and then discuss how to add tests to your software project using CMake. Finally we will look at additional tools that support creating centralized software status dashboards.

The tests for a software package may take a number of forms. At the most basic level there are smoke tests, such as one that simply verifies that the software compiles. While this may seem like a simple test, with the wide variety of platforms and configurations available, smoke tests catch more problems than any other type of test. Another form of smoke test is to verify that a test runs without crashing. This can be handy for situations where the developer does not want to spend the time creating more complex tests, but is willing to run some simple tests. Most of the time these simple tests can be small example programs. Running them verifies not only that the build was successful, but that any required shared libraries can be loaded (for projects that use them) and that at least some of the code can be executed without crashing.

Moving beyond basic smoke tests leads to more specific tests such as regression, black, and white box testing. Each of these has its strengths. Regression testing verifies that the results of a test do not change over time, or platform. This is very useful when performed frequently, as it provides a quick check that the behavior and results of the software have not changed. When a regression test fails a quick look at recent code changes can usually identify the culprit. Unfortunately, regression tests typically require more effort to create than other tests.

White and black box testing refer to tests written to exercise units of code (at various levels of integration) with and without knowledge of how those units are implemented respectively. White box testing is designed to stress potential failure points in the code knowing how that code was written, and hence its weaknesses. As with regression testing this can take a substantial amount of effort to create good tests. Black box testing typically knows little or nothing about the implementation of the software other than its public API. Black box testing can provide a lot of code coverage without too much effort in developing the tests. This is especially true for libraries of object oriented software where the APIs are well defined. A black box test can be written to go through and invoke a number of typical methods on all the classes in the software.

The final type of testing we will discuss is software standard compliance testing. While the other test types we have discussed are focused on determining if the code works properly, compliance testing tries to determine if the code adheres to the coding standards of the software project. This could be a check to verify that all classes have implemented some key method, or that all functions have a common prefix. The options for this type of test are limitless and there are a number of ways to perform such testing. There are software analysis tools that can be used, or specialized test programs (maybe python scripts etc) could be written. The key point to realize is that the tests do not necessarily have to involve running some part of the software. The tests might run some other tool on the source code itself.

There are a number of reasons why it helps to have testing support integrated into the build process. First, complex software projects may have a number of configuration or platform-dependent options. The build system knows what options can be enabled and can then enable the appropriate tests for those options. For example, the Visualization Toolkit (VTK) includes support for a parallel processing library called MPI. If VTK is built with MPI support then additional tests are enabled that make use of MPI and verify that the MPI-specific code in VTK works as expected. Secondly, the build system knows where the executables will be placed and it has tools for finding other required executables (such as perl, python etc). The third reason is that with UNIX Makefiles it is common to have a test target in the Makefile so that developers can type make test and have the test(s) run. In order for this to work, the build system must have some knowledge of the testing process.

10.2 How Does CMake Facilitate Testing?

CMake facilitates testing your software through special testing commands and the CTest executable. First we will discuss the key testing commands in CMake. To add testing to a CMake-based project is fairly simple using the `add_test` command. The `add_test` command has a simple syntax as follows:

```
add_test (TestName ExecutableToRun arg1 arg2 arg3 ...)
```

The first argument is simply a string name for the test. This is the name that will be displayed by testing programs. The second argument is the executable to run. The executable can be built as part of the project or it can be a standalone executable such as python, perl, etc. The remaining arguments will be passed to the running executable. A typical example of testing using the `add_test` command would look like this:

```
add_executable (TestInstantiator TestInstantiator.cxx)
target_link_libraries (TestInstantiator vtkCommon)
add_test (TestInstantiator
          ${EXECUTABLE_OUTPUT_PATH}/TestInstantiator)
```

The `add_test` command is typically placed in the CMakeLists file for the directory that has the test in it. For large projects there may be multiple CMakeLists files with `add_test` commands in them. Once the `add_test` commands are present in the project, the user can run the tests by invoking the “test” target of Makefile, or the `RUN_TESTS` target of Visual Studio or Xcode. An example of running tests on the CMake tests using the Makefile generator on Linux would be:

```
$ make test
Running tests...
Test project
    Start 2: kwsys.testEncode
    1/20 Test #2: kwsys.testEncode ..... Passed      0.02 sec
    Start 3: kwsys.testTerminal
    2/20 Test #3: kwsys.testTerminal ..... Passed      0.02 sec
    Start 4: kwsys.testAutoPtr
    3/20 Test #4: kwsys.testAutoPtr ..... Passed      0.02 sec
```

10.3 Additional Test Properties

By default a test passes if all of the following conditions are true:

- The test executable was found
- The test ran without exception
- The test exited with return code 0

That said, these behaviors can be modified using the `set_property` command:

```
set_property (TEST test_name
              PROPERTY prop1 value1 value2 ...)
```

This command will set additional properties for the specified tests. Example properties are:

ENVIRONMENT

Specify environment variables that should be defined for running a test. If set to a list of environment variables and values of the form `MYVAR=value` those environment variables will be defined while the test is running. The environment is restored to its previous state after the test is done.

LABELS

Specify a list of text labels associated with a test. These labels can be used to group tests together based on what they test. For example you could add a label of MPI to all tests that exercise MPI code.

WILL_FAIL

If this option is set to true, then the test will pass if the return code is not 0 and fail if it is. This reverses the third condition of the pass requirements.

PASS_REGULAR_EXPRESSION

If this option is specified, then the output of the test is checked against the regular expression provided (a list of regular expressions may be passed in as well). If none of the regular expressions match, then the test will fail. If at least one of them matches, then the test will pass.

FAIL_REGULAR_EXPRESSION

If this option is specified, then the output of the test is checked against the regular expression provided (a list of regular expressions may be passed in as well). If none of the regular expressions match, then the test will pass. If at least one of them matches, then the test will fail.

If both `PASS_REGULAR_EXPRESSION` and `FAIL_REGULAR_EXPRESSION` are specified, then the `FAIL_REGULAR_EXPRESSION` takes precedence. The following example illustrates using the `PASS_REGULAR_EXPRESSION` and `FAIL_REGULAR_EXPRESSION`:

```
add_test (outputTest ${EXECUTABLE_OUTPUT_PATH}/outputTest)

set (passRegex "^Test passed" "^All ok")
set (failRegex "Error" "Fail")

set_property (TEST outputTest
              PROPERTY PASS_REGULAR_EXPRESSION "${passRegex}")
set_property (TEST outputTest
              PROPERTY FAIL_REGULAR_EXPRESSION "${failRegex}")
```

10.4 Testing Using CTest

When you run the tests from your build environment what really happens is that the build environment runs CTest. CTest is an executable that comes with CMake, it handles running the tests for the project. While CTest works well with CMake, you do not have to use CMake in order to use CTest. The main input file for CTest is called `CTestTestfile.cmake`. This file will be created in each directory that was processed by CMake (typically every directory with a `CMakeLists` file). The syntax of `CTestTestfile.cmake` is like the regular CMake syntax, with a subset of the commands available. If CMake is used to generate testing files, they will list any subdirectories that need to be processed as well as any `add_test` calls. The subdirectories are those that were added by `subdirs` or `add_subdirectory` commands. CTest can then parse these files to determine what tests to run. An example of such a file is shown below:

```
# CMake generated Testfile for
# Source directory: C:/CMake
# Build directory: C:/CMakeBin
#
# This file includes the relevant testing commands required
# for testing this directory and lists subdirectories to
# be tested as well.

ADD_TEST (SystemInformationNew ...)

SUBDIRS (Source/kwsys)
SUBDIRS (Utilities/cmzlib)
...
```

When CTest parses the `CTestTestfile.cmake` files it will extract the list of tests from them. These tests will be run, and for each test CTest will display the name of the test and its status. Consider the following sample output:

```
$ ctest
Test project C:/CMake-build26
      Start 1: SystemInformationNew
1/21 Test #1: SystemInformationNew ..... Passed    5.78 sec
          Start 2: kwsys.testEncode
      Start 3: kwsys.testTerminal
2/21 Test #2: kwsys.testEncode .....
          Start 4: kwsys.testAutoPtr
3/21 Test #3: kwsys.testTerminal .....
          Start 5: kwsys.testHashSTL
4/21 Test #4: kwsys.testAutoPtr .....
          Start 6: kwsys.testHashSTL
```

```
5/21 Test #5: kwsys.testHashSTL ..... Passed 0.02 sec
...
100% tests passed, 0 tests failed out of 21
Total Test time (real) = 59.22 sec
```

CTest is run from within your build tree. It will run all the tests found in the current directory as well as any subdirectories listed in the `CTestTestfile.cmake`. For each test that is run CTest will report if the test passed and how long it took to run the test.

The CTest executable includes some handy command line options to make testing a little easier. We will start by looking at the options you would typically use from the command line.

<code>-R <regex></code>	Run tests matching regular expression
<code>-E <regex></code>	Exclude tests matching regular expression
<code>-L <regex></code>	Run tests with labels matching the regex
<code>-LE <regex></code>	Run tests with labels not matching regexp
<code>-C <config></code>	Choose the configuration to test
<code>-V,--verbose</code>	Enable verbose output from tests.
<code>-N,--show-only</code>	Disable actual execution of tests.
<code>-I [Start,End,Stride,test#,test# Test file]</code>	Run specific tests by range and number.
<code>-H</code>	Display a help message

The `-R` option is probably the most commonly used. It allows you to specify a regular expression, only the tests with names matching the regular expression will be run. Using the `-R` option with the name (or part of the name) of a test is a quick way to run a single test. The `-E` option is similar except that it excludes all tests matching the regular expression. The `-L` and `-LE` options are similar to `-R` and `-E` except that they apply to test labels that were set using the `set_property` command as described in section 10.3. The `-C` option is mainly for IDE builds where you might have multiple configurations, such as Release and Debug in the same tree. The argument following the `-C` determines which configuration will be tested. The `-V` argument is useful when you are trying to determine why a test is failing, with `-V` CTest will print out the command line used to run the test as well as any output from the test itself. The `-V` option can be used with any invocation of CTest to provide more verbose output. The `-N` option is useful if you want to see what tests CTest would run without actually running them.

Running the tests and making sure they all pass before committing any changes to the software is a sure fire way to improve your software quality and development process. Unfortunately, for large projects the number of tests and the time required to run them may be prohibitive. In these situations the `-I` option of CTest can be used. The `-I` option allows you to

flexibly specify a subset of the tests to run. For example, the following invocation of CTest will run every seventh test.

```
ctest -I , ,7
```

While this is not as good as running every test, it is better than not running any and it may be a more practical solution for many developers. Note that if the start and end arguments are not specified, as in this example, then they will default to the first and last tests. In another example, assume that you always want to run a few tests plus a subset of the others. In this case you can explicitly add those tests to the end of the arguments for `-I`. For example:

```
ctest -I , ,5,1,2,3,10
```

will run tests 1, 2, 3, and 10, plus every fifth test. You can pass as many test numbers as you want after the stride argument.

10.5 Using CTest to Drive Complex Tests

Sometimes to properly test a project you need to actually compile code during the testing phase. There are several reasons for this. First, if test programs are compiled as part of the main project, they can end up taking up a significant amount of the build time. Also, if a test fails to build, the main build should not fail as well. Finally, IDE projects can quickly become too large to load and work with. The CTest command supports a group of command line options that allow it to be used as the test executable to run. When used as the test executable, CTest can run CMake, run the compile step, and finally run a compiled test. We will now look at the command line options to CTest that support building and running tests.

```
--build-and-test src_directory build_directory
Run cmake on the given source directory using the specified
build directory.
--test-command           Name of the program to run.
--build-target           Specify a specific target to build.
--build-nocmake          Run the build without running cmake first.
--build-run-dir          Specify directory to run programs from.
--build-two-config       Run CMake twice before the build.
--build-exe-dir          Specify the directory for the executable.
--build-generator        Specify the generator to use.
--build-project          Specify the name of the project to build.
--build-makeprogram      Specify the make program to use.
--build-noclean          Skip the make clean step.
--build-options          Add extra options to the build step.
```

For an example, consider the following `add_test` command taken from the `CMakeLists.txt` file of CMake itself. It shows how CTest can be used both to compile and run a test.

```
add_test (simple ${CMAKE_CTEST_COMMAND}
          --build-and-test "${CMake_SOURCE_DIR}/Tests/Simple"
                           "${CMake_BINARY_DIR}/Tests/Simple"
          --build-generator ${CMAKE_GENERATOR}
          --build-makeprogram ${CMAKE_MAKE_PROGRAM}
          --build-project Simple
          --test-command simple)
```

In this example, the `add_test` command is first passed the name of the test, "simple". After the name of the test, the command to be run is specified. In this case, the test command to be run is CTest. The CTest command is referenced via the `CMAKE_CTEST_COMMAND` variable. This variable is always set by CMake to the CTest command that came from the CMake installation used to build the project. Next, the source and binary directories are specified. The next options to CTest are the `--build-generator`, and `--build-makeprogram` options. These are specified using the `cmake` variables `CMAKE_MAKE_PROGRAM` and `CMAKE_GENERATOR`. Both `CMAKE_MAKE_PROGRAM` and `CMAKE_GENERATOR` are defined by CMake. This is an important step as it makes sure that the same generator is used for building the test as was used for building the project itself. The `--build-project` option is passed `Simple` which corresponds to the `project` command used in the `Simple` test. The final argument is the `--test-command` which tells CTest the command to run once it gets a successful build. That should be the name of the executable that will be compiled by the test.

10.6 Handling a Large Number of Tests

When a large number of tests exist in a single project, it is cumbersome to have individual executables available for each test. That said, the developer of the project should not be required to create tests with complex argument parsing. This is why CMake provides a convenience command for creating a test driver program. This command is called `create_test_sourcelist`. A test driver is a program that links together many small tests into a single executable. This is useful when building static executables with large libraries to shrink the total required size. The signature for `create_test_sourcelist` is as follows:

```
create_test_sourcelist (SourceListName
                        DriverName
                        test1 test2 test3
                        EXTRA_INCLUDE include.h
                        FUNCTION function
                        )
```

The first argument is the variable which will contain the list of source files that must be compiled to make the test executable. The `DriverName` is the name of the test driver program (e.g. the name of the resulting executable). The rest of the arguments consist of a list of test source files. Each test source file should have a function in it that has the same name as the file with no extension (`foo.cxx` should have `int foo(argc, argv);`) The resulting executable will be able to invoke each of the tests by name on the command line. The `EXTRA_INCLUDE` and `FUNCTION` arguments support additional customization of the test driver program. Consider the following CMakeLists file fragment to see how this command can be used:

```
# create the testing file and list of tests
create_test_sourcelist(Tests
    CommonCxxTests.cxx
    ObjectFactory.cxx
    otherArrays.cxx
    otherEmptyCell.cxx
    TestSmartPointer.cxx
    SystemInformation.cxx
    ...
)

# add the executable
add_executable(CommonCxxTests ${Tests})

# remove the test driver source file
set(TestsToRun ${Tests})
remove(TestsToRun CommonCxxTests.cxx)

# Add all the ADD_TEST for each test
foreach(test ${TestsToRun})
    get_filename_component(TName ${test} NAME_WE)
    add_test(${TName} CommonCxxTests ${TName})
endforeach()
```

The `create_test_sourcelist` command is invoked to create a test driver. In this case it creates and writes `CommonCXXTests.cxx` into the binary tree of the project using the rest of the arguments to determine its contents. Next the `add_executable` command is used to add that executable to the build. Then a new variable called `TestsToRun` is created with an initial value of the sources required for the test driver. The `remove` command is used to remove the driver program itself from the list. Then a `foreach` command is used to loop over the remaining sources. For each source its name without a file extension is extracted and put in the variable `TName`, then a new test is added for `TName`. The end result is that for each source file in the `create_test_sourcelist` an `add_test` command is called with the name of

the test. As more tests are added to the `create_test_sourcelist` command, the `foreach` loop will automatically call `add_test` for each one.

10.7 Producing Test Dashboards

As your project's testing needs grow, keeping track of the test results can become overwhelming. This is especially true for projects that are tested nightly on a number of different platforms. In these cases we recommend using a test dashboard to summarize the test results. (see Figure 30)

The screenshot shows a Mozilla Firefox browser window displaying the CDash - CMake dashboard at <http://www.cdash.org/CDash/index.php?project=CMake>. The dashboard has a dark theme with a banner at the top. Below the banner, there are tabs for DASHBOARD, CALENDAR, PREVIOUS, CURRENT, and PROJECT. A message indicates "6 files changed by 2 authors as of Monday, August 03 2009 21:00:00 EDT". The main area contains two tables: "Style" and "Nightly Expected".

Site	Build Name	Update			Configure			Build			Test			Build Time
		Files	Min	Error	Warn	Min	Error	Warn	Min	NotRun	Fail	Pass	Min	
dash21.kitware	KWStyle	5	0.2	0	0	0.4	0	0	0	0	0	0	2009-08-03T22:10:58 EDT	
Totals	1 Builds	5	0.2	0	0	0.4	0	0	0	0	0	0		

Site	Build Name	Update			Configure			Build			Test			Build Time
		Files	Min	Error	Warn	Min	Error	Warn	Min	NotRun	Fail	Pass	Min	
v20n17.ppm.ihost.com	AIX53-xlc	14	0.5	0	0	5	0	0	4	0	0	110	21.2 2009-08-04T01:06:34 EDT	
dash22	CVS-Win32-x86-hcc32	5	0	0	0	0.7	0	0	2	0	0	108	7.1 2009-08-04T00:14:00 EDT	
dash22	CYGWIN2008-01.5.25i686-gcc	5	0	0	0	3.8	0	0	2.1	0	0	109	69.5 2009-08-03T22:48:00 EDT	
krondor.kitware	Darwin-c++	5	0.4	0	0	0.1	0	0	26.5	0	0	113	39.3 2009-08-03T21:07:31 EDT	
dashmacmini3.kitware	Darwin-Leopard-Xcode21-univ	0	0	0	0	0	0	0	10.5	0	0	108	16.4 2009-08-04T02:01:00 EDT	
dashmacmini3.kitware	Darwin-Leopardintel-p++	5	0	0	0	0	0	0	2.2	0	0	112	12.7 2009-08-04T01:29:00 EDT	

At the bottom left, there is a "Done" button.

Figure 30 - Sample Testing Dashboard

A test dashboard summarizes the results for many tests on many platforms, and its hyperlinks allow people to drill down into additional levels of detail quickly. The CTest executable includes support for producing test dashboards. When run with the correct options, CTest will

produce XML-based output recording the build and test results and post them to a dashboard server. The dashboard server runs an open source software package called CDash. CDash collects the XML results and produces HTML web pages from them.

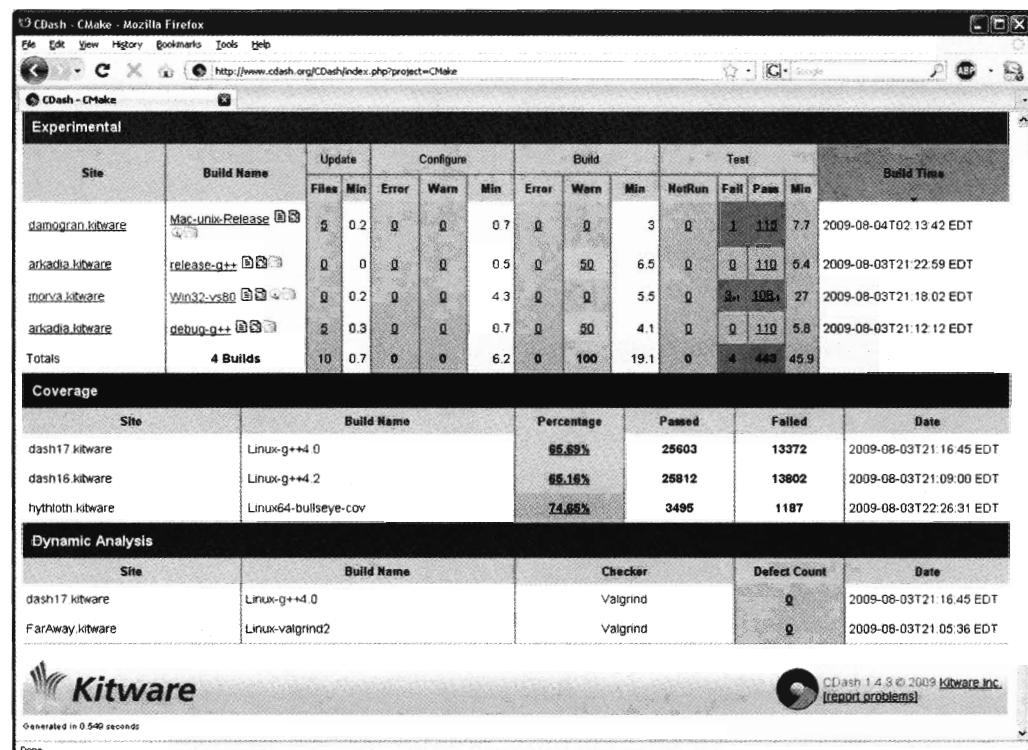


Figure 31 - Experimental, Coverage, and Dynamic Analysis Results

Before discussing how to use CTest to produce a dashboard let us consider the main parts of a testing dashboard. Each night at a specified time the dashboard server will open up a new dashboard, so each day there is a new web page showing the results of tests for that twenty-four hour period. There are links on the main page that allow you to quickly navigate through different days. Looking at the main page for a project (such as CMake's dashboard off of www.cmake.org) you will see that it is divided into a few main components. Near the top you will find a set of links that allow you to step to previous dashboards as well as links to project pages such as the bug tracker, documentation etc.

Below that you will find groups that you will find include Nightly, Experimental, Continuous, Coverage, and Dynamic Analysis (see Figure 31). The category into which a dashboard entry will be placed depends on how it was generated. The simplest are Experimental entries which represent dashboard results for someone's current

copy of the project's source code. With an experimental dashboard the source code is not guaranteed to be up to date. In contrast a Nightly dashboard entry is one where CTest tries to update the source code to a specific date and time. The expected result is that all nightly dashboard entries for a given day should be based on the same source code.

A continuous dashboard entry is one that is designed to run every time new files are checked in. Depending on how frequently new files are checked in a single day's dashboard could have many continuous entries. Continuous dashboards are particularly helpful for cross platform projects where a problem may only show up on some platforms. In those cases a developer can commit a change that works for them on their platform and then another platform running a continuous build could catch the error, allowing the developer to correct the problem promptly.

Dynamic Analysis and Coverage dashboards are designed to test the memory safety and code coverage of a project. A Dynamic Analysis dashboard entry is one where all the tests are run with a memory access/leak checking program enabled. Any resulting errors or warnings are parsed, summarized and displayed. This is important to verify that your software is not leaking memory, reading from un-initialized memory, etc. Coverage dashboard entries are similar in that all the tests are run, but as the tests are run the lines of code executed are tracked. When all the tests have been run, a listing of how many times each line of code was executed is produced and displayed on the dashboard.

Adding CDash Dashboard Support to a Project

In this section we show how to submit results to the CDash dashboard. You can either use the Kitware CDash servers at my.cdash.org or you can setup your own CDash server as described in section 10.11. If you are using my.cdash.org you can click on the “Start My Project” button which will ask you to create an account (or login if you already have one), and then bring you to a page to start creating your project. If you have installed your own CDash server then you should login to your CDash server as administrator and select “Create New Project” from the administration panel. Regardless of which approach you use the next few steps will be to fill in information about your project as shown in Figure 32. Many of the items below are optional, so do not be concerned if you do not have a value for them.

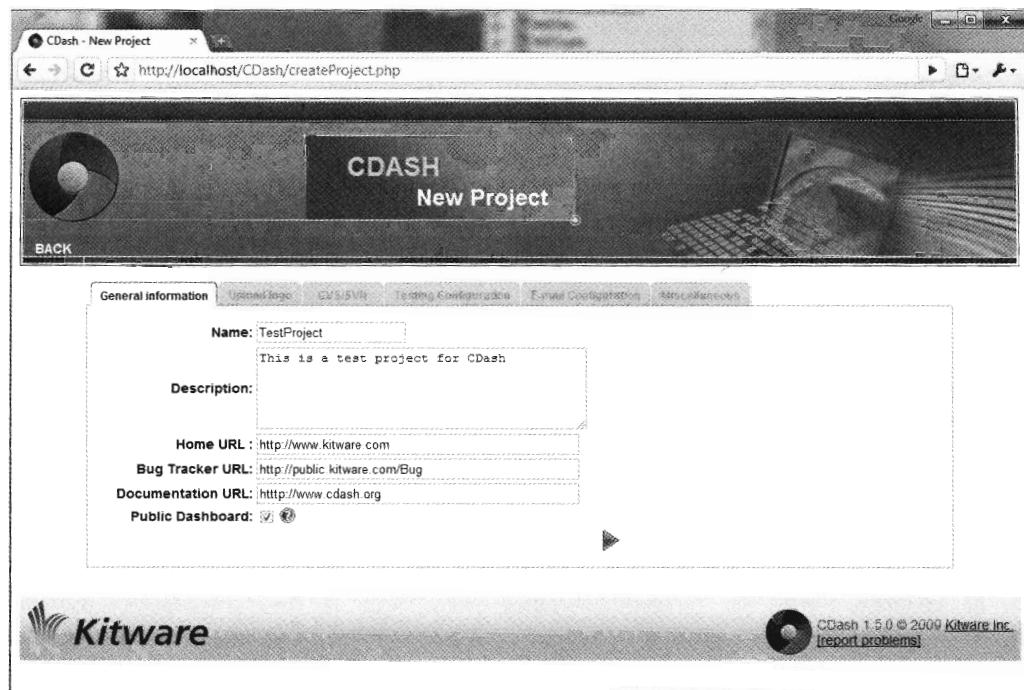


Figure 32 – Creating a new project in CDash

Name: what you want to call the project.

Description: description of the project to be shown on the first page.

Home URL: home URL of the project to appear in the main menu of the dashboard.

Bug Tracker URL: URL to the bug tracker. Currently CDash supports Mantis (<http://www.mantisbt.org/>), and if a bug is entered in the repository with the message “BUG: 132456”, CDash will automatically link to the appropriate bug.

Documentation URL: URL to where the project’s documentation is kept. This will appear in the main menu of the dashboard.

Public Dashboard: if checked, the dashboard is public and anybody can see the results of the dashboard. If unchecked, only users assigned to this project can access the dashboard.

Logo: logo of the project to be displayed on the main dashboard. Optimal size for a logo is 100x100 pixels. Transparent GIFs work best as they can blend in with the CDash background.

Repository Viewer URL: URL of the web repository browser. CDash currently supports: ViewCVS, Trac, Fisheye, ViewVC, WebSVN and Loggerhead. For instance a URL for ViewVC is <http://public.kitware.com/cgi-bin/viewvc.cgi/?cvsroot=CMak>, or for WebSVN: <https://www.kitware.com/websvn/listing.php?repname=MyRepository>

Repositories: in order to display the daily updates CDash gets a diff version of the modified files. Current CDash supports only anonymous repository access. A typical URL is `:pserver:anoncvs@myproject.org:/cvsroot/MyProject`.

Nightly Start Time: CDash displays the current dashboard using a 24 hour window. The nightly start time defines the beginning of this window. Note that the start time is expressed in the form HH:MM:SS TZ, e.g. 21:00:00 EDT.

Coverage Threshold: CDash marks that coverage has passed (green) if the global coverage for a build or specific files is above this threshold. It is recommended to set the coverage threshold to a high value and decrease it as you focus on improving your coverage.

Enable Test Timing: enable/Disable test timing for this project See “Test timing” in the next section for more information.

Test Time Standard Deviation: set a multiplier for the standard deviation of a test time. If the time for a test is higher than the mean + multiplier*standard deviation, the test time status is marked as failed. The default value is 4 if not specified. Note that changing this value does not affect previous builds, only builds submitted after the modification.

Test Time Standard Deviation Threshold: set a minimum standard deviation for a test time. If the current standard deviation for a test is lower than this threshold then the threshold is used instead. This is particularly important for tests that have a very low standard deviation but still some variability. The default threshold is set to 2 if not specified. Note that changing this value does not affect previous builds, only builds submitted after the modification.

Test Time # Max Failures Before Flag: some tests might take longer from one day to another depending on the client machine load. This variable defines the number of times a test should fail because of timing issues before being flagged.

Email Submission Failures: enable/disable sending email when a build fails (configure, error, warnings, update, test failings) for this project. This is a general feature.

Email Redundant Failure: by default CDash does not send email for the same failures. For instance if a build continues to fail over time, only one email would be sent. If the email redundant failures is checked then CDash will send an email every time a build has a failure. (not all versions of CDash have this option).

Email Administrator: enable/disable sending email when submission to CDash is invalid. This can help tracking down misconfigured clients. This is particularly useful when CTest is not used to submit to CDash..

Email Build Missing: enable/disable sending email when a build has not submitted.

Email Low Coverage: enable/disable sending email when the coverage for files is lower than the default threshold value specified above.

Email Test Timing Changed: enable/disable sending email when a test's timing has changed.

Maximum Number of Items in Email: how many failures should be sent in an email.

Maximum Number of Characters in Email: how many characters from the log should be sent in the email.

Google Analytics Tracker: CDash supports visitor tracking through Google analytics. See “Adding Google Analytics” for more information.

Show Site IP Addresses: Enable/disable the display of IP addresses of the sites submitting to this project.

Display Labels: as of CDash 1.4, and recent versions of CTest, labels can be attached to build files, these can be displayed and retrieved on CDash.

AutoRemove Timeframe: set the number of days to keep for this project. If the timeframe is less than 2 days, CDash will not remove any builds (not all versions of CDash have this option).

AutoRemove Max Builds: set the maximum number of builds to remove when performing the auto removal of builds (not all versions of CDash have this option).

After providing this information you can click on “Create Project” to create the project in CDash. At this point the server is ready to accept dashboard submissions. The next step is to provide the dashboard server information to your software project. This information is kept in a file named `CTestConfig.cmake` at the top level of your source tree. You can download this file by clicking on the “Edit Project” button for your dashboard (it looks like a pie chart with a wrench underneath it), then click on the miscellaneous tab and select “Download CTestConfig” and save the `CTestConfig.cmake` in your source tree. In the next section we review this file in more detail.

Client Setup

To support dashboards in your project you need to include the CTest module as follows.

```
# Include CDash dashboard testing module
include (CTest)
```

The CTest module will then read settings from the `CTestConfig.cmake` file you downloaded from CDash. If you have added `add_test` command calls to your project creating a dashboard entry is as simple as running:

```
ctest -D Experimental
```

The `-D` option tells CTest to create a dashboard entry. The next argument indicates what type of dashboard entry to create. Creating a dashboard entry involves quite a few steps that can be run independently or as one command. In this example the `Experimental` argument will cause CTest to perform a number of different steps as one command. The different steps of creating a dashboard entry are summarized below.

Start

Prepare a new dashboard entry. This creates a `Testing` subdirectory in the build directory. The `Testing` subdirectory will contain a subdirectory for the dashboard results with a name that corresponds to the dashboard time. The `Testing` subdirectory will also contain a subdirectory for the temporary testing results called `Temporary`.

Update

Perform a source control update of the source code (typically used for nightly or continuous runs). Currently CTest supports Concurrent Versions System (CVS), Subversion, Git, Mercurial and Bazaar.

Configure

Run CMake on the project to make sure the Makefiles or project files are up to date.

Build

Build the software using the specified generator.

Test

Run all the tests and record results.

MemoryCheck

Perform memory checks using Purify or valgrind.

Coverage

Collect source code coverage information.

Submit

Submit the testing results as a dashboard entry to the server.

Each of these steps can be run independently for a Nightly or Experimental entry using the following syntax:

```
ctest -D NightlyStart  
ctest -D NightlyBuild  
ctest -D NightlyCoverage -D NightlySubmit
```

or

```
ctest -D ExperimentalStart  
ctest -D ExperimentalConfigure  
ctest -D ExperimentalCoverage -D ExperimentalSubmit
```

etc. Alternatively you can use shortcuts that perform the most common combinations all at once. The shortcuts that CTest has defined include:

ctest -D Experimental

performs a start, configure, build, test, coverage, submit

ctest -D Nightly

performs a start, update, configure, build, test, coverage, submit

ctest -D Continuous

performs a start, update, configure, build, test, coverage, submit

ctest -D MemoryCheck

performs a start, configure, build, memory check, coverage, submit

When first setting up a dashboard it is often useful to combine the **-D** option with the **-V** option. This will allow you to see the output of all the different stages of the dashboard process. Likewise, CTest maintains log files in the `Testing/Temporary` directory it creates in your binary tree. There you will find log files for the most recent build, update, test, etc. The actual dashboard results are stored in the `Testing` directory too.

10.8 Customizing Dashboards for a Project

CTest has a few options that can be used to control how it processes a project. If, when CTest runs a dashboard, it finds `CTestCustom.cmake` files in the binary tree, it will load these files and use the settings from them to control its behavior. The syntax of a `CTestCustom` file is the same as regular CMake syntax. That said, only set commands are normally used in this file. These commands specify properties that CTest will consider when performing the testing.

Dashboard Submissions Settings

A number of the basic dashboard settings are provided in the file that you download from CDash. You can edit these initial values and provide additional values if you wish. The first value that is set is the nightly start time. This is the time that dashboards all around the world will use for checking out their copy of the nightly source code. This time also controls how dashboard submissions will be grouped together. All submissions from the nightly start time until the next nightly start time will be included on the same "day".

```
# Dashboard is opened for submissions for a 24 hour period
# starting at the specified NIGHTLY_START_TIME. Time is
# specified in 24 hour format.
set (CTEST_NIGHTLY_START_TIME "21:00:00 EDT")
```

The next group of settings control where to submit the testing results. This is the location of the CDash server.

```
# CDash server to submit results (used by client)
set (CTEST_DROP_METHOD http)
set (CTEST_DROP_SITE "my.cdash.org")
set (CTEST_DROP_LOCATION "/submit.php?project=KensTest")
set (CTEST_DROP_SITE_CDASH TRUE)
```

The `CTEST_DROP_SITE` specifies the location of the CDash server. Build and test results generated by CDash clients are sent to this location. The `CTEST_DROP_LOCATION` is the directory or the HTTP URL on the server where CDash clients leave their build and test reports. The `CTEST_DROP_SITE_CDASH` specifies that the current server is CDash which prevents CTest from trying to trigger the submission (backwards compatibility with Dart and Dart 2).

Currently CDash supports only the HTTP drop submission method; however CTest supports other submission types. The `CTEST_DROP_METHOD` specifies the method used to submit testing results. The most common setting for this will be HTTP which uses the Hyper Text Transfer Protocol (HTTP) to transfer the test data to the server. Other drop methods are supported for special cases such as FTP and SCP. In the example below, clients that are

submitting their results using the HTTP protocol use a web address as their drop site. If the submission is via FTP, this location is relative to where the `CTEST_DROP_SITE_USER` will log in by default. The `CTEST_DROP_SITE_USER` specifies the FTP username the client will use on the server. For FTP submissions this user will typically be "anonymous". However, any username that can communicate with the server can be used. For FTP servers that require a password this can be stored in the `CTEST_DROP_SITE_PASSWORD` variable. The `CTEST_DROP_SITE_MODE` (not used in this example) is an optional variable that you can use to specify the FTP mode. Most FTP servers will handle the default passive mode but you can set the mode explicitly to active if your server does not.

CTest can also be run from behind a firewall. If the firewall allows FTP or HTTP traffic, then no additional settings are required. If the firewall requires an FTP/HTTP proxy or uses a SOCKS4 or SOCKS5 type proxy, some environment variables need to be set. `HTTP_PROXY` and `FTP_PROXY` specify the servers that service HTTP and FTP proxy requests. `HTTP_PROXY_PORT` and `FTP_PROXY_PORT` specify the port on which the HTTP and FTP proxies reside. `HTTP_PROXY_TYPE` specifies the type of the HTTP proxy used. The three different types of proxies supported are the default, which is a generic HTTP/FTP proxy, "SOCKS4", and "SOCKS5", which specify SOCKS4 and SOCKS5 compatible proxies.

Filtering Errors and Warnings

By default CTest has a list of regular expressions that it matches for finding the errors and warnings from the output of the build process. You can override these settings in your `CTestCustom.ctest` files using several variables as shown below.

```
set (CTEST_CUSTOM_WARNING_MATCH
${CTEST_CUSTOM_WARNING_MATCH}
"^{standard input}:[0-9][0-9]*: Warning: "
)

set (CTEST_CUSTOM_WARNING_EXCEPTION
${CTEST_CUSTOM_WARNING_EXCEPTION}
"tk8.4.5/[^\n]+/[^\n]+.c[:\"]"
"xtree.[0-9]+. : warning C4702: unreachable code"
"warning LNK4221"
"variable .var_args[2]*. is used before its value is set"
"jobserver unavailable"
)
```

Another useful feature of the `CTestCustom` files is that you can use it to limit the tests that are run for memory checking dashboards. Memory checking using `purify` or `valgrind` is a CPU intensive process that can take twenty hours for a dashboard that normally takes one hour. To

help alleviate this problem CTest allows you to exclude some of the tests from the memory checking process as follows:

```
set (CTEST_CUSTOM_MEMCHECK_IGNORE
    ${CTEST_CUSTOM_MEMCHECK_IGNORE}
TestSetGet
otherPrint-ParaView
Example-vtkLocal
Example-vtkMy
)
```

The format for excluding tests is simply a list of test names as specified when the tests were added in your CMakeLists file with `add_test`.

In addition to the demonstrated settings, such as `CTEST_CUSTOM_WARNING_MATCH`, `CTEST_CUSTOM_WARNING_EXCEPTION`, and `CTEST_CUSTOM_MEMCHECK_IGNORE`, CTest also checks several other variables.

`CTEST_CUSTOM_ERROR_MATCH`

Additional regular expressions to consider a build line as an error line

`CTEST_CUSTOM_ERROR_EXCEPTION`

Additional regular expressions to consider a build line not as an error line

`CTEST_CUSTOM_WARNING_MATCH`

Additional regular expressions to consider a build line as a warning line

`CTEST_CUSTOM_WARNING_EXCEPTION`

Additional regular expressions to consider a build line not as a warning line

`CTEST_CUSTOM_MAXIMUM_NUMBER_OF_ERRORS`

Maximum number of errors before CTest stops reporting errors (default 50)

`CTEST_CUSTOM_MAXIMUM_NUMBER_OF_WARNINGS`

Maximum number of warnings before CTest stops reporting warnings (default 50)

`CTEST_CUSTOM_COVERAGE_EXCLUDE`

Regular expressions for files to be excluded from the coverage analysis

CTEST_CUSTOM_PRE_MEMCHECK**CTEST_CUSTOM_POST_MEMCHECK**

List of commands to execute before/after performing memory checking

CTEST_CUSTOM_MEMCHECK_IGNORE

List of tests to exclude from the memory checking step

CTEST_CUSTOM_PRE_TEST**CTEST_CUSTOM_POST_TEST**

List of commands to execute before/after performing testing

CTEST_CUSTOM_TESTS_IGNORE

List of tests to exclude from the testing step

CTEST_CUSTOM_MAXIMUM_PASSED_TEST_OUTPUT_SIZE

Maximum size of test output for the passed test (default 1k)

CTEST_CUSTOM_MAXIMUM_FAILED_TEST_OUTPUT_SIZE

Maximum size of test output for the failed test (default 300k)

Commands specified in `CTEST_CUSTOM_PRE_TEST` and `CTEST_CUSTOM_POST_TEST`, as well as the equivalent memory checking ones, are executed once per CTest run. These commands can be used for example if all tests require some initial setup and cleanup to be performed.

Adding Notes to a Dashboard

CTest and CDash support adding notes to a dashboard submission. These notes will appear on the dashboard as a graphical icon that when clicked expands to the text of the notes. Using CTest you can use the `-A` option followed by a semicolon separated list of filenames. The contents of these files will be submitted as notes for the dashboard. For example:

```
ctest -D Continuous -A C:/MyNotes.txt;C:/OtherNotes.txt
```

Another way to submit notes with a dashboard is to copy or write the notes as files into a Notes directory under the Testing directory of your binary tree. Any files found there when CTest submits a dashboard will also be uploaded as notes.

10.9 Setting up Automated Dashboard Clients

CTest has a built-in scripting mode to help make the process of setting up dashboard clients even easier. CTest scripts will handle most of the common tasks and options that CTest -D Nightly does not. The dashboard script is written using CMake syntax and mainly involves setting up different variables or options, or creating an elaborate procedure, depending on the complexity of testing. Once you have written the script you can run the nightly dashboard on Windows, Mac OS X or UNIX systems as follows:

```
ctest -S myScript.cmake
```

First we will consider the most basic script you can use, and then we will cover the different options you can make use of. There are four variables that you must always set in your scripts. The first two variables are the names of the source and binary directories on disk, `CTEST_SOURCE_DIRECTORY` and `CTEST_BINARY_DIRECTORY`. These should be fully specified paths. The next variable, `CTEST_COMMAND`, specifies which CTest command to use for running the dashboard. This may seem a bit confusing at first. The `-S` option of CTest is provided to do all the setup and customization for a dashboard, but the actual running of the dashboard is done with another invocation of CTest -D. Basically once the CTest script has done what it needs to do to setup the dashboard, it invokes CTest -D to actually generate the results. You can adjust the value of `CTEST_COMMAND` to control what type of dashboard to generate (Nightly, Experimental, Continuous) as well as to pass other options to the internal CTest process such as `-I ,,7` to run every 7th test. To refer to the CTest that is running the script, use the variable: `CTEST_EXECUTABLE_NAME`. The last required variable is `CTEST_CMAKE_COMMAND` which specifies the full path to the `cmake` executable that will be used to configure the dashboard. To refer to the CMake command that corresponds to the CTest command running the script, use the variable: `CMAKE_EXECUTABLE_NAME`. The CTest script does an initial configuration with CMake in order to generate the `CTestConfig.cmake` file that CTest will use for the dashboard. The following example demonstrates the use of these four variables and is an example of the simplest script you can have.

```
# these are the source and binary directories on disk
set (CTEST_SOURCE_DIRECTORY C:/martink/test/CMake)
set (CTEST_BINARY_DIRECTORY C:/martink/test/CMakeBin)

# which CTest command to use for running the dashboard
set (CTEST_COMMAND
    "\"${CTEST_EXECUTABLE_NAME}\" -D Nightly"
)

# what cmake command to use for configuring this dashboard
set (CTEST_CMAKE_COMMAND
```

```
"\"${CMAKE_EXECUTABLE_NAME}\"
)
```

The script above is not that different to running CTest -D from the command line yourself. All it adds is that it verifies that the binary directory exists and creates it if it does not. Where CTest scripting really shines is in the optional features it supports. We will consider these options one by one, starting with one of the most commonly used options CTEST_START_WITH_EMPTY_BINARY_DIRECTORY. When this variable is set to true it will delete the binary directory and then recreate it as an empty directory prior to running the dashboard. This guarantees that you are testing a clean build every time the dashboard is run. To use this option you simply set it in your script. In the example above we would simply add the following lines:

```
# should CTest wipe the binary tree before running
set (CTEST_START_WITH_EMPTY_BINARY_DIRECTORY TRUE)
```

Another commonly used option is the CTEST_INITIAL_CACHE variable. Whatever values you set this to will be written into the CMakeCache file prior to running the dashboard. This is an effective and simple way to initialize a cache with some preset values. The syntax is the same as what is in the cache with the exception that you must escape any quotes. Consider the following example:

```
# this is the initial cache to use for the binary tree, be
# careful to escape any quotes inside of this string
set (CTEST_INITIAL_CACHE "
//Command used to build entire project from the command line.
MAKECOMMAND:STRING=\"C:/PROGRA~1/MICROS~1.NET/Common7/IDE/devenv
.com\" CMake.sln /build Debug /project ALL_BUILD

//make program
CMAKE_MAKE_PROGRAM:FILEPATH=C:/PROGRA~1/MICROS~1.NET/Common7/IDE
/devenv.com

//Name of generator.
CMAKE_GENERATOR:INTERNAL=Visual Studio 7 .NET 2003

//Path to a program.
CVSCOMMAND:FILEPATH=C:/cygwin/bin/cvs.exe

//Name of the build
BUILDNAME:STRING=Win32-vs71
```

```
//Name of the computer/site where compile is being run  
SITE:STRING=DASH1.kitware  
")
```

Note that the above code is basically just one set command setting the value of CTEST_INITIAL_CACHE to a multiline string value. For Windows builds these are the most common cache entries that need to be set prior to running the dashboard. The first three values control what compiler will be used to build this dashboard (Visual Studio 7.1 in this example). CVSCOMMAND might be found automatically, but if not it can be set here. The last two cache entries are the names that will be used to identify this dashboard submission on the dashboard.

The next two variables work together to support additional directories and projects. For example, imagine that you had a separate data directory that you needed to keep up to date with your source directory. Setting the variables CTEST_CVS_COMMAND and CTEST_EXTRA_UPDATES_1 tells CTest to perform a cvs update on the specified directory, with the specified arguments prior to running the dashboard. For example:

```
# what cvs command to use for configuring this dashboard  
set (CTEST_CVS_COMMAND "C:/cygwin/bin/cvs.exe")  
  
# set any extra directories to do an update on  
set (CTEST_EXTRA_UPDATES_1  
"C:/Dashboards/My Tests/VTKData" "-dAP")
```

If you have more than one directory that needs to be updated you can use CTEST_EXTRA_UPDATES_2 through CTEST_EXTRA_UPDATES_9 in the same manner. The next variable you can set is called CTEST_ENVIRONMENT. This variable consolidates several set commands into a single command. Setting this variable allows you to set environment variables that will be used by the process running the dashboards. You can set as many environment variables as you want using the syntax shown below.

```
# set any extra environment variables here  
set (CTEST_ENVIRONMENT  
"DISPLAY=:0"  
"USE_GCC_MALLOC=1"  
)  
# is the same as  
set (ENV{DISPLAY} ":0")  
set (ENV{USE_GCC_MALLOC} "1")
```

The final general purpose option we will discuss is CTest's support for restoring a bad dashboard. In some cases you might want to make sure that you always have a working build of the software. Or you might use the resulting executables or libraries from one dashboard in the build process of another dashboard. In these cases if the first dashboard fails it is best to drop back to the last previously working dashboard. You can do this in CTest by setting `CTEST_BACKUP_AND_RESTORE` to true. When this is set to true CTest will first backup the source and binary directories, it will then check out a new source directory and create a new binary directory. After that it will run a full dashboard. If the dashboard is successful the backup directories are removed, if for some reason the new dashboard fails the new directories will be removed and the old directories restored. To make this work you must also set the `CTEST_CVS_CHECKOUT` variable. This should be set to the command required to check out your source tree. This doesn't actually have to be cvs but it must result in a source tree in the correct location. Consider the following example:

```
# do a backup and should the build fail restore,  
# if this is true you must set the CTEST_CVS_CHECKOUT  
# variable below.  
set (CTEST_BACKUP_AND_RESTORE TRUE)  
  
# this is the full cvs command to checkout the source dir  
# this will be run from the directory above the source dir  
set (CTEST_CVS_CHECKOUT  
    "/usr/bin/cvs -d /cvsroot/FOO co -d FOO FOO"  
)
```

Note that whatever checkout command you specify will be run from the directory above the source directory. A typical nightly dashboard client script will look like this:

```
set (CTEST_SOURCE_NAME CMake)  
set (CTEST_BINARY_NAME CMake-gcc)  
set (CTEST_DASHBOARD_ROOT "$ENV{HOME}/Dashboards/My Tests")  
  
set (CTEST_SOURCE_DIRECTORY  
    "${CTEST_DASHBOARD_ROOT}/${CTEST_SOURCE_NAME}")  
set (CTEST_BINARY_DIRECTORY  
    "${CTEST_DASHBOARD_ROOT}/${CTEST_BINARY_NAME}")  
  
# which ctest command to use for running the dashboard  
set (CTEST_COMMAND  
    "\"${CTEST_EXECUTABLE_NAME}\""  
    "-D Nightly  
    -A \"${CTEST_SCRIPT_DIRECTORY}/${CTEST_SCRIPT_NAME}\")
```

```
# what cmake command to use for configuring this dashboard
set (CTEST_CMAKE_COMMAND "\"${CMAKE_EXECUTABLE_NAME}\"")

# should ctest wipe the binary tree before running
set (CTEST_START_WITH_EMPTY_BINARY_DIRECTORY TRUE)
# this is the initial cache to use for the binary tree
set (CTEST_INITIAL_CACHE "
SITE:STRING=midworld.kitware
BUILDNAME:STRING=DarwinG5-g++
MAKECOMMAND:STRING=make -i -j2
")

# set any extra environment variables here
set (CTEST_ENVIRONMENT
    "CC=gcc"
    "CXX=g++"
)
```

Settings for Continuous Dashboards

The next three variables are used for setting up continuous dashboards. As mentioned earlier a continuous dashboard is designed to run continuously throughout the day providing quick feedback on the state of the software. If you are doing a continuous dashboard you can use `CTEST_CONTINUOUS_DURATION` and `CTEST_CONTINUOUS_MINIMUM_INTERVAL` to run the continuous repeatedly. The duration controls how long the script should run continuous dashboards, and the minimum interval specifies the shortest allowed time between continuous dashboards. For example, say that you want to run a continuous dashboard from 9AM until 7PM and that you want no more than one dashboard every twenty minutes. To do this you would set the duration to 600 minutes (ten hours) and the minimum interval to 20 minutes. If you run the test script at 9AM it will start a continuous dashboard. When that dashboard finishes it will check to see how much time has elapsed. If less than 20 minutes has elapsed CTest will sleep until the 20 minutes are up. If 20 or more minutes have elapsed then it will immediately start another continuous dashboard. Do not be concerned that you will end up with 300 dashboards a day (10 hours * three times an hour). If there have been no changes to the source code, CTest will not build and submit a dashboard. It will instead start waiting until the next interval is up and then check again. Using this feature just involves setting the following variables to the values you desire.

```
set (CTEST_CONTINUOUS_DURATION 600)
set (CTEST_CONTINUOUS_MINIMUM_INTERVAL 20)
```

Earlier we introduced the `CTEST_START_WITH_EMPTY_BINARY_DIRECTORY` variable that can be set to start the dashboards with an empty binary directory. If this is set to true for a

continuous dashboard then every continuous where there has been a change in the source code will result in a complete build from scratch. For larger projects this can significantly limit the number of continuous dashboards that can be generated in a day, while not using it can result in build errors or omissions because it is not a clean build. Fortunately there is a compromise, if you set `CTEST_START_WITH_EMPTY_BINARY_DIRECTORY_ONCE` to true CTest will start with a clean binary directory for the first continuous build but not subsequent ones. Based on your settings for the duration this is an easy way to start with a clean build every morning, but use existing builds for the rest of the day.

Another useful feature to use with a continuous dashboard is the `-I` option. A large project may have so many tests that running all the tests limits how frequently a continuous dashboard can be generated. By adding `-I ,,7` (or `-I ,,5` etc) to the `CTEST_COMMAND` value the continuous dashboard will only run every seventh test significantly reducing the time required between continuous dashboards. For example:

```
# which ctest command to use for running the dashboard
set (CTEST_COMMAND
    "\"${CTEST_EXECUTABLE_NAME}\" -D Continuous -I ,,7"
)
```

As you can imagine there is a compromise to be made between the coverage of the continuous and the frequency of its updates. Depending on the size of your project and the compute resources at your disposal these variables can be used to fine tune a continuous dashboard to meet your needs. An example of a CTest script for a continuous dashboard looks like this:

```
# these are the names of the source and binary directories
set (CTEST_SOURCE_NAME CMake-cont)
set (CTEST_BINARY_NAME CMakeBCC-cont)
set (CTEST_DASHBOARD_ROOT "c:/Dashboards/My Tests")
set (CTEST_SOURCE_DIRECTORY
    "${CTEST_DASHBOARD_ROOT}/${CTEST_SOURCE_NAME}")
set (CTEST_BINARY_DIRECTORY
    "${CTEST_DASHBOARD_ROOT}/${CTEST_BINARY_NAME}")

# which ctest command to use for running the dashboard
set (CTEST_COMMAND
    "\"${CTEST_EXECUTABLE_NAME}\""
    -D Continuous
    -A "\"${CTEST_SCRIPT_DIRECTORY}/${CTEST_SCRIPT_NAME}\\"")

# what cmake command to use for configuring this dashboard
set (CTEST_CMAKE_COMMAND "\"${CMAKE_EXECUTABLE_NAME}\\"")
```

```
# this is the initial cache to use for the binary tree
set (CTEST_INITIAL_CACHE "
SITE:STRING=dash14.kitware
BUILDNAME:STRING=Win32-bcc5.6
CMAKE_GENERATOR:INTERNAL=Borland Makefiles
CVSCOMMAND:FILEPATH=C:/Program Files/TortoiseCVS/cvs.exe
CMAKE_CXX_FLAGS:STRING=-w- -whid -waus -wpar -tWM
CMAKE_C_FLAGS:STRING=-w- -whid -waus -tWM
")

# set any extra environment variables here
set (ENV{PATH} "C:/Program
Files/Borland/CBuilder6/Bin\;C:/Program
Files/Borland/CBuilder6/Projects/Bpl"
)
```

Variables Available in CTest Scripts

There are a few variables that will be set before your script executes. The first two variables are the directory the script is in, `CTEST_SCRIPT_DIRECTORY`, and name of the script itself `CTEST_SCRIPT_NAME`. These two variables can be used to make your scripts more portable. For example if you wanted to include the script itself as a note for the dashboard you could do the following:

```
set (CTEST_COMMAND
"\\"${CTEST_EXECUTABLE_NAME}\\" -D Continuous
-A \\"${CTEST_SCRIPT_DIRECTORY}/${CTEST_SCRIPT_NAME}\\""
)
```

Another variable you can use is `CTEST_SCRIPT_ARG`. This variable can be set by providing a comma separated argument after the script name when invoking `CTest -S`. For example `CTest -S foo.cmake,21` would result in `CTEST_SCRIPT_ARG` being set to 21.

10.10 Advanced CTest Scripting

CTest scripting described in section [Setting up Automated Dashboard Clients](#) 10.9 should provide support for most dashboards. That said it has some limitations that advanced users may want to circumvent. This section describes how to write CTest scripts that allow the dashboard maintainer to have much more control.

Limitations of Traditional CTest Scripting

Let us start with the limitations of traditional CTest scripting. The first limitation is that the dashboard will always fail if the Configure step fails. The reason for that is that the input files for CTest are actually generated by the Configure step. To make things worse, the update step will not happen and the dashboard will be stuck. To prevent this, an additional update step is necessary. This can be achieved by adding `CTEST_EXTRA_UPDATES_1` variable with “-D yesterday” or similar flag. This will update the repository prior to doing a dashboard. Since it will update to yesterday’s time stamp, the actual update step of CTest will find the files that were modified since the previous day.

The second limitation of traditional CTest scripting is that it is not actually scripting. We only have control over what happens before the actual CTest run, but not what happens during or after. For example, if we want to run the testing and then move the binaries somewhere, or if we want to build the project, do some extra tasks and then run tests or something similar, we need to perform several complicated tasks, such as run CMake with `-P` option as a part of `CTEST_COMMAND`.

Extended CTest Scripting

To overcome the limitations of traditional CTest scripting, CTest provides an extended scripting mode. In this mode, the dashboard maintainer has access to individual CTest handlers. By running these handlers individually, she can develop relatively elaborate testing schemes. An example of an extended CTest script would be something like this:

```
cmake_minimum_required (VERSION 2.2)

set (CTEST_SITE           "andoria.kitware")
set (CTEST_BUILD_NAME     "Linux-g++")
set (CTEST_NOTES_FILES    "${CTEST_SCRIPT_DIRECTORY}/${CTEST_SCRIPT_NAME}")

set (CTEST_DASHBOARD_ROOT  "$ENV{HOME}/Dashboards/My Tests")
set (CTEST_SOURCE_DIRECTORY "${CTEST_DASHBOARD_ROOT}/CMake")
set (CTEST_BINARY_DIRECTORY "${CTEST_DASHBOARD_ROOT}/CMake-gcc
")

set (CTEST_UPDATE_COMMAND   "/usr/bin/cvs")
set (CTEST_CONFIGURE_COMMAND
      "\"${CTEST_SOURCE_DIRECTORY}/bootstrap\"")
set (CTEST_BUILD_COMMAND     "/usr/bin/make -j 2")

ctest_empty_binary_directory (${CTEST_BINARY_DIRECTORY})
```

```
ctest_start (Nightly)
ctest_update (SOURCE "${CTEST_SOURCE_DIRECTORY}")
ctest_configure (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_build (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_test (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_submit ()
```

The first line is there to make sure an appropriate version of CTest is used, the advanced scripting was introduced in CTest 2.2. The CMake parser is used, and so all scriptable commands from CMake are available. This includes the `cmake_minimum_required` command:

```
cmake_minimum_required (VERSION 2.2)
```

Overall the layout of the rest of this script is similar to the traditional one. There are several settings that CTest will use to perform its tasks. Then, unlike with traditional CTest, there are the actual tasks that CTest will perform. Instead of providing information in the project's CMake cache, in this scripting mode, all the information is provided to CTest. For compatibility reasons we may choose to write the information to the cache, but that is up to the dashboard maintainer. The first block contains the variables about the submission.

```
set (CTEST_SITE "andoria.kitware")
set (CTEST_BUILD_NAME "Linux-g++")
set (CTEST_NOTES_FILES
    "${CTEST_SCRIPT_DIRECTORY}/${CTEST_SCRIPT_NAME}")
```

These variables serve the same role as the `SITE` and `BUILD_NAME` cache variables. They are used to identify the system once it submits the results to the dashboard. `CTEST_NOTES_FILES` is a list of files that should be submitted as the notes of the dashboard submission. This variable corresponds to the `-A` flag of CTest.

The second block describes the information that CTest handlers will use to perform the tasks:

```
set (CTEST_DASHBOARD_ROOT "$ENV{HOME}/Dashboards/My Tests")
set (CTEST_SOURCE_DIRECTORY "${CTEST_DASHBOARD_ROOT}/CMake")
set (CTEST_BINARY_DIRECTORY "${CTEST_DASHBOARD_ROOT}/CMake-gcc")
set (CTEST_UPDATE_COMMAND "/usr/bin/cvs")
set (CTEST_CONFIGURE_COMMAND
    "\"${CTEST_SOURCE_DIRECTORY}/bootstrap\"")
set (CTEST_BUILD_COMMAND "/usr/bin/make -j 2")
```

The `CTEST_SOURCE_DIRECTORY` and `CTEST_BINARY_DIRECTORY` serve the same purpose as in the traditional CTest script. The only difference is that we will be able to overwrite these variables later on when calling the CTest handlers. The `CTEST_UPDATE_COMMAND` is the path to the command used to update the source directory from the repository. Currently CTest supports Concurrent Versions System (CVS), Subversion, Git, Mercurial and Bazaar.

Both the configure and build handlers support two modes. One mode is to provide the full command that will be invoked during that stage, this is designed to support projects that do not use CMake as their configuration or build tool. In this case you specify the full command lines to configure and build your project by setting the `CTEST_CONFIGURE_COMMAND` and `CTEST_BUILD_COMMAND` variables respectively. This is similar to specifying `CTEST_CMAKE_COMMAND` in the traditional CTest scripting.

For projects that use CMake for their configuration and build steps you do not need to specify the command lines for configuring and building your project. Instead you will specify the CMake generator to use by setting the `CTEST_CMAKE_GENERATOR` variable. This way CMake will be run with the appropriate generator. One example of this is:

```
set (CTEST_CMAKE_GENERATOR "Visual Studio 8 2005")
```

For the build step you should also set the variables `CTEST_PROJECT_NAME` and `CTEST_BUILD_CONFIGURATION` to specify how to build the project. In this case `CTEST_PROJECT_NAME` will match the top level CMakeLists file's `PROJECT` command. The `CTEST_BUILD_CONFIGURATION` should be one of `Release`, `Debug`, `MinSizeRel`, or `RelWithDebInfo`. Additionally `CTEST_BUILD_FLAGS` can be provided as a hint to the build command. An example of testing for a CMake based project would be:

```
set (CTEST_CMAKE_GENERATOR "Visual Studio 8 2005")
set (CTEST_PROJECT_NAME "Grommit")
set (CTEST_BUILD_CONFIGURATION "Debug")
```

The final block performs the actual testing and submission:

```
ctest_empty_binary_directory (${CTEST_BINARY_DIRECTORY})

ctest_start (Nightly)

ctest_update (SOURCE
              "${CTEST_SOURCE_DIRECTORY}" RETURN_VALUE res)
ctest_configure (BUILD
                  "${CTEST_BINARY_DIRECTORY}" RETURN_VALUE res)
ctest_build (BUILD "${CTEST_BINARY_DIRECTORY}" RETURN_VALUE res)
```

```
ctest_test (BUILD "${CTEST_BINARY_DIRECTORY}" RETURN_VALUE res)
ctest_submit (RETURN_VALUE res)
```

The `ctest_empty_binary_directory` command empties the directory and all subdirectories. Please note that this command has a safety measure built in, which is that it will only remove the directory if there is a `CMakeCache.txt` file in the top level directory. This is to prevent CMake from mistakenly removing a directory.

The rest of the block contains the calls to the actual CTest handlers. Each of them corresponds to a CTest `-D` option. For example, instead of:

```
ctest -D ExperimentalBuild
```

The script would contain:

```
ctest_start (Experimental)
ctest_build (BUILD "${CTEST_BINARY_DIRECTORY}" RETURN_VALUE res)
```

Each step can return a return value, which can then be used to determine if the step was successful. For example the return value of the Update stage can be used in a continuous dashboard to determine if the rest of the dashboard should be run.

To demonstrate some advantages of using extended CTest scripting, let us examine a more advanced CTest script. This script drives testing of an application called Slicer. Slicer uses CMake internally, but it drives the build process through a series of Tcl scripts. One of the problems of this approach is that it does not support out-of-source builds. Also, on Windows, certain modules come pre-built, so they have to be copied to the build directory. To test a project like that, we would use a script like this:

```
cmake_minimum_required (VERSION 2.2)

# set the dashboard specific variables -- name and notes
set (CTEST_SITE "dash11.kitware")
set (CTEST_BUILD_NAME "Win32-VS71")
set (CTEST_NOTES_FILES
    "${CTEST_SCRIPT_DIRECTORY}/${CTEST_SCRIPT_NAME}")

# do not let the test run for more than 1500 seconds
set (CTEST_TIMEOUT "1500")

# set the source and binary directories
```

```
set (CTEST_SOURCE_DIRECTORY "C:/Dashboards/MyTests/slicer2")
set (CTEST_BINARY_DIRECTORY "${CTEST_SOURCE_DIRECTORY}-build")

set (SLICER_SUPPORT
    "//Dash11/Shared/Support/SlicerSupport/Lib")
set (TCLSH  "${SLICER_SUPPORT}/win32/bin/tclsh84.exe")
# set the complete update, configure and build commands
set (CTEST_UPDATE_COMMAND
    "C:/Program Files/TortoiseCVS/cvs.exe")
set (CTEST_CONFIGURE_COMMAND
    "\\"${TCLSH}\"
    \\"${CTEST_BINARY_DIRECTORY}/Scripts/genlib.tcl\"")
set (CTEST_BUILD_COMMAND
    "\\"${TCLSH}\"
    \\"${CTEST_BINARY_DIRECTORY}/Scripts/cmaker.tcl\"")

# clear out the binary tree
file (WRITE "${CTEST_BINARY_DIRECTORY}/CMakeCache.txt"
    "// Dummy cache just so that ctest will wipe binary dir")
ctest_empty_binary_directory (${CTEST_BINARY_DIRECTORY})

# special variables for the Slicer build process
set (ENV{MSVC6}           "0")
set (ENV{GENERATOR}        "Visual Studio 7 .NET 2003")
set (ENV{MAKE}              "devenv.exe ")
set (ENV{COMPILER_PATH}
    "C:/Program Files/Microsoft Visual Studio .NET
2003/Common7/Vc7/bin")
set (ENV{CVS}                "${CTEST_UPDATE_COMMAND}")

# start and update the dashboard
ctest_start (Nightly)
ctest_update (SOURCE "${CTEST_SOURCE_DIRECTORY}")

# define a macro to copy a directory
macro (COPY_DIR srccdir destdir)
    exec_program ("${CMAKE_EXECUTABLE_NAME}" ARGS
        "-E copy_directory \"${srccdir}\" \"${destdir}\\"")
endmacro ()

# Slicer does not support out of source builds so we
# first copy the source directory to the binary directory
# and then build it
copy_dir ("${CTEST_SOURCE_DIRECTORY}"
    "${CTEST_BINARY_DIRECTORY}")
```

```
# copy support libraries that slicer needs into the binary tree
copy_dir ("${SLICER_SUPPORT}"
          "${CTEST_BINARY_DIRECTORY}/Lib")

# finally do the configure, build, test and submit steps
ctest_configure (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_build (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_test (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_submit ()
```

With extended CTest scripting we have full control over the flow, so we can perform arbitrary commands at any point. For example, after performing an update of the project, the script copies the source tree into the build directory. This allows it to do an “out-of-source” build.

10.11 Setting up a Dashboard Server

For many people making use of Kitware’s my.cdash.org dashboard hosting will be sufficient, if that is the case for you then you can skip this section. If you wish to setup your own server, then this section will walk you through the process. There are a few options for what to run on the server to process the dashboard results. The preferred option is to use CDash, a new dashboard server based on PHP, MySQL, CSS, and XSLT. Predecessors to CDash such as DART 1 and DART 2 can also be used, information on them can be found at <http://www.itk.org/Dart/HTML/Index.shtml>.

CDash Server

CDash is a new dashboard server developed by Kitware that is based on the common LAMP platform. It makes use of PHP, CSS, XSL, MySQL/PostgreSQL and of course your web server, normally Apache. CDash takes the dashboard submissions as XML and stores them in an SQL database (currently MySQL and PostgreSQL are supported). When the web server receives requests for pages the PHP scripts extract the relevant data from the database and produce XML that is sent to XSL templates, that in turn convert it into HTML. CSS is used to provide the overall look and feel for the pages. CDash can handle large projects ,and has been hosting up to 20 projects on a typical server (dual-core PC running Ubuntu), with about 300M records stored in the database.

Server requirements

- MySQL (5.x and higher) or PostgreSQL (8.3 and higher)
- PHP (5.0 recommended)
- XSL module for PHP (`apt-get install php5-xsl`)

- cURL module for PHP
- GD module for PHP

Getting CDash

You can get CDash from the www.cdash.org website, or you can get the latest code from SVN using the following command:

```
svn co https://www.kitware.com:8443/svn/CDash/trunk CDash
```

Quick installation

1. Unzip or checkout CDash in your webroot directory on the server. Make sure the web server has read permission to the files
2. Create a `cdash/config.local.php` and add the following lines, adapted for your server configuration

```
// Hostname of the database server  
$CDASH_DB_HOST = 'localhost';  
// Login for database access  
$CDASH_DB_LOGIN = 'root';  
// Password for database access  
$CDASH_DB_PASS = '';  
// Name of the database  
$CDASH_DB_NAME = 'cdash';  
// Database type  
$CDASH_DB_TYPE = 'mysql';
```

3. Point your web browser to the `install.php` script

```
http://mywebsite.com/CDash/install.php
```

4. Follow the installation instructions
5. When the installation is done, add the following line in the `config.local.php` to ensure the installation script is no longer accessible

```
$CDASH_PRODUCTION_MODE = true;
```

Testing the installation

In order to test the installation of the CDash server, you can download a small test project and test the submission to CDash, by following these steps:

1. Download the test project at

```
• http://www.cdash.org/download/CDashTest.zip
```

2. Create a CDash project named “test” on your CDash server (see 10.7 Producing Test Dashboards)
3. Download the `CTestConfig.cmake` file from the CDash server
4. Run CMake on `CDashTest` to configure the project
5. Run

```
make experimental
```

6. Go to the dashboard page for the “test” project, you should see the submission in the experimental section.

Advanced Server Management

Project Roles: CDash supports three role levels for users:

- Normal users are regular users with read and/or write access to the project’s code repository.
- Site maintainers are responsible for periodic submissions to CDash.
- Project administrators have reserved privileges to administer the project in CDash.

The first two levels can be defined by the users themselves. Project administrators access must be granted by another administrator of the project, or a CDash administrator.

In order to change the current role for a user:

1. Select [Manage project roles] in the administration section
2. If you have more than one project, select the appropriate project
3. In the “current users” section change the role for a user
4. Click “update” to update the current role
5. In order to completely remove a user from a project, click “remove”

6. If the CVS login is not correct it can be changed from this page. Note that users can also change their CVS login manually from their profile

In order to add a current role for a user:

1. Select [Manage project roles] in the administration section
2. Then, if you have more than one project, select the appropriate project
3. In the “Add new user” section type the first letters of the first name, last name or email address of the user you want to add. Or type ‘%’ in order to show all the users registered in CDash
4. Select the appropriate user’s role
5. Optionally enter the user’s CVS login
6. Click on “add user”

The screenshot shows the 'Project Roles' management page for the 'CDASH' project. At the top, there's a navigation bar with a 'BACK' button. Below it, a table lists 'Current users' for the 'CMake' project. The table has columns for Firstname, Lastname, Email, Role, CVS Login, and Action. One row is shown with the values: administrator, admin@cdash.org, Project Administrator, jjomier, update, remove. Below the table, there's a section for 'Add new user' with a search input containing 'test2'. Underneath, there's a list of users: françois test (test2@test.com) and a dropdown for 'role' set to 'Normal User'. There's also a 'cvslogin:' field and a 'add user' button. At the bottom, there's a 'CVS Users File' input with 'Browse...' and 'import' buttons. The footer features the Kitware logo and a link to 'report problems'.

Figure 33 – Project Role management page in CDash

Importing users: to batch import a list of current users for a given project

1. Click on [manage project role] in the administration section
2. Select the appropriate project
3. Click “Browse” to select a CVS users file.
4. The current file should be formatted as follows:

```
cvsuser:email:first_name last_name
```

5. Click “import”
6. Make sure the reported names and email addresses are correct, deselect any that should not be imported
7. Click on “Register and send email”. This will automatically register the users, set a random password and send a registration request to the appropriate email addresses

Google Analytics

Usage statistics of the CDash server can be assessed using Google Analytics. In order to setup google analytics:

1. Go to <http://www.google.com/analytics/index.html>
2. Setup an account if necessary
3. Add a website project
4. Login into CDash as the administrator of a project
5. Click on “Edit Project”
6. Add the code from Google into the Google Analytics Tracker (i.e. UA-43XXXX-X) for your project

Submission backup

CDash backups all the incoming XML submissions and places them in the `backup` directory by default. The default timeframe is 48 hours. The timeframe can be changed in the `config.local.php` as follows:

```
$CDASH_BACKUP_TIMEFRAME=72;
```

If projects are private it is recommended to set the backup directory outside of the apache root directory to make sure that nobody can access the XML files, or to add the following lines to the `.htaccess` in the backup directory:

```
<Files *>
order allow,deny
deny from all
</files>
```

Note that the backup directory is emptied only when a new submission arrives. If necessary, CDash can also import builds from the backup directory:

1. Log into CDash as administrator
2. Click on [Import from backups] in the administration section

3. Click on “Import backups”

Build Groups

Builds can be organized by groups. In CDash, three groups are defined automatically and cannot be removed: `Nightly`, `Continuous` and `Experimental`. These groups are the same as the ones imposed by CTest. Each group has an associated description that is displayed when clicking on the name of the group on the main dashboard.

To add a new group:

1. Click on [manage project groups] in the administration section
2. Select the appropriate project
3. Under the “create new group” section enter the name of the new group
4. Click on “create group”. The newly created group appears at the bottom of the current dashboard

To order groups:

1. Click on [manage project groups] in the administration section
2. Select the appropriate project
3. Under the “Current Groups” section, click on the [up] or [down] links. The order displayed in this page is exactly the same as the order on the dashboard

To update group description:

1. Click on [manage project groups] in the administration section
2. Select the appropriate project
3. Under the “Current Groups” section, update or add a description in the field next to the [up][down] links
4. Click “Update Description” in order to commit your changes

By default a build belongs to the group associated with the build type defined by CTest, i.e. a nightly build will go in the nightly section. CDash matches a build by its name, site, and build type. For instance a nightly build named “Linux-gcc-4.3” from the site “midworld.kitware” will be moved to the nightly section unless a rule on “Linux-gcc-4.3”-“midworld.kitware”-“Nightly” is defined. There are two ways to move a build into a given group by defining a rule: Global Move and Single Move.

Global move allows moving builds in batch.

1. Click on [manage project groups] in the administration section .
2. Select the appropriate project (if more than one) .

3. Under “Global Move” you will see a list of the builds submitted in the past 7 days. (without duplicates). Note that expected builds are also shown, even if they have not been submitting in the past 7 days.
4. You can narrow your search by selecting a specific group (default is All).
5. Select the builds to move. Hold “shift” in order to select multiple builds.
6. Select the target group. This is mandatory.
7. Optionally check the “expected” box if you expect the builds to be submitted on a daily basis. For more information on expected builds see the “Expected builds” section below.
8. Click “Selected Builds to Group” to move the groups.

Single move allows modifying only a particular build. From the main dashboard page, if logged in as an administrator of the project, a small folder icon is displayed next to each build. Clicking on the icon shows some options for each build. In particular, project administrators can mark a build as expected, move a build to a specific group, or delete a bogus build.

Expected builds: Project administrators can mark certain builds as expected. That means builds are expected to submit daily. This allows you to quickly check if a build has not been submitting on today's dashboard, or to quickly assess how long the build has been missing by clicking on the info icon on the main dashboard.

midworld.kitware	DarwinG5-g++ 
	This build has not been submitting since <u>2008-04-23 15:56:00 (9 days)</u>

Figure 34 –Information regarding a build from the main dashboard page

If an expected build was not submitted the previous day and the option “Email Build Missing” is checked for the project, an email will be sent to the site maintainer and project administrator to alert them (see the Sites section for more information).

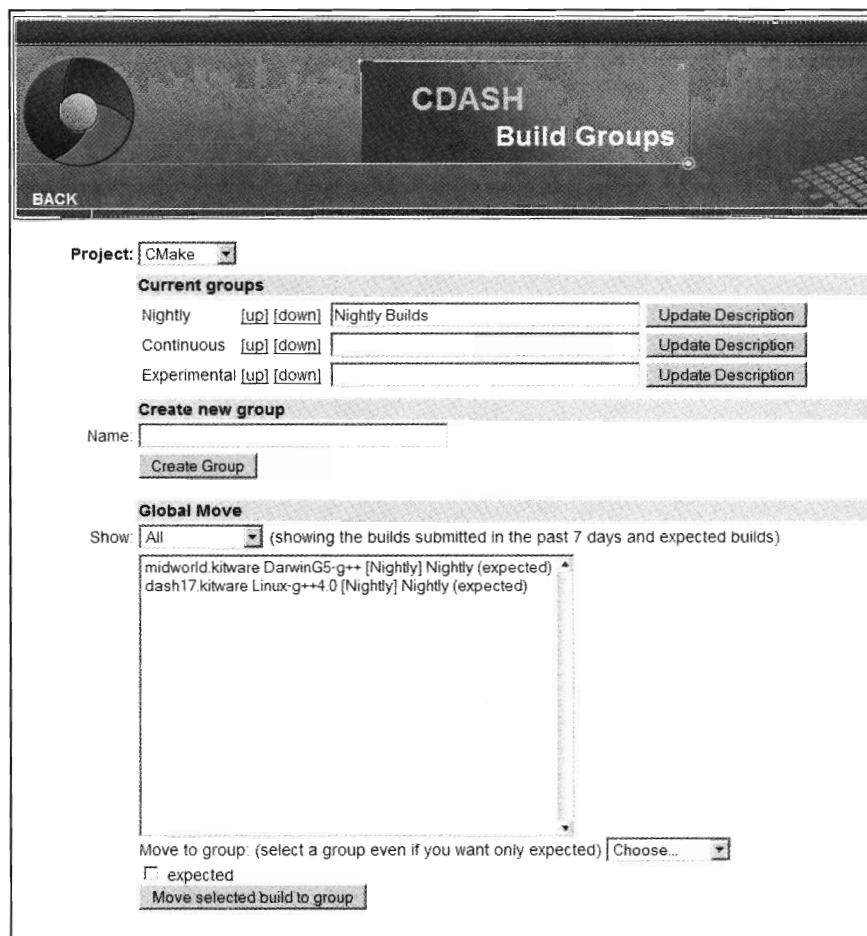


Figure 35 –Build Group Configuration Page

Email

CDash sends email to developers and project administrators when a failure occurs for a given build. The configuration of the email feature is located in three places: the config.local.php file, the project administration section and the project's groups section.

In the config.local.php, two variables are defined to specify the email address from which email is sent and the reply address. Note that the SMTP server cannot be defined in the current version of CDash, it is assumed that a local email server is running on the machine.

```
$CDASH_EMAIL_FROM = 'admin@mywebsite.com';
$CDASH_EMAIL_REPLY = 'noreply@mywebsite.com';
```

In the email configuration section of the project, several parameters can be tuned to control the email feature. These parameters were described in the previous section, “Adding CDash Support to a Project”.

In the “build groups” administration section of a project, an administrator can decide if emails are sent to a specific group, or if only a summary email should be sent. The summary email is sent for a given group when at least one build is failing on the current day.

Sites

CDash refers to a site as an individual machine submitting at least one build to a given project. A site might submit multiple builds (e.g. nightly and continuous) to multiple projects stored in CDash.

In order to see the site description, click on the name of the site from the main dashboard page for a project. The description of a site includes information regarding the processor type and speed as well as the amount of memory available on the given machine. The description of a site is automatically sent by CTest, however in some cases it might be required to manually edit it. Moreover, if the machine is upgraded, i.e. the memory is upgraded; CDash keeps track of the history of the description, allowing users to compare performance before and after the upgrade.

Sites usually belong to one maintainer, responsible for the submissions to CDash. It is important for site maintainers to be warned when a site is not submitting as it could be related to a configuration issue. In order to claim a site, a maintainer should

1. Log into CDash
2. Click on a dashboard containing at least one build for the site
3. Click on the site name to open the description of the site
4. Click on [claim this site]

Once a site is claimed, its maintainer will receive emails if the client machine does not submit for an unknown reason, assuming that the site is expected to submit nightly. Furthermore, the site will appear in the “My Sites” section of the maintainer’s profile, facilitating a quick check of the site’s status.

Another feature of the site page is the pie chart showing the load of the machine. Assuming that a site submits to multiple projects, it is usually useful to know if the machine has room for other submissions to CDash. The pie chart gives an overview of the machine submission time for each project.

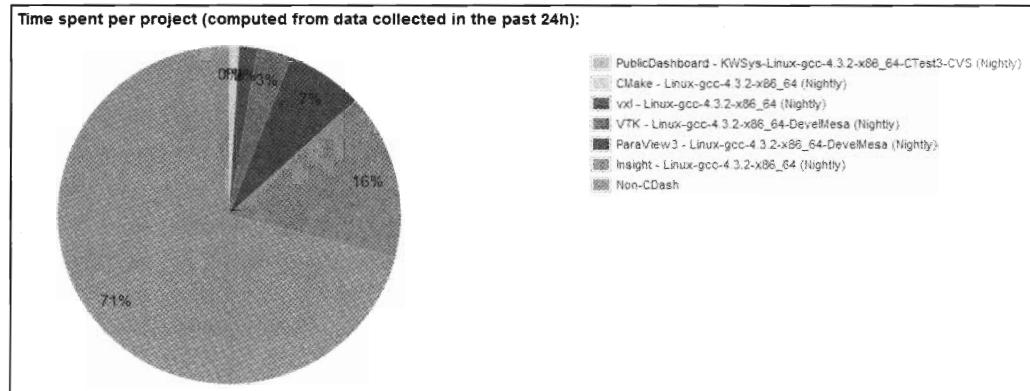


Figure 36 –Pie chart showing how much time is spent by a given site on building CDash projects

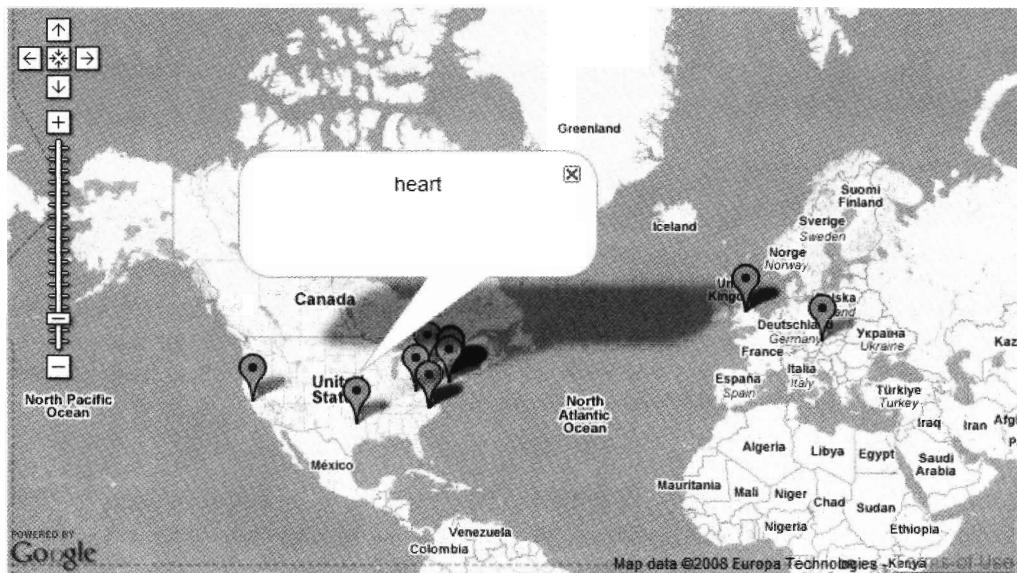


Figure 37 –Map showing the location of the different sites building a project

Graphs

CDash currently plots three types of graph. The graphs are generated dynamically from the database records and are interactive.

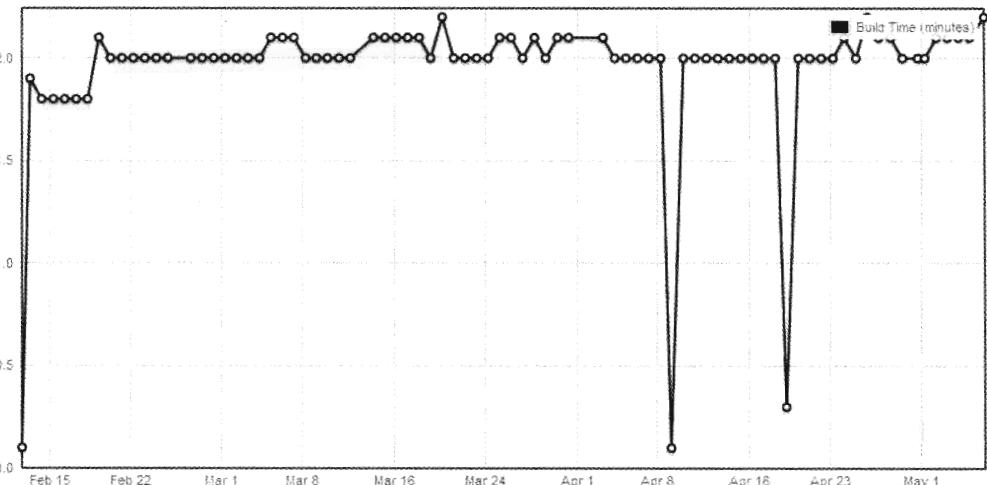


Figure 38 – Example of build time over time

The build time graph displays the time required to build a project over time. In order to display the graph you need to:

1. Go to the main dashboard for the project.
2. Click on the build name you want to track.
3. On the build summary page, click on [Show Build Time Graph].

The test time graphs display the time to run a specific test as well as its status (passed/failed) over time. To display it:

1. Go to the main dashboard for a project.
2. Click on the number of test passed or failed.
3. From the list of tests, click on the status of the test.
4. Click on [Show Test Time Graph] and/or [Show Failing/Passing Graph].

Adding Notes to a Build

In some cases, it is useful to inform other developers that someone is currently looking at the errors for a build. CDash implements a simple note mechanism for that purpose:

1. Login to CDash.
2. On the dashboard project page, click on the build name that you would like to add the note to.
3. Click on the [Add a Note to this Build] link, located next to the current build matrix (see thumbnail).

4. Enter a short message that will be added as a note.
5. Select the status of the note: Simple note, Fix in progress, Fixed.
6. Click on "Add Note".

Logging

CDash supports internal logging mechanism using the `error_log()` function from PHP. Most of the critical SQL queries are logged. By default the CDash log file is located in the `backup` directory under the name `cdash.log`. The location of the log file can be modified by changing the variable in the `config.local.php` configuration file.

```
$CDASH_BACKUP_DIRECTORY='/var/temp/cdashbackup/log';
```

The log file can be accessed directly from CDash if the log file is in the standard location:

1. Log into CDash as administrator.
2. Click on [CDash logs] in the administration section.
3. Click on `cdash.log` to see the log file.

Note that CDash does not perform any log rotation.

Test Timing

CDash supports checks on the duration of tests. CDash keeps the current weighted average of the mean and standard deviation for the time each test takes to run in the database. In order to keep the computation as efficient as possible, the following formula is used, which only involves the previous build.

```
newMean = (1-alpha)*oldMean + alpha*currentTime  
  
newSD = sqrt((1-alpha)*SD*SD + alpha*(currentTime-newMean)*(currentTime-newMean))
```

A test is defined as failing based on the following logic:

```
if previousSD < thresholdSD then previousSD = thresholdSD  
if currentTime > previousMean+multiplier*previousSD then fail
```

It should be noted that alpha defines the current “window” for the computation. By default alpha is set to 0.3.

Mobile Support

Since CDash is written using template layers via XSLT, developing new layouts is as simple as adding new rendering templates. As a demonstration, an iPhone web template is provided with the current version of CDash.

```
http://mycdashserver/CDash/iphone
```

The main page shows a list of the public projects hosted on the server. Clicking on the name of a project loads its current dashboard. In the same manner, clicking on a given build displays more detailed information about that build. For the moment, the ability to login and to access private sections of CDash are not supported with this layout.



Figure 39 –Example of dashboard on the iPhone

Backing up CDash

All of the data (except the logs) used by CDash is stored in its database. It is important to backup the database regularly, especially so before performing a CDash upgrade. There are a couple of ways to backup a MySQL database. The easiest is to use the `mysqldump` (<http://dev.mysql.com/doc/refman/5.1/en/mysqldump.html>) command:

```
mysqldump -r cdashbackup.sql cdash
```

If you are using MyISAM tables exclusively you can copy the CDash directory in your MySQL data directory. Note that you need to shutdown MySQL before doing the copy so that

no file could be changed during the copy. Similarly to MySQL, PostGreSQL has a pg_dump utility:

```
pg_dump -U posgreSQL_user cdash > cdashbackup.sql
```

Upgrading CDash

When a new version of CDash is released or if you decide to update from the SVN repository, CDash will warn you on the front page if the current database needs to be upgraded. When upgrading to a new release version the following steps should be taken:

1. Backup your SQL database (see previous section).
2. Backup your config.local.php (or config.php) configuration files.
3. Replace your current cdash directory with the latest version and copy the config.local.php in the cdash directory.
4. Navigate your browser to your CDash page. (e.g.<http://localhost/CDash>).
5. Note the version number on the main page, it should match the version that you are upgrading to.
6. The following message may appear: "The current database schema doesn't match the version of CDash you are running, upgrade your database structure in the Administration panel of CDash." This is a helpful reminder to perform the following steps.
7. Login to CDash as administrator.
8. In the 'Administration' section, click on '[CDash Maintenance]'.
9. Click on 'Upgrade CDash': this process might take some time depending on the size of your database (do not close your browser).
 - Progress messages may appear while CDash performs the upgrade.
 - If the upgrade process takes too long you can check in the backup/cdash.log file to see where the process is taking a long time and/or failing.
 - It has been reported that on some systems the spinning icon never turns into a check mark. Please check the cdash.log for the "Upgrade done." string if you feel that the upgrade is taking too long.
 - On a 50GB database the upgrade might take up to 2 hours.
10. Some web browsers might have issues when upgrading (with some javascript variables not being passed correctly), in that case you can perform individual updates. For example, upgrading from CDash 1-2 to 1-4:
<http://mywebsite.com/CDash/backwardCompatibilityTools.php?upgrade=1-4=1>

CDash Maintenance

Database maintenance: we recommend that you perform database optimization (reindexing, purging, etc.) regularly to maintain a stable database. MySQL has a utility called `mysqlcheck`, and PostgreSQL has several utilities such as `vacuumdb`.

Deleting builds with incorrect dates: some builds might be submitted to CDash with the wrong date, either because the date in the XML file is incorrect or the timezone was not recognized by CDash (mainly by PHP). These builds will not show up in any dashboard because the start time is bogus. In order to remove these builds:

1. Login to CDash as administrator.
2. Click on [CDash maintenance] in the administration section.
3. Click on ‘Delete builds with wrong start date’.

Recompute test timing: if you just upgraded CDash you might notice that the current submissions are showing a high number of failing test due to time defects. This is because CDash does not have enough sample points to compute the mean and standard deviation for each test, in particular the standard deviation might be very small (probably zero for the first few samples). You should turn the “enable test timing” off for about a week, or until you get enough build submissions and CDash has calculated an approximate mean and standard deviation for each test time.

The other option is to force CDash to compute the mean and standard deviation for each test for the past few days. Be warned that this process may take a long time, depending on the number of test and projects involved. In order to recompute the test timing:

1. Login to CDash as administrator.
2. Click on [CDash maintenance] in the administration section.
3. Specify the number of days (default is 4) to recompute the test timings for.
4. Click on “Compute test timing”. When the process is done the new mean, standard deviation and status should be updated for the tests submitted during this period.

Automatic build removal

In order to keep the database at a reasonable size, CDash can automatically purge old builds. There are currently two ways to setup automatic removal of builds: without a cronjob edit the `config.local.php` and add/edit the following line

```
$CDASH_AUTOREMOVE_BUILD=1;
```

CDash will automatically remove builds on the first submission of the day. Note that removing builds might add an extra load on the database, or slow down the current

submission process if your database is large and the number of submissions is high. If you can use a cronjob the PHP command line tool can be used to trigger build removals at a convenient time. For example, removing the builds for all the projects at 6am every Sunday:

```
0 6 * * 0 php5 /var/www/CDash/autoRemoveBuilds.php all
```

Note that the ‘all’ parameter can be changed to a specific project name in order to purge builds from a single project.

CDash XML Schema

The XML parsers in CDash can be easily extended to support new features. The current XML schemas generated by CTest, and their features as described in the book, are located at:

```
http://public.kitware.com/Wiki/CDash:XML
```

10.12 Subprojects

CDash (versions 1.4 and later) supports splitting projects into subprojects. Some of the subprojects may in turn depend on other subprojects. A typical real life project consists of libraries, executables, test suites, documentation, web pages and installers. Organizing your project into well-defined subprojects and presenting the results of nightly builds on a CDash dashboard can help identify where the problems are at different levels of granularity.

A project with subprojects has a different view for its top level CDash page than a project without any. It contains a summary row for the project as a whole, and then one summary row for each subproject.

The screenshot shows the TRILINOS Dashboard interface. At the top, there's a navigation bar with links for 'Login | Dashboards' and the date 'Monday, September 14 2009 10:05:05 MDT'. Below the header, the title 'TRILINOS Dashboard' is centered above a dark background image of a computer monitor displaying code. The main content area has a light gray header labeled 'Project'. Below it is a table for the 'Trilinos' project with columns for Error, Warning, Pass, and various build/test metrics. To the right, a 'Last submission' timestamp is shown. A second table below, titled 'SubProjects', lists several subprojects (Teuchos, RTOp, Kokkos, etc.) with their own detailed build/test statistics and last submission times.

Project	Configure			Build			Test			Last submission
	Error	Warning	Pass	Error	Warning	Pass	Not Run	Fail	Pass	
Trilinos	0	6	120	0	29	149	0	1	825	2009-09-14 10:05:05

Project	Configure			Build			Test			Last submission
	Error	Warning	Pass	Error	Warning	Pass	Not Run	Fail	Pass	
Teuchos	0	1	6	0	0	9	0	0	2	2009-09-14 16:03:47
RTOp	0	1	7	0	0	7				2009-09-14 15:02:49
Kokkos	0	1	7	0	0	7	0	0	19	2009-09-14 15:03:00
Epetra	0	1	7	0	5	1	0	0	4	2009-09-14 15:03:17
Zoltan	0	1	7	0	0	7	0	0	4	2009-09-14 15:03:46
Shards	0	0	4	0	0	4				2009-09-14 15:05:20
Intrepid	0	0	1	0	1	3	0	0	80	2009-09-14 15:36:54
GlobiPack	0	0	4	0	0	4	0	0	9	2009-09-14 15:05:32
Trilutils	0	0	4	0	1	3	0	0	2	2009-09-14 15:05:48
Tpetra	0	0	4	0	0	4	0	0	5	2009-09-14 15:07:42

Figure 40 –Main project page with subprojects

Organizing and defining subprojects

To add subproject organization to your project, you must: (1) define the subprojects for CDash, so that it knows how to display them properly and (2) use build scripts with CTest to submit subproject builds of your project. Some (re-)organization of your project's CMakeLists.txt files may also be necessary to allow building of your project by subprojects.

There are two ways to define subprojects and their dependencies: interactively in the CDash GUI when logged in as a project administrator, or by submitting a Project.xml file describing the subprojects and dependencies.

Adding Subprojects Interactively

As a project administrator, a “Manage subprojects” button will appear for each of your projects on the My CDash page. Clicking the Manage Subprojects button opens the manage subproject page where you may add new subprojects, or establish dependencies between existing subprojects for any project that you are an administrator of. There are two tabs on this page, one for viewing the current subprojects along with their dependencies, and one for creating new subprojects.

To add subprojects, for instance two subprojects called Exes and Libs, and to make Exes depend on Libs, the following steps are necessary:

- Click the “Add a subplot” tab.
- Type “Exes” in the “Add a subplot” edit field.
- Click the “Add subplot” button.
- Click the “Add a subplot” tab.
- Type “Libs” in the “Add a subplot” edit field.
- Click the “Add Subproject” button.
- In the “Exes” row of the “Current Subprojects” tab, choose “Libs” from the “Add dependency” drop down list and click the “Add dependency” button.

To remove a dependency or a subplot, click on the “X” next to the item you wish to delete.

Adding Subprojects Automatically

Another way to define CDash subplots and their dependencies is to submit a “Project.xml” file along with the usual submission files that CTest sends when it submits a build to CDash. To define the same two subplots as in the interactive example above (Exes and Libs) with the same dependency (Exes depend on Libs) the `Project.xml` file would look like the following example:

```
<Project name="Tutorial">
    <SubProject name="Libs"></SubProject>
    <SubProject name="Exes">
        <Dependency name="Libs">
        </Dependency>
    </SubProject>
</Project>
```

Once the `Project.xml` file is written or generated, it can be submitted to CDash from a `ctest -S` script using the `new FILES` argument to the `ctest_submit` command, or directly from the `ctest` command line in a build tree configured for dashboard submission.

From inside a `ctest -S` script:

```
ctest_submit(FILES "${CTEST_BINARY_DIRECTORY}/Project.xml")
```

From the command line:

```
cd ..../Project-build
ctest --extra-submit Project.xml
```

CDash will automatically add subprojects and dependencies according to the `Project.xml` file. CDash will also remove any subprojects or dependencies not defined in the `Project.xml` file. Additionally, if the `Project.xml` is submitted multiple times, the second and subsequent submissions will have no observable effect: the first submission adds/modifies the data, the second and later submissions send the same data, so no changes are necessary. CDash tracks changes to the subproject definitions over time to allow for projects to evolve. If you view dashboards from a past date, CDash will present the project/subproject views according to the subproject definitions in effect on that date.

Using `ctest_submit` with PARTS and FILES

In `ctest` version 2.8 and later, the `ctest_submit` command supports new `PARTS` and `FILES` arguments. With `PARTS`, you can send any subset of the XML files with each `ctest_submit` call. Previously, all parts would be sent with any call to `ctest_submit`. Typically, the script would wait until all dashboard stages were complete and then call `ctest_submit` once to send the results of all stages at the end of the run. Now, a script may call `ctest_submit` with `PARTS` to do partial submissions of subsets of the results. For example, you can submit configure results after `ctest_configure`, build results after `ctest_build` and test results after `ctest_test`. This allows for information to be posted as the builds progress.

With `FILES`, you can send arbitrary XML files to CDash. In addition to the standard build result XML files that CTest sends, CDash also handles the new `Project.xml` file that describes subprojects and dependencies, as described previously. Prior to the addition of the `ctest_submit` `PARTS` handling, a typical dashboard script would contain a single `ctest_submit()` call on its last line:

```
ctest_start (Experimental)
ctest_update (SOURCE "${CTEST_SOURCE_DIRECTORY}")
ctest_configure (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_build (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_test (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_submit ()
```

Now, submissions can occur incrementally, with each part of the submission sent piecemeal as it becomes available:

```
ctest_start (Experimental)
ctest_update (SOURCE "${CTEST_SOURCE_DIRECTORY}")
ctest_configure (BUILD "${CTEST_BINARY_DIRECTORY}")
ctest_submit (PARTS Update Configure Notes)

ctest_build (BUILD "${CTEST_BINARY_DIRECTORY}" APPEND)
ctest_submit (PARTS Build)
```

```
ctest_test (BUILD "${CTEST_BINARY_DIRECTORY}" APPEND)
ctest_submit (PARTS Test)
```

Submitting incrementally by parts means that you can inspect the results of the configure stage live on the CDash dashboard while the build is still in progress. Likewise, you can inspect the results of the build stage live while the tests are still running.

Splitting Your Project into Multiple Subprojects

One `ctest_build` invocation that builds everything followed by one `ctest_test` invocation that tests everything is sufficient for a project that has no subprojects. But if you want to submit results on a per-subproject basis to CDash, you will have to make some changes to your project and test scripts. For your project you need to identify what targets are part of what sub-projects. If you organize your CMakeLists files such that you have a target to build for each subproject, and you can derive (or look up) the name of that target based on the subproject name, then revising your script to separate it into multiple smaller configure/build/test chunks should be relatively painless. To do this you can modify your CMakeLists files in various ways depending on your needs. The most common changes are listed below.

CMakeLists.txt modifications

- Name targets the same as subprojects, or base target names on subproject names, or provide a look up mechanism to map from subproject name to target name.
- Possibly add custom targets to aggregate existing targets into subprojects, using `add_dependencies` to say which existing targets the custom target depends on.
- Add the `LABELS` target property to targets with a value of the subproject name.
- Add the `LABELS` test property to tests with a value of the subproject name.

Next you need to modify your CTest scripts that run your dashboards. To split your one large monolithic build into smaller subproject builds you can use a `foreach` loop in your CTest driver script. To help you iterate over your subprojects, CDash provides a variable named `CTEST_PROJECT_SUBPROJECTS` in `CTestConfig.cmake`. Given the above example, CDash produces a variable like this:

```
set (CTEST_PROJECT_SUBPROJECTS Libs Exes)
```

CDash orders the elements in this list such that the independent subprojects (that do not depend on any other subprojects) are first, followed by subprojects that depend only on the independent subprojects. After that subprojects that depend on those. The same logic continues until all subprojects are listed exactly once in this list in an order that makes sense for building them sequentially, one after the other.

To facilitate building just the targets associated with a subproject, use the variable named `CTEST_BUILD_TARGET` to tell `ctest_build` what to build. To facilitate running just the tests associated with a subproject, assign the `LABELS` test property to your tests and use the new `INCLUDE_LABEL` argument to `ctest_test`.

ctest driver script modifications

- Iterate over the subprojects in dependency order (from independent to most dependent...).
- Set the SubProject and Label global properties – CTest uses these properties to submit the results to the correct subproject on the CDash server.
- Build the target(s) for this subproject: compute the name of the target to build from the subproject name, set `CTEST_BUILD_TARGET`, call `ctest_build`.
- Run the tests for this subproject using the `INCLUDE` or `INCLUDE_LABEL` arguments to `ctest_ctest`.
- Use `ctest_submit` with the `PARTS` argument to submit partial results as they complete.

To illustrate this, the following example shows the changes required to split a build into smaller pieces. Assume that the subproject name is the same as the target name required to build the subproject's components. For example, here is a snippet from `CMakeLists.txt`, in the hypothetical Tutorial project. The only additions necessary (since the target names are the same as the subproject names) are the calls to `set_property` for each target and each test.

```
# "Libs" is the library name (therefore a target name) and
# the subproject name
add_library (Libs ...)
set_property (TARGET Libs PROPERTY LABELS Libs)
add_test (LibsTest1 ...)
add_test (LibsTest2 ...)
set_property (TEST LibsTest1 LibsTest2 PROPERTY LABELS Libs)

# "Exes" is the executable name (therefore a target name)
# and the subproject name
add_executable (Exes ...)
target_link_libraries (Exes Libs)
set_property (TARGET Exes PROPERTY LABELS Exes)
add_test (ExesTest1 ...)
add_test (ExesTest2 ...)
set_property (TEST ExesTest1 ExesTest2 PROPERTY LABELS Exes)
```

Here is an example of what the CTest driver script might look like before and after organizing this project into subprojects. Before the changes:

```
ctest_start (Experimental)
ctest_update (SOURCE "${CTEST_SOURCE_DIRECTORY}")
ctest_configure (BUILD "${CTEST_BINARY_DIRECTORY}")
# builds *all* targets: Libs and Exes
ctest_build (BUILD "${CTEST_BINARY_DIRECTORY}")
# runs *all* tests
ctest_test (BUILD "${CTEST_BINARY_DIRECTORY}")
# submits everything all at once at the end
ctest_submit ()
```

After the changes (new lines emphasized):

```
ctest_start (Experimental)
ctest_update (SOURCE "${CTEST_SOURCE_DIRECTORY}")
ctest_submit (PARTS Update Notes)

# to get CTEST_PROJECT_SUBPROJECTS definition:
include ("${CTEST_SOURCE_DIRECTORY}/CTestConfig.cmake")

foreach (subproject ${CTEST_PROJECT_SUBPROJECTS})
    set_property (GLOBAL PROPERTY SubProject ${subproject})
    set_property (GLOBAL PROPERTY Label ${subproject})

    ctest_configure (BUILD "${CTEST_BINARY_DIRECTORY}")
    ctest_submit (PARTS Configure)

    set (CTEST_BUILD_TARGET "${subproject}")
    ctest_build (BUILD "${CTEST_BINARY_DIRECTORY}")
        # builds target ${CTEST_BUILD_TARGET}
    ctest_submit (PARTS Build)

    ctest_test (BUILD "${CTEST_BINARY_DIRECTORY}"
        INCLUDE_LABEL "${subproject}"
    )
# runs only tests that have a LABELS property matching
# "${subproject}"
    ctest_submit (PARTS Test)
endforeach ()
```

In some projects, more than one `ctest_build` step may be required to build all the pieces of the subproject. For example, in Trilinos, each subproject builds the `subproject_libs` target and then builds the `all` target to build all the configured executables in the test suite. They also

configure dependencies such that only the executables that need to be built for the currently configured packages build when the all target is built.

Normally, if you submit multiple `Build.xml` files to CDash with the same exact build stamp, it will delete the existing entry and add the new entry in its place. In the case where multiple `ctest_build` steps are required, each with their own `ctest_submit(PARTS Build)` call, use the `APPEND` keyword argument in all of the `ctest_build` calls that belong together. The `APPEND` flag tells CDash to accumulate the results from multiple submissions and display the aggregation of all of them in one row on the dashboard. From CDash's perspective, multiple `ctest_build` calls (with the same build stamp and subproject and `APPEND` turned on) result in a single CDash build.

Adopt some of these tips and techniques in your favorite CMake-based project:

- `LABELS` is a new CMake/CTest property that applies to source files, targets and tests. Labels are sent to CDash inside the resulting xml files.
- Use `ctest_submit (PARTS ...)` to do incremental submissions. Results are available for viewing on the dashboards sooner.
- Use `INCLUDE_LABEL` with `ctest_test` to run only the tests with labels that match the regular expression.
- Use `CTEST_BUILD_TARGET` to build your subprojects one at a time, submitting subproject dashboards along the way.

Porting CMake to New Platforms and Languages

In order to generate build files for a particular system, CMake needs to determine what system it is running on, and what compiler tools to use for enabled languages. To do this CMake loads a series of files containing CMake code from the Modules directory. This all has to happen before the first try-compile or try-run is executed. To avoid having to re-compute all of this information for each try-compile and for subsequent runs of CMake, the discovered values are stored in several configured files that are read each time CMake is run. These files are also copied into the try-compile and try-run directories. This chapter will describe how this process of system and tool discovery works. An understanding of the process is necessary to extend CMake to run on new platforms, and to add support for new languages.

11.1 The Determine System Process

The first thing CMake needs to do is to determine what platform it is running on and what the target platform is. Except for when you are cross compiling the host platform and the target platform are identical. The host platform is determined by loading the `CMakeDetermineSystem.cmake` file. On POSIX systems, "uname" is used to get the name of the system. `CMAKE_HOST_SYSTEM_NAME` is set to the result of `uname -s`, and `CMAKE_HOST_SYSTEM_VERSION` is set to the result of `uname -r`. On Windows systems, `CMAKE_HOST_SYSTEM_NAME` is set to Windows and `CMAKE_HOST_SYSTEM_VERSION` is set to the value returned by the system function `GetVersionEx`. The variable `CMAKE_HOST_SYSTEM` is set to a combination of `CMAKE_HOST_SYSTEM_NAME` and `CMAKE_HOST_SYSTEM_VERSION` as follows:

```
${CMAKE_HOST_SYSTEM_NAME}-${CMAKE_HOST_SYSTEM_VERSION}
```

Additionally CMake tries to figure out the processor of the host, on POSIX systems it uses `uname -m` or `uname -p` to retrieve this information, on Windows it uses the environment variable `PROCESSOR_ARCHITECTURE`. `CMAKE_HOST_SYSTEM_PROCESSOR` holds the value of the result.

Now that CMake has the information about the host that it is running on, it needs to find this information for the target platform. The results will be stored in the `CMAKE_SYSTEM_NAME`, `CMAKE_SYSTEM_VERSION`, `CMAKE_SYSTEM` and `CMAKE_SYSTEM_PROCESSOR` variables, corresponding to the `CMAKE_HOST_SYSTEM_*` variables described above. See the "Cross compiling with CMake" chapter on how this is done when cross compiling. In all other cases the `CMAKE_SYSTEM_*` variables will be set to the value of their corresponding `CMAKE_HOST_SYSTEM_*` variable.

Once the `CMAKE_SYSTEM` information has been determined, `CMakeSystem.cmake.in` is configured into `$(CMAKE_BINARY_DIR)/CMakeFiles/CMakeSystem.cmake`. CMake versions prior to 2.6.0 did not support cross compiling, and so only the `CMAKE_SYSTEM_*` set of variables was available.

11.2 The Enable Language Process

After the platform has been determined, the next step is to enable all languages specified in the `project` command. For each language specified CMake loads `CMakeDetermine(LANG)Compiler.cmake` where `LANG` is the name of the language specified in the `project` command. For example with project (`f` Fortran) the file is called `CMakeDetermineFortranCompiler.cmake`. This file discovers the compiler and tools that will be used to compile files for the particular language. Starting with version 2.6.0 CMake tries to identify the compiler for C, C++ and Fortran not only by its filename, but by compiling some source code, which is named `CMake(LANG)CompilerId.(LANG_SUFFIX)`. If this succeeds, it will return a unique id for every compiler supported by CMake. Once the compiler has been determined for a language, CMake configures the file `CMake(LANG)Compiler.cmake.in` into `CMake(LANG)Compiler.cmake`.

After the platform and compiler tools have been determined, CMake loads `CMakeSystemSpecificInformation.cmake` which in turn will load `$(CMAKE_SYSTEM_NAME).cmake` from the platform subdirectory of modules if it exists for the platform. An example would be `SunOS.cmake`. This file contains OS specific information about compiler flags, creation of executables, libraries, and object files.

Next, CMake loads `CMake(LANG)Information.cmake` for each `LANG` that was enabled. This file in turn loads two files; `$(CMAKE_SYSTEM_NAME)-$(COMPILER_ID)-LANG-`

`$(CMAKE_SYSTEM_PROCESSOR).cmake` if it exists, and after that `$(CMAKE_SYSTEM_NAME)-${COMPILER_ID}-LANG.cmake`. In these file names `COMPILER_ID` references the compiler identification determined as described above. The `CMake(LANG)Information.cmake` file contains default rules for creating executables, libraries, and object files on most UNIX systems. The defaults can be overridden by setting values in either `$(CMAKE_SYSTEM_NAME).cmake` or `$(CMAKE_SYSTEM_NAME)-${COMPILER_ID}-LANG.cmake`.

`$(CMAKE_SYSTEM_NAME)-${COMPILER_ID}-LANG-`

`$(CMAKE_SYSTEM_PROCESSOR).cmake` is intended to be used only for cross compiling, it is loaded before `$(CMAKE_SYSTEM_NAME)-${COMPILER_ID}-LANG.cmake`, so variables can be set up which can then be used in the rule variables.

In addition to the files with the `COMPILER_ID` in their name, CMake also supports these files using the `COMPILER_BASE_NAME`. `COMPILER_BASE_NAME` is the name of the compiler with no path information. For example `cl` would be the `COMPILER_BASE_NAME` for the Microsoft Windows compiler, and `Windows-cl.cmake` would be loaded. If a `COMPILER_ID` exists, it will be preferred over the `COMPILER_BASE_NAME`, since on one side the same compiler can have different names, but there can be also different compilers all with the same name. This means, if

```
$(CMAKE_SYSTEM_NAME)-${COMPILER_ID}-LANG-
$(CMAKE_SYSTEM_PROCESSOR).cmake
```

was not found, CMake tries

```
$(CMAKE_SYSTEM_NAME)-${COMPILER_BASE_NAME}.cmake
```

and if

```
$(CMAKE_SYSTEM_NAME)-${COMPILER_ID}-LANG.cmake
```

was not found, CMake tries

```
$(CMAKE_SYSTEM_NAME)-${COMPILER_BASE_NAME}.cmake.
```

`CMake(LANG)Information.cmake` and associated Platform files define special CMake variables, called rule variables. A rule variable consists of a list of commands separated by spaces. The commands are enclosed by quotes. In addition to the normal variable expansion performed by CMake, some special tag variables are expanded by the Makefile generator. Tag

variables have the syntax of <NAME> where NAME is the name of the variable. An example rule variable is `CMAKE_CXX_CREATE_SHARED_LIBRARY`, and the default setting is

```
set (CMAKE_CXX_CREATE_SHARED_LIBRARY
    "<CMAKE_CXX_COMPILER> <CMAKE_SHARED_LIBRARY_CXX_FLAGS>
     <LINK_FLAGS> <CMAKE_SHARED_LIBRARY_CREATE_CXX_FLAGS>
     <CMAKE_SHARED_LIBRARY SONAME_CXX_FLAG><TARGET SONAME> -o
     <TARGET> <OBJECTS> <LINK_LIBRARIES>")
```

At this point, CMake has determined the system it is running on, the tools it will be using to compile the enabled languages, and the rules to use the tools. This means there is enough information for CMake to perform a try-compile. CMake will now test the detected compilers for each enabled language by loading `CMakeTest(LANG)Compiler.cmake`. This file will usually run a try-compile on a simple source file for the given language to make sure the chosen compiler actually works.

Once the platform has been determined, and the compilers have been tested, CMake loads a few more files that can be used to change some of the computed values. The first file that is loaded is `CMake(PROJECTNAME)Compatibility.cmake` where `PROJECTNAME` is the name given to the top level `PROJECT` command in the project. The project compatibility file is used to add backwards compatibility fixes into CMake. For example, if a new version of CMake fails to build a project that the previous version of CMake could build, then fixes can be added on a per project basis to CMake. The last file that is loaded is `$(CMAKE_USER_MAKE_RULES_OVERRIDE)`. This file is an optionally user supplied variable, that can allow a project to make very specific platform based changes to the build rules.

11.3 Porting to a New Platform

Many common platforms are already supported by CMake. However, you may come across a compiler or platform that has not yet been used. If the compiler uses an Integrated Development Environment (IDE), then you will have to extend CMake from the C++ level. However, if the compiler supports a standard make program, then you can specify in CMake the rules to use to compile object code and build libraries by creating CMake configuration files. These files are written using the CMake language with a few special tags that are expanded when the Makefiles are created by CMake. If you run CMake on your system and get a message like the following, you will want to read how to create platform specific settings.

```
System is unknown to cmake, create:
Modules/Platform/MySystem.cmake
to use this system, please send your config file to
cmake@www.cmake.org so it can be added to cmake
```

At a minimum you will need to create the `Platform/${CMAKE_SYSTEM_NAME}.cmake` file for the new platform. Depending on the tools for the platform, you may also want to create `Platform/${CMAKE_SYSTEM_NAME}-${COMPILER_BASE_NAME}.cmake`. On most systems, there is a vendor compiler and the GNU compiler. The rules for both of these compilers can be put in `Platform/${CMAKE_SYSTEM_NAME}.cmake` instead of creating separate files for each of the compilers. For most new systems or compilers, if they follow the basic UNIX compiler flags, you will only need to specify the system specific flags for shared library and module creation.

The following example is from `Platform/IRIX.cmake`. This file specifies several flags, and also one CMake rule variable. The rule variable tells CMake how to use the IRIX CC compiler to create a static library, which is required for template instantiation to work with IRIX CC.

```
# there is no -ldl required on this system
set (CMAKE_DL_LIBS "")

# Specify the flag to create a shared c library
set (CMAKE_SHARED_LIBRARY_CREATE_C_FLAGS
    "-shared -rdata_shared")

# Specify the flag to create a shared c++ library
set (CMAKE_SHARED_LIBRARY_CREATE_CXX_FLAGS
    "-shared -rdata_shared")

# specify the flag to specify run time paths for shared
# libraries -rpath
set (CMAKE_SHARED_LIBRARY_RUNTIME_C_FLAG "-Wl,-rpath,")

# specify a separator for paths on the -rpath, if empty
# then -rpath will be repeated.
set (CMAKE_SHARED_LIBRARY_RUNTIME_C_FLAG_SEP "")

# if the compiler is not GNU, then specify the initial flags
if (NOT CMAKE_COMPILER_IS_GNUCXX)
    # use the CC compiler to create static library
    set (CMAKE_CXX_CREATE_STATIC_LIBRARY
        "<CMAKE_CXX_COMPILER> -ar -o <TARGET> <OBJECTS>")

    # initializes flags for the native compiler
    set (CMAKE_CXX_FLAGS_INIT "")
    set (CMAKE_CXX_FLAGS_DEBUG_INIT "-g")
    set (CMAKE_CXX_FLAGS_MINSIZEREL_INIT "-O3 -DNDEBUG")
    set (CMAKE_CXX_FLAGS_RELEASE_INIT "-O2 -DNDEBUG")
```

```
set (CMAKE_CXX_FLAGS_RELWITHDEBINFO_INIT "-O2")
endif (NOT CMAKE_COMPILER_IS_GNUCXX)
```

11.4 Adding a New Language

In addition to porting CMake to new platforms a user may want to add a new language. This can be done either through the use of custom commands, or by defining a new language for CMake. Once a new language is defined, the standard `add_library` and `add_executable` commands can be used to create libraries and executables for the new language. To add a new language, you need to create four files. The name `LANG` has to match in exact case the name used in the `PROJECT` command to enable the language. For example Fortran has the file `CMakeDetermineFortranCompiler.cmake`, and it is enabled with a call like this `project(f Fortran)`. The four files are as follows:

CMakeDetermine(LANG)Compiler.cmake

This file will find the path to the compiler for `LANG` and then configure `CMake(LANG)Compiler.cmake.in`.

CMake(LANG)Compiler.cmake.in

This file should be used as input to a `configure` file call in the `CMakeDetermine(LANG)Compiler.cmake` file. It is used to store compiler information and is copied down into try-compile directories so that try compiles do not need to re-determine and test the `LANG`.

CMakeTest(LANG)Compiler.cmake

This should use a `try compile` command to make sure the compiler and tools are working. If the tools are working, the following variable should be set in this way:

```
set (CMAKE_(LANG)_COMPILER_WORKS 1 CACHE INTERNAL "")
```

CMake(LANG)Information.cmake

Set values for the following rule variables for `LANG`:

```
CMAKE_(LANG)_CREATE_SHARED_LIBRARY
CMAKE_(LANG)_CREATE_SHARED_MODULE
CMAKE_(LANG)_CREATE_STATIC_LIBRARY
CMAKE_(LANG)_COMPILE_OBJECT
CMAKE_(LANG)_LINK_EXECUTABLE
```

11.5 Rule Variable Listing

For each language that CMake supports, the following rule variables are expanded into build Makefiles at generation time. `LANG` is the name used in the `PROJECT` (name `LANG`) command. CMake currently supports CXX, C, Fortran, and Java as values for `LANG`.

General Tag Variables

The following set of variables will be expanded by CMake.

<TARGET>

The name of the target being built (this may be a full path).

<TARGET_QUOTED>

The name of the target being built (this may be a full path) double quoted.

<TARGET_BASE>

This is replaced by the name of the target without a suffix.

<TARGET SONAME>

This is replaced by

`CMAKE_SHARED_LIBRARY SONAME_(LANG)_FLAG`

<OBJECTS>

This is the list of object files to be linked into the target.

<OBJECTS_QUOTED>

This is the list of object files to be linked into the target double quoted.

<OBJECT>

This is the name of the object file to be built.

<LINK_LIBRARIES>

This is the list of libraries that are linked into an executable or shared object.

<FLAGS>

This is the command line flags for the linker or compiler.

<LINK_FLAGS>

This is the flags used at link time.

<SOURCE>

The source file name.

Language Specific Information

The following set of variables related to the compiler tools will also be expanded.

<CMAKE_(LANG)_COMPILER>

This is the (LANG) compiler command.

<CMAKE_SHARED_LIBRARY_CREATE_(LANG)_FLAGS>

This is the flags used to create a shared library for (LANG) code.

<CMAKE_SHARED_MODULE_CREATE_(LANG)_FLAGS>

This is the flags used to create a shared module for (LANG) code.

<CMAKE_(LANG)_LINK_FLAGS>

This is the flags used to link a (LANG) program.

<CMAKE_AR>

This is the command to create a .a archive file.

<CMAKE_RANLIB>

This is the command to ranlib a .a archive file.

11.6 Compiler and Platform Examples

Como Compiler

A good example to look at is the como compiler on Linux found in Modules/Platforms/Linux-como.cmake. This compiler requires several non-standard commands when creating libraries and executables in order to instantiate C++ templates.

```
# create a shared C++ library

set (CMAKE_CXX_CREATE_SHARED_LIBRARY
    "<CMAKE_CXX_COMPILER> --prelink_objects <OBJECTS>"
    "<CMAKE_CXX_COMPILER>
<CMAKE_SHARED_LIBRARY_CREATE_CXX_FLAGS> <LINK_FLAGS> -o <TARGET>
<OBJECTS> <LINK_LIBRARIES>")

# create a C++ static library
```

```
set (CMAKE_CXX_CREATE_STATIC_LIBRARY
    "<CMAKE_CXX_COMPILER> --prelink_objects <OBJECTS>"
    "<CMAKE_AR> cr <TARGET> <LINK_FLAGS> <OBJECTS> "
    "<CMAKE_RANLIB> <TARGET> ")

set (CMAKE_CXX_LINK_EXECUTABLE
    "<CMAKE_CXX_COMPILER> --prelink_objects <OBJECTS>"
    "<CMAKE_CXX_COMPILER> <CMAKE_CXX_LINK_FLAGS> <LINK_FLAGS>
<FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>")

set (CMAKE_SHARED_LIBRARY_RUNTIME_FLAG "")
set (CMAKE_SHARED_LIBRARY_C_FLAGS "")
set (CMAKE_SHARED_LIBRARY_LINK_FLAGS "")
```

This overrides the creation of libraries (shared and static), and the linking of executable C++ programs. You can see that the linking process of executables and shared libraries requires an extra command that calls the compiler with the flag `--prelink_objects`, and gets all of the object files passed to it.

Borland Compiler

The full Borland compiler rules can be found in `Platforms/Windows-bcc32.cmake`. The following code is an excerpt from that file, showing some of the features used to define rules for the Borland compiler set.

```
set (CMAKE_CXX_CREATE_SHARED_LIBRARY
    "<CMAKE_CXX_COMPILER> ${CMAKE_START_TEMP_FILE}-e<TARGET>
-tWD <LINK_FLAGS> -tWR <LINK_LIBRARIES>
<OBJECTS>${CMAKE_END_TEMP_FILE}"
    "implib -c -w <TARGET_BASE>.lib <TARGET_BASE>.dll"
)

set (CMAKE_CXX_CREATE_SHARED_MODULE
    ${CMAKE_CXX_CREATE_SHARED_LIBRARY})

# create a C shared library
set (CMAKE_C_CREATE_SHARED_LIBRARY
    "<CMAKE_C_COMPILER> ${CMAKE_START_TEMP_FILE}-e<TARGET> -tWD
<LINK_FLAGS> -tWR <LINK_LIBRARIES>
<OBJECTS>${CMAKE_END_TEMP_FILE}"
    "implib -c -w <TARGET_BASE>.lib <TARGET_BASE>.dll"
)

# create a C++ static library
```

```
set (CMAKE_CXX_CREATE_STATIC_LIBRARY "tlib"
${CMAKE_START_TEMP_FILE}/p512 <LINK_FLAGS> /a <TARGET_QUOTED>
<OBJECTS_QUOTED>${CMAKE_END_TEMP_FILE}")

# compile a C++ file into an object file
set (CMAKE_CXX_COMPILE_OBJECT
    "<CMAKE_CXX_COMPILER> ${CMAKE_START_TEMP_FILE}-DWIN32 -P
<FLAGS> -o<OBJECT> -c <SOURCE>${CMAKE_END_TEMP_FILE}")
```

11.7 Extending CMake

Occasionally you will come across a situation where you want to do something during your build process that CMake cannot seem to handle. Examples of this include creating wrappers for C++ classes to make them available to other languages, or creating bindings for C++ classes to support runtime introspection. In these cases you may want to extend CMake by adding your own commands. CMake supports this capability through its C plugin API. Using this API a project can extend CMake to add specialized commands to handle project specific tasks.

A loaded command in CMake is essentially a C code plugin that is compiled into a shared library (a.k.a. DLL). This shared library can then be loaded into the running CMake to provide the functionality of the loaded command. Creating a loaded command is a two step process. You must first write the C code and CMakeLists file for the command, and then place it in your source tree. Next you must modify your project's CMakeLists file to compile the loaded command and load it. We will start by looking at writing the plugin. Before resorting to creating a loaded command you should first see if you can accomplish what you want with a macro. With the commands in CMake a macro/function has almost the same level of flexibility as a loaded command, but does not require compilation or as much complexity. You can almost always, and should, use a macro/function instead of a loaded command.

Creating a Loaded Command

While CMake itself is written in C++ we suggest that you write your plugins using only C code. This avoids a number of portability and compiler issues that can plague C++ plugins being loaded into CMake executables. The API for a plugin is defined in the header file `cmCPluginAPI.h`. This file defines all of the CMake functions that you can invoke from your plugin. It also defines the `cmLoadedCommandInfo` structure that is passed to a plugin. Before going into detail about these functions, consider the following simple plugin:

```
#include "cmCPluginAPI.h"

static int InitialPass(void *inf, void *mf,
                      int argc, char *argv[])
```

```
{  
    cmLoadedCommandInfo *info = (cmLoadedCommandInfo *)inf;  
    info->CAPI->AddDefinition(mf, "FOO", "BAR");  
  
    return 1;  
}  
void CM_PLUGIN_EXPORT  
HELLO_WORLDInit(cmLoadedCommandInfo *info)  
{  
    info->InitialPass = InitialPass;  
    info->Name = "HELLO_WORLD";  
}
```

First this plugin includes the `cmCPluginAPI.h` file to get the definitions and structures required for a plugin. Next it defines a static function called `InitialPass` that will be called whenever this loaded command is invoked. This function is always passed four parameters: the `cmLoadedCommandInfo` structure, the Makefile, the number of arguments, and the list of arguments. Inside this function we typecast the `inf` argument to its actual type and then use it to invoke the C API (CAPI) `AddDefinition` function. This function will set the variable `FOO` to the value of `BAR` in the current `cmMakefile` instance.

The second function is called `HELLO_WORLDInit` and it will be called when the plugin is loaded. The name of this function must exactly match the name of the loaded command with `Init` appended. In this example the name of the command is `HELLO_WORLD` so the function is named `HELLO_WORLDInit`. This function will be called as soon as your command is loaded. It is responsible for initializing the elements of the `cmLoadedCommandInfo` structure. In this example it sets the `InitialPass` member to the address of the `InitialPass` function defined above. It will then set the name of the command by setting the `Name` member to `"HELLO_WORLD"`.

Using a Loaded Command

Now let us consider how to use this new `HELLO_WORLD` command in a project. The basic process is that CMake will have to compile the plugin into a shared library and then dynamically load it. To do this you first create a subdirectory in your project's source tree called CMake or CMakeCommands (by convention, any name can be used). Place the source code to your plugin in that directory. We recommend naming the file with the prefix `cm` and then the name of the command. For example, `cmHELLO_WORLD.c`. Then you must create a simple `CMakeLists.txt` file for this directory that includes instructions to build the shared library. Typically this will be the following:

```
project (HELLO_WORLD)  
  
set (CMAKE_CXX_FLAGS "${CMAKE_CXX_FLAGS}"
```

```
"${CMAKE_ANSI_CXXFLAGS}"
)

set (CMAKE_C_FLAGS "${CMAKE_C_FLAGS}"
"${CMAKE_ANSI_CFLAGS}"
)
include_directories (${CMAKE_ROOT}/include
${CMAKE_ROOT}/Source
)

add_library (cmHELLO_WORLD MODULE cmHELLO_WORLD.c)
```

It is critical that you name the library `cm` followed by the name of the command as shown in the `add_library` call in the above example (e.g. `cmHELLO_WORLD`). When CMake loads a command it assumes that the command is in a library named using that pattern. The next step is to modify your project's main CMakeLists file to compile and load the plugin. This can be accomplished with the following code:

```
# if the command has not been loaded, compile and load it
if (NOT COMMAND HELLO_WORLD)

# try compiling it first
try_compile (COMPILE_OK
${PROJECT_BINARY_DIR}/CMake
${PROJECT_SOURCE_DIR}/CMake
HELLO_WORLD
)

# if it compiled OK then load it
if (COMPILE_OK)
load_command (HELLO_WORLD
${PROJECT_BINARY_DIR}/CMake
${PROJECT_BINARY_DIR}/CMake/Debug
)

# if it did not compile OK, then display an error
else (COMPILE_OK)
message ("error compiling HELLO_WORLD extension")
endif (COMPILE_OK)

endif (NOT COMMAND HELLO_WORLD)
```

In the above example you would simply replace `HELLO_WORLD` with the name of your command and replace `${PROJECT_SOURCE_DIR}/CMake` with the actual name of the subdirectory where you placed your loaded command. Now let us look at creating loaded commands in more detail. We will start by looking at the `cmLoadedCommandInfo` structure.

```
typedef const char* (*CM_DOC_FUNCTION) () ;

typedef int (*CM_INITIAL_PASS_FUNCTION) (
    void *info, void *mf, int argc, char *[]);

typedef void (*CM_FINAL_PASS_FUNCTION) (
    void *info, void *mf);

typedef void (*CM_DESTRUCTOR_FUNCTION) (void *info);

typedef struct {
    unsigned long reserved1;
    unsigned long reserved2;
    cmCAPI *CAPI;
    int m_Inherited;
    CM_INITIAL_PASS_FUNCTION InitialPass;
    CM_FINAL_PASS_FUNCTION FinalPass;
    CM_DESTRUCTOR_FUNCTION Destructor;
    CM_DOC_FUNCTION GetTerseDocumentation;
    CM_DOC_FUNCTION GetFullDocumentation;
    const char *Name;
    char *Error;
    void *ClientData;
} cmLoadedCommandInfo;
```

The first two entries of the structure are reserved for future use. The next entry, `CAPI`, is a pointer to a structure containing pointers to all the CMake functions you can invoke from a plugin. The `m_Inherited` member only applies to CMake versions 2.0 and earlier. It can be set to indicate if this command should be inherited by subdirectories or not. If you are creating a command that will work with versions of CMake prior to 2.2 then you probably want to set this to zero. The next five members are pointers to functions that your plugin may provide. The `InitialPass` function must be provided and it is invoked whenever your loaded command is invoked from a CMakeLists file. The `FinalPass` function is optional and is invoked after configuration but before generation of the output. The `Destructor` function is optional and will be invoked when your command is destroyed by CMake (typically on exit). It can be used to clean up any memory that you have allocated in the `InitialPass` or `FinalPass`. The next two functions are optional and are used to provide documentation for your command. The `Name` member is used to store the name of your command. This is what will be compared against when parsing a CMakeLists file, it should be in all caps in keeping

with CMake's naming conventions. The `Error` and `ClientData` members are used internally by CMake, you should not directly access them. Instead you can use the CAPI functions to manipulate them.

Now let us consider some of the common CAPI functions you will use from within a loaded command. First we will consider some utility functions that are provided specifically for loaded commands. Since loaded commands use a C interface they will receive arguments as `(int argc, char *argv[])`, for convenience you can call `GetTotalArgumentSize(argc, argv)` which will return the total length of all the arguments. Likewise some CAPI methods will return an `(argc, argv)` pair that you will be responsible for freeing. The `FreeArguments(argc, argv)` function can be used to free such return values. If your loaded command has a `FinalPass()` then you might want to pass data from the `InitialPass()` to the `FinalPass()` invocation. This can be accomplished using the `SetClientData(void *info, void *data)` and `void *GetClientData(void *info)` functions. Since the client data is passed as a `void *` argument, any client data larger than a pointer must be allocated and then finally freed in your `Destructor()` function. Be aware that CMake will create multiple instances of your loaded command so using global variables or static variables is not recommended. If you should encounter an error in executing your loaded command, you can call `SetError(void *info, const char *errorMessage)` to pass an error message on to the user.

Another group of CAPI functions worth noting are the `cmSourceFile` functions. `cmSourceFile` is a C++ object that represents information about a single file including its full path, file extension, special compiler flags, etc. Some loaded commands will need to either create or access `cmSourceFile` instances. This can be done using the `void *CreateSourceFile()` and `void * GetSource (void *mf, const char *sourceName)` functions. Both of these functions return a pointer to a `cmSourceFile` as a `void *` return value. This pointer can then be passed into other functions that manipulate `cmSourceFiles` such as `SourceFileGetProperty()` or `SourceFileSetProperty()`.

Tutorials

This chapter provides a step by step tutorial that covers common build system issues that CMake helps address. Many of these topics have been introduced in prior chapters as separate issues but seeing how they all work together in an example project can be very helpful. This tutorial can be found in the Tests/Tutorial directory of the CMake source code tree. Each step has its own subdirectory containing a complete copy of the tutorial for that step.

12.1 A Basic Starting Point (Step 1)

The most basic project is an executable built from source code files. For simple projects a two line CMakeLists file is all that is required. This will be the starting point for our tutorial. The CMakeLists file looks like:

```
cmake_minimum_required (2.6)
project (Tutorial)

add_executable(Tutorial tutorial.cxx)
```

Note that this example uses lower case commands in the CMakeLists file. Upper, lower, and mixed case commands are supported by CMake. The source code for `tutorial.cxx` will compute the square root of a number and the first version of it is very simple, as follows:

```
// A simple program that computes the square root of a number
#include <stdio.h>
#include <math.h>
int main (int argc, char *argv[])
{
    if (argc < 2)
    {
        fprintf(stdout,"Usage: %s number\n", argv[0]);
        return 1;
    }
    double inputValue = atof(argv[1]);
    double outputValue = sqrt(inputValue);
    fprintf(stdout,"The square root of %g is %g\n",
            inputValue, outputValue);
    return 0;
}
```

Adding a Version Number and Configured Header File

The first feature we will add is to provide our executable and project with a version number. While you can do this exclusively in the source code, doing it in the CMakeLists file provides more flexibility. To add a version number we modify the CMakeLists file as follows:

```
cmake_minimum_required (2.6)
project (Tutorial)

# The version number.
set (Tutorial_VERSION_MAJOR 1)
set (Tutorial_VERSION_MINOR 0)

# configure a header file to pass some of the CMake settings
# to the source code
configure_file (
    "${PROJECT_SOURCE_DIR}/TutorialConfig.h.in"
    "${PROJECT_BINARY_DIR}/TutorialConfig.h"
)

# add the binary tree to the search path for include files
# so that we will find TutorialConfig.h
include_directories("${PROJECT_BINARY_DIR}")

# add the executable
add_executable(Tutorial tutorial.cxx)
```

Since the configured file will be written into the binary tree we must add that directory to the list of paths to search for include files. We then create a `TutorialConfig.h.in` file in the source tree with the following contents:

```
// the configured options and settings for Tutorial
#define Tutorial_VERSION_MAJOR @Tutorial_VERSION_MAJOR@
#define Tutorial_VERSION_MINOR @Tutorial_VERSION_MINOR@
```

When CMake configures this header file the values for `@Tutorial_VERSION_MAJOR@` and `@Tutorial_VERSION_MINOR@` will be replaced by the values from the `CMakeLists` file. Next we modify `tutorial.cxx` to include the configured header file and to make use of the version numbers. The resulting source code is listed below.

```
// A simple program that computes the square root of a number
#include <stdio.h>
#include <math.h>
#include "TutorialConfig.h"

int main (int argc, char *argv[])
{
    if (argc < 2)
    {
        fprintf(stdout,"%s Version %d.%d\n",
                argv[0],
                Tutorial_VERSION_MAJOR,
                Tutorial_VERSION_MINOR);
        fprintf(stdout,"Usage: %s number\n", argv[0]);
        return 1;
    }
    double inputValue = atof(argv[1]);
    double outputValue = sqrt(inputValue);
    printf(stdout,"The square root of %g is %g\n",
           inputValue, outputValue);
    return 0;
}
```

The main changes are the inclusion of the `TutorialConfig.h` header file and printing out a version number as part of the usage message.

12.2 Adding a Library (Step 2)

Now we will add a library to our project. This library will contain our own implementation for computing the square root of a number. The executable can then use this library instead of the standard square root function provided by the compiler. For this tutorial we will put the library into a subdirectory called MathFunctions. It will have the following one line CMakeLists file:

```
add_library(MathFunctions mysqrt.cxx)
```

The source file mysqrt.cxx has one function called mysqrt that provides similar functionality to the compiler's sqrt function. To make use of the new library we add an add_subdirectory call in the top level CMakeLists file so that the library will get built. We also add another include directory so that the MathFunctions/mysqrt.h header file can be found for the function prototype. The last change is to add the new library to the executable. The last few lines of the top level CMakeLists file now look like:

```
include_directories ("${PROJECT_SOURCE_DIR}/MathFunctions")
add_subdirectory (MathFunctions)

# add the executable
add_executable (Tutorial tutorial.cxx)
target_link_libraries (Tutorial MathFunctions)
```

Now let us consider making the MathFunctions library optional. In this tutorial there really isn't any reason to do so, but with larger libraries or libraries that rely on third party code you might want to. The first step is to add an option to the top level CMakeLists file.

```
# should we use our own math functions?
option (USE_MYMATH
        "Use tutorial provided math implementation" ON)
```

This will show up in the CMake GUI with a default value of ON that the user can change as desired. This setting will be stored in the cache so that the user does not need to keep setting it each time they run CMake on this project. The next change is to make the build and linking of the MathFunctions library conditional. To do this we change the end of the top level CMakeLists file to look like the following:

```
# add the MathFunctions library?  
#  
if (USE_MYMATH)  
    include_directories ("${PROJECT_SOURCE_DIR}/MathFunctions")  
    add_subdirectory (MathFunctions)  
    set (EXTRA_LIBS ${EXTRA_LIBS} MathFunctions)  
endif (USE_MYMATH)  
  
# add the executable  
add_executable (Tutorial tutorial.cxx)  
target_link_libraries (Tutorial ${EXTRA_LIBS})
```

This uses the setting of `USE_MYMATH` to determine if the `MathFunctions` should be compiled and used. Note the use of a variable (`EXTRA_LIBS` in this case) to collect up any optional libraries to later be linked into the executable. This is a common approach used to keep larger projects with many optional components clean. The corresponding changes to the source code are fairly straight forward and leave us with:

```
// A simple program that computes the square root of a number  
#include <stdio.h>  
#include <math.h>  
#include "TutorialConfig.h"  
  
#ifdef USE_MYMATH  
#include "MathFunctions.h"  
#endif  
  
int main (int argc, char *argv[]){  
    if (argc < 2)  
    {  
        fprintf(stdout,"%s Version %d.%d\n", argv[0],  
                Tutorial_VERSION_MAJOR,  
                Tutorial_VERSION_MINOR);  
        fprintf(stdout,"Usage: %s number\n", argv[0]);  
        return 1;  
    }  
  
    double inputValue = atof(argv[1]);  
  
    #ifdef USE_MYMATH  
        double outputValue = mysqrt(inputValue);  
    #else
```

```

    double outputValue = sqrt(inputValue);
#endif

    fprintf(stdout,"The square root of %g is %g\n",
            inputValue, outputValue);
    return 0;
}

```

In the source code we make use of `USE_MYMATH` as well. This is provided from CMake to the source code through the `TutorialConfig.h.in` configured file by adding the following line to it:

```
#cmakedefine USE_MYMATH
```

12.3 Installing and Testing (Step 3)

For the next step we will add install rules and testing support to our project. The install rules are fairly straight forward. For the `MathFunctions` library we setup the library and the header file to be installed by adding the following two lines to `MathFunctions'` `CMakeLists` file:

```

install (TARGETS MathFunctions DESTINATION bin)
install (FILES MathFunctions.h DESTINATION include)

```

For the application the following lines are added to the top level `CMakeLists` file to install the executable and the configured header file:

```

# add the install targets
install (TARGETS Tutorial DESTINATION bin)
install (FILES "${PROJECT_BINARY_DIR}/TutorialConfig.h"
        DESTINATION include)

```

That is all there is to it. At this point you should be able to build the tutorial, then type `make install` (or build the `INSTALL` target from an IDE) and it will install the appropriate header files, libraries, and executables. The CMake variable `CMAKE_INSTALL_PREFIX` is used to determine the root of where the files will be installed. Adding testing is also a fairly straight forward process. At the end of the top level `CMakeLists` file we can add a number of basic tests to verify that the application is working correctly.

```
# does the application run
add_test (TutorialRuns Tutorial 25)

# does it sqrt of 25
add_test (TutorialComp25 Tutorial 25)

set_tests_properties (TutorialComp25
    PROPERTIES PASS_REGULAR_EXPRESSION "25 is 5")

# does it handle negative numbers
add_test (TutorialNegative Tutorial -25)
set_tests_properties (TutorialNegative
    PROPERTIES PASS_REGULAR_EXPRESSION "-25 is 0")

# does it handle small numbers
add_test (TutorialSmall Tutorial 0.0001)
set_tests_properties (TutorialSmall
    PROPERTIES PASS_REGULAR_EXPRESSION "0.0001 is 0.01")

# does the usage message work?
add_test (TutorialUsage Tutorial)
set_tests_properties (TutorialUsage
    PROPERTIES
    PASS_REGULAR_EXPRESSION "Usage:.*number")
```

The first test simply verifies that the application runs, does not segfault or otherwise crash, and has a zero return value. This is the basic form of a CTest test. The next few tests all make use of the `PASS_REGULAR_EXPRESSION` test property to verify that the output of the test contains certain strings. In this case verifying that the computed square root is what it should be and that the usage message is printed when an incorrect number of arguments are provided. If you wanted to add a lot of tests to test different input values you might consider creating a macro like the following:

```
#define a macro to simplify adding tests, then use it
macro (do_test arg result)
    add_test (TutorialComp${arg} Tutorial ${arg})
    set_tests_properties (TutorialComp${arg}
        PROPERTIES PASS_REGULAR_EXPRESSION ${result})
endmacro (do_test)

# do a bunch of result based tests
do_test (25 "25 is 5")
do_test (-25 "-25 is 0")
```

For each invocation of `do_test`, another test is added to the project with a name, input, and results based on the passed arguments.

12.4 Adding System Introspection (Step 4)

Next let us consider adding some code to our project that depends on features the target platform may not have. For this example we will add some code that depends on whether or not the target platform has the `log` and `exp` functions. Of course almost every platform has these functions but for this tutorial assume that they are less common. If the platform has `log` then we will use that to compute the square root in the `mysqrt` function. We first test for the availability of these functions using the `CheckFunctionExists.cmake` macro in the top level CMakeLists file as follows:

```
# does this system provide the log and exp functions?  
include (CheckFunctionExists.cmake)  
check_function_exists (log HAVE_LOG)  
check_function_exists (exp HAVE_EXP)
```

Next we modify the `TutorialConfig.h.in` to define those values if CMake found them on the platform as follows:

```
// does the platform provide exp and log functions?  
#cmakedefine HAVE_LOG  
#cmakedefine HAVE_EXP
```

It is important that the tests for `log` and `exp` are done before the `configure_file` command for `TutorialConfig.h`. The `configure_file` command immediately configures the file using the current settings in CMake. Finally in the `mysqrt` function we can provide an alternate implementation based on `log` and `exp` if they are available on the system using the following code:

```
// if we have both log and exp then use them  
#if defined (HAVE_LOG) && defined (HAVE_EXP)  
    result = exp(log(x)*0.5);  
#else // otherwise use an iterative approach  
    . . .
```

12.5 Adding a Generated File and Generator (Step 5)

In this section we will show how you can add a generated source file into the build process of an application. For this example we will create a table of precomputed square roots as part of the build process, and then compile that table into our application. To accomplish this we first need a program that will generate the table. In the MathFunctions subdirectory a new source file named `MakeTable.cxx` will do just that.

```
// A simple program that builds a sqrt table
#include <stdio.h>
#include <math.h>

int main (int argc, char *argv[])
{
    int i;
    double result;

    // make sure we have enough arguments
    if (argc < 2)
    {
        return 1;
    }

    // open the output file
    FILE *fout = fopen(argv[1],"w");
    if (!fout)
    {
        return 1;
    }

    // create a source file with a table of square roots
    fprintf(fout,"double sqrtTable[] = {\n");
    for (i = 0; i < 10; ++i)
    {
        result = sqrt(static_cast<double>(i));
        fprintf(fout,"%g,\n",result);
    }

    // close the table with a zero
    fprintf(fout,"0};\n");
    fclose(fout);
    return 0;
}
```

Note that the table is produced as valid C++ code and that the name of the file to write the output to is passed in as an argument. The next step is to add the appropriate commands to MathFunctions' CMakeLists file to build the MakeTable executable, and then run it as part of the build process. A few commands are needed to accomplish this, as shown below.

```
# first we add the executable that generates the table
add_executable(MakeTable MakeTable.cxx)

# add the command to generate the source code
add_custom_command (
    OUTPUT ${CMAKE_CURRENT_BINARY_DIR}/Table.h
    COMMAND MakeTable ${CMAKE_CURRENT_BINARY_DIR}/Table.h
    DEPENDS MakeTable
)

# add the binary tree directory to the search path for
# include files
include_directories( ${CMAKE_CURRENT_BINARY_DIR} )

# add the main library
add_library(MathFunctions mysqrt.cxx
${CMAKE_CURRENT_BINARY_DIR}/Table.h )
```

First the executable for MakeTable is added as any other executable would be added. Then we add a custom command that specifies how to produce Table.h by running MakeTable. Next we have to let CMake know that mysqrt.cxx depends on the generated file Table.h. This is done by adding the generated Table.h to the list of sources for the library MathFunctions. We also have to add the current binary directory to the list of include directories so that Table.h can be found and included by mysqrt.cxx.

When this project is built it will first build the MakeTable executable. It will then run MakeTable to produce Table.h. Finally, it will compile mysqrt.cxx which includes Table.h to produce the MathFunctions library.

At this point the top level CMakeLists file with all the features we have added looks like the following:

```
cmake_minimum_required (2.6)
project (Tutorial)

# The version number.
set (Tutorial_VERSION_MAJOR 1)
set (Tutorial_VERSION_MINOR 0)
```

```
# does this system provide the log and exp functions?
include (${CMAKE_ROOT}/Modules/CheckFunctionExists.cmake)

check_function_exists (log HAVE_LOG)
check_function_exists (exp HAVE_EXP)

# should we use our own math functions
option(USE_MYMATH
    "Use tutorial provided math implementation" ON)

# configure a header file to pass some of the CMake settings
# to the source code
configure_file (
    "${PROJECT_SOURCE_DIR}/TutorialConfig.h.in"
    "${PROJECT_BINARY_DIR}/TutorialConfig.h"
)

# add the binary tree to the search path for include files
# so that we will find TutorialConfig.h
include_directories ("${PROJECT_BINARY_DIR}")

# add the MathFunctions library?
if (USE_MYMATH)
    include_directories ("${PROJECT_SOURCE_DIR}/MathFunctions")
    add_subdirectory (MathFunctions)
    set (EXTRA_LIBS ${EXTRA_LIBS} MathFunctions)
endif (USE_MYMATH)

# add the executable
add_executable (Tutorial tutorial.cxx)
target_link_libraries (Tutorial ${EXTRA_LIBS})

# add the install targets
install (TARGETS Tutorial DESTINATION bin)
install (FILES "${PROJECT_BINARY_DIR}/TutorialConfig.h"
        DESTINATION include)

# does the application run
add_test (TutorialRuns Tutorial 25)

# does the usage message work?
add_test (TutorialUsage Tutorial)
set_tests_properties (TutorialUsage
    PROPERTIES
```

```
PASS_REGULAR_EXPRESSION "Usage:.*number"
)

#define a macro to simplify adding tests
macro (do_test arg result)
    add_test (TutorialComp${arg} Tutorial ${arg})
    set_tests_properties (TutorialComp${arg}
        PROPERTIES PASS_REGULAR_EXPRESSION ${result}
    )
endmacro (do_test)

# do a bunch of result based tests
do_test (4 "4 is 2")
do_test (9 "9 is 3")
do_test (5 "5 is 2.236")
do_test (7 "7 is 2.645")
do_test (25 "25 is 5")
do_test (-25 "-25 is 0")
do_test (0.0001 "0.0001 is 0.01")
```

TutorialConfig.h looks like:

```
// the configured options and settings for Tutorial
#define Tutorial_VERSION_MAJOR @Tutorial_VERSION_MAJOR@
#define Tutorial_VERSION_MINOR @Tutorial_VERSION_MINOR@
#define USE_MYMath

// does the platform provide exp and log functions?
#define HAVE_LOG
#define HAVE_EXP
```

And the CMakeLists file for MathFunctions looks like:

```
# first we add the executable that generates the table
add_executable(MakeTable MakeTable.cxx)
# add the command to generate the source code
add_custom_command (
    OUTPUT ${CMAKE_CURRENT_BINARY_DIR}/Table.h
    DEPENDS MakeTable
    COMMAND MakeTable ${CMAKE_CURRENT_BINARY_DIR}/Table.h
)
# add the binary tree directory to the search path
# for include files
include_directories( ${CMAKE_CURRENT_BINARY_DIR} )

# add the main library
add_library(MathFunctions mysqrt.cxx
${CMAKE_CURRENT_BINARY_DIR}/Table.h)

install (TARGETS MathFunctions DESTINATION bin)
install (FILES MathFunctions.h DESTINATION include)
```

12.6 Building an Installer (Step 6)

Next suppose that we want to distribute our project to other people so that they can use it. We want to provide both binary and source distributions on a variety of platforms. This is a little different from the instal we did previously in section 12.3, where we were installing the binaries that we had built from the source code. In this example we will be building installation packages that support binary installations and package management features as found in cygwin, debian, RPMs etc. To accomplish this we will use CPack to create platform specific installers as described in Chapter 9. Specifically we need to add a few lines to the bottom of our toplevel CMakeLists.txt file.

```
# build a CPack driven installer package
include (InstallRequiredSystemLibraries)
set (CPACK_RESOURCE_FILE_LICENSE
    "${CMAKE_CURRENT_SOURCE_DIR}/License.txt")
set (CPACK_PACKAGE_VERSION_MAJOR "${Tutorial_VERSION_MAJOR}")
set (CPACK_PACKAGE_VERSION_MINOR "${Tutorial_VERSION_MINOR}")
include (CPack)
```

That is all there is to it. We start by including InstallRequiredSystemLibraries. This module will include any runtime libraries that are needed by the project for the current platform. Next

we set some CPack variables to where we have stored the license and version information for this project. The version information makes use of the variables we set earlier in this tutorial. Finally we include the CPack module which will use these variables and some other properties of the system you are on to setup an installer.

The next step is to build the project in the usual manner and then run CPack on it. To build a binary distribution you would run:

```
cpack -C CPackConfig.cmake
```

To create a source distribution you would type

```
cpack -C CPackSourceConfig.cmake
```

12.7 Adding Support for a Dashboard (Step 7)

Adding support for submitting our test results to a dashboard is very easy. We already defined a number of tests for our project in the earlier steps of this tutorial. We just have to run those tests and submit them to a dashboard. To include support for dashboards we include the CTest module in our toplevel CMakeLists file.

```
# enable dashboard scripting
include (CTest)
```

We also create a CTestConfig.cmake file where we can specify the name of this project for the dashboard.

```
set (CTEST_PROJECT_NAME "Tutorial")
```

CTest will read in this file when it runs. To create a simple dashboard you can run CMake on your project, change directory to the binary tree, and then run ctest -D Experimental. The results of your dashboard will be uploaded to Kitware's public dashboard at:

```
http://www.cdash.org/CDash/index.php?project=PublicDashboard
```

Appendix A - Variables

Variables That Change Behavior

BUILD_SHARED_LIBS: Global flag to cause add_library to create shared libraries if on.

If present and true, this will cause all libraries to be built shared unless the library was explicitly added as a static library. This variable is often added to projects as an OPTION so that each user of a project can decide if they want to build the project using shared or static libraries.

CMAKE_BACKWARDS_COMPATIBILITY: Version of cmake required to build project

From the point of view of backwards compatibility, this specifies what version of CMake should be supported. By default this value is the version number of CMake that you are running. You can set this to an older version of CMake to support deprecated commands of CMake in projects that were written to use older versions of CMake. This can be set by the user or set at the beginning of a CMakeLists file.

CMAKE_BUILD_TYPE: Specifies the build type for make based generators.

This specifies what build type will be built in this tree. Possible values are empty, Debug, Release, RelWithDebInfo and MinSizeRel. This variable is only supported for make based generators. If this variable is supported, then CMake will also provide initial values for the variables with the name CMAKE_C_FLAGS_[Debug|Release|RelWithDebInfo|MinSizeRel]. For example, if CMAKE_BUILD_TYPE is Debug, then CMAKE_C_FLAGS_DEBUG will be added to the CMAKE_C_FLAGS.

CMAKE_COLOR_MAKEFILE: Enables color output when using the Makefile generator.

When enabled, the generated Makefiles will produce colored output. Default is ON.

CMAKE_CONFIGURATION_TYPES: Specifies the available build types.

This specifies what build types will be available such as Debug, Release, RelWithDebInfo etc. This has reasonable defaults on most platforms. But can be extended to provide other build types. See also CMAKE_BUILD_TYPE.

CMAKE_FIND_LIBRARY_PREFIXES: Prefixes to prepend when looking for libraries.

This specifies what prefixes to add to library names when the find_library command looks for libraries. On UNIX systems this is typically lib, meaning that when trying to find the foo library it will look for libfoo.

CMAKE_FIND_LIBRARY_SUFFIXES: Suffixes to append when looking for libraries.

This specifies what suffixes to add to library names when the find_library command looks for libraries. On Windows systems this is typically .lib and .dll, meaning that when trying to find the foo library it will look for foo.dll etc.

CMAKE_INCLUDE_PATH: Path used for searching by FIND_FILE() and FIND_PATH().

Specifies a path which will be used both by FIND_FILE() and FIND_PATH(). Both commands will check each of the contained directories for the existence of the file which is currently searched. By default it is empty, it is intended to be set by the project. See also CMAKE_SYSTEM_INCLUDE_PATH, CMAKE_PREFIX_PATH.

CMAKE_INSTALL_PREFIX: Install directory used by install.

If "make install" is invoked or INSTALL is built, this directory is pre-pended onto all install directories. This variable defaults to /usr/local on UNIX and c:/Program Files on Windows.

CMAKE_LIBRARY_PATH: Path used for searching by FIND_LIBRARY().

Specifies a path which will be used by FIND_LIBRARY(). FIND_LIBRARY() will check each of the contained directories for the existence of the library which is currently searched. By default it is empty, it is intended to be set by the project. See also CMAKE_SYSTEM_LIBRARY_PATH, CMAKE_PREFIX_PATH.

CMAKE_MFC_FLAG: Tell CMake to use MFC for an executable or dll.

This can be set in a CMakeLists.txt file and will enable MFC in the application. It should be set to 1 for static the static MFC library, and 2 for the shared MFC library. This is used in visual studio 6 and 7 project files. The CMakeSetup dialog used MFC and the CMakeLists.txt looks like this:

```
add_definitions(-D_AFXDLL)
set(CMAKE_MFC_FLAG 2)
add_executable(CMakeSetup WIN32 ${SRCS})
```

CMAKE_MODULE_PATH: Path to look for cmake modules to load.

Specifies a path to override the default search path for CMake modules. For example include commands will look in this path first for modules to include.

CMAKE_NOT_USING_CONFIG_FLAGS: Skip _BUILD_TYPE flags if true.

This is an internal flag used by the generators in CMake to tell CMake to skip the _BUILD_TYPE flags.

CMAKE_PREFIX_PATH: Path used for searching by FIND_XXX(), with appropriate suffixes added.

Specifies a path which will be used by the FIND_XXX() commands. It contains the "base" directories, the FIND_XXX() commands append appropriate subdirectories to the base directories. So FIND_PROGRAM() adds /bin to each of the directories in the path, FIND_LIBRARY() appends /lib to each of the directories, and FIND_PATH() and FIND_FILE() append /include . By default it is empty, it is intended to be set by the project. See also CMAKE_SYSTEM_PREFIX_PATH, CMAKE_INCLUDE_PATH, CMAKE_LIBRARY_PATH, CMAKE_PROGRAM_PATH.

CMAKE_PROGRAM_PATH: Path used for searching by FIND_PROGRAM().

Specifies a path which will be used by FIND_PROGRAM(). FIND_PROGRAM() will check each of the contained directories for the existence of the program which is currently searched. By default it is empty, it is intended to be set by the project. See also CMAKE_SYSTEM_PROGRAM_PATH, CMAKE_PREFIX_PATH.

CMAKE_SKIP_INSTALL_ALL_DEPENDENCY: Don't make the install target depend on the all target.

By default, the "install" target depends on the "all" target. This has the effect, that when "make install" is invoked or INSTALL is built, first the "all" target is built, then the installation starts. If CMAKE_SKIP_INSTALL_ALL_DEPENDENCY is set to TRUE, this dependency is not created, so the installation process will start immediately, independent from whether the project has been completely built or not.

CMAKE_SYSTEM_INCLUDE_PATH: Path used for searching by FIND_FILE() and FIND_PATH().

Specifies a path which will be used both by FIND_FILE() and FIND_PATH(). Both commands will check each of the contained directories for the existence of the file which is currently searched. By default it contains the standard directories for the current system. It is

NOT intended to be modified by the project, use CMAKE_INCLUDE_PATH for this. See also CMAKE_SYSTEM_PREFIX_PATH.

CMAKE_SYSTEM_LIBRARY_PATH: Path used for searching by FIND_LIBRARY().

Specifies a path which will be used by FIND_LIBRARY(). FIND_LIBRARY() will check each of the contained directories for the existence of the library which is currently searched. By default it contains the standard directories for the current system. It is NOT intended to be modified by the project, use CMAKE_SYSTEM_LIBRARY_PATH for this. See also CMAKE_SYSTEM_PREFIX_PATH.

CMAKE_SYSTEM_PREFIX_PATH: Path used for searching by FIND_XXX(), with appropriate suffixes added.

Specifies a path which will be used by the FIND_XXX() commands. It contains the "base" directories, the FIND_XXX() commands append appropriate subdirectories to the base directories. So FIND_PROGRAM() adds /bin to each of the directories in the path, FIND_LIBRARY() appends /lib to each of the directories, and FIND_PATH() and FIND_FILE() append /include . By default this contains the standard directories for the current system. It is NOT intended to be modified by the project, use CMAKE_PREFIX_PATH for this. See also CMAKE_SYSTEM_INCLUDE_PATH, CMAKE_SYSTEM_LIBRARY_PATH, CMAKE_SYSTEM_PROGRAM_PATH.

CMAKE_SYSTEM_PROGRAM_PATH: Path used for searching by FIND_PROGRAM().

Specifies a path which will be used by FIND_PROGRAM(). FIND_PROGRAM() will check each of the contained directories for the existence of the program which is currently searched. By default it contains the standard directories for the current system. It is NOT intended to be modified by the project, use CMAKE_PROGRAM_PATH for this. See also CMAKE_SYSTEM_PREFIX_PATH.

CMAKE_USER_MAKE_RULES_OVERRIDE: Specify a file that can change the build rule variables.

If this variable is set, it should point to a CMakeLists.txt file that will be read in by CMake after all the system settings have been set, but before they have been used. This would allow you to override any variables that need to be changed for some special project.

Variables That Describe the System

APPLE: True if running on Mac OS X.

BORLAND: True if the Borland compiler is being used.

CMAKE_CL_64: Using the 64 bit compiler from Microsoft

Set to true when using the 64 bit cl compiler from Microsoft.

CMAKE_COMPILER_2005: Using the Visual Studio 2005 compiler from Microsoft

Set to true when using the Visual Studio 2005 compiler from Microsoft.

CMAKE_HOST_APPLE: True for Apple OS X operating systems.

CMAKE_HOST_SYSTEM: Name of system CMake is being run on.

The same as CMAKE_SYSTEM but for the host system instead of the target system when cross compiling.

CMAKE_HOST_SYSTEM_NAME: Name of the OS CMake is running on.

The same as CMAKE_SYSTEM_NAME but for the host system instead of the target system when cross compiling.

CMAKE_HOST_SYSTEM_PROCESSOR: The name of the CPU CMake is running on.

The same as CMAKE_SYSTEM_PROCESSOR but for the host system instead of the target system when cross compiling.

CMAKE_HOST_SYSTEM_VERSION: OS version CMake is running on.

The same as CMAKE_SYSTEM_VERSION but for the host system instead of the target system when cross compiling.

CMAKE_HOST_UNIX: True for UNIX and UNIX like operating systems.

Set to true when the host system is UNIX or UNIX like (i.e. APPLE and CYGWIN).

CMAKE_HOST_WIN32: True on windows systems, including win64.

Set to true when the host system is Windows and on cygwin.

CMAKE_OBJECT_PATH_MAX: Maximum object file full-path length allowed by native build tools.

CMake computes for every source file an object file name that is unique to the source file and deterministic with respect to the full path to the source file. This allows multiple source files in a target to share the same name if they lie in different directories without rebuilding when one is added or removed. However, it can produce long full paths in a few cases, so CMake shortens the path using a hashing scheme when the full path to an object file exceeds a limit. CMake has a built-in limit for each platform that is sufficient for common tools, but some native tools may have a lower limit. This variable may be set to specify the limit explicitly. The value must be an integer no less than 128.

CMAKE_SYSTEM: Name of system CMake is compiling for.

This variable is the composite of CMAKE_SYSTEM_NAME and CMAKE_SYSTEM_VERSION, like this \${CMAKE_SYSTEM_NAME}-

`${CMAKE_SYSTEM_VERSION}`. If `CMAKE_SYSTEM_VERSION` is not set, then `CMAKE_SYSTEM` is the same as `CMAKE_SYSTEM_NAME`.

CMAKE_SYSTEM_NAME: Name of the OS CMake is building for.

This is the name of the operating system on which CMake is targeting. On systems that have the `uname` command, this variable is set to the output of `uname -s`. Linux, Windows, and Darwin for Mac OS X are the values found on the big three operating systems.

CMAKE_SYSTEM_PROCESSOR: The name of the CPU CMake is building for.

On systems that support `uname`, this variable is set to the output of `uname -p`, on Windows it is set to the value of the environment variable `PROCESSOR_ARCHITECTURE`

CMAKE_SYSTEM_VERSION: OS version CMake is building for.

A numeric version string for the system, on systems that support `uname`, this variable is set to the output of `uname -r`. On other systems this is set to major-minor version numbers.

CYGWIN: True for cygwin.

MSVC: True when using Microsoft Visual C

MSVC80: True when using Microsoft Visual C 8.0

MSVC_IDE: True when using the Microsoft Visual C IDE

Set to true when the target platform is the Microsoft Visual C IDE, as opposed to the command line compiler.

MSVC_VERSION: The version of Microsoft Visual C/C++ being used if any.

The version of Microsoft Visual C/C++ being used if any. For example 1300 is MSVC 6.0.

UNIX: True for UNIX and UNIX like operating systems.

Set to true when the target system is UNIX or UNIX like (i.e. APPLE and CYGWIN).

WIN32: True on windows systems, including win64.

Set to true when the target system is Windows and on cygwin.

XCODE_VERSION: Version of Xcode (Xcode generator only).

Under the Xcode generator, this is the version of Xcode as specified in "Xcode.app/Contents/version.plist" (such as "3.1.2").

Variables for Languages

CMAKE_<LANG>_ARCHIVE_APPEND: Rule variable to append to a static archive.

This is a rule variable that tells CMake how to append to a static archive. It is used in place of `CMAKE_<LANG>_CREATE_STATIC_LIBRARY` on some platforms in order to support large object counts. See also `CMAKE_<LANG>_ARCHIVE_CREATE` and `CMAKE_<LANG>_ARCHIVE_FINISH`.

`CMAKE_<LANG>_ARCHIVE_CREATE`: Rule variable to create a new static archive.

This is a rule variable that tells CMake how to create a static archive. It is used in place of `CMAKE_<LANG>_CREATE_STATIC_LIBRARY` on some platforms in order to support large object counts. See also `CMAKE_<LANG>_ARCHIVE_APPEND` and `CMAKE_<LANG>_ARCHIVE_FINISH`.

`CMAKE_<LANG>_ARCHIVE_FINISH`: Rule variable to finish an existing static archive.

This is a rule variable that tells CMake how to finish a static archive. It is used in place of `CMAKE_<LANG>_CREATE_STATIC_LIBRARY` on some platforms in order to support large object counts. See also `CMAKE_<LANG>_ARCHIVE_CREATE` and `CMAKE_<LANG>_ARCHIVE_APPEND`.

`CMAKE_<LANG>_COMPILER`: The full path to the compiler for LANG.

This is the command that will be used as the `<LANG>` compiler. Once set, you cannot change this variable.

`CMAKE_<LANG>_COMPILER_ABI`: An internal variable subject to change.

This is used in determining the compiler ABI and is subject to change.

`CMAKE_<LANG>_COMPILER_ID`: An internal variable subject to change.

This is used in determining the compiler and is subject to change.

`CMAKE_<LANG>_COMPILER_LOADED`: Defined to true if the language is enabled.

When language `<LANG>` is enabled by `project()` or `enable_language()` this variable is defined to 1.

`CMAKE_<LANG>_COMPILE_OBJECT`: Rule variable to compile a single object file.

This is a rule variable that tells CMake how to compile a single object file for the language `<LANG>`.

`CMAKE_<LANG>_CREATE_SHARED_LIBRARY`: Rule variable to create a shared library.

This is a rule variable that tells CMake how to create a shared library for the language `<LANG>`.

`CMAKE_<LANG>_CREATE_SHARED_MODULE`: Rule variable to create a shared module.

This is a rule variable that tells CMake how to create a shared library for the language <LANG>.

CMAKE_<LANG>_CREATE_STATIC_LIBRARY: Rule variable to create a static library.

This is a rule variable that tells CMake how to create a static library for the language <LANG>.

CMAKE_<LANG>_FLAGS_DEBUG: Flags for Debug build type or configuration.

<LANG> flags used when CMAKE_BUILD_TYPE is Debug.

CMAKE_<LANG>_FLAGS_MINSIZEREL: Flags for MinSizeRel build type or configuration.

<LANG> flags used when CMAKE_BUILD_TYPE is MinSizeRel. Short for minimum size release.

CMAKE_<LANG>_FLAGS_RELEASE: Flags for Release build type or configuration.

<LANG> flags used when CMAKE_BUILD_TYPE is Release

CMAKE_<LANG>_FLAGS_RELWITHDEBINFO: Flags for RelWithDebInfo type or configuration.

<LANG> flags used when CMAKE_BUILD_TYPE is RelWithDebInfo. Short for Release With Debug Information.

CMAKE_<LANG>_IGNORE_EXTENSIONS: File extensions that should be ignored by the build.

This is a list of file extensions that may be part of a project for a given language but are not compiled.

CMAKE_<LANG>_IMPLICIT_INCLUDE_DIRECTORIES: Directories implicitly searched by the compiler for header files.

CMake does not explicitly specify these directories on compiler command lines for language <LANG>. This prevents system include directories from being treated as user include directories on some compilers.

CMAKE_<LANG>_IMPLICIT_LINK_DIRECTORIES: Implicit linker search path detected for language <LANG>.

Compilers typically pass directories containing language runtime libraries and default library search paths when they invoke a linker. These paths are implicit linker search directories for the compiler's language. CMake automatically detects these directories for each language and reports the results in this variable.

CMAKE_<LANG>_IMPLICIT_LINK_LIBRARIES: Implicit link libraries and flags detected for language <LANG>.

Compilers typically pass language runtime library names and other flags when they invoke a linker. These flags are implicit link options for the compiler's language. CMake automatically detects these libraries and flags for each language and reports the results in this variable.

CMAKE_<LANG>_LINKER_PREFERENCE: Preference value for linker language selection.

The "linker language" for executable, shared library, and module targets is the language whose compiler will invoke the linker. The LINKER_LANGUAGE target property sets the language explicitly. Otherwise, the linker language is that whose linker preference value is highest among languages compiled and linked into the target. See also the CMAKE_<LANG>_LINKER_PREFERENCE_PROPAGATES variable.

CMAKE_<LANG>_LINKER_PREFERENCE_PROPAGATES: True if CMAKE_<LANG>_LINKER_PREFERENCE propagates across targets.

This is used when CMake selects a linker language for a target. Languages compiled directly into the target are always considered. A language compiled into static libraries linked by the target is considered if this variable is true.

CMAKE_<LANG>_LINK_EXECUTABLE : Rule variable to link and executable.

Rule variable to link and executable for the given language.

CMAKE_<LANG>_OUTPUT_EXTENSION: Extension for the output of a compile for a single file.

This is the extension for an object file for the given <LANG>. For example .obj for C on Windows.

CMAKE_<LANG>_PLATFORM_ID: An internal variable subject to change.

This is used in determining the platform and is subject to change.

CMAKE_<LANG>_SIZEOF_DATA_PTR: Size of pointer-to-data types for language <LANG>.

This holds the size (in bytes) of pointer-to-data types in the target platform ABI. It is defined for languages C and CXX (C++).

CMAKE_<LANG>_SOURCE_FILE_EXTENSIONS: Extensions of source files for the given language.

This is the list of extensions for a given languages source files.

CMAKE_COMPILER_IS_GNU<LANG>: True if the compiler is GNU.

If the selected <LANG> compiler is the GNU compiler then this is TRUE, if not it is FALSE.

CMAKE_INTERNAL_PLATFORM_ABI: An internal variable subject to change.

This is used in determining the compiler ABI and is subject to change.

CMAKE_USER_MAKE_RULES_OVERRIDE_<LANG>: Specify a file that can change the build rule variables.

If this variable is set, it should point to a CMakeLists.txt file that will be read in by CMake after all the system settings have been set, but before they have been used. This would allow you to override any variables that need to be changed for some language.

Variables That Control the Build

CMAKE_<CONFIG>_POSTFIX: Default filename postfix for libraries under configuration <CONFIG>.

When a non-executable target is created its <CONFIG>_POSTFIX target property is initialized with the value of this variable if it is set.

CMAKE_ARCHIVE_OUTPUT_DIRECTORY: Where to put all the ARCHIVE targets when built.

This variable is used to initialize the ARCHIVE_OUTPUT_DIRECTORY property on all the targets. See that target property for additional information.

CMAKE_BUILD_WITH_INSTALL_RPATH: Use the install path for the RPATH

Normally CMake uses the build tree for the RPATH when building executables etc on systems that use RPATH. When the software is installed the executables etc are relinked by CMake to have the install RPATH. If this variable is set to true then the software is always built with the install path for the RPATH and does not need to be relinked when installed.

CMAKE_DEBUG_POSTFIX: See variable CMAKE_<CONFIG>_POSTFIX.

This variable is a special case of the more-general CMAKE_<CONFIG>_POSTFIX variable for the DEBUG configuration.

CMAKE_EXE_LINKER_FLAGS: Linker flags used to create executables.

Flags used by the linker when creating an executable.

CMAKE_EXE_LINKER_FLAGS_[CMAKE_BUILD_TYPE]: Flag used when linking an executable.

Same as CMAKE_C_FLAGS_* but used by the linker when creating executables.

CMAKE_Fortran_MODULE_DIRECTORY: Fortran module output directory.

This variable is used to initialize the Fortran_MODULE_DIRECTORY property on all the targets. See that target property for additional information.

CMAKE_INCLUDE_CURRENT_DIR: Automatically add the current source- and build directories to the include path.

If this variable is enabled, CMake automatically adds in each directory `${CMAKE_CURRENT_SOURCE_DIR}` and `${CMAKE_CURRENT_BINARY_DIR}` to the include path for this directory. These additional include directories do not propagate down to subdirectories. This is useful mainly for out-of-source builds, where files generated into the build tree are included by files located in the source tree.

By default `CMAKE_INCLUDE_CURRENT_DIR` is OFF.

`CMAKE_INSTALL_NAME_DIR`: Mac OS X directory name for installed targets.

`CMAKE_INSTALL_NAME_DIR` is used to initialize the `INSTALL_NAME_DIR` property on all targets. See that target property for more information.

`CMAKE_INSTALL_RPATH`: The rpath to use for installed targets.

A semicolon-separated list specifying the rpath to use in installed targets (for platforms that support it). This is used to initialize the target property `INSTALL_RPATH` for all targets.

`CMAKE_INSTALL_RPATH_USE_LINK_PATH`: Add paths to linker search and installed rpath.

`CMAKE_INSTALL_RPATH_USE_LINK_PATH` is a boolean that if set to true will append directories in the linker search path and outside the project to the `INSTALL_RPATH`. This is used to initialize the target property `INSTALL_RPATH_USE_LINK_PATH` for all targets.

`CMAKE_LIBRARY_OUTPUT_DIRECTORY`: Where to put all the LIBRARY targets when built.

This variable is used to initialize the `LIBRARY_OUTPUT_DIRECTORY` property on all the targets. See that target property for additional information.

`CMAKE_LIBRARY_PATH_FLAG`: The flag used to add a library search path to a compiler.

The flag used to specify a library directory to the compiler. On most compilers this is "-L".

`CMAKE_LINK_DEF_FILE_FLAG` : Linker flag used to specify a .def file for dll creation.

The flag used to add a .def file when creating a dll on Windows, this is only defined on Windows.

`CMAKE_LINK_LIBRARY_FILE_FLAG`: Flag used to link a library specified by a path to its file.

The flag used before a library file path is given to the linker. This is needed only on very few platforms.

`CMAKE_LINK_LIBRARY_FLAG`: Flag used to link a library into an executable.

The flag used to specify a library to link to an executable. On most compilers this is "-l".

`CMAKE_NO_BUILTIN_CHRPATH`: Do not use the builtin ELF editor to fix RPATHS on installation.

When an ELF binary needs to have a different RPATH after installation than it does in the build tree, CMake uses a builtin editor to change the RPATH in the installed copy. If this variable is set to true then CMake will relink the binary before installation instead of using its builtin editor.

CMAKE_RUNTIME_OUTPUT_DIRECTORY: Where to put all the RUNTIME targets when built.

This variable is used to initialize the RUNTIME_OUTPUT_DIRECTORY property on all the targets. See that target property for additional information.

CMAKE_SKIP_BUILD_RPATH: Do not include RPATHs in the build tree.

Normally CMake uses the build tree for the RPATH when building executables etc on systems that use RPATH. When the software is installed the executables etc are relinked by CMake to have the install RPATH. If this variable is set to true then the software is always built with no RPATH.

CMAKE_USE_RELATIVE_PATHS: Use relative paths (May not work!).

If this is set to TRUE, then the CMake will use relative paths between the source and binary tree. This option does not work for more complicated projects, and relative paths are used when possible. In general, it is not possible to move CMake generated Makefiles to a different location regardless of the value of this variable.

EXECUTABLE_OUTPUT_PATH: Old executable location variable.

The target property RUNTIME_OUTPUT_DIRECTORY supercedes this variable for a target if it is set. Executable targets are otherwise placed in this directory.

LIBRARY_OUTPUT_PATH: Old library location variable.

The target properties ARCHIVE_OUTPUT_DIRECTORY, LIBRARY_OUTPUT_DIRECTORY, and RUNTIME_OUTPUT_DIRECTORY supercede this variable for a target if they are set. Library targets are otherwise placed in this directory.

Variables That Provide Information

variables defined by CMake, that give information about the project, and CMake

CMAKE_AR: Name of archiving tool for static libraries.

This specifies name of the program that creates archive or static libraries.

CMAKE_BINARY_DIR: The path to the top level of the build tree.

This is the full path to the top level of the current CMake build tree. For an in-source build, this would be the same as CMAKE_SOURCE_DIR.

CMAKE_BUILD_TOOL: Tool used for the actual build process.

This variable is set to the program that will be needed to build the output of CMake. If the generator selected was Visual Studio 6, the CMAKE_MAKE_PROGRAM will be set to msdev, for Unix Makefiles it will be set to make or gmake, and for Visual Studio 7 it set to devenv. For NMake Makefiles the value is nmake. This can be useful for adding special flags and commands based on the final build environment.

CMAKE_CACHEFILE_DIR: The directory with the CMakeCache.txt file.

This is the full path to the directory that has the CMakeCache.txt file in it. This is the same as CMAKE_BINARY_DIR.

CMAKE_CACHE_MAJOR_VERSION: Major version of CMake used to create the CMakeCache.txt file

This stores the major version of CMake used to write a CMake cache file. It is only different when a different version of CMake is run on a previously created cache file.

CMAKE_CACHE_MINOR_VERSION: Minor version of CMake used to create the CMakeCache.txt file

This stores the minor version of CMake used to write a CMake cache file. It is only different when a different version of CMake is run on a previously created cache file.

CMAKE_CACHE_PATCH_VERSION: Patch version of CMake used to create the CMakeCache.txt file

This stores the patch version of CMake used to write a CMake cache file. It is only different when a different version of CMake is run on a previously created cache file.

CMAKE_CFG_INDIR: Build-time reference to per-configuration output subdirectory.

For native build systems supporting multiple configurations in the build tree (such as Visual Studio and Xcode), the value is a reference to a build-time variable specifying the name of the per-configuration output subdirectory. On Makefile generators this evaluates to "." because there is only one configuration in a build tree. Example values:

```
$ (IntDir)      = Visual Studio 6  
$ (OutDir)      = Visual Studio 7, 8, 9  
$ (Configuration) = Visual Studio 10  
$ (CONFIGURATION) = Xcode  
.               = Make-based tools
```

Since these values are evaluated by the native build system, this variable is suitable only for use in command lines that will be evaluated at build time. Example of intended usage:

```

add_executable(mytool mytool.c)
add_custom_command(
    OUTPUT out.txt
    COMMAND
        ${CMAKE_CURRENT_BINARY_DIR}/${CMAKE_CFG_INTDIR}/mytool
        ${CMAKE_CURRENT_SOURCE_DIR}/in.txt out.txt
    DEPENDS mytool in.txt
)
add_custom_target(drive ALL DEPENDS out.txt)

```

Note that CMAKE_CFG_INTDIR is no longer necessary for this purpose but has been left for compatibility with existing projects. Instead add_custom_command() recognizes executable target names in its COMMAND option, so "\${CMAKE_CURRENT_BINARY_DIR}/\${CMAKE_CFG_INTDIR}/mytool" can be replaced by just "mytool".

This variable is read-only. Setting it is undefined behavior. In multi-configuration build systems the value of this variable is passed as the value of preprocessor symbol "CMAKE_INDIR" to the compilation of all source files.

CMAKE_COMMAND: The full path to the cmake executable.

This is the full path to the CMake executable cmake which is useful from custom commands that want to use the cmake -E option for portable system commands, e.g. /usr/local/bin/cmake

CMAKE_CROSSCOMPILING: Is CMake currently cross compiling.

This variable will be set to true by CMake if CMake is cross compiling. Specifically if the build platform is different from the target platform.

CMAKE_CTEST_COMMAND: Full path to ctest command installed with cmake.

This is the full path to the CTest executable ctest which is useful from custom commands that want to use the cmake -E option for portable system commands.

CMAKE_CURRENT_BINARY_DIR: The path to the binary directory currently being processed.

This is the full path to the build directory that is currently being processed by cmake. Each directory added by add_subdirectory will create a binary directory in the build tree, and as it is being processed this variable will be set. For in-source builds this is the current source directory being processed.

CMAKE_CURRENT_LIST_FILE: Full path to the listfile currently being processed.

As CMake processes the listfiles in your project this variable will always be set to the one currently being processed. See also CMAKE_PARENT_LIST_FILE.

CMAKE_CURRENT_LIST_LINE: The line number of the current file being processed.

This is the line number of the file currently being processed by cmake.

CMAKE_CURRENT_SOURCE_DIR: The path to the source directory currently being processed.

This is the full path to the source directory that is currently being processed by cmake.

CMAKE_DL_LIBS: Name of library containing dlopen and dlclose.

The name of the library that has dlopen and dlclose in it, usually -ldl on most UNIX machines.

CMAKE_EDIT_COMMAND: Full path to cmake-gui or ccmake.

This is the full path to the CMake executable that can graphically edit the cache. For example, cmake-gui, ccmake, or cmake -i.

CMAKE_EXECUTABLE_SUFFIX: The suffix for executables on this platform.

The suffix to use for the end of an executable if any, .exe on Windows.

CMAKE_EXECUTABLE_SUFFIX_<LANG> overrides this for language <LANG>.

CMAKE_EXTRA_SHARED_LIBRARY_SUFFIXES: Additional suffixes for shared libraries.

Extensions for shared libraries other than that specified by CMAKE_SHARED_LIBRARY_SUFFIX, if any. CMake uses this to recognize external shared library files during analysis of libraries linked by a target.

CMAKE_GENERATOR: The generator used to build the project.

The name of the generator that is being used to generate the build files. (e.g. "Unix Makefiles", "Visual Studio 6", etc.)

CMAKE_HOME_DIRECTORY: Path to top of source tree.

This is the path to the top level of the source tree.

CMAKE_IMPORT_LIBRARY_PREFIX: The prefix for import libraries that you link to.

The prefix to use for the name of an import library if used on this platform.

CMAKE_IMPORT_LIBRARY_PREFIX_<LANG> overrides this for language <LANG>.

CMAKE_IMPORT_LIBRARY_SUFFIX: The suffix for import libraries that you link to.

The suffix to use for the end of an import library if used on this platform.

CMAKE_IMPORT_LIBRARY_SUFFIX_<LANG> overrides this for language <LANG>.

CMAKE_LINK_LIBRARY_SUFFIX: The suffix for libraries that you link to.

The suffix to use for the end of a library, .lib on Windows.

CMAKE_MAJOR_VERSION: The Major version of cmake (i.e. the 2 in 2.X.X)

This specifies the major version of the CMake executable being run.

CMAKE_MAKE_PROGRAM: See CMAKE_BUILD_TOOL.

This variable is around for backwards compatibility, see CMAKE_BUILD_TOOL.

CMAKE_MINOR_VERSION: The Minor version of cmake (i.e. the 4 in X.4.X).

This specifies the minor version of the CMake executable being run.

CMAKE_PARENT_LIST_FILE: Full path to the parent listfile of the one currently being processed.

As CMake processes the listfiles in your project this variable will always be set to the listfile that included or somehow invoked the one currently being processed. See also CMAKE_CURRENT_LIST_FILE.

CMAKE_PATCH_VERSION: The patch version of cmake (i.e. the 3 in X.X.3).

This specifies the patch version of the CMake executable being run.

CMAKE_PROJECT_NAME: The name of the current project.

This specifies name of the current project from the closest inherited PROJECT command.

CMAKE_RANLIB: Name of randomizing tool for static libraries.

This specifies name of the program that randomizes libraries on UNIX, not used on Windows, but may be present.

CMAKE_ROOT: Install directory for running cmake.

This is the install root for the running CMake and the Modules directory can be found here. This is commonly used in this format: \${CMAKE_ROOT}/Modules

CMAKE_SHARED_LIBRARY_PREFIX: The prefix for shared libraries that you link to.

The prefix to use for the name of a shared library, lib on UNIX.

CMAKE_SHARED_LIBRARY_PREFIX_<LANG> overrides this for language <LANG>.

CMAKE_SHARED_LIBRARY_SUFFIX: The suffix for shared libraries that you link to.

The suffix to use for the end of a shared library, .dll on Windows.

CMAKE_SHARED_LIBRARY_SUFFIX_<LANG> overrides this for language <LANG>.

CMAKE_SHARED_MODULE_PREFIX: The prefix for loadable modules that you link to.

The prefix to use for the name of a loadable module on this platform.

CMAKE_SHARED_MODULE_PREFIX_<LANG> overrides this for language <LANG>.

CMAKE_SHARED_MODULE_SUFFIX: The suffix for shared libraries that you link to.

The suffix to use for the end of a loadable module on this platform

CMAKE_SHARED_MODULE_SUFFIX_<LANG> overrides this for language <LANG>.

CMAKE_SIZEOF_VOID_P: Size of a void pointer.

This is set to the size of a pointer on the machine, and is determined by a try compile. If a 64 bit size is found, then the library search path is modified to look for 64 bit libraries first.

CMAKE_SKIP_RPATH: If true, do not add run time path information.

If this is set to TRUE, then the rpath information is not added to compiled executables. The default is to add rpath information if the platform supports it. This allows for easy running from the build tree.

CMAKE_SOURCE_DIR: The path to the top level of the source tree.

This is the full path to the top level of the current CMake source tree. For an in-source build, this would be the same as **CMAKE_BINARY_DIR**.

CMAKE_STANDARD_LIBRARIES: Libraries linked into every executable and shared library.

This is the list of libraries that are linked into all executables and libraries.

CMAKE_STATIC_LIBRARY_PREFIX: The prefix for static libraries that you link to.

The prefix to use for the name of a static library, lib on UNIX.

CMAKE_STATIC_LIBRARY_PREFIX_<LANG> overrides this for language <LANG>.

CMAKE_STATIC_LIBRARY_SUFFIX: The suffix for static libraries that you link to.

The suffix to use for the end of a static library, .lib on Windows.

CMAKE_STATIC_LIBRARY_SUFFIX_<LANG> overrides this for language <LANG>.

CMAKE_USING_VC_FREE_TOOLS: True if free visual studio tools being used.

This is set to true if the compiler is Visual Studio free tools.

CMAKE_VERBOSE_MAKEFILE: Create verbose Makefiles if on.

This variable defaults to false. You can set this variable to true to make CMake produce verbose Makefiles that show each command line as it is used.

CMAKE_VERSION: The full version of cmake in major.minor.patch format.

This specifies the full version of the CMake executable being run. This variable is defined by versions 2.6.3 and higher. See variables CMAKE_MAJOR_VERSION, CMAKE_MINOR_VERSION, and CMAKE_PATCH_VERSION for individual version components.

PROJECT_BINARY_DIR: Full path to build directory for project.

This is the binary directory of the most recent PROJECT command.

PROJECT_NAME: Name of the project given to the project command.

This is the name given to the most recent PROJECT command.

PROJECT_SOURCE_DIR: Top level source directory for the current project.

This is the source directory of the most recent PROJECT command.

[Project name]_BINARY_DIR: Top level binary directory for the named project.

A variable is created with the name used in the PROJECT command, and is the binary directory for the project. This can be useful when SUBDIR is used to connect several projects.

[Project name]_SOURCE_DIR: Top level source directory for the named project.

A variable is created with the name used in the PROJECT command, and is the source directory for the project. This can be useful when add_subdirectory is used to connect several projects.

Appendix B – Command Line Reference

CMake Command Line Options

```
cmake [options] <path-to-source>
cmake [options] <path-to-existing-build>
```

The "cmake" executable is the CMake command-line interface. It may be used to configure projects in scripts. Project configuration settings may be specified on the command line with the -D option. The -i option will cause cmake to interactively prompt for such settings.

-C <initial-cache>: Pre-load a script to populate the cache.

When cmake is first run in an empty build tree, it creates a CMakeCache.txt file and populates it with customizable settings for the project. This option may be used to specify a file from which to load cache entries before the first pass through the project's cmake listfiles. The loaded entries take priority over the project's default values. The given file should be a CMake script containing SET commands that use the CACHE option, not a cache-format file.

-D <var>:<type>=<value>: Create a cmake cache entry.

When cmake is first run in an empty build tree, it creates a CMakeCache.txt file and populates it with customizable settings for the project. This option may be used to specify a setting that takes priority over the project's default value. The option may be repeated for as many cache entries as desired.

-U <globbing_expr>: Remove matching entries from CMake cache.

This option may be used to remove one or more variables from the CMakeCache.txt file, globbing expressions using * and ? are supported. The option may be repeated for as many cache entries as desired.

Use with care, you can make your CMakeCache.txt non-working.

-G <generator-name>: Specify a Makefile generator.

CMake may support multiple native build systems on certain platforms. A Makefile generator is responsible for generating a particular build system. Possible generator names are specified in the Generators section.

-Wno-dev: Suppress developer warnings.

Suppress warnings that are meant for the author of the CMakeLists.txt files.

-Wdev: Enable developer warnings.

Enable warnings that are meant for the author of the CMakeLists.txt files.

-E: CMake command mode.

For true platform independence, CMake provides a list of commands that can be used on all systems. Run with -E help for the usage information. Commands available are: chdir, copy, copy_if_different, copy_directory, compare_files, echo, echo_append, environment, make_directory, md5sum, remove_directory, remove, tar, time, touch, touch_nocreate, write_regv, delete_regv, comspec, create_symlink.

-i: Run in wizard mode.

Wizard mode runs cmake interactively without a GUI. The user is prompted to answer questions about the project configuration. The answers are used to set cmake cache values.

-L [A] [H] : List non-advanced cached variables.

List cache variables will run CMake and list all the variables from the CMake cache that are not marked as INTERNAL or ADVANCED. This will effectively display current CMake settings, which can be then changed with -D option. Changing some of the variable may result in more variables being created. If A is specified, then it will display also advanced variables. If H is specified, it will also display help for each variable.

--build <dir>: Build a CMake-generated project binary tree.

This abstracts a native build tool's command-line interface with the following options:

```
<dir>          = Project binary directory to be built.  
--target <tgt> = Build <tgt> instead of default targets.  
--config <cfg> = For multi-configuration tools, choose <cfg>.  
--clean-first  = Build target 'clean' first, then build.  
                  (To clean only, use --target 'clean'.)  
--           = Pass remaining options to the native tool.
```

Run cmake --build with no options for quick help.

-N: View mode only.

Only load the cache. Do not actually run configure and generate steps.

-P <file>: Process script mode.

Process the given cmake file as a script written in the CMake language. No configure or generate step is performed and the cache is not modified. If variables are defined using -D, this must be done before the -P argument.

--graphviz=[file]: Generate graphviz of dependencies.

Generate a graphviz input file that will contain all the library and executable dependencies in the project.

--system-information [file]: Dump information about this system.

Dump a wide range of information about the current system. If run from the top of a binary tree for a CMake project it will dump additional information such as the cache, log files etc.

--debug-trycompile: Do not delete the try compile directories..

Do not delete the files and directories created for try_compile calls. This is useful in debugging failed try_compiles. It may however change the results of the try-compiles as old junk from a previous try-compile may cause a different test to either pass or fail incorrectly. This option is best used for one try-compile at a time, and only when debugging.

--debug-output: Put cmake in a debug mode.

Print extra stuff during the cmake run like stack traces with message(send_error) calls.

--trace: Put cmake in trace mode.

Print a trace of all calls made and from where with message(send_error) calls.

--help-command cmd [file]: Print help for a single command and exit.

Full documentation specific to the given command is displayed. If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-command-list [file]: List available listfile commands and exit.

The list contains all commands for which help may be obtained by using the --help-command argument followed by a command name. If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-commands [file]: Print help for all commands and exit.

Full documentation specific for all current command is displayed. If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-compatcommands [file]: Print help for compatibility commands.

Full documentation specific for all compatibility commands is displayed. If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-module module [file]: Print help for a single module and exit.

Full documentation specific to the given module is displayed. If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-module-list [file]: List available modules and exit.

The list contains all modules for which help may be obtained by using the --help-module argument followed by a module name. If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-modules [file]: Print help for all modules and exit.

Full documentation for all modules is displayed. If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-custom-modules [file]: Print help for all custom modules and exit.

Full documentation for all custom modules is displayed. If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-policy *cmp* [file]: Print help for a single policy and exit.

Full documentation specific to the given policy is displayed.If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-policies [file]: Print help for all policies and exit.

Full documentation for all policies is displayed.If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-property *prop* [file]: Print help for a single property and exit.

Full documentation specific to the given property is displayed.If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-property-list [file]: List available properties and exit.

The list contains all properties for which help may be obtained by using the --help-property argument followed by a property name. If a file is specified, the help is written into it.If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-properties [file]: Print help for all properties and exit.

Full documentation for all properties is displayed.If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-variable *var* [file]: Print help for a single variable and exit.

Full documentation specific to the given variable is displayed.If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-variable-list [file]: List documented variables and exit.

The list contains all variables for which help may be obtained by using the --help-variable argument followed by a variable name. If a file is specified, the help is written into it.If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--help-variables [file]: Print help for all variables and exit.

Full documentation for all variables is displayed.If a file is specified, the documentation is written into and the output format is determined depending on the filename suffix. Supported are man page, HTML, DocBook and plain text.

--copyright [file]: Print the CMake copyright and exit.

If a file is specified, the copyright is written into it.

--help: Print usage information and exit.

Usage describes the basic command line interface and its options.

--help-full [file]: Print full help and exit.

Full help displays most of the documentation provided by the UNIX man page. It is provided for use on non-UNIX platforms, but is also convenient if the man page is not installed. If a file is specified, the help is written into it.

--help-html [file]: Print full help in HTML format.

This option is used by CMake authors to help produce web pages. If a file is specified, the help is written into it.

--help-man [file]: Print full help as a UNIX man page and exit.

This option is used by the cmake build to generate the UNIX man page. If a file is specified, the help is written into it.

--version [file]: Show program name/version banner and exit.

If a file is specified, the version is written into it.

CMake Generators

The following generators are available within CMake:

Borland Makefiles: Generates Borland Makefiles.

MSYS Makefiles: Generates MSYS Makefiles.

The Makefiles use /bin/sh as the shell. They require msys to be installed on the machine.

MinGW Makefiles: Generates a Makefile for use with mingw32-make.

The Makefiles generated use cmd.exe as the shell. They do not require msys or a UNIX shell.

NMake Makefiles: Generates NMake Makefiles.

Unix Makefiles: Generates standard UNIX Makefiles.

A hierarchy of UNIX Makefiles is generated into the build tree. Any standard UNIX-style make program can build the project through the default make target. A "make install" target is also provided.

Visual Studio 10: Generates Visual Studio 10 project files.

Visual Studio 6: Generates Visual Studio 6 project files.

Visual Studio 7: Generates Visual Studio .NET 2002 project files.

Visual Studio 7 .NET 2003: Generates Visual Studio .NET 2003 project files.

Visual Studio 8 2005: Generates Visual Studio .NET 2005 project files.

Visual Studio 8 2005 Win64: Generates Visual Studio .NET 2005 Win64 project files.

Visual Studio 9 2008: Generates Visual Studio 9 2008 project files.

Visual Studio 9 2008 Win64: Generates Visual Studio 9 2008 Win64 project files.

Watcom WMake: Generates Watcom WMake Makefiles.

CodeBlocks - MinGW Makefiles: Generates CodeBlocks project files.

Project files for CodeBlocks will be created in the top directory and in every subdirectory which features a CMakeLists.txt file containing a PROJECT() call. Additionally a hierarchy of Makefiles is generated into the build tree. The appropriate make program can build the project through the default make target. A "make install" target is also provided.

CodeBlocks - NMake Makefiles: Generates CodeBlocks project files.

Project files for CodeBlocks will be created in the top directory and in every subdirectory which features a CMakeLists.txt file containing a PROJECT() call. Additionally a hierarchy of Makefiles is generated into the build tree. The appropriate make program can build the project through the default make target. A "make install" target is also provided.

CodeBlocks - Unix Makefiles: Generates CodeBlocks project files.

Project files for CodeBlocks will be created in the top directory and in every subdirectory which features a CMakeLists.txt file containing a PROJECT() call. Additionally a hierarchy of Makefiles is generated into the build tree. The appropriate make program can build the project through the default make target. A "make install" target is also provided.

Eclipse CDT4 - MinGW Makefiles: Generates Eclipse CDT 4.0 project files.

Project files for Eclipse will be created in the top directory and will have a linked resource to every subdirectory which features a CMakeLists.txt file containing a PROJECT() call. Additionally a hierarchy of Makefiles is generated into the build tree. The appropriate make program can build the project through the default make target. A "make install" target is also provided.

Eclipse CDT4 - NMake Makefiles: Generates Eclipse CDT 4.0 project files.

Project files for Eclipse will be created in the top directory and will have a linked resource to every subdirectory which features a CMakeLists.txt file containing a PROJECT() call. Additionally a hierarchy of Makefiles is generated into the build tree. The appropriate

make program can build the project through the default make target. A "make install" target is also provided.

Eclipse CDT4 - Unix Makefiles: Generates Eclipse CDT 4.0 project files.

Project files for Eclipse will be created in the top directory and will have a linked resource to every subdirectory which features a CMakeLists.txt file containing a PROJECT() call. Additionally a hierarchy of Makefiles is generated into the build tree. The appropriate make program can build the project through the default make target. A "make install" target is also provided.

CTest Command Line Options

```
ctest [options]
```

The "ctest" executable is the CMake test driver program. CMake-generated build trees created for projects that use the ENABLE_TESTING and ADD_TEST commands have testing support. This program will run the tests and report results.

-C <cfg>, --build-config <cfg>: Choose configuration to test.

Some CMake-generated build trees can have multiple build configurations in the same tree. This option can be used to specify which one should be tested. Example configurations are "Debug" and "Release".

-v, --verbose: Enable verbose output from tests.

Test output is normally suppressed and only summary information is displayed. This option will show all test output.

-vv, --extra-verbose: Enable more verbose output from tests.

Test output is normally suppressed and only summary information is displayed. This option will show even more test output.

--debug: Displaying more verbose internals of CTest.

This feature will result in large number of output that is mostly useful for debugging dashboard problems.

--output-on-failure: Output anything outputted by the test program if the test should fail. This option can also be enabled by setting the environment variable
CTEST_OUTPUT_ON_FAILURE

-F: Enable failover.

This option allows ctest to resume a test set execution that was previously interrupted. If no interruption occurred, the -F option will have no effect.

-Q, --quiet: Make ctest quiet.

This option will suppress all the output. The output log file will still be generated if the --output-log is specified. Options such as --verbose, --extra-verbose, and --debug are ignored if --quiet is specified.

-O <file>, --output-log <file>: Output to log file

This option tells ctest to write all its output to a log file.

-N, --show-only: Disable actual execution of tests.

This option tells ctest to list the tests that would be run but not actually run them. Useful in conjunction with the -R and -E options.

-L <regex>, --label-regex <regex>: Run tests with labels matching regular expression.

This option tells ctest to run only the tests whose labels match the given regular expression.

-R <regex>, --tests-regex <regex>: Run tests matching regular expression.

This option tells ctest to run only the tests whose names match the given regular expression.

-E <regex>, --exclude-regex <regex>: Exclude tests matching regular expression.

This option tells ctest to NOT run the tests whose names match the given regular expression.

-LE <regex>, --label-exclude <regex>: Exclude tests with labels matching regular expression.

This option tells ctest to NOT run the tests whose labels match the given regular expression.

-D <dashboard>, --dashboard <dashboard>: Execute dashboard test

This option tells ctest to perform act as a Dart client and perform a dashboard test. All tests are <Mode><Test>, where Mode can be Experimental, Nightly, and Continuous, and Test can be Start, Update, Configure, Build, Test, Coverage, and Submit.

-M <model>, --test-model <model>: Sets the model for a dashboard

This option tells ctest to act as a Dart client where the TestModel can be Experimental, Nightly, and Continuous. Combining -M and -T is similar to -D

-T <action>, --test-action <action>: Sets the dashboard action to perform

This option tells ctest to act as a Dart client and perform some action such as start, build, test etc. Combining -M and -T is similar to -D

--track <track>: Specify the track to submit dashboard to

Submit dashboard to specified track instead of default one. By default, the dashboard is submitted to Nightly, Experimental, or Continuous track, but by specifying this option, the track can be arbitrary.

-S <script>, --script <script>: Execute a dashboard for a configuration

This option tells ctest to load in a configuration script which sets a number of parameters such as the binary and source directories. Then ctest will do what is required to create and run a dashboard. This option basically sets up a dashboard and then runs ctest -D with the appropriate options.

-SP <script>, --script-new-process <script>: Execute a dashboard for a configuration

This option does the same operations as -S but it will do them in a separate process. This is primarily useful in cases where the script may modify the environment and you do not want the modified environment to impact other -S scripts.

-A <file>, --add-notes <file>: Add a notes file with submission

This option tells ctest to include a notes file when submitting dashboard.

-I [Start,End,Stride,test#,test#|Test file], --tests-information: Run a specific number of tests by number.

This option causes ctest to run tests starting at number Start, ending at number End, and incrementing by Stride. Any additional numbers after Stride are considered individual test numbers. Start, End, or stride can be empty. Optionally a file can be given that contains the same syntax as the command line.

-U, --union: Take the Union of -I and -R

When both -R and -I are specified by default the intersection of tests are run. By specifying -U the union of tests is run instead.

--max-width <width>: Set the max width for a test name to output

Set the maximum width for each test name to show in the output. This allows the user to widen the output to avoid clipping the test name which can be very annoying.

--interactive-debug-mode [0|1]: Set the interactive mode to 0 or 1.

This option causes ctest to run tests in either an interactive mode or a non-interactive mode. On Windows this means that in non-interactive mode, all system debug pop up windows are blocked. In dashboard mode (Experimental, Nightly, Continuous), the default is non-interactive. When just running tests not for a dashboard the default is to allow popups and interactive debugging.

--no-label-summary: Disable timing summary information for labels.

This option tells ctest to not print summary information for each label associated with the tests run. If there are no labels on the tests, nothing extra is printed.

--build-and-test: Configure, build and run a test.

This option tells ctest to configure (i.e. run cmake on), build, and or execute a test. The configure and test steps are optional. The arguments to this command line are the source and binary directories. By default this will run CMake on the Source/Bin directories specified unless --build-nocmake is specified. Both --build-makeprogram and --build-generator MUST be provided to use --built-and-test. If --test-command is specified then that will be run after the build is complete. Other options that affect this mode are --build-target --build-nocmake, -build-run-dir, --build-two-config, --build-exe-dir, --build-project,--build-noclean, --build-options

--build-target: Specify a specific target to build.

This option goes with the --build-and-test option, if left out the all target is built.

--build-nocmake: Run the build without running cmake first.

Skip the cmake step.

--build-run-dir: Specify directory to run programs from.

Directory where programs will be after it has been compiled.

--build-two-config: Run CMake twice

--build-exe-dir: Specify the directory for the executable.

--build-generator: Specify the generator to use.

--build-project: Specify the name of the project to build.

--build-makeprogram: Specify the make program to use.

--build-noclean: Skip the make clean step.

--build-config-sample: A sample executable to use to determine the configuraiton

A sample executable to use to determine the configuraiton that should be used. e.g. Debug/Release/etc

--build-options: Add extra options to the build step.

This option must be the last option with the exception of --test-command

--test-command: The test to run with the --build-and-test option.

--test-timeout: The time limit in seconds, internal use only.

--tomorrow-tag: Nightly or experimental starts with next day tag.

This is useful if the build will not finish in one day.

--ctest-config: The configuration file used to initialize CTest state when submitting dashboards.

This option tells CTest to use different initialization file instead of CTestConfiguration.tcl. This way multiple initialization files can be used for example to submit to multiple dashboards.

--overwrite: Overwrite CTest configuration option.

By default `ctest` uses configuration options from configuration file. This option will overwrite the configuration option.

--extra-submit <file>[;<file>]: Submit extra files to the dashboard.

This option will submit extra files to the dashboard.

--force-new-ctest-process: Run child CTest instances as new processes

By default CTest will run child CTest instances within the same process. If this behavior is not desired, this argument will enforce new processes for child CTest processes.

--submit-index: Submit individual dashboard tests with specific index

This option allows performing the same CTest action (such as `test`) multiple times and submit all stages to the same dashboard (`Dart2` required). Each execution requires different index.

CPack Command Line Options

```
cpack -G <generator> [options]
```

The "cpack" executable is the CMake packaging program. CMake-generated build trees created for projects that use the `INSTALL_*` commands have packaging support. This program will generate the package.

-G <generator>: Use the specified generator to generate package.

CPack may support multiple native packaging systems on certain platforms. A generator is responsible for generating input files for particular system and invoking that systems. Possible generator names are specified in the Generators section.

-C <Configuration>: Specify the project configuration

This option specifies the configuration that the project was build with, for example 'Debug', 'Release'.

-D <var>=<value>: Set a CPack variable.

Set a variable that can be used by the generator.

--config <config file>: Specify the config file.

Specify the config file to use to create the package. By default CPackConfig.cmake in the current directory will be used.

CPack Generators

NSIS: Null Soft Installer

STGZ: Self extracting Tar GZip compression

TBZ2: Tar BZip2 compression

TGZ: Tar GZip compression

TZ: Tar Compress compression

ZIP: ZIP file format

Appendix C – Listfile Commands

Current Commands

The following is an alphabetical listing of the commands that are currently used by CMake. Later in the appendix there is a section on commands that are deprecated but may still be handled by CMake for backwards compatibility.

add_custom_command: Add a custom build rule to the generated build system.

There are two main signatures for `add_custom_command`. The first signature is for adding a custom command to produce an output.

```
add_custom_command(OUTPUT output1 [output2 ...]
                    COMMAND command1 [ARGS] [args1...]
                    [COMMAND command2 [ARGS] [args2...] ...]
                    [MAIN_DEPENDENCY depend]
                    [DEPENDS [depends...]]
                    [IMPLICIT_DEPENDS <lang1> depend1 ...]
                    [WORKING_DIRECTORY dir]
                    [COMMENT comment] [VERBATIM] [APPEND])
```

This defines a command to generate specified OUTPUT file(s). A target created in the same directory (CMakeLists.txt file) that specifies any output of the custom command as a source file is given a rule to generate the file using the command at build time. If an output name is a relative path it will be interpreted relative to the build tree directory corresponding to the

current source directory. Note that MAIN_DEPENDENCY is completely optional and is used as a suggestion to Visual Studio about where to hang the custom command. In Makefile terms this creates a new target of the following form:

```
OUTPUT: MAIN_DEPENDENCY DEPENDS  
COMMAND
```

If more than one command is specified they will be executed in order. The optional ARGS argument is for backward compatibility and will be ignored.

The second signature adds a custom command to a target such as a library or executable. This is useful for performing an operation before or after building the target. The command becomes part of the target and will only execute when the target itself is built. If the target is already built, the command will not execute.

```
add_custom_command(TARGET target  
                   PRE_BUILD | PRE_LINK | POST_BUILD  
                   COMMAND command1 [ARGS] [args1...]  
                   [COMMAND command2 [ARGS] [args2...] ...]  
                   [WORKING_DIRECTORY dir]  
                   [COMMENT comment] [VERBATIM])
```

This defines a new command that will be associated with building the specified target. When the command will happen is determined by which of the following is specified:

```
PRE_BUILD - run before all other dependencies  
PRE_LINK  - run after other dependencies  
POST_BUILD - run after the target has been built
```

Note that the PRE_BUILD option is only supported on Visual Studio 7 or later. For all other generators PRE_BUILD will be treated as PRE_LINK.

If WORKING_DIRECTORY is specified the command will be executed in the directory given. If COMMENT is set, the value will be displayed as a message before the commands are executed at build time. If APPEND is specified the COMMAND and DEPENDS option values are appended to the custom command for the first output specified. There must have already been a previous call to this command with the same output. The COMMENT, WORKING_DIRECTORY, and MAIN_DEPENDENCY options are currently ignored when APPEND is given, but may be used in the future.

If VERBATIM is given then all arguments to the commands will be escaped properly for the build tool so that the invoked command receives each argument unchanged. Note that one

level of escapes is still used by the CMake language processor before add_custom_command even sees the arguments. Use of VERBATIM is recommended as it enables correct behavior. When VERBATIM is not given the behavior is platform specific because there is no protection of tool-specific special characters.

If the output of the custom command is not actually created as a file on disk it should be marked as SYMBOLIC with SET_SOURCE_FILES_PROPERTIES.

The IMPLICIT_DEPENDS option requests scanning of implicit dependencies of an input file. The language given specifies the programming language whose corresponding dependency scanner should be used. Currently only C, C++ and Fortran language scanners are supported. Dependencies discovered from the scanning are added to those of the custom command at build time. Note that the IMPLICIT_DEPENDS option is currently supported only for Makefile generators and will be ignored by other generators.

If COMMAND specifies an executable target (created by ADD_EXECUTABLE) it will automatically be replaced by the location of the executable created at build time. Additionally a target-level dependency will be added so that the executable target will be built before any target using this custom command. However this does NOT add a file-level dependency that would cause the custom command to re-run whenever the executable is recompiled.

The DEPENDS option specifies files on which the command depends. If any dependency is an OUTPUT of another custom command in the same directory (CMakeLists.txt file) CMake automatically brings the other custom command into the target in which this command is built. If DEPENDS specifies any target (created by an ADD_* command) a target-level dependency is created to make sure the target is built before any target using this custom command. Additionally, if the target is an executable or library a file-level dependency is created to cause the custom command to re-run whenever the target is recompiled.

add_custom_target: Add a target with no output so it will always be built.

```
add_custom_target(Name [ALL] [command1 [args1...]]  
                  [COMMAND command2 [args2...] ...]  
                  [DEPENDS depend depend depend ... ]  
                  [WORKING_DIRECTORY dir]  
                  [COMMENT comment] [VERBATIM]  
                  [SOURCES src1 [src2...]])
```

Adds a target with the given name that executes the given commands. The target has no output file and is ALWAYS CONSIDERED OUT OF DATE even if the commands try to create a file with the name of the target. Use ADD_CUSTOM_COMMAND to generate a file with dependencies. By default nothing depends on the custom target. Use ADD_DEPENDENCIES to add dependencies to or from other targets. If the ALL option is

specified it indicates that this target should be added to the default build target so that it will be run every time (the command cannot be called ALL). The command and arguments are optional and if not specified an empty target will be created. If WORKING_DIRECTORY is set, then the command will be run in that directory. If COMMENT is set, the value will be displayed as a message before the commands are executed at build time. Dependencies listed with the DEPENDS argument may reference files and outputs of custom commands created with add_custom_command() in the same directory (CMakeLists.txt file).

If VERBATIM is given then all arguments to the commands will be escaped properly for the build tool so that the invoked command receives each argument unchanged. Note that one level of escapes is still used by the CMake language processor before add_custom_target even sees the arguments. Use of VERBATIM is recommended as it enables correct behavior. When VERBATIM is not given the behavior is platform specific because there is no protection of tool-specific special characters.

The SOURCES option specifies additional source files to be included in the custom target. Specified source files will be added to IDE project files for convenience in editing even if they have not build rules.

add_definitions: Adds -D define flags to the compilation of source files.

```
add_definitions(-DFOO -DBAR ...)
```

Adds flags to the compiler command line for sources in the current directory and below. This command can be used to add any flags, but it was originally intended to add preprocessor definitions. Flags beginning in -D or /D that look like preprocessor definitions are automatically added to the COMPILE_DEFINITIONS property for the current directory. Definitions with non-trivial values may be left in the set of flags instead of being converted for reasons of backwards compatibility. See documentation of the directory, target, and source file COMPILE_DEFINITIONS properties for details on adding preprocessor definitions to specific scopes and configurations.

add_dependencies: Add a dependency between top-level targets.

```
add_dependencies(target-name depend-target1  
                depend-target2 ...)
```

Make a top-level target depend on other top-level targets. A top-level target is one created by ADD_EXECUTABLE, ADD_LIBRARY, or ADD_CUSTOM_TARGET. Adding dependencies with this command can be used to make sure one target is built before another target. See the DEPENDS option of ADD_CUSTOM_TARGET and ADD_CUSTOM_COMMAND for adding file-level dependencies in custom rules. See the

OBJECT_DEPENDS option in SET_SOURCE_FILES_PROPERTIES to add file-level dependencies to object files.

add_executable: Add an executable to the project using the specified source files.

```
add_executable(<name> [WIN32] [MACOSX_BUNDLE]
               [EXCLUDE_FROM_ALL]
               source1 source2 ... sourceN)
```

Adds an executable target called <name> to be built from the source files listed in the command invocation. The <name> corresponds to the logical target name and must be globally unique within a project. The actual file name of the executable built is constructed based on conventions of the native platform (such as <name>.exe or just <name>).

By default the executable file will be created in the build tree directory corresponding to the source tree directory in which the command was invoked. See documentation of the RUNTIME_OUTPUT_DIRECTORY target property to change this location. See documentation of the OUTPUT_NAME target property to change the <name> part of the final file name.

If WIN32 is given the property WIN32_EXECUTABLE will be set on the target created. See documentation of that target property for details.

If MACOSX_BUNDLE is given the corresponding property will be set on the created target. See documentation of the MACOSX_BUNDLE target property for details.

If EXCLUDE_FROM_ALL is given the corresponding property will be set on the created target. See documentation of the EXCLUDE_FROM_ALL target property for details.

The add_executable command can also create IMPORTED executable targets using this signature:

```
add_executable(<name> IMPORTED)
```

An IMPORTED executable target references an executable file located outside the project. No rules are generated to build it. The target name has scope in the directory in which it is created and below. It may be referenced like any target built within the project. IMPORTED executables are useful for convenient reference from commands like add_custom_command. Details about the imported executable are specified by setting properties whose names begin in "IMPORTED_". The most important such property is IMPORTED_LOCATION (and its per-configuration version IMPORTED_LOCATION_<CONFIG>) which specifies the

location of the main executable file on disk. See documentation of the IMPORTED_* properties for more information.

add_library: Add a library to the project using the specified source files.

```
add_library(<name> [STATIC | SHARED | MODULE]
            [EXCLUDE_FROM_ALL]
            source1 source2 ... sourceN)
```

Adds a library target called <name> to be built from the source files listed in the command invocation. The <name> corresponds to the logical target name and must be globally unique within a project. The actual file name of the library built is constructed based on conventions of the native platform (such as lib<name>.a or <name>.lib).

STATIC, SHARED, or MODULE may be given to specify the type of library to be created. STATIC libraries are archives of object files for use when linking other targets. SHARED libraries are linked dynamically and loaded at runtime. MODULE libraries are plugins that are not linked into other targets but may be loaded dynamically at runtime using dlopen-like functionality. If no type is given explicitly the type is STATIC or SHARED based on whether the current value of the variable BUILD_SHARED_LIBS is true.

By default the library file will be created in the build tree directory corresponding to the source tree directory in which the command was invoked. See documentation of the ARCHIVE_OUTPUT_DIRECTORY, LIBRARY_OUTPUT_DIRECTORY, and RUNTIME_OUTPUT_DIRECTORY target properties to change this location. See documentation of the OUTPUT_NAME target property to change the <name> part of the final file name.

If EXCLUDE_FROM_ALL is given the corresponding property will be set on the created target. See documentation of the EXCLUDE_FROM_ALL target property for details.

The add_library command can also create IMPORTED library targets using this signature:

```
add_library(<name> <SHARED|STATIC|MODULE|UNKNOWN> IMPORTED)
```

An IMPORTED library target references a library file located outside the project. No rules are generated to build it. The target name has scope in the directory in which it is created and below. It may be referenced like any target built within the project. IMPORTED libraries are useful for convenient reference from commands like target_link_libraries. Details about the imported library are specified by setting properties whose names begin in "IMPORTED_". The most important such property is IMPORTED_LOCATION (and its per-configuration

version IMPORTED_LOCATION_<CONFIG>) which specifies the location of the main library file on disk. See documentation of the IMPORTED_* properties for more information.

add_subdirectory: Add a subdirectory to the build.

```
add_subdirectory(source_dir [binary_dir]
                 [EXCLUDE_FROM_ALL])
```

Add a subdirectory to the build. The source_dir specifies the directory in which the source CmakeLists.txt and code files are located. If it is a relative path it will be evaluated with respect to the current directory (the typical usage), but it may also be an absolute path. The binary_dir specifies the directory in which to place the output files. If it is a relative path it will be evaluated with respect to the current output directory, but it may also be an absolute path. If binary_dir is not specified, the value of source_dir, before expanding any relative path, will be used (the typical usage). The CMakeLists.txt file in the specified source directory will be processed immediately by CMake before processing in the current input file continues beyond this command.

If the EXCLUDE_FROM_ALL argument is provided then targets in the subdirectory will not be included in the ALL target of the parent directory by default, and will be excluded from IDE project files. Users must explicitly build targets in the subdirectory. This is meant for use when the subdirectory contains a separate part of the project that is useful but not necessary, such as a set of examples. Typically the subdirectory should contain its own project() command invocation so that a full build system will be generated in the subdirectory (such as a VS IDE solution file). Note that inter-target dependencies supercede this exclusion. If a target built by the parent project depends on a target in the subdirectory, the dependee target will be included in the parent project build system to satisfy the dependency.

add_test: Add a test to the project with the specified arguments.

```
add_test(testname Exename arg1 arg2 ...)
```

If the ENABLE_TESTING command has been run, this command adds a test target to the current directory. If ENABLE_TESTING has not been run, this command does nothing. The tests are run by the testing subsystem by executing Exename with the specified arguments. Exename can be either an executable built by this project or an arbitrary executable on the system (like tcsh). The test will be run with the current working directory set to the CMakeList.txt files corresponding directory in the binary tree.

```
add_test(NAME <name> [CONFIGURATIONS [Debug|Release|...]]
         COMMAND <command> [arg1 [arg2 ...]])
```

If COMMAND specifies an executable target (created by add_executable) it will automatically be replaced by the location of the executable created at build time. If a CONFIGURATIONS option is given then the test will be executed only when testing under one of the named configurations.

Arguments after COMMAND may use "generator expressions" with the syntax "\$<...>". These expressions are evaluated during build system generation and produce information specific to each generated build configuration. Valid expressions are:

\$<CONFIGURATION>	= configuration name
\$<TARGET_FILE:tgt>	= main file (.exe, .so.1.2, .a)
\$<TARGET_LINKER_FILE:tgt>	= file used to link (.a, .lib, .so)
\$<TARGET SONAME FILE:tgt>	= file with soname (.so.3)

where "tgt" is the name of a target. Target file expressions produce a full path, but _DIR and _NAME versions can produce the directory and file name components:

\$<TARGET_FILE_DIR:tgt>/	\$<TARGET_FILE_NAME:tgt>
\$<TARGET_LINKER_FILE_DIR:tgt>/	\$<TARGET_LINKER_FILE_NAME:tgt>
\$<TARGET SONAME FILE DIR:tgt>/	\$<TARGET SONAME FILE NAME:tgt>

Example usage:

```
add_test(NAME mytest
        COMMAND testDriver --config $<CONFIGURATION>
                    --exe $<TARGET_FILE:myexe>)
```

This creates a test "mytest" whose command runs a testDriver tool passing the configuration name and the full path to the executable file produced by target "myexe".

aux_source_directory: Find all source files in a directory.

aux_source_directory(<dir> <variable>)
--

Collects the names of all the source files in the specified directory and stores the list in the <variable> provided. This command is intended to be used by projects that use explicit template instantiation. Template instantiation files can be stored in a "Templates" subdirectory and collected automatically using this command to avoid manually listing all instantiations.

It is tempting to use this command to avoid writing the list of source files for a library or executable target. While this seems to work, there is no way for CMake to generate a build system that knows when a new source file has been added. Normally the generated build system knows when it needs to rerun CMake because the CMakeLists.txt file is modified to add a new source. When the source is just added to the directory without modifying this file, one would have to manually rerun CMake to generate a build system incorporating the new file.

break: Break from an enclosing foreach or while loop.

```
break()
```

Breaks from an enclosing foreach loop or while loop

build_command: Get the command line that will build this project.

```
build_command(<variable> <makecommand>)
```

Sets the given <variable> to a string containing the command that will build this project from the root of the build tree using the build tool given by <makecommand>. <makecommand> should be msdev, nmake, make or one of the end user build tools. This is useful for configuring testing systems.

cmake_minimum_required: Set the minimum required version of cmake for a project.

```
cmake_minimum_required(VERSION major[.minor[.patch]]  
                      [FATAL_ERROR])
```

If the current version of CMake is lower than that required it will stop processing the project and report an error. When a version higher than 2.4 is specified the command implicitly invokes

```
cmake_policy(VERSION major[.minor[.patch]])
```

which sets the cmake policy version level to the version specified. When version 2.4 or lower is given the command implicitly invokes

```
cmake_policy(VERSION 2.4)
```

which enables compatibility features for CMake 2.4 and lower. The FATAL_ERROR option is accepted but ignored by CMake 2.6 and higher. It should be specified so CMake versions 2.4 and lower fail with an error instead of just a warning.

`cmake_policy`: Manage CMake Policy settings.

As CMake evolves it is sometimes necessary to change existing behavior in order to fix bugs or improve implementations of existing features. The CMake Policy mechanism is designed to help keep existing projects building as new versions of CMake introduce changes in behavior. Each new policy (behavioral change) is given an identifier of the form "CMP<NNNN>" where "<NNNN>" is an integer index. Documentation associated with each policy describes the OLD and NEW behavior and the reason the policy was introduced. Projects may set each policy to select the desired behavior. When CMake needs to know which behavior to use it checks for a setting specified by the project. If no setting is available the OLD behavior is assumed and a warning is produced requesting that the policy be set.

The `cmake_policy` command is used to set policies to OLD or NEW behavior. While setting policies individually is supported, we encourage projects to set policies based on CMake versions.

```
cmake_policy(VERSION major.minor[.patch])
```

Specify that the current CMake list file is written for the given version of CMake. All policies introduced in the specified version or earlier will be set to use NEW behavior. All policies introduced after the specified version will be unset. This effectively requests behavior preferred as of a given CMake version and tells newer CMake versions to warn about their new policies. The policy version specified must be at least 2.4 or the command will report an error. In order to get compatibility features supporting versions earlier than 2.4 see documentation of policy CMP0001.

```
cmake_policy(SET CMP<NNNN> NEW)
cmake_policy(SET CMP<NNNN> OLD)
```

Tell CMake to use the OLD or NEW behavior for a given policy. Projects depending on the old behavior of a given policy may silence a policy warning by setting the policy state to OLD. Alternatively one may fix the project to work with the new behavior and set the policy state to NEW.

```
cmake_policy(GET CMP<NNNN> <variable>)
```

Check whether a given policy is set to OLD or NEW behavior. The output variable value will be "OLD" or "NEW" if the policy is set, and empty otherwise.

CMake keeps policy settings on a stack, so changes made by the `cmake_policy` command affect only the top of the stack. A new entry on the policy stack is managed automatically for each subdirectory to protect its parents and siblings. CMake also manages a new entry for scripts loaded by `include()` and `find_package()` commands except when invoked with the `NO_POLICY_SCOPE` option (see also policy CMP0011). The `cmake_policy` command provides an interface to manage custom entries on the policy stack:

```
cmake_policy(PUSH)
cmake_policy(POP)
```

Each PUSH must have a matching POP to erase any changes. This is useful to make temporary changes to policy settings.

Functions and macros record policy settings when they are created and use the pre-record policies when they are invoked. If the function or macro implementation sets policies, the changes automatically propagate up through callers until they reach the closest nested policy stack entry.

configure_file: Copy a file to another location and modify its contents.

```
configure_file(<input> <output>
              [COPYONLY] [ESCAPE_QUOTES] [@ONLY])
```

Copies a file `<input>` to file `<output>` and substitutes variable values referenced in the file content. If `<input>` is a relative path it is evaluated with respect to the current source directory. The `<input>` must be a file, not a directory. If `<output>` is a relative path it is evaluated with respect to the current binary directory. If `<output>` names an existing directory the input file is placed in that directory with its original name.

This command replaces any variables in the input file referenced as `${VAR}` or `@VAR@` with their values as determined by CMake. If a variable is not defined, it will be replaced with nothing. If `COPYONLY` is specified, then no variable expansion will take place. If `ESCAPE_QUOTES` is specified then any substituted quotes will be C-style escaped. The file will be configured with the current values of CMake variables. If `@ONLY` is specified, only variables of the form `@VAR@` will be replaced and `${VAR}` will be ignored. This is useful for configuring scripts that use `${VAR}` . Any occurrences of `#cmakedefine VAR` will be replaced with either `#define VAR` or `/* #undef VAR */` depending on the setting of `VAR` in CMake. Any occurrences of `#cmakedefine01 VAR` will be replaced with either `#define VAR 1` or `#define VAR 0` depending on whether `VAR` evaluates to TRUE or FALSE in CMake.

create_test_sourcelist: Create a test driver and source list for building test programs.

```
create_test_sourcelist(sourceListName driverName
                      test1 test2 test3
                      EXTRA_INCLUDE include.h
                      FUNCTION function)
```

A test driver is a program that links together many small tests into a single executable. This is useful when building static executables with large libraries to shrink the total required size. The list of source files needed to build the test driver will be in sourceListName. DriverName is the name of the test driver program. The rest of the arguments consist of a list of test source files, can be semicolon separated. Each test source file should have a function in it that is the same name as the file with no extension (foo.cxx should have int foo(int, char*[]);) DriverName will be able to call each of the tests by name on the command line. If EXTRA_INCLUDE is specified, then the next argument is included into the generated file. If FUNCTION is specified, then the next argument is taken as a function name that is passed a pointer to ac and av. This can be used to add extra command line processing to each test. The cmake variable CMAKE_TESTDRIVER_BEFORE_TESTMAIN can be set to have code that will be placed directly before calling the test main function. CMAKE_TESTDRIVER_AFTER_TESTMAIN can be set to have code that will be placed directly after the call to the test main function.

define_property: Define and document custom properties.

```
define_property(<GLOBAL | DIRECTORY | TARGET | SOURCE |
                TEST | VARIABLE | CACHED_VARIABLE>
                PROPERTY <name> [INHERITED]
                BRIEF_DOCS <brief-doc> [docs...]
                FULL_DOCS <full-doc> [docs...])
```

Define one property in a scope for use with the set_property and get_property commands. This is primarily useful to associate documentation with property names that may be retrieved with the get_property command. The first argument determines the kind of scope in which the property should be used. It must be one of the following:

GLOBAL	= associated with the global namespace
DIRECTORY	= associated with one directory
TARGET	= associated with one target
SOURCE	= associated with one source file
TEST	= associated with a test named with add_test
VARIABLE	= documents a CMake language variable
CACHED_VARIABLE	= documents a CMake cache variable

Note that unlike set_property and get_property no actual scope needs to be given; only the kind of scope is important. The required PROPERTY option is immediately followed by the name of the property being defined.

If the INHERITED option then the get_property command will chain up to the next higher scope when the requested property is not set in the scope given to the command. DIRECTORY scope chains to GLOBAL. TARGET, SOURCE, and TEST chain to DIRECTORY.

The BRIEF_DOCS and FULL_DOCS options are followed by strings to be associated with the property as its brief and full documentation. Corresponding options to the get_property command will retrieve the documentation.

else: Starts the else portion of an if block. See the if command.

elseif: Starts the elseif portion of an if block. See the if command.

enable_language: Enable a language (CXX/C/Fortran/etc)

```
enable_language(languageName [OPTIONAL] )
```

This command enables support for the named language in CMake. This is the same as the project command but does not create any of the extra variables that are created by the project command. Example languages are CXX, C, Fortran. If OPTIONAL is used, use the CMAKE_<languageName>_COMPILER_WORKS variable to check whether the language has been enabled successfully.

enable_testing: Enable testing for current directory and below.

```
enable_testing()
```

Enables testing for this directory and below. See also the add_test command. Note that ctest expects to find a test file in the build directory root. Therefore, this command should be in the source directory root.

endforeach: Ends a list of commands in a foreach block. See the foreach command.

endfunction: Ends a list of commands in a function block. See the function command.

endif: Ends a list of commands in an if block. See the if command.

endmacro: Ends a list of commands in a macro block. See the macro command.

endwhile: Ends a list of commands in a while block. See the while command.

execute_process: Execute one or more child processes.

```
execute_process(COMMAND <cmd1> [args1...])
    [COMMAND <cmd2> [args2...] ...]
    [WORKING_DIRECTORY <directory>]
    [TIMEOUT <seconds>]
    [RESULT_VARIABLE <variable>]
    [OUTPUT_VARIABLE <variable>]
    [ERROR_VARIABLE <variable>]
    [INPUT_FILE <file>]
    [OUTPUT_FILE <file>]
    [ERROR_FILE <file>]
    [OUTPUT_QUIET]
    [ERROR_QUIET]
    [OUTPUT_STRIP_TRAILING_WHITESPACE]
    [ERROR_STRIP_TRAILING_WHITESPACE])
```

Runs the given sequence of one or more commands with the standard output of each process piped to the standard input of the next. A single standard error pipe is used for all processes. If WORKING_DIRECTORY is given the named directory will be set as the current working directory of the child processes. If TIMEOUT is given the child processes will be terminated if they do not finish in the specified number of seconds (fractions are allowed). If RESULT_VARIABLE is given the variable will be set to contain the result of running the processes. This will be an integer return code from the last child or a string describing an error condition. If OUTPUT_VARIABLE or ERROR_VARIABLE are given the variable named will be set with the contents of the standard output and standard error pipes respectively. If the same variable is named for both pipes their output will be merged in the order produced. If INPUT_FILE, OUTPUT_FILE, or ERROR_FILE is given the file named will be attached to the standard input of the first process, standard output of the last process, or standard error of all processes respectively. If OUTPUT_QUIET or ERROR_QUIET is given then the standard output or standard error results will be quietly ignored. If more than one OUTPUT_* or ERROR_* option is given for the same pipe the precedence is not specified. If no OUTPUT_* or ERROR_* options are given the output will be shared with the corresponding pipes of the CMake process itself.

The execute_process command is a newer more powerful version of exec_program, but the old command has been kept for compatibility.

export: Export targets from the build tree for use by outside projects.

```
export(TARGETS [target1 [target2 [...]]] [NAMESPACE <namespace>]
    [APPEND] FILE <filename>)
```

Create a file <filename> that may be included by outside projects to import targets from the current project's build tree. This is useful during cross-compiling to build utility executables

that can run on the host platform in one project and then import them into another project being compiled for the target platform. If the NAMESPACE option is given the <namespace> string will be prepended to all target names written to the file. If the APPEND option is given the generated code will be appended to the file instead of overwriting it. If a library target is included in the export but a target to which it links is not included the behavior is unspecified.

The file created by this command is specific to the build tree and should never be installed. See the `install(EXPORT)` command to export targets from an installation tree.

```
export(PACKAGE <name>)
```

Store the current build directory in the CMake user package registry for package <name>. The `find_package` command may consider the directory while searching for package <name>. This helps dependent projects find and use a package from the current project's build tree without help from the user. Note that the entry in the package registry that this command creates works only in conjunction with a package configuration file (<name>Config.cmake) that works with the build tree.

file: File manipulation command.

```
file(WRITE filename "message to write"...)
file(APPEND filename "message to write"...)
file(READ filename variable [LIMIT numBytes]
    [OFFSET offset]
    [HEX])
file(STRINGS filename variable [LIMIT_COUNT num]
    [LIMIT_INPUT numBytes] [LIMIT_OUTPUT numBytes]
    [LENGTH_MINIMUM numBytes] [LENGTH_MAXIMUM numBytes]
    [NEWLINE_CONSUME] [REGEX regex]
    [NO_HEX_CONVERSION])
file(GLOB variable [RELATIVE path] [globbing expressions]...)
file(GLOB_RECURSE variable [RELATIVE path]
    [FOLLOW_SYMLINKS] [globbing expressions]...)
file(RENAME <oldname> <newname>)
file(REMOVE [file1 ...])
file(REMOVE_RECURSE [file1 ...])
file(MAKE_DIRECTORY [directory1 directory2 ...])
file(RELATIVE_PATH variable directory file)
file(TO_CMAKE_PATH path result)
file(TO_NATIVE_PATH path result)
file(DOWNLOAD url file [TIMEOUT timeout]
    [STATUS status] [LOG log])
```

WRITE will write a message into a file called 'filename'. It overwrites the file if it already exists, and creates the file if it does not exist.

APPEND will write a message into a file same as WRITE, except it will append it to the end of the file

READ will read the content of a file and store it into the variable. It will start at the given offset and read up to numBytes. If the argument HEX is given, the binary data will be converted to hexadecimal representation and this will be stored in the variable.

STRINGS will parse a list of ASCII strings from a file and store it in a variable. Binary data in the file are ignored. Carriage return (CR) characters are ignored. It works also for Intel Hex and Motorola S-record files, which are automatically converted to binary format when reading them. Disable this using NO_HEX_CONVERSION.

LIMIT_COUNT sets the maximum number of strings to return. LIMIT_INPUT sets the maximum number of bytes to read from the input file. LIMIT_OUTPUT sets the maximum number of bytes to store in the output variable. LENGTH_MINIMUM sets the minimum length of a string to return. Shorter strings are ignored. LENGTH_MAXIMUM sets the maximum length of a string to return. Longer strings are split into strings no longer than the maximum length. NEWLINE_CONSUME allows newlines to be included in strings instead of terminating them.

REGEX specifies a regular expression that a string must match to be returned. Typical usage

```
file(STRINGS myfile.txt myfile)
```

stores a list in the variable "myfile" in which each item is a line from the input file.

GLOB will generate a list of all files that match the globbing expressions and store it into the variable. Globbing expressions are similar to regular expressions, but much simpler. If RELATIVE flag is specified for an expression, the results will be returned as a relative path to the given path. Examples of globbing expressions include:

*.cxx	- match all files with extension cxx
*.vt?	- match all files with extension vta,...,vtz
f[3-5].txt	- match files f3.txt, f4.txt, f5.txt

GLOB_RECURSE will generate a list similar to the regular GLOB, except it will traverse all the subdirectories of the matched directory and match the files. Subdirectories that are symlinks are only traversed if FOLLOW_SYMLINKS is given or cmake policy CMP0009 is not set to NEW. See cmake --help-policy CMP0009 for more information. Examples of recursive globbing include:

```
/dir/*.py - match all python files in /dir and subdirectories
```

`MAKE_DIRECTORY` will create the given directories, also if their parent directories don't exist yet

`RENAME` moves a file or directory within a filesystem, replacing the destination atomically.

`REMOVE` will remove the given files, also in subdirectories

`REMOVE_RECURSE` will remove the given files and directories, also non-empty directories

`RELATIVE_PATH` will determine relative path from directory to the given file.

`TO_CMAKE_PATH` will convert path into a CMake style path with UNIX /. The input can be a single path or a system path like "\$ENV{PATH}". Note the double quotes around the ENV call `TO_CMAKE_PATH` only takes one argument.

`TO_NATIVE_PATH` works just like `TO_CMAKE_PATH`, but will convert from a cmake style path into the native path style \ for windows and / for UNIX.

`DOWNLOAD` will download the given URL to the given file. If LOG var is specified a log of the download will be put in var. If STATUS var is specified the status of the operation will be put in var. The status is returned in a list of length 2. The first element is the numeric return value for the operation, and the second element is a string value for the error. A 0 numeric error means no error in the operation. If TIMEOUT time is specified, the operation will timeout after time seconds, time can be specified as a float.

The `file()` command also provides `COPY` and `INSTALL` signatures:

```
file(<COPY|INSTALL> files... DESTINATION <dir>
    [FILE_PERMISSIONS permissions...]
    [DIRECTORY_PERMISSIONS permissions...]
    [NO_SOURCE_PERMISSIONS] [USE_SOURCE_PERMISSIONS]
    [FILES_MATCHING]
    [[PATTERN <pattern> | REGEX <regex>]
    [EXCLUDE] [PERMISSIONS permissions...]] [...])
```

The `COPY` signature copies files, directories, and symlinks to a destination folder. Relative input paths are evaluated with respect to the current source directory, and a relative destination is evaluated with respect to the current build directory. Copying preserves input file timestamps, and optimizes out a file if it exists at the destination with the same timestamp. Copying preserves input permissions unless explicit permissions or `NO_SOURCE_PERMISSIONS` are given (default is `USE_SOURCE_PERMISSIONS`). See

the `install(DIRECTORY)` command for documentation of permissions, PATTERN, REGEX, and EXCLUDE options.

The `INSTALL` signature differs slightly from `COPY`: it prints status messages, and `NO_SOURCE_PERMISSIONS` is default. Installation scripts generated by the `install()` command use this signature (with some undocumented options for internal use).

`find_file`: Find the full path to a file.

```
find_file(<VAR> name1 [path1 path2 ...])
```

This is the short-hand signature for the command that is sufficient in many cases. It is the same as `find_file(<VAR> name1 [PATHS path1 path2 ...])`

```
find_file(  
    <VAR>  
    name | NAMES name1 [name2 ...]  
    [HINTS path1 [path2 ... ENV var]]  
    [PATHS path1 [path2 ... ENV var]  
    [PATH_SUFFIXES suffix1 [suffix2 ...]]  
    [DOC "cache documentation string"]  
    [NO_DEFAULT_PATH]  
    [NO_CMAKE_ENVIRONMENT_PATH]  
    [NO_CMAKE_PATH]  
    [NO_SYSTEM_ENVIRONMENT_PATH]  
    [NO_CMAKE_SYSTEM_PATH]  
    [CMAKE_FIND_ROOT_PATH_BOTH |  
     ONLY_CMAKE_FIND_ROOT_PATH |  
     NO_CMAKE_FIND_ROOT_PATH]  
)
```

This command is used to find a full path to named file. A cache entry named by `<VAR>` is created to store the result of this command. If the full path to a file is found the result is stored in the variable and the search will not be repeated unless the variable is cleared. If nothing is found, the result will be `<VAR>-NOTFOUND`, and the search will be attempted again the next time `find_file` is invoked with the same variable. The name of the full path to a file that is searched for is specified by the names listed after the `NAMES` argument. Additional search locations can be specified after the `PATHS` argument. If `ENV var` is found in the `HINTS` or `PATHS` section the environment variable `var` will be read and converted from a system environment variable to a `cmake` style list of paths. For example `ENV PATH` would be a way to list the system path variable. The argument after `DOC` will be used for the documentation

string in the cache. PATH_SUFFIXES specifies additional subdirectories to check below each search path.

If NO_DEFAULT_PATH is specified, then no additional paths are added to the search. If NO_DEFAULT_PATH is not specified, the search process is as follows:

1. Search paths specified in cmake-specific cache variables. These are intended to be used on the command line with a -DVAR=value. This can be skipped if NO_CMAKE_PATH is passed.

```
<prefix>/include for each <prefix> in CMAKE_PREFIX_PATH,  
CMAKE_INCLUDE_PATH and CMAKE_FRAMEWORK_PATH
```

2. Search paths specified in cmake-specific environment variables. These are intended to be set in the user's shell configuration. This can be skipped if NO_CMAKE_ENVIRONMENT_PATH is passed.

```
<prefix>/include for each <prefix> in CMAKE_PREFIX_PATH,  
CMAKE_INCLUDE_PATH and CMAKE_FRAMEWORK_PATH
```

3. Search the paths specified by the HINTS option. These should be paths computed by system introspection, such as a hint provided by the location of another item already found. Hard-coded guesses should be specified with the PATHS option.

4. Search the standard system environment variables. This can be skipped if NO_SYSTEM_ENVIRONMENT_PATH is an argument. Specifically PATH and INCLUDE

5. Search cmake variables defined in the Platform files for the current system. This can be skipped if NO_CMAKE_SYSTEM_PATH is passed.

```
<prefix>/include for each <prefix> in CMAKE_SYSTEM_PREFIX_PATH,  
CMAKE_SYSTEM_INCLUDE_PATH and CMAKE_SYSTEM_FRAMEWORK_PATH
```

6. Search the paths specified by the PATHS option or in the short-hand version of the command. These are typically hard-coded guesses.

On Darwin or systems supporting OS X Frameworks, the cmake variable CMAKE_FIND_FRAMEWORK can be set to empty or one of the following:

"FIRST"	- Try to find frameworks before standard libraries or headers. This is the default on Darwin.
"LAST"	- Try to find frameworks after standard libraries or headers.
"ONLY"	- Only try to find frameworks.
"NEVER"	- Never try to find frameworks.

On Darwin or systems supporting OS X Application Bundles, the cmake variable `CMAKE_FIND_APPBUNDLE` can be set to empty or one of the following:

"FIRST"	- Try to find application bundles before standard programs. This is the default on Darwin.
"LAST"	- Try to find application bundles after standard programs.
"ONLY"	- Only try to find application bundles.
"NEVER"	- Never try to find application bundles.

The CMake variable `CMAKE_FIND_ROOT_PATH` specifies one or more directories to be prepended to all other search directories. This effectively "re-roots" the entire search under given locations. By default it is empty. It is especially useful when cross-compiling to point to the root directory of the target environment and CMake will search there too. By default at first the directories listed in `CMAKE_FIND_ROOT_PATH` and then the non-rooted directories will be searched. The default behavior can be adjusted by setting `CMAKE_FIND_ROOT_PATH_MODE_INCLUDE`. This behavior can be manually overridden on a per-call basis. By using `CMAKE_FIND_ROOT_PATH_BOTH` the search order will be as described above. If `NO_CMAKE_FIND_ROOT_PATH` is used then `CMAKE_FIND_ROOT_PATH` will not be used. If `ONLY_CMAKE_FIND_ROOT_PATH` is used then only the re-rooted directories will be searched.

The default search order is designed to be most-specific to least-specific for common use cases. Projects may override the order by simply calling the command multiple times and using the `NO_*` options:

```
find_file(<VAR> NAMES name PATHS paths... NO_DEFAULT_PATH)
find_file(<VAR> NAMES name)
```

Once one of the calls succeeds the result variable will be set and stored in the cache so that no call will search again.

find_library: Find a library.

```
find_library(<VAR> name1 [path1 path2 ...])
```

This is the short-hand signature for the command that is sufficient in many cases. It is the same as `find_library(<VAR> name1 [PATHS path1 path2 ...])`

```
find_library(
    <VAR>
    name | NAMES name1 [name2 ...]
    [HINTS path1 [path2 ... ENV var]]
    [PATHS path1 [path2 ... ENV var]]
    [PATH_SUFFIXES suffix1 [suffix2 ...]]
    [DOC "cache documentation string"]
    [NO_DEFAULT_PATH]
    [NO_CMAKE_ENVIRONMENT_PATH]
    [NO_CMAKE_PATH]
    [NO_SYSTEM_ENVIRONMENT_PATH]
    [NO_CMAKE_SYSTEM_PATH]
    [CMAKE_FIND_ROOT_PATH_BOTH |
     ONLY_CMAKE_FIND_ROOT_PATH |
     NO_CMAKE_FIND_ROOT_PATH]
)
```

This command is used to find a library. A cache entry named by `<VAR>` is created to store the result of this command. If the library is found the result is stored in the variable and the search will not be repeated unless the variable is cleared. If nothing is found, the result will be `<VAR>-NOTFOUND`, and the search will be attempted again the next time `find_library` is invoked with the same variable. The name of the library that is searched for is specified by the names listed after the `NAMES` argument. Additional search locations can be specified after the `PATHS` argument. If `ENV var` is found in the `HINTS` or `PATHS` section the environment variable `var` will be read and converted from a system environment variable to a cmake style list of paths. For example `ENV PATH` would be a way to list the system path variable. The argument after `DOC` will be used for the documentation string in the cache. `PATH_SUFFIXES` specifies additional subdirectories to check below each search path.

If `NO_DEFAULT_PATH` is specified, then no additional paths are added to the search. If `NO_DEFAULT_PATH` is not specified, the search process is as follows:

1. Search paths specified in cmake-specific cache variables. These are intended to be used on the command line with a `-DVAR=value`. This can be skipped if `NO_CMAKE_PATH` is passed.

```
<prefix>/lib for each <prefix> in CMAKE_PREFIX_PATH  
CMAKE_LIBRARY_PATH and CMAKE_FRAMEWORK_PATH
```

2. Search paths specified in cmake-specific environment variables. These are intended to be set in the user's shell configuration. This can be skipped if NO_CMAKE_ENVIRONMENT_PATH is passed.

```
<prefix>/lib for each <prefix> in CMAKE_PREFIX_PATH  
CMAKE_LIBRARY_PATH and CMAKE_FRAMEWORK_PATH
```

3. Search the paths specified by the HINTS option. These should be paths computed by system introspection, such as a hint provided by the location of another item already found. Hard-coded guesses should be specified with the PATHS option.

4. Search the standard system environment variables. This can be skipped if NO_SYSTEM_ENVIRONMENT_PATH is an argument. Specifically PATH and LIB.

5. Search cmake variables defined in the Platform files for the current system. This can be skipped if NO_CMAKE_SYSTEM_PATH is passed.

```
<prefix>/lib for each <prefix> in CMAKE_SYSTEM_PREFIX_PATH  
CMAKE_SYSTEM_LIBRARY_PATH and CMAKE_SYSTEM_FRAMEWORK_PATH
```

6. Search the paths specified by the PATHS option or in the short-hand version of the command. These are typically hard-coded guesses.

On Darwin or systems supporting OS X Frameworks, the cmake variable CMAKE_FIND_FRAMEWORK can be set to empty or one of the following:

```
"FIRST" - Try to find frameworks before standard  
           libraries or headers. This is the default on Darwin.  
"LAST"   - Try to find frameworks after standard  
           libraries or headers.  
"ONLY"   - Only try to find frameworks.  
"NEVER"  - Never try to find frameworks.
```

On Darwin or systems supporting OS X Application Bundles, the cmake variable CMAKE_FIND_APPBUNDLE can be set to empty or one of the following:

```
"FIRST" - Try to find application bundles before standard
           programs. This is the default on Darwin.
"LAST"   - Try to find application bundles after standard
           programs.
"ONLY"   - Only try to find application bundles.
"NEVER"  - Never try to find application bundles.
```

The CMake variable `CMAKE_FIND_ROOT_PATH` specifies one or more directories to be prepended to all other search directories. This effectively "re-roots" the entire search under given locations. By default it is empty. It is especially useful when cross-compiling to point to the root directory of the target environment and CMake will search there too. By default at first the directories listed in `CMAKE_FIND_ROOT_PATH` and then the non-rooted directories will be searched. The default behavior can be adjusted by setting `CMAKE_FIND_ROOT_PATH_MODE_LIBRARY`. This behavior can be manually overridden on a per-call basis. By using `CMAKE_FIND_ROOT_PATH_BOTH` the search order will be as described above. If `NO_CMAKE_FIND_ROOT_PATH` is used then `CMAKE_FIND_ROOT_PATH` will not be used. If `ONLY_CMAKE_FIND_ROOT_PATH` is used then only the re-rooted directories will be searched.

The default search order is designed to be most-specific to least-specific for common use cases. Projects may override the order by simply calling the command multiple times and using the `NO_*` options:

```
find_library(<VAR> NAMES name PATHS paths... NO_DEFAULT_PATH)
find_library(<VAR> NAMES name)
```

Once one of the calls succeeds the result variable will be set and stored in the cache so that no call will search again.

If the library found is a framework, then `VAR` will be set to the full path to the framework <fullPath>/A.framework. When a full path to a framework is used as a library, CMake will use a -framework A, and a -F<fullPath> to link the framework to the target.

find_package: Load settings for an external project.

```
find_package(<package> [version] [EXACT] [QUIET]
            [[REQUIRED|COMPONENTS] [components...]]
            [NO_POLICY_SCOPE])
```

Finds and loads settings from an external project. `<package>_FOUND` will be set to indicate whether the package was found. When the package is found package-specific information is provided through variables documented by the package itself. The `QUIET` option disables

messages if the package cannot be found. The REQUIRED option stops processing with an error message if the package cannot be found. A package-specific list of components may be listed after the REQUIRED option or after the COMPONENTS option if no REQUIRED option is given. The [version] argument requests a version with which the package found should be compatible (format is major[.minor[.patch[.tweak]]]). The EXACT option requests that the version be matched exactly. If no [version] is given to a recursive invocation inside a find-module, the [version] and EXACT arguments are forwarded automatically from the outer call. Version support is currently provided only on a package-by-package basis (details below).

User code should generally look for packages using the above simple signature. The remainder of this command documentation specifies the full command signature and details of the search process. Project maintainers wishing to provide a package to be found by this command are encouraged to read on.

The command has two modes by which it searches for packages: "Module" mode and "Config" mode. Module mode is available when the command is invoked with the above reduced signature. CMake searches for a file called "Find<package>.cmake" in the CMAKE_MODULE_PATH followed by the CMake installation. If the file is found, it is read and processed by CMake. It is responsible for finding the package, checking the version, and producing any needed messages. Many find-modules provide limited or no support for versioning; check the module documentation. If no module is found the command proceeds to Config mode.

The complete Config mode command signature is:

```
find_package(<package> [version] [EXACT] [QUIET]
             [[REQUIRED|COMPONENTS] [components...]] [NO_MODULE]
             [NO_POLICY_SCOPE]
             [NAMES name1 [name2 ...]]
             [CONFIGS config1 [config2 ...]]
             [HINTS path1 [path2 ... ]]
             [PATHS path1 [path2 ... ]]
             [PATH_SUFFIXES suffix1 [suffix2 ... ]]
             [NO_DEFAULT_PATH]
             [NO_CMAKE_ENVIRONMENT_PATH]
             [NO_CMAKE_PATH]
             [NO_SYSTEM_ENVIRONMENT_PATH]
             [NO_CMAKE_PACKAGE_REGISTRY]
             [NO_CMAKE_BUILDS_PATH]
             [NO_CMAKE_SYSTEM_PATH]
             [CMAKE_FIND_ROOT_PATH_BOTH |
              ONLY_CMAKE_FIND_ROOT_PATH |
              NO_CMAKE_FIND_ROOT_PATH])
```

The NO_MODULE option may be used to skip Module mode explicitly. It is also implied by use of options not specified in the reduced signature.

Config mode attempts to locate a configuration file provided by the package to be found. A cache entry called <package>_DIR is created to hold the directory containing the file. By default the command searches for a package with the name <package>. If the NAMES option is given the names following it are used instead of <package>. The command searches for a file called "<name>Config.cmake" or "<lower-case-name>-config.cmake" for each name specified. A replacement set of possible configuration file names may be given using the CONFIGS option. The search procedure is specified below. Once found, the configuration file is read and processed by CMake. Since the file is provided by the package it already knows the location of package contents. The full path to the configuration file is stored in the cmake variable <package>_CONFIG.

If the package configuration file cannot be found CMake will generate an error describing the problem unless the QUIET argument is specified. If REQUIRED is specified and the package is not found a fatal error is generated and the configure step stops executing. If <package>_DIR has been set to a directory not containing a configuration file CMake will ignore it and search from scratch.

When the [version] argument is given Config mode will only find a version of the package that claims compatibility with the requested version (format is major[.minor[.patch[.tweak]]]]. If the EXACT option is given only a version of the package claiming an exact match of the requested version may be found. CMake does not establish any convention for the meaning of version numbers. Package version numbers are checked by "version" files provided by the packages themselves. For a candidate package configuration file "<config-file>.cmake" the corresponding version file is located next to it and named either "<config-file>-version.cmake" or "<config-file>Version.cmake". If no such version file is available then the configuration file is assumed to not be compatible with any requested version. When a version file is found it is loaded to check the requested version number. The version file is loaded in a nested scope in which the following variables have been defined:

PACKAGE_FIND_NAME	= the <package> name
PACKAGE_FIND_VERSION	= full requested version string
PACKAGE_FIND_VERSION_MAJOR	= major version if requested, else 0
PACKAGE_FIND_VERSION_MINOR	= minor version if requested, else 0
PACKAGE_FIND_VERSION_PATCH	= patch version if requested, else 0
PACKAGE_FIND_VERSION_TWEAK	= tweak version if requested, else 0
PACKAGE_FIND_VERSION_COUNT	= number of version components 0 to 4

The version file checks whether it satisfies the requested version and sets these variables:

PACKAGE_VERSION	= full provided version string
PACKAGE_VERSION_EXACT	= true if version is exact match
PACKAGE_VERSION_COMPATIBLE	= true if version is compatible
PACKAGE_VERSION_UNSUITABLE	= true if unsuitable as any version

These variables are checked by the `find_package` command to determine whether the configuration file provides an acceptable version. They are not available after the `find_package` call returns. If the version is acceptable the following variables are set:

<package>_VERSION	= full provided version string
<package>_VERSION_MAJOR	= major version if provided, else 0
<package>_VERSION_MINOR	= minor version if provided, else 0
<package>_VERSION_PATCH	= patch version if provided, else 0
<package>_VERSION_TWEAK	= tweak version if provided, else 0
<package>_VERSION_COUNT	= number of version components, 0 to 4

and the corresponding package configuration file is loaded. When multiple package configuration files are available whose version files claim compatibility with the version requested it is unspecified which one is chosen. No attempt is made to choose a highest or closest version number.

Config mode provides an elaborate interface and search procedure. Much of the interface is provided for completeness and for use internally by find-modules loaded by Module mode. Most user code should simply call

```
find_package(<package> [major[.minor]] [EXACT] [REQUIRED|QUIET])
```

in order to find a package. Package maintainers providing CMake package configuration files are encouraged to name and install them such that the procedure outlined below will find them without requiring use of additional options.

CMake constructs a set of possible installation prefixes for the package. Under each prefix several directories are searched for a configuration file. The tables below show the directories searched. Each entry is meant for installation trees following Windows (W), UNIX (U), or Apple (A) conventions.

<prefix>/	(W)
<prefix>/ (cmake CMake) /	(W)
<prefix>/<name>*/	(W)
<prefix>/<name>*/ (cmake CMake) /	(W)
<prefix>/ (share lib) /cmake/<name>*/	(U)
<prefix>/ (share lib) /<name>*/	(U)
<prefix>/ (share lib) /<name>*/ (cmake CMake) /	(U)

On systems supporting OS X Frameworks and Application Bundles the following directories are searched for frameworks or bundles containing a configuration file:

<prefix>/<name>.framework/Resources/	(A)
<prefix>/<name>.framework/Resources/CMake/	(A)
<prefix>/<name>.framework/Versions/*/Resources/	(A)
<prefix>/<name>.framework/Versions/*/Resources/CMake/	(A)
<prefix>/<name>.app/Contents/Resources/	(A)
<prefix>/<name>.app/Contents/Resources/CMake/	(A)

In all cases the <name> is treated as case-insensitive and corresponds to any of the names specified (<package> or names given by NAMES). If PATH_SUFFIXES is specified the suffixes are appended to each (W) or (U) directory entry one-by-one.

This set of directories is intended to work in cooperation with projects that provide configuration files in their installation trees. Directories above marked with (W) are intended for installations on Windows where the prefix may point at the top of an application's installation directory. Those marked with (U) are intended for installations on UNIX platforms where the prefix is shared by multiple packages. This is merely a convention, so all (W) and (U) directories are still searched on all platforms. Directories marked with (A) are intended for installations on Apple platforms. The cmake variables CMAKE_FIND_FRAMEWORK and CMAKE_FIND_APPBUNDLE determine the order of preference as specified below.

The set of installation prefixes is constructed using the following steps. If NO_DEFAULT_PATH is specified all NO_* options are enabled.

1. Search paths specified in cmake-specific cache variables. These are intended to be used on the command line with a -DVAR=value. This can be skipped if NO_CMAKE_PATH is passed. Specifically CMAKE_PREFIX_PATH CMAKE_FRAMEWORK_PATH and CMAKE_APPBUNDLE_PATH.
2. Search paths specified in cmake-specific environment variables. These are intended to be set in the user's shell configuration. This can be skipped if

`NO_CMAKE_ENVIRONMENT_PATH` is passed. Specifically `CMAKE_PREFIX_PATH`, `CMAKE_FRAMEWORK_PATH` and `CMAKE_APPBUNDLE_PATH`.

3. Search paths specified by the `HINTS` option. These should be paths computed by system introspection, such as a hint provided by the location of another item already found. Hard-coded guesses should be specified with the `PATHS` option.
4. Search the standard system environment variables. This can be skipped if `NO_SYSTEM_ENVIRONMENT_PATH` is passed. Path entries ending in `"/bin"` or `"/sbin"` are automatically converted to their parent directories. Specifically: `PATH`.
5. Search project build trees recently configured in a CMake GUI. This can be skipped if `NO_CMAKE_BUILDS_PATH` is passed. It is intended for the case when a user is building multiple dependent projects one after another.
6. Search paths stored in the CMake user package registry. This can be skipped if `NO_CMAKE_PACKAGE_REGISTRY` is passed. Paths are stored in the registry when CMake configures a project that invokes `export(PACKAGE <name>)`. See the `export(PACKAGE)` command documentation for more details.
7. Search `cmake` variables defined in the Platform files for the current system. This can be skipped if `NO_CMAKE_SYSTEM_PATH` is passed. Specifically: `CMAKE_SYSTEM_PREFIX_PATH`, `CMAKE_SYSTEM_FRAMEWORK_PATH` and `CMAKE_SYSTEM_APPBUNDLE_PATH`
8. Search paths specified by the `PATHS` option. These are typically hard-coded guesses.

On Darwin or systems supporting OS X Frameworks, the `cmake` variable `CMAKE_FIND_FRAMEWORK` can be set to empty or one of the following:

"FIRST"	- Try to find frameworks before standard libraries or headers. This is the default on Darwin.
"LAST"	- Try to find frameworks after standard libraries or headers.
"ONLY"	- Only try to find frameworks.
"NEVER"	- Never try to find frameworks.

On Darwin or systems supporting OS X Application Bundles, the `cmake` variable `CMAKE_FIND_APPBUNDLE` can be set to empty or one of the following:

"FIRST"	- Try to find application bundles before standard programs. This is the default on Darwin.
"LAST"	- Try to find application bundles after standard programs.
"ONLY"	- Only try to find application bundles.
"NEVER"	- Never try to find application bundles.

The CMake variable `CMAKE_FIND_ROOT_PATH` specifies one or more directories to be prepended to all other search directories. This effectively "re-roots" the entire search under given locations. By default it is empty. It is especially useful when cross-compiling to point to the root directory of the target environment and CMake will search there too. By default at first the directories listed in `CMAKE_FIND_ROOT_PATH` and then the non-rooted directories will be searched. The default behavior can be adjusted by setting `CMAKE_FIND_ROOT_PATH_MODE_PACKAGE`. This behavior can be manually overridden on a per-call basis. By using `CMAKE_FIND_ROOT_PATH_BOTH` the search order will be as described above. If `NO_CMAKE_FIND_ROOT_PATH` is used then `CMAKE_FIND_ROOT_PATH` will not be used. If `ONLY_CMAKE_FIND_ROOT_PATH` is used then only the re-rooted directories will be searched.

The default search order is designed to be most-specific to least-specific for common use cases. Projects may override the order by simply calling the command multiple times and using the `NO_*` options:

```
find_package(<package> PATHS paths... NO_DEFAULT_PATH)
find_package(<package>)
```

Once one of the calls succeeds the result variable will be set and stored in the cache so that no call will search again. See the `cmake_policy()` command documentation for discussion of the `NO_POLICY_SCOPE` option.

find_path: Find the directory containing a file.

```
find_path(<VAR> name1 [path1 path2 ...])
```

This is the short-hand signature for the command that is sufficient in many cases. It is the same as `find_path(<VAR> name1 [PATHS path1 path2 ...])`

```

find_path(
    <VAR>
    NAMES name1 [name2 ...]
    [HINTS path1 [path2 ... ENV var]]
    [PATHS path1 [path2 ... ENV var]]
    [PATH_SUFFIXES suffix1 [suffix2 ...]]
    [DOC "cache documentation string"]
    [NO_DEFAULT_PATH]
    [NO_CMAKE_ENVIRONMENT_PATH]
    [NO_CMAKE_PATH]
    [NO_SYSTEM_ENVIRONMENT_PATH]
    [NO_CMAKE_SYSTEM_PATH]
    [CMAKE_FIND_ROOT_PATH_BOTH |
     ONLY_CMAKE_FIND_ROOT_PATH |
     NO_CMAKE_FIND_ROOT_PATH]
)

```

This command is used to find a directory containing the named file. A cache entry named by <VAR> is created to store the result of this command. If the file in a directory is found the result is stored in the variable and the search will not be repeated unless the variable is cleared. If nothing is found, the result will be <VAR>-NOTFOUND, and the search will be attempted again the next time `find_path` is invoked with the same variable. The name of the file in a directory that is searched for is specified by the names listed after the NAMES argument. Additional search locations can be specified after the PATHS argument. If ENV var is found in the HINTS or PATHS section the environment variable var will be read and converted from a system environment variable to a cmake style list of paths. For example ENV PATH would be a way to list the system path variable. The argument after DOC will be used for the documentation string in the cache. PATH_SUFFIXES specifies additional subdirectories to check below each search path.

If NO_DEFAULT_PATH is specified, then no additional paths are added to the search. If NO_DEFAULT_PATH is not specified, the search process is as follows:

1. Search paths specified in cmake-specific cache variables. These are intended to be used on the command line with a -DVAR=value. This can be skipped if NO_CMAKE_PATH is passed.

```

<prefix>/include for each <prefix> in CMAKE_PREFIX_PATH
CMAKE_INCLUDE_PATH and CMAKE_FRAMEWORK_PATH

```

2. Search paths specified in cmake-specific environment variables. These are intended to be set in the user's shell configuration. This can be skipped if NO_CMAKE_ENVIRONMENT_PATH is passed.

```
<prefix>/include for each <prefix> in CMAKE_PREFIX_PATH
CMAKE_INCLUDE_PATH and CMAKE_FRAMEWORK_PATH
```

3. Search the paths specified by the HINTS option. These should be paths computed by system introspection, such as a hint provided by the location of another item already found. Hard-coded guesses should be specified with the PATHS option.
4. Search the standard system environment variables. This can be skipped if NO_SYSTEM_ENVIRONMENT_PATH is an argument. Specifically: PATH and INCLUDE
5. Search cmake variables defined in the Platform files for the current system. This can be skipped if NO_CMAKE_SYSTEM_PATH is passed.

```
<prefix>/include for each <prefix> in CMAKE_SYSTEM_PREFIX_PATH
CMAKE_SYSTEM_INCLUDE_PATH and CMAKE_SYSTEM_FRAMEWORK_PATH
```

6. Search the paths specified by the PATHS option or in the short-hand version of the command. These are typically hard-coded guesses.

On Darwin or systems supporting OS X Frameworks, the cmake variable CMAKE_FIND_FRAMEWORK can be set to empty or one of the following:

<ul style="list-style-type: none"> "FIRST" - Try to find frameworks before standard libraries or headers. This is the default on Darwin.
<ul style="list-style-type: none"> "LAST" - Try to find frameworks after standard libraries or headers.
<ul style="list-style-type: none"> "ONLY" - Only try to find frameworks.
<ul style="list-style-type: none"> "NEVER" - Never try to find frameworks.

On Darwin or systems supporting OS X Application Bundles, the cmake variable CMAKE_FIND_APPBUNDLE can be set to empty or one of the following:

<ul style="list-style-type: none"> "FIRST" - Try to find application bundles before standard programs. This is the default on Darwin.
<ul style="list-style-type: none"> "LAST" - Try to find application bundles after standard programs.
<ul style="list-style-type: none"> "ONLY" - Only try to find application bundles.
<ul style="list-style-type: none"> "NEVER" - Never try to find application bundles.

The CMake variable CMAKE_FIND_ROOT_PATH specifies one or more directories to be prepended to all other search directories. This effectively "re-roots" the entire search under

given locations. By default it is empty. It is especially useful when cross-compiling to point to the root directory of the target environment and CMake will search there too. By default at first the directories listed in `CMAKE_FIND_ROOT_PATH` and then the non-rooted directories will be searched. The default behavior can be adjusted by setting `CMAKE_FIND_ROOT_PATH_MODE_INCLUDE`. This behavior can be manually overridden on a per-call basis. By using `CMAKE_FIND_ROOT_PATH_BOTH` the search order will be as described above. If `NO_CMAKE_FIND_ROOT_PATH` is used then `CMAKE_FIND_ROOT_PATH` will not be used. If `ONLY_CMAKE_FIND_ROOT_PATH` is used then only the re-rooted directories will be searched.

The default search order is designed to be most-specific to least-specific for common use cases. Projects may override the order by simply calling the command multiple times and using the `NO_*` options:

```
find_path(<VAR> NAMES name PATHS paths... NO_DEFAULT_PATH)
find_path(<VAR> NAMES name)
```

Once one of the calls succeeds the result variable will be set and stored in the cache so that no call will search again.

When searching for frameworks, if the file is specified as `A/b.h`, then the framework search will look for `A.framework/Headers/b.h`. If that is found the path will be set to the path to the framework. CMake will convert this to the correct `-F` option to include the file.

find_program: Find an executable program.

```
find_program(<VAR> name1 [path1 path2 ...])
```

This is the short-hand signature for the command that is sufficient in many cases. It is the same as `find_program(<VAR> name1 [PATHS path1 path2 ...])`

```
find_program(
    <VAR>
    name | NAMES name1 [name2 ...]
    [HINTS path1 [path2 ... ENV var]]
    [PATHS path1 [path2 ... ENV var]]
    [PATH_SUFFIXES suffix1 [suffix2 ...]]
    [DOC "cache documentation string"]
    [NO_DEFAULT_PATH]
    [NO_CMAKE_ENVIRONMENT_PATH]
    [NO_CMAKE_PATH]
    [NO_SYSTEM_ENVIRONMENT_PATH]
```

```
[NO_CMAKE_SYSTEM_PATH]
[CMAKE_FIND_ROOT_PATH_BOTH | 
 ONLY_CMAKE_FIND_ROOT_PATH | 
 NO_CMAKE_FIND_ROOT_PATH]
)
```

This command is used to find a program. A cache entry named by <VAR> is created to store the result of this command. If the program is found the result is stored in the variable and the search will not be repeated unless the variable is cleared. If nothing is found, the result will be <VAR>-NOTFOUND, and the search will be attempted again the next time `find_program` is invoked with the same variable. The name of the program that is searched for is specified by the names listed after the NAMES argument. Additional search locations can be specified after the PATHS argument. If ENV var is found in the HINTS or PATHS section the environment variable var will be read and converted from a system environment variable to a cmake style list of paths. For example ENV PATH would be a way to list the system path variable. The argument after DOC will be used for the documentation string in the cache. PATH_SUFFIXES specifies additional subdirectories to check below each search path.

If NO_DEFAULT_PATH is specified, then no additional paths are added to the search. If NO_DEFAULT_PATH is not specified, the search process is as follows:

1. Search paths specified in cmake-specific cache variables. These are intended to be used on the command line with a -DVAR=value. This can be skipped if NO_CMAKE_PATH is passed.

```
<prefix>/[s]bin for each <prefix> in CMAKE_PREFIX_PATH
CMAKE_PROGRAM_PATH and CMAKE_APPBUNDLE_PATH
```

2. Search paths specified in cmake-specific environment variables. These are intended to be set in the user's shell configuration. This can be skipped if NO_CMAKE_ENVIRONMENT_PATH is passed.

```
<prefix>/[s]bin for each <prefix> in CMAKE_PREFIX_PATH
CMAKE_PROGRAM_PATH and CMAKE_APPBUNDLE_PATH
```

3. Search the paths specified by the HINTS option. These should be paths computed by system introspection, such as a hint provided by the location of another item already found. Hard-coded guesses should be specified with the PATHS option.

4. Search the standard system environment variables. This can be skipped if NO_SYSTEM_ENVIRONMENT_PATH is an argument. Specifically: PATH.

5. Search cmake variables defined in the Platform files for the current system. This can be skipped if NO_CMAKE_SYSTEM_PATH is passed.

```
<prefix>/[s]bin for each <prefix> in CMAKE_SYSTEM_PREFIX_PATH  
CMAKE_SYSTEM_PROGRAM_PATH and CMAKE_SYSTEM_APPBUNDLE_PATH
```

6. Search the paths specified by the PATHS option or in the short-hand version of the command. These are typically hard-coded guesses.

On Darwin or systems supporting OS X Frameworks, the cmake variable CMAKE_FIND_FRAMEWORK can be set to empty or one of the following:

"FIRST"	- Try to find frameworks before standard libraries or headers. This is the default on Darwin.
"LAST"	- Try to find frameworks after standard libraries or headers.
"ONLY"	- Only try to find frameworks.
"NEVER"	- Never try to find frameworks.

On Darwin or systems supporting OS X Application Bundles, the cmake variable CMAKE_FIND_APPBUNDLE can be set to empty or one of the following:

"FIRST"	- Try to find application bundles before standard programs. This is the default on Darwin.
"LAST"	- Try to find application bundles after standard programs.
"ONLY"	- Only try to find application bundles.
"NEVER"	- Never try to find application bundles.

The CMake variable CMAKE_FIND_ROOT_PATH specifies one or more directories to be prepended to all other search directories. This effectively "re-roots" the entire search under given locations. By default it is empty. It is especially useful when cross-compiling to point to the root directory of the target environment and CMake will search there too. By default at first the directories listed in CMAKE_FIND_ROOT_PATH and then the non-rooted directories will be searched. The default behavior can be adjusted by setting CMAKE_FIND_ROOT_PATH_MODE_PROGRAM. This behavior can be manually overridden on a per-call basis. By using CMAKE_FIND_ROOT_PATH_BOTH the search order will be as described above. If NO_CMAKE_FIND_ROOT_PATH is used then CMAKE_FIND_ROOT_PATH will not be used. If ONLY_CMAKE_FIND_ROOT_PATH is used then only the re-rooted directories will be searched.

The default search order is designed to be most-specific to least-specific for common use cases. Projects may override the order by simply calling the command multiple times and using the NO_* options:

```
find_program(<VAR> NAMES name PATHS paths... NO_DEFAULT_PATH)
find_program(<VAR> NAMES name)
```

Once one of the calls succeeds the result variable will be set and stored in the cache so that no call will search again.

fltk_wrap_ui: Create FLTK user interfaces Wrappers.

```
fltk_wrap_ui(resultingLibraryName source1
              source2 ... sourceN )
```

Produce .h and .cxx files for all the .fl and .fld files listed. The resulting .h and .cxx files will be added to a variable named resultingLibraryName_FLTK_UI_SRCS which should be added to your library.

foreach: Evaluate a group of commands for each value in a list.

```
foreach(loop_var arg1 arg2 ...)
    COMMAND1 (ARGS ...)
    COMMAND2 (ARGS ...)
    ...
endforeach([loop_var])
```

All commands between foreach and the matching endforeach are recorded without being invoked. Once the endforeach is evaluated, the recorded list of commands is invoked once for each argument listed in the original foreach command. Before each iteration of the loop "\${loop_var}" will be set as a variable with the current value in the list.

```
foreach(loop_var RANGE total)
foreach(loop_var RANGE start stop [step])
```

Foreach can also iterate over a generated range of numbers. There are three types of this iteration:

* When specifying single number, the range will have elements 0 to "total".

* When specifying two numbers, the range will have elements from the first number to the second number.

* The third optional number is the increment used to iterate from the first number to the second number.

```
foreach(loop_var IN [LISTS [list1 [...]]]
      [ITEMS [item1 [...]]])
```

Iterates over a precise list of items. The LISTS option names list-valued variables to be traversed, including empty elements (an empty string is a zero-length list). The ITEMS option ends argument parsing and includes all arguments following it in the iteration.

function: Start recording a function for later invocation as a command.

```
function(<name> [arg1 [arg2 [arg3 ...]])  
  COMMAND1 (ARGS ...)  
  COMMAND2 (ARGS ...)  
  ...  
endfunction ([<name>])
```

Define a function named <name> that takes arguments named arg1 arg2 arg3 (...). Commands listed after function, but before the matching endfunction, are not invoked until the function is invoked. When it is invoked, the commands recorded in the function are first modified by replacing formal parameters (\${arg1}) with the arguments passed, and then invoked as normal commands. In addition to referencing the formal parameters you can reference the variable ARGC which will be set to the number of arguments passed into the function as well as ARGV0 ARGV1 ARGV2 ... which will have the actual values of the arguments passed in. This facilitates creating functions with optional arguments. Additionally ARGV holds the list of all arguments given to the function and ARGN holds the list of argument past the last expected argument.

See the cmake_policy() command documentation for the behavior of policies inside functions.

get_cmake_property: Get a property of the CMake instance.

```
get_cmake_property(VAR property)
```

Get a property from the CMake instance. The value of the property is stored in the variable VAR. If the property is not found, CMake will report an error. Some supported properties

include: VARIABLES, CACHE_VARIABLES, COMMANDS, MACROS, and COMPONENTS.

get_directory_property: Get a property of DIRECTORY scope.

```
get_directory_property(<variable> [DIRECTORY <dir>] <prop-name>)
```

Store a property of directory scope in the named variable. If the property is not defined the empty-string is returned. The DIRECTORY argument specifies another directory from which to retrieve the property value. The specified directory must have already been traversed by CMake.

```
get_directory_property(<variable> [DIRECTORY <dir>]  
                      DEFINITION <var-name>)
```

Get a variable definition from a directory. This form is useful to get a variable definition from another directory.

get_filename_component: Get a specific component of a full filename.

```
get_filename_component(VarName FileName  
                      PATH|ABSOLUTE|NAME|EXT|NAME_WE|REALPATH  
                      [CACHE])
```

Set VarName to be the path (PATH), file name (NAME), file extension (EXT), file name without extension (NAME_WE) of FileName, the full path (ABSOLUTE), or the full path with all symlinks resolved (REALPATH). Note that the path is converted to UNIX slashes format and has no trailing slashes. The longest file extension is always considered. If the optional CACHE argument is specified, the result variable is added to the cache.

```
get_filename_component(VarName FileName  
                      PROGRAM [PROGRAM_ARGS ArgVar]  
                      [CACHE])
```

The program in FileName will be found in the system search path or left as a full path. If PROGRAM_ARGS is present with PROGRAM, then any command-line arguments present in the FileName string are split from the program name and stored in ArgVar. This is used to separate a program name from its arguments in a command line string.

get_property: Get a property.

```
get_property(<variable>
    <GLOBAL           |
    DIRECTORY [dir]   |
    TARGET     <target> |
    SOURCE      <source> |
    TEST        <test>  |
    CACHE       <entry>  |
    VARIABLE>
PROPERTY <name>
[SET | DEFINED | BRIEF_DOCS | FULL_DOCS])
```

Get one property from one object in a scope. The first argument specifies the variable in which to store the result. The second argument determines the scope from which to get the property. It must be one of the following:

GLOBAL scope is unique and does not accept a name.

DIRECTORY scope defaults to the current directory but another directory (already processed by CMake) may be named by full or relative path.

TARGET scope must name one existing target.

SOURCE scope must name one source file.

TEST scope must name one existing test.

CACHE scope must name one cache entry.

VARIABLE scope is unique and does not accept a name.

The required PROPERTY option is immediately followed by the name of the property to get. If the property is not set an empty value is returned. If the SET option is given the variable is set to a boolean value indicating whether the property has been set. If the DEFINED option is given the variable is set to a boolean value indicating whether the property has been defined such as with define_property. If BRIEF_DOCS or FULL_DOCS is given then the variable is set to a string containing documentation for the requested property. If documentation is requested for a property that has not been defined NOTFOUND is returned.

get_source_file_property: Get a property for a source file.

```
get_source_file_property(VAR file property)
```

Get a property from a source file. The value of the property is stored in the variable VAR. If the property is not found, VAR will be set to "NOTFOUND". Use set_source_files_properties to set property values. Source file properties usually control how the file is built. One property that is always there is LOCATION

get_target_property: Get a property from a target.

```
get_target_property(VAR target property)
```

Get a property from a target. The value of the property is stored in the variable VAR. If the property is not found, VAR will be set to "NOTFOUND". Use set_target_properties to set property values. Properties are usually used to control how a target is built, but some query the target instead. This command can get properties for any target so far created. The targets do not need to be in the current CMakeLists.txt file.

get_test_property: Get a property of the test.

```
get_test_property(test VAR property)
```

Get a property from the Test. The value of the property is stored in the variable VAR. If the property is not found, CMake will report an error. For a list of standard properties you can type cmake --help-property-list

if: Conditionally execute a group of commands.

```
if(expression)
  # then section.
  COMMAND1 (ARGS ...)
  COMMAND2 (ARGS ...)
  ...
elseif(expression2)
  # elseif section.
  COMMAND1 (ARGS ...)
  COMMAND2 (ARGS ...)
  ...
else([expression2])
  # else section.
  COMMAND1 (ARGS ...)
  COMMAND2 (ARGS ...)
  ...
endif([expression2])
```

Evaluates the given expression. If the result is true, the commands in the THEN section are invoked. Otherwise, the commands in the else section are invoked. The elseif and else sections are optional. You may have multiple elseif clauses. Note that the expression in the else and endif clause is optional. Long expressions can be used and there is a traditional order of precedence. Parenthetical expressions are evaluated first followed by unary operators such as EXISTS, COMMAND, and DEFINED. Then any EQUAL, LESS, GREATER, STRLESS, STRGREATER, STREQUAL, MATCHES will be evaluated. Then NOT operators and finally AND, OR operators will be evaluated. Possible expressions are:

```
if(variable)
```

True if the variable's value is not empty, 0, N, NO, OFF, FALSE, NOTFOUND, or <variable>-NOTFOUND.

```
if(NOT variable)
```

True if the variable's value is empty, 0, N, NO, OFF, FALSE, NOTFOUND, or <variable>-NOTFOUND.

```
if(variable1 AND variable2)
```

True if both variables would be considered true individually.

```
if(variable1 OR variable2)
```

True if either variable would be considered true individually.

```
if(COMMAND command-name)
```

True if the given name is a command, macro or function that can be invoked.

```
if(POLICY policy-id)
```

True if the given name is an existing policy (of the form CMP<NNNN>).

```
if(TARGET target-name)
```

True if the given name is an existing target, built or imported.

```
if(EXISTS file-name)
if(EXISTS directory-name)
```

True if the named file or directory exists. Behavior is well-defined only for full paths.

```
if(file1 IS_NEWER_THAN file2)
```

True if file1 is newer than file2 or if one of the two files doesn't exist. Behavior is well-defined only for full paths.

```
if(IS_DIRECTORY directory-name)
```

True if the given name is a directory. Behavior is well-defined only for full paths.

```
if(IS_ABSOLUTE path)
```

True if the given path is an absolute path.

```
if(variable MATCHES regex)
if(string MATCHES regex)
```

True if the given string or variable's value matches the given regular expression.

```
if(variable LESS number)
if(string LESS number)
if(variable GREATER number)
if(string GREATER number)
if(variable EQUAL number)
if(string EQUAL number)
```

True if the given string or variable's value is a valid number and the inequality or equality is true.

```
if(variable STRLESS string)
if(string STRLESS string)
if(variable STRGREATER string)
if(string STRGREATER string)
```

```
if(variable STREQUAL string)
if(string STREQUAL string)
```

True if the given string or variable's value is lexicographically less (or greater, or equal) than the string or variable on the right.

```
if(version1 VERSION_LESS version2)
if(version1 VERSION_EQUAL version2)
if(version1 VERSION_GREATER version2)
```

Component-wise integer version number comparison (version format is major[minor[.patch[.tweak]]]).

```
if(DEFINED variable)
```

True if the given variable is defined. It does not matter if the variable is true or false just if it has been set.

```
if((expression) AND (expression OR (expression)))
```

The expressions inside the parenthesis are evaluated first and then the remaining expression is evaluated as in the previous examples. Where there are nested parenthesis the innermost are evaluated as part of evaluating the expression that contains them.

The if statement was written fairly early in CMake's history and it has some convenience features that are worth covering. The if statement reduces operations until there is a single remaining value, at that point if the case insensitive value is: ON, 1, YES, TRUE, Y it returns true, if it is OFF, 0, NO, FALSE, N, NOTFOUND, *-NOTFOUND, IGNORE it will return false.

This is fairly reasonable. The convenience feature that sometimes throws new authors is how CMake handles values that do not match the true or false list. Those values are treated as variables and are dereferenced even though they do not have the required \${} syntax. This means that if you write

```
if (boobah)
```

CMake will treat it as if you wrote

```
if (${boobah})
```

likewise if you write

```
if (fubar AND sol)
```

CMake will conveniently treat it as

```
if ("${fubar}" AND "${sol}")
```

The later is really the correct way to write it, but the former will work as well. Only some operations in the if statement have this special handling of arguments. The specific details follow:

- 1) The left hand argument to MATCHES is first checked to see if it is a defined variable, if so the variable's value is used, otherwise the original value is used.
- 2) If the left hand argument to MATCHES is missing it returns false without error
- 3) Both left and right hand arguments to LESS GREATER EQUAL are independently tested to see if they are defined variables, if so their defined values are used otherwise the original value is used.
- 4) Both left and right hand arguments to STRLESS STREQUAL STRGREATER are independently tested to see if they are defined variables, if so their defined values are used otherwise the original value is used.
- 5) Both left and right hand arguments to VERSION_LESS VERSION_EQUAL VERSION_GREATER are independently tested to see if they are defined variables, if so their defined values are used otherwise the original value is used.
- 6) The right hand argument to NOT is tested to see if it is a boolean constant, if so the value is used, otherwise it is assumed to be a variable and it is dereferenced.
- 7) The left and right hand arguments to AND OR are independently tested to see if they are boolean constants, if so they are used as such, otherwise they are assumed to be variables and are dereferenced.

include: Read CMake listfile code from the given file.

```
include(<file|module> [OPTIONAL] [RESULT_VARIABLE <VAR>]
       [NO_POLICY_SCOPE])
```

Reads CMake listfile code from the given file. Commands in the file are processed immediately as if they were written in place of the include command. If OPTIONAL is present, then no error is raised if the file does not exist. If RESULT_VARIABLE is given the variable will be set to the full filename which has been included or NOTFOUND if it failed.

If a module is specified instead of a file, the file with name <modulename>.cmake is searched in the CMAKE_MODULE_PATH.

See the cmake_policy() command documentation for discussion of the NO_POLICY_SCOPE option.

include_directories: Add include directories to the build.

```
include_directories([AFTER|BEFORE] [$SYSTEM] dir1 dir2 ...)
```

Add the given directories to those searched by the compiler for include files. By default the directories are appended onto the current list of directories. This default behavior can be changed by setting CMAKE_include_directories_BEFORE to ON. By using BEFORE or AFTER you can select between appending and prepending, independent from the default. If the SYSTEM option is given the compiler will be told that the directories are meant as system include directories on some platforms.

include_external_msproject: Include an external Microsoft project file in a workspace.

```
include_external_msproject(projectname location
                           dep1 dep2 ...)
```

Includes an external Microsoft project in the generated workspace file. Currently does nothing on UNIX. This will create a target named INCLUDE_EXTERNAL_MSPROJECT_[projectname]. This can be used in the add_dependencies command to make things depend on the external project.

include_regular_expression: Set the regular expression used for dependency checking.

```
include_regular_expression(regex_match [regex_complain])
```

Set the regular expressions used in dependency checking. Only files matching regex_match will be traced as dependencies. Only files matching regex_complain will generate warnings if they cannot be found (standard header paths are not searched). The defaults are:

```
regex_match      = "^.*$" (match everything)
regex_complain = "^\$" (match empty string only)
```

install: Specify rules to run at install time.

This command generates installation rules for a project. Rules specified by calls to this command within a source directory are executed in order during installation. The order across directories is not defined.

There are multiple signatures for this command. Some of them define installation properties for files and targets. Properties common to multiple signatures are covered here but they are valid only for signatures that specify them.

DESTINATION arguments specify the directory on disk to which a file will be installed. If a full path (with a leading slash or drive letter) is given it is used directly. If a relative path is given it is interpreted relative to the value of CMAKE_INSTALL_PREFIX.

PERMISSIONS arguments specify permissions for installed files. Valid permissions are OWNER_READ, OWNER_WRITE, OWNER_EXECUTE, GROUP_READ, GROUP_WRITE, GROUP_EXECUTE, WORLD_READ, WORLD_WRITE, WORLD_EXECUTE, SETUID, and SETGID. Permissions that do not make sense on certain platforms are ignored on those platforms.

The CONFIGURATIONS argument specifies a list of build configurations for which the install rule applies (Debug, Release, etc.).

The COMPONENT argument specifies an installation component name with which the install rule is associated, such as "runtime" or "development". During component-specific installation only install rules associated with the given component name will be executed. During a full installation all components are installed.

The RENAME argument specifies a name for an installed file that may be different from the original file. Renaming is allowed only when a single file is installed by the command.

The OPTIONAL argument specifies that it is not an error if the file to be installed does not exist.

The TARGETS signature:

```
install(TARGETS targets... [EXPORT <export-name>]
        [ [ARCHIVE|LIBRARY|RUNTIME|FRAMEWORK|BUNDLE|
          PRIVATE_HEADER|PUBLIC_HEADER|RESOURCE]
        [DESTINATION <dir>]
        [PERMISSIONS permissions...]
        [CONFIGURATIONS [Debug|Release|...]]
        [COMPONENT <component>]
        [OPTIONAL] [NAMELINK_ONLY|NAMELINK_SKIP]
      ] [...])
```

The TARGETS form specifies rules for installing targets from a project. There are five kinds of target files that may be installed: ARCHIVE, LIBRARY, RUNTIME, FRAMEWORK, and BUNDLE. Executables are treated as RUNTIME targets, except that those marked with the MACOSX_BUNDLE property are treated as BUNDLE targets on OS X. Static libraries are always treated as ARCHIVE targets. Module libraries are always treated as LIBRARY targets. For non-DLL platforms shared libraries are treated as LIBRARY targets, except that those marked with the FRAMEWORK property are treated as FRAMEWORK targets on OS X. For DLL platforms the DLL part of a shared library is treated as a RUNTIME target and the corresponding import library is treated as an ARCHIVE target. All Windows-based systems including Cygwin are DLL platforms. The ARCHIVE, LIBRARY, RUNTIME, and FRAMEWORK arguments change the type of target to which the subsequent properties apply. If none is given the installation properties apply to all target types. If only one is given then only targets of that type will be installed (which can be used to install just a DLL or just an import library).

The PRIVATE_HEADER, PUBLIC_HEADER, and RESOURCE arguments cause subsequent properties to be applied to installing a FRAMEWORK shared library target's associated files on non-Apple platforms. Rules defined by these arguments are ignored on Apple platforms because the associated files are installed into the appropriate locations inside the framework folder. See documentation of the PRIVATE_HEADER, PUBLIC_HEADER, and RESOURCE target properties for details.

Either NAMELINK_ONLY or NAMELINK_SKIP may be specified as a LIBRARY option. On some platforms a versioned shared library has a symbolic link such as

```
lib<name>.so -> lib<name>.so.1
```

where "lib<name>.so.1" is the soname of the library and "lib<name>.so" is a "namelink" allowing linkers to find the library when given "-l<name>". The NAMELINK_ONLY option causes installation of only the namelink when a library target is installed. The NAMELINK_SKIP option causes installation of library files other than the namelink when a library target is installed. When neither option is given both portions are installed. On platforms where versioned shared libraries do not have namelinks or when a library is not

versioned the NAMELINK_SKIP option installs the library and the NAMELINK_ONLY option installs nothing. See the VERSION and SOVERSION target properties for details on creating versioned shared libraries.

One or more groups of properties may be specified in a single call to the TARGETS form of this command. A target may be installed more than once to different locations. Consider hypothetical targets "myExe", "mySharedLib", and "myStaticLib". The code

```
install(TARGETS myExe mySharedLib myStaticLib
        RUNTIME DESTINATION bin
        LIBRARY DESTINATION lib
        ARCHIVE DESTINATION lib/static)
install(TARGETS mySharedLib DESTINATION /some/full/path)
```

will install myExe to <prefix>/bin and myStaticLib to <prefix>/lib/static. On non-DLL platforms mySharedLib will be installed to <prefix>/lib and /some/full/path. On DLL platforms the mySharedLib DLL will be installed to <prefix>/bin and /some/full/path and its import library will be installed to <prefix>/lib/static and /some/full/path. On non-DLL platforms mySharedLib will be installed to <prefix>/lib and /some/full/path.

The EXPORT option associates the installed target files with an export called <export-name>. It must appear before any RUNTIME, LIBRARY, or ARCHIVE options. See documentation of the install(EXPORT ...) signature below for details.

Installing a target with EXCLUDE_FROM_ALL set to true has undefined behavior.

The FILES signature:

```
install(FILES files... DESTINATION <dir>
        [PERMISSIONS permissions...]
        [CONFIGURATIONS [Debug|Release|...]]
        [COMPONENT <component>]
        [RENAME <name>] [OPTIONAL])
```

The FILES form specifies rules for installing files for a project. File names given as relative paths are interpreted with respect to the current source directory. Files installed by this form are by default given permissions OWNER_WRITE, OWNER_READ, GROUP_READ, and WORLD_READ if no PERMISSIONS argument is given.

The PROGRAMS signature:

```
install(PROGRAMS files... DESTINATION <dir>
        [PERMISSIONS permissions...]
        [CONFIGURATIONS [Debug|Release|...]]
        [COMPONENT <component>]
        [RENAME <name>] [OPTIONAL])
```

The PROGRAMS form is identical to the FILES form except that the default permissions for the installed file also include OWNER_EXECUTE, GROUP_EXECUTE, and WORLD_EXECUTE. This form is intended to install programs that are not targets, such as shell scripts. Use the TARGETS form to install targets built within the project.

The DIRECTORY signature:

```
install(DIRECTORY dirs... DESTINATION <dir>
        [FILE_PERMISSIONS permissions...]
        [DIRECTORY_PERMISSIONS permissions...]
        [USE_SOURCE_PERMISSIONS] [OPTIONAL]
        [CONFIGURATIONS [Debug|Release|...]]
        [COMPONENT <component>]
        [FILES_MATCHING
            [[PATTERN <pattern> | REGEX <regex>]
            [EXCLUDE] [PERMISSIONS permissions...]] [...]])
```

The DIRECTORY form installs contents of one or more directories to a given destination. The directory structure is copied verbatim to the destination. The last component of each directory name is appended to the destination directory but a trailing slash may be used to avoid this because it leaves the last component empty. Directory names given as relative paths are interpreted with respect to the current source directory. If no input directory names are given the destination directory will be created but nothing will be installed into it. The FILE_PERMISSIONS and DIRECTORY_PERMISSIONS options specify permissions given to files and directories in the destination. If USE_SOURCE_PERMISSIONS is specified and FILE_PERMISSIONS is not, file permissions will be copied from the source directory structure. If no permissions are specified files will be given the default permissions specified in the FILES form of the command, and the directories will be given the default permissions specified in the PROGRAMS form of the command.

Installation of directories may be controlled with fine granularity using the PATTERN or REGEX options. These "match" options specify a globbing pattern or regular expression to match directories or files encountered within input directories. They may be used to apply certain options (see below) to a subset of the files and directories encountered. The full path to each input file or directory (with forward slashes) is matched against the expression. A PATTERN will match only complete file names: the portion of the full path matching the pattern must occur at the end of the file name and be preceded by a slash. A REGEX will

match any portion of the full path but it may use '/' and '\$' to simulate the PATTERN behavior. By default all files and directories are installed whether or not they are matched. The FILES_MATCHING option may be given before the first match option to disable installation of files (but not directories) not matched by any expression. For example, the code

```
install(DIRECTORY src/ DESTINATION include/myproj  
        FILES_MATCHING PATTERN "*.h")
```

will extract and install header files from a source tree.

Some options may follow a PATTERN or REGEX expression and are applied only to files or directories matching them. The EXCLUDE option will skip the matched file or directory. The PERMISSIONS option overrides the permissions setting for the matched file or directory. For example the code

```
install(DIRECTORY icons scripts/ DESTINATION share/myproj  
        PATTERN "CVS" EXCLUDE  
        PATTERN "scripts/*"  
        PERMISSIONS OWNER_EXECUTE OWNER_WRITE OWNER_READ  
                  GROUP_EXECUTE GROUP_READ)
```

will install the icons directory to share/myproj/icons and the scripts directory to share/myproj. The icons will get default file permissions, the scripts will be given specific permissions, and any CVS directories will be excluded.

The SCRIPT and CODE signatures:

```
install([[SCRIPT <file>] [CODE <code>]] [...])
```

The SCRIPT form will invoke the given CMake script files during installation. If the script file name is a relative path it will be interpreted with respect to the current source directory. The CODE form will invoke the given CMake code during installation. Code is specified as a single argument inside a double-quoted string. For example, the code

```
install(CODE "MESSAGE(\"Sample install message.\")")
```

will print a message during installation.

The EXPORT signature:

```
install(EXPORT <export-name> DESTINATION <dir>
       [NAMESPACE <namespace>] [FILE <name>.cmake]
       [PERMISSIONS permissions...]
       [CONFIGURATIONS [Debug|Release|...]]
       [COMPONENT <component>])
```

The EXPORT form generates and installs a CMake file containing code to import targets from the installation tree into another project. Target installations are associated with the export <export-name> using the EXPORT option of the install(TARGETS ...) signature documented above. The NAMESPACE option will prepend <namespace> to the target names as they are written to the import file. By default the generated file will be called <export-name>.cmake but the FILE option may be used to specify a different name. The value given to the FILE option must be a file name with the ".cmake" extension. If a CONFIGURATIONS option is given then the file will only be installed when one of the named configurations is installed. Additionally, the generated import file will reference only the matching target configurations. If a COMPONENT option is specified that does not match that given to the targets associated with <export-name> the behavior is undefined. If a library target is included in the export but a target to which it links is not included the behavior is unspecified.

The EXPORT form is useful to help outside projects use targets built and installed by the current project. For example, the code

```
install(TARGETS myexe EXPORT myproj DESTINATION bin)
install(EXPORT myproj NAMESPACE mp_ DESTINATION lib/myproj)
```

will install the executable myexe to <prefix>/bin and code to import it in the file "<prefix>/lib/myproj/myproj.cmake". An outside project may load this file with the include command and reference the myexe executable from the installation tree using the imported target name mp_myexe as if the target were built in its own tree.

NOTE: This command supercedes the INSTALL_TARGETS command and the target properties PRE_INSTALL_SCRIPT and POST_INSTALL_SCRIPT. It also replaces the FILES forms of the INSTALL_FILES and INSTALL_PROGRAMS commands. The processing order of these install rules relative to those generated by INSTALL_TARGETS, INSTALL_FILES, and INSTALL_PROGRAMS commands is not defined.

link_directories: Specify directories in which the linker will look for libraries.

```
link_directories(directory1 directory2 ...)
```

Specify the paths in which the linker should search for libraries. The command will apply only to targets created after it is called. For historical reasons, relative paths given to this command are passed to the linker unchanged (unlike many CMake commands which interpret them relative to the current source directory).

list: List operations.

```
list(LENGTH <list> <output variable>)
list(GET <list> <element index> [<element index> ...]
      <output variable>)
list(APPEND <list> <element> [<element> ...])
list(FIND <list> <value> <output variable>)
list(INSERT <list> <element_index> <element> [<element> ...])
list(REMOVE_ITEM <list> <value> [<value> ...])
list(REMOVE_AT <list> <index> [<index> ...])
list(REMOVE_DUPLICATES <list>)
list(REVERSE <list>)
list(SORT <list>)
```

LENGTH will return a given list's length.

GET will return list of elements specified by indices from the list.

APPEND will append elements to the list.

FIND will return the index of the element specified in the list or -1 if it wasn't found.

INSERT will insert elements to the list to the specified location.

REMOVE_AT and REMOVE_ITEM will remove items from the list. The difference is that REMOVE_ITEM will remove the given items, while REMOVE_AT will remove the items at the given indices.

REMOVE_DUPLICATES will remove duplicated items in the list.

REVERSE reverses the contents of the list in-place.

SORT sorts the list in-place alphabetically.

NOTES: A list in cmake is a ; separated group of strings. To create a list the set command can be used. For example, set(var a b c d e) creates a list with a;b;c;d;e, and set(var "a b c d e") creates a string or a list with one item in it.

When specifying index values, if <element index> is 0 or greater, it is indexed from the beginning of the list, with 0 representing the first list element. If <element index> is -1 or lesser, it is indexed from the end of the list, with -1 representing the last list element. Be careful when counting with negative indices: they do not start from 0. -0 is equivalent to 0, the first list element.

load_cache: Load in the values from another project's CMake cache.

```
load_cache(pathToCacheFile READ_WITH_PREFIX  
           prefix entry1...)
```

Read the cache and store the requested entries in variables with their name prefixed with the given prefix. This only reads the values, and does not create entries in the local project's cache.

```
load_cache(pathToCacheFile [EXCLUDE entry1...]  
           [INCLUDE_INTERNALS entry1...])
```

Load in the values from another cache and store them in the local project's cache as internal entries. This is useful for a project that depends on another project built in a different tree. EXCLUDE option can be used to provide a list of entries to be excluded. INCLUDE_INTERNALS can be used to provide a list of internal entries to be included. Normally, no internal entries are brought in. Use of this form of the command is strongly discouraged, but it is provided for backward compatibility.

load_command: Load a command into a running CMake.

```
load_command(COMMAND_NAME <loc1> [loc2 ...])
```

The given locations are searched for a library whose name is cmCOMMAND_NAME. If found, it is loaded as a module and the command is added to the set of available CMake commands. Usually, TRY_COMPILE is used before this command to compile the module. If the command is successfully loaded a variable named

```
CMAKE_LOADED_COMMAND_<COMMAND_NAME>
```

will be set to the full path of the module that was loaded. Otherwise the variable will not be set.

macro: Start recording a macro for later invocation as a command.

```
macro(<name> [arg1 [arg2 [arg3 ...]]])
    COMMAND1 (ARGS ...)
    COMMAND2 (ARGS ...)
    ...
endmacro ([<name>])
```

Define a macro named <name> that takes arguments named arg1 arg2 arg3 (...). Commands listed after macro, but before the matching endmacro, are not invoked until the macro is invoked. When it is invoked, the commands recorded in the macro are first modified by replacing formal parameters (`${arg1}`) with the arguments passed, and then invoked as normal commands. In addition to referencing the formal parameters you can reference the values `${ARGC}` which will be set to the number of arguments passed into the function as well as `${ARGV0}` `${ARGV1}` `${ARGV2}` ... which will have the actual values of the arguments passed in. This facilitates creating macros with optional arguments. Additionally `${ARGV}` holds the list of all arguments given to the macro and `${ARGN}` holds the list of argument past the last expected argument. Note that the parameters to a macro and values such as ARGN are not variables in the usual CMake sense. They are string replacements much like the c preprocessor would do with a macro. If you want true CMake variables you should look at the function command.

See the `cmake_policy()` command documentation for the behavior of policies inside macros.

mark_as_advanced: Mark cmake cached variables as advanced.

```
mark_as_advanced([CLEAR|FORCE] VAR VAR2 VAR...)
```

Mark the named cached variables as advanced. An advanced variable will not be displayed in any of the cmake GUIs unless the show advanced option is on. If CLEAR is the first argument advanced variables are changed back to unadvanced. If FORCE is the first argument, then the variable is made advanced. If neither FORCE nor CLEAR is specified, new values will be marked as advanced, but if the variable already has an advanced/non-advanced state, it will not be changed. This command does nothing in script mode.

math: Mathematical expressions.

```
math(EXPR <output variable> <math expression>)
```

EXPR evaluates mathematical expression and return result in the output variable. Example mathematical expression is `'5 * (10 + 13)'`. Supported operators are `+ - * / % | & ^ ~ << >> *` `/ %`. They have the same meaning as they do in c code.

message: Display a message to the user.

```
message([STATUS|WARNING|AUTHOR_WARNING|FATAL_ERROR|SEND_ERROR]
       "message to display" ...)
```

The optional keyword determines the type of message:

(none)	= Important information
STATUS	= Incidental information
WARNING	= CMake Warning, continue processing
AUTHOR_WARNING	= CMake Warning (dev), continue processing
SEND_ERROR	= CMake Error, continue but skip generation
FATAL_ERROR	= CMake Error, stop all processing

The CMake command-line tool displays STATUS messages on stdout and all other message types on stderr. The CMake GUI displays all messages in its log area. The interactive dialogs (ccmake and CMakeSetup) show STATUS messages one at a time on a status line and other messages in interactive pop-up boxes.

CMake Warning and Error message text displays using a simple markup language. Non-indented text is formatted in line-wrapped paragraphs delimited by newlines. Indented text is considered pre-formatted.

option: Provides an option that the user can optionally select.

```
option(<option_variable> "help string describing option"
      [initial value])
```

Provide an option for the user to select as ON or OFF. If no initial value is provided, OFF is used.

output_required_files: Output a list of required source files for a specified source file.

```
output_required_files(srcfile outfile)
```

Outputs a list of all the source files that are required by the specified srcfile. This list is written into outfile. This is similar to writing out the dependencies for srcfile except that it jumps from .h files into .cxx, .c and .cpp files if possible.

project: Set a name for the entire project.

```
project(<projectname> [languageName1 languageName2 ... ] )
```

Sets the name of the project. Additionally this sets the variables <projectName>_BINARY_DIR and <projectName>_SOURCE_DIR to the respective values.

Optionally you can specify which languages your project supports. Example languages are CXX (i.e. C++), C, Fortran, etc. By default C and CXX are enabled. E.g. if you do not have a C++ compiler, you can disable the check for it by explicitly listing the languages you want to support, e.g. C. By using the special language "NONE" all checks for any language can be disabled.

qt_wrap_cpp: Create Qt Wrappers.

```
qt_wrap_cpp(resultingLibraryName DestName  
           SourceLists ...)
```

Produce moc files for all the .h files listed in the SourceLists. The moc files will be added to the library using the DestName source list.

qt_wrap_ui: Create Qt user interfaces Wrappers.

```
qt_wrap_ui(resultingLibraryName HeadersDestName  
           SourcesDestName SourceLists ...)
```

Produce .h and .cxx files for all the .ui files listed in the SourceLists. The .h files will be added to the library using the HeadersDestNamesource list. The .cxx files will be added to the library using the SourcesDestNamesource list.

remove_definitions: Removes -D define flags added by add_definitions.

```
remove_definitions (-DFOO -DBAR ...)
```

Removes flags (added by add_definitions) from the compiler command line for sources in the current directory and below.

return: Return from a file, directory or function.

```
return()
```

Returns from a file, directory or function. When this command is encountered in an included file (via include() or find_package()), it causes processing of the current file to stop and control is returned to the including file. If it is encountered in a file which is not included by

another file, e.g. a CMakeLists.txt, control is returned to the parent directory if there is one. If return is called in a function, control is returned to the caller of the function. Note that a macro is not a function and does not handle return like a function does.

separate_arguments: Parse space-separated arguments into a semicolon-separated list.

```
separate_arguments(<var> <UNIX|WINDOWS>_COMMAND "<args>")
```

Parses a UNIX- or Windows-style command-line string "<args>" and stores a semicolon-separated list of the arguments in <var>. The entire command line must be given in one "<args>" argument.

The UNIX_COMMAND mode separates arguments by unquoted whitespace. It recognizes both single-quote and double-quote pairs. A backslash escapes the next literal character (\\" is "); there are no special escapes (\n is just n).

The WINDOWS_COMMAND mode parses a windows command-line using the same syntax the runtime library uses to construct argv at startup. It separates arguments by whitespace that is not double-quoted. Backslashes are literal unless they precede double-quotes. See the MSDN article "Parsing C Command-Line Arguments" for details.

```
separate_arguments(VARIABLE)
```

Convert the value of VARIABLE to a semi-colon separated list. All spaces are replaced with ':'. This helps with generating command lines.

set: Set a CMAKE variable to a given value.

```
set(<variable> <value>
    [ [CACHE <type> <docstring> [FORCE]] | PARENT_SCOPE])
```

Within CMake sets <variable> to the value <value>. <value> is expanded before <variable> is set to it. If CACHE is present and <variable> is not yet in the cache, then <variable> is put in the cache. If it is already in the cache, <variable> is assigned the value stored in the cache. If CACHE is present, also <type> and <docstring> are required. <type> is used by the CMake GUI to choose a widget with which the user sets a value. The value for <type> may be one of

```
FILEPATH = File chooser dialog.
PATH      = Directory chooser dialog.
STRING    = Arbitrary string.
```

```
BOOL      = Boolean ON/OFF checkbox.  
INTERNAL = No GUI entry (used for persistent variables).
```

If <type> is INTERNAL, then the <value> is always written into the cache, replacing any values existing in the cache. If it is not a cache variable, then this always writes into the current Makefile. The FORCE option will overwrite the cache value removing any changes by the user.

If PARENT_SCOPE is present, the variable will be set in the scope above the current scope. Each new directory or function creates a new scope. This command will set the value of a variable into the parent directory or calling function (whichever is applicable to the case at hand).

If <value> is not specified then the variable is removed instead of set. See also: the unset() command.

```
set(<variable> <value1> ... <valueN>)
```

In this case <variable> is set to a semicolon separated list of values. <variable> can be an environment variable such as:

```
set( ENV{PATH} /home/martink )
```

in which case the environment variable will be set.

set_directory_properties: Set a property of the directory.

```
set_directory_properties(PROPERTIES prop1 value1 prop2 value2)
```

Set a property for the current directory and subdirectories. If the property is not found, CMake will report an error. The properties include: INCLUDE_DIRECTORIES, LINK_DIRECTORIES, INCLUDE_REGULAR_EXPRESSION, and ADDITIONAL_MAKE_CLEAN_FILES. ADDITIONAL_MAKE_CLEAN_FILES is a list of files that will be cleaned as a part of "make clean" stage.

set_property: Set a named property in a given scope.

```
set_property(<GLOBAL  
            DIRECTORY [dir]  
            TARGET     [target1 [target2 ...]] |
```

```
SOURCE      [src1 [src2 ...]]      |
TEST        [test1 [test2 ...]]      |
CACHE       [entry1 [entry2 ...]]>
[APPEND]
PROPERTY <name> [value1 [value2 ...]])
```

Set one property on zero or more objects of a scope. The first argument determines the scope in which the property is set. It must be one of the following:

GLOBAL scope is unique and does not accept a name.

DIRECTORY scope defaults to the current directory but another directory (already processed by CMake) may be named by full or relative path.

TARGET scope may name zero or more existing targets.

SOURCE scope may name zero or more source files.

TEST scope may name zero or more existing tests.

CACHE scope must name zero or more cache existing entries.

The required PROPERTY option is immediately followed by the name of the property to set. Remaining arguments are used to compose the property value in the form of a semicolon-separated list. If the APPEND option is given the list is appended to any existing property value.

set_source_files_properties: Source files can have properties that affect how they are built.

```
set_source_files_properties(file1 file2 ...
                           PROPERTIES prop1 value1
                           prop2 value2 ...)
```

Set properties on a file. The syntax for the command is to list all the files you want to change, and then provide the values you want to set next. You can make up your own properties as well. The following are used by CMake. The ABSTRACT flag (boolean) is used by some class wrapping commands. If WRAP_EXCLUDE (boolean) is true then many wrapping commands will ignore this file. If GENERATED (boolean) is true then it is not an error if this source file does not exist when it is added to a target. Obviously, it must be created (presumably by a custom command) before the target is built. If the HEADER_FILE_ONLY (boolean) property is true then the file is not compiled. This is useful if you want to add extra non build files to an IDE. OBJECT_DEPENDS (string) adds dependencies to the object file.

COMPILE_FLAGS (string) is passed to the compiler as additional command line arguments when the source file is compiled. LANGUAGE (string) CXX|C will change the default compiler used to compile the source file. The languages used need to be enabled in the PROJECT command. If SYMBOLIC (boolean) is set to true the build system will be informed that the source file is not actually created on disk but instead used as a symbolic name for a build rule.

set_target_properties: Targets can have properties that affect how they are built.

```
set_target_properties(target1 target2 ...
                      PROPERTIES prop1 value1
                                prop2 value2 ...)
```

Set properties on a target. The syntax for the command is to list all the files you want to change, and then provide the values you want to set next. You can use any prop value pair you want and extract it later with the GET_TARGET_PROPERTY command. See Appendix E - Properties for a full list of target properties.

set_tests_properties: Set a property of the tests.

```
set_tests_properties(test1 [test2...]
                      PROPERTIES prop1 value1 prop2 value2)
```

Set a property for the tests. If the property is not found, CMake will report an error. See Appendix E - Properties for a list of test properties.

site_name: Set the given variable to the name of the computer.

```
site_name(variable)
```

source_group: Define a grouping for sources in the Makefile.

```
source_group(name [REGULAR_EXPRESSION regex]
              [FILES src1 src2 ...])
```

Defines a group into which sources will be placed in project files. This is mainly used to setup file tabs in Visual Studio. Any file whose name is listed or matches the regular expression will be placed in this group. If a file matches multiple groups, the LAST group that explicitly lists the file will be favored, if any. If no group explicitly lists the file, the LAST group whose regular expression matches the file will be favored.

The name of the group may contain backslashes to specify subgroups:

```
source_group(outer\\inner ...)
```

For backwards compatibility, this command is also supports the format:

```
source_group(name regex)
```

string: String operations.

```
string(REGEX MATCH <regular_expression>
      <output variable> <input> [<input>...])
string(REGEX MATCHALL <regular_expression>
      <output variable> <input> [<input>...])
string(REGEX REPLACE <regular_expression>
      <replace_expression> <output variable>
      <input> [<input>...])
string(REPLACE <match_string>
      <replace_string> <output variable>
      <input> [<input>...])
string(COMPARE EQUAL <string1> <string2> <output variable>)
string(COMPARE NOTEQUAL <string1> <string2> <output variable>)
string(COMPARE LESS <string1> <string2> <output variable>)
string(COMPARE GREATER <string1> <string2> <output variable>)
string(ASCII <number> [<number> ...] <output variable>)
string(CONFIGURE <string1> <output variable>
      [&ONLY] [ESCAPE_QUOTES])
string(TOUPPER <string1> <output variable>)
string(TOLOWER <string1> <output variable>)
string(LENGTH <string> <output variable>)
string(SUBSTRING <string> <begin> <length> <output variable>)
string(STRIPE <string> <output variable>)
string(RANDOM [LENGTH <length>] [ALPHABET <alphabet>]
      <output variable>)
```

REGEX MATCH will match the regular expression once and store the match in the output variable.

REGEX MATCHALL will match the regular expression as many times as possible and store the matches in the output variable as a list.

REGEX REPLACE will match the regular expression as many times as possible and substitute the replacement expression for the match in the output. The replace expression may refer to paren-delimited subexpressions of the match using \1, \2, ..., \9. Note that two backslashes (\\\) are required in CMake code to get a backslash through argument parsing.

REPLACE will replace all occurrences of match_string in the input with replace_string and store the result in the output.

COMPARE EQUAL/NOTEQUAL/LESS/GREATER will compare the strings and store true or false in the output variable.

ASCII will convert all numbers into corresponding ASCII characters.

CONFIGURE will transform a string like CONFIGURE_FILE transforms a file.

TOUPPER/TOLOWER will convert string to upper/lower characters.

LENGTH will return a given string's length.

SUBSTRING will return a substring of a given string.

STRIP will return a substring of a given string with leading and trailing spaces removed.

RANDOM will return a random string of given length consisting of characters from the given alphabet. Default length is 5 characters and default alphabet is all numbers and upper and lower case letters.

The following characters have special meaning in regular expressions:

^	Matches at beginning of a line
\$	Matches at end of a line
.	Matches any single character
[]	Matches any character(s) inside the brackets
[^]	Matches any character(s) not inside the brackets
-	Matches characters in range on either side of a dash
*	Matches preceding pattern zero or more times
+	Matches preceding pattern one or more times
?	Matches preceding pattern zero or once only
	Matches a pattern on either side of the
()	Saves a matched subexpression, which can be referenced in the REGEX REPLACE operation. Additionally it is Saved by all regular expression-related commands, including e.g. if(MATCHES), in the variables CMAKE_MATCH_(0..9).

target_link_libraries: Link a target to given libraries.

```
target_link_libraries(<target> [item1 [item2 [...]]]
                      [[debug|optimized|general] <item>] ...)
```

Specify libraries or flags to use when linking a given target. If a library name matches that of another target in the project a dependency will automatically be added in the build system to make sure the library being linked is up-to-date before the target links. Item names starting with '`-`', but not '`-l`' or '`-framework`', are treated as linker flags.

A "`debug`", "`optimized`", or "`general`" keyword indicates that the library immediately following it is to be used only for the corresponding build configuration. The "`debug`" keyword corresponds to the Debug configuration (or to configurations named in the `DEBUG_CONFIGURATIONS` global property if it is set). The "`optimized`" keyword corresponds to all other configurations. The "`general`" keyword corresponds to all configurations, and is purely optional (assumed if omitted). Higher granularity may be achieved for per-configuration rules by creating and linking to IMPORTED library targets. See the IMPORTED mode of the `add_library` command for more information.

Library dependencies are transitive by default. When this target is linked into another target then the libraries linked to this target will appear on the link line for the other target too. See the `LINK_INTERFACE_LIBRARIES` target property to override the set of transitive link dependencies for a target.

```
target_link_libraries(<target> LINK_INTERFACE_LIBRARIES
                      [[debug|optimized|general] <lib>] ...)
```

The `LINK_INTERFACE_LIBRARIES` mode appends the libraries to the `LINK_INTERFACE_LIBRARIES` and its per-configuration equivalent target properties instead of using them for linking. Libraries specified as "`debug`" are appended to the `LINK_INTERFACE_LIBRARIES_DEBUG` property (or to the properties corresponding to configurations listed in the `DEBUG_CONFIGURATIONS` global property if it is set). Libraries specified as "`optimized`" are appended to the `LINK_INTERFACE_LIBRARIES` property. Libraries specified as "`general`" (or without any keyword) are treated as if specified for both "`debug`" and "`optimized`".

The library dependency graph is normally acyclic (a DAG), but in the case of mutually-dependent STATIC libraries CMake allows the graph to contain cycles (strongly connected components). When another target links to one of the libraries CMake repeats the entire connected component. For example, the code

```
add_library(A STATIC a.c)
add_library(B STATIC b.c)
target_link_libraries(A B)
target_link_libraries(B A)
add_executable(main main.c)
target_link_libraries(main A)
```

links 'main' to 'A B A B'. (While one repetition is usually sufficient, pathological object file and symbol arrangements can require more. One may handle such cases by manually repeating the component in the last target_link_libraries call. However, if two archives are really so interdependent they should probably be combined into a single archive.)

try_compile: Try compiling some code.

```
try_compile(RESULT_VAR bindir srccdir
            projectName <targetname> [CMAKE_FLAGS <Flags>]
            [OUTPUT_VARIABLE var])
```

Try compiling a program. In this form, srccdir should contain a complete CMake project with a CMakeLists.txt file and all sources. The bindir and srccdir will not be deleted after this command is run. If <target name> is specified then build just that target otherwise the all or ALL_BUILD target is built.

```
try_compile(RESULT_VAR bindir srcfile
            [CMAKE_FLAGS <Flags>]
            [COMPILE_DEFINITIONS <flags> ...]
            [OUTPUT_VARIABLE var]
            [COPY_FILE <filename> ])
```

Try compiling a srcfile. In this case, the user need only supply a source file. CMake will create the appropriate CMakeLists.txt file to build the source. If COPY_FILE is used, the compiled file will be copied to the given file.

In this version all files in bindir/CMakeFiles/CMakeTmp, will be cleaned automatically, for debugging a --debug-trycompile can be passed to cmake to avoid the clean. Some extra flags that can be included are, INCLUDE_DIRECTORIES, LINK_DIRECTORIES, and LINK_LIBRARIES. COMPILE_DEFINITIONS are -Ddefinition that will be passed to the compile line. try_compile creates a CMakeList.txt file on the fly that looks like this:

```
add_definitions(<expanded COMPILE_DEFINITIONS from cmake>)
include_directories(${INCLUDE_DIRECTORIES})
```

```
link_directories(${LINK_DIRECTORIES})
add_executable(cmTryCompileExec sources)
target_link_libraries(cmTryCompileExec ${LINK_LIBRARIES})
```

In both versions of the command, if OUTPUT_VARIABLE is specified, then the output from the build process is stored in the given variable. Return the success or failure in RESULT_VAR. CMAKE_FLAGS can be used to pass -DVAR:TYPE=VALUE flags to the cmake that is run during the build.

try_run: Try compiling and then running some code.

```
try_run(RUN_RESULT_VAR COMPILE_RESULT_VAR
        bindir srcfile [CMAKE_FLAGS <Flags>]
        [COMPILE_DEFINITIONS <flags>]
        [COMPILE_OUTPUT_VARIABLE comp]
        [RUN_OUTPUT_VARIABLE run]
        [OUTPUT_VARIABLE var]
        [ARGS <arg1> <arg2>...])
```

Try compiling a srcfile. Return TRUE or FALSE for success or failure in COMPILE_RESULT_VAR. Then if the compile succeeded, run the executable and return its exit code in RUN_RESULT_VAR. If the executable was built, but failed to run, then RUN_RESULT_VAR will be set to FAILED_TO_RUN. COMPILE_OUTPUT_VARIABLE specifies the variable where the output from the compile step goes. RUN_OUTPUT_VARIABLE specifies the variable where the output from the running executable goes.

For compatibility reasons OUTPUT_VARIABLE is still supported, which gives you the output from the compile and run step combined.

Cross compiling issues

When cross compiling, the executable compiled in the first step usually cannot be run on the build host. try_run() checks the CMAKE_CROSSCOMPILING variable to detect whether CMake is in crosscompiling mode. If that's the case, it will still try to compile the executable, but it will not try to run the executable. Instead it will create cache variables which must be filled by the user or by presetting them in some CMake script file to the values the executable would have produced if it would have been run on its actual target platform. These variables are RUN_RESULT_VAR (explanation see above) and if RUN_OUTPUT_VARIABLE (or OUTPUT_VARIABLE) was used, an additional cache variable RUN_RESULT_VAR__COMPILE_RESULT_VAR__TRYRUN_OUTPUT. This is intended to hold stdout and stderr from the executable.

In order to make cross compiling your project easier, use `try_run` only if really required. If you use `try_run`, use `RUN_OUTPUT_VARIABLE` (or `OUTPUT_VARIABLE`) only if really required. Using them will require that when crosscompiling, the cache variables will have to be set manually to the output of the executable. You can also "guard" the calls to `try_run` with `if(CMAKE_CROSSCOMPILING)` and provide an easy-to-preset alternative for this case.

unset: Unset a variable, cache variable, or environment variable.

```
unset(<variable> [CACHE])
```

Removes the specified variable causing it to become undefined. If `CACHE` is present then the variable is removed from the cache instead of the current scope. `<variable>` can be an environment variable such as:

```
unset(ENV{LD_LIBRARY_PATH})
```

in which case the variable will be removed from the current environment.

variable_watch: Watch the CMake variable for change.

```
variable_watch(<variable name> [<command to execute>])
```

If the specified variable changes, the message will be printed about the variable being changed. If the command is specified, the command will be executed. The command will receive the following arguments: `COMMAND(<variable> <access> <value> <current list file> <stack>)`

while: Evaluate a group of commands while a condition is true.

```
while(condition)
  COMMAND1(ARGS ...)
  COMMAND2(ARGS ...)
  ...
endwhile([condition])
```

All commands between `while` and the matching `endwhile` are recorded without being invoked. Once the `endwhile` is evaluated, the recorded list of commands is invoked as long as the condition is true. The condition is evaluated using the same logic as the `if` command.

Compatibility Commands

This is the documentation for now obsolete listfile commands from previous CMake versions, which are still supported for compatibility reasons. You should instead use the newer, preferred commands.

`build_name`: Deprecated. Use `${CMAKE_SYSTEM}` and `${CMAKE_CXX_COMPILER}` instead.

```
build_name(variable)
```

Sets the specified variable to a string representing the platform and compiler settings. These values are now available through the `CMAKE_SYSTEM` and `CMAKE_CXX_COMPILER` variables.

`exec_program`: Deprecated. Use the `execute_process()` command instead.

Run an executable program during the processing of the `CMakeList.txt` file.

```
exec_program(Executable [directory in which to run]
              [ARGS <arguments to executable>]
              [OUTPUT_VARIABLE <var>]
              [RETURN_VALUE <var>])
```

The executable is run in the optionally specified directory. The executable can include arguments if it is double quoted, but it is better to use the optional ARGS argument to specify arguments to the program. This is because `cmake` will then be able to escape spaces in the executable path. An optional argument `OUTPUT_VARIABLE` specifies a variable in which to store the output. To capture the return value of the execution, provide a `RETURN_VALUE`. If `OUTPUT_VARIABLE` is specified, then no output will go to the `stdout/stderr` of the console running `cmake`.

`export_library_dependencies`: Deprecated. Use `INSTALL(EXPORT)` or `EXPORT` command.

This command generates an old-style library dependencies file. Projects requiring CMake 2.6 or later should not use the command. Use instead the `install(EXPORT)` command to help export targets from an installation tree and the `export()` command to export targets from a build tree.

The old-style library dependencies file does not take into account per-configuration names of libraries or the `LINK_INTERFACE_LIBRARIES` target property.

```
export_library_dependencies(<file> [APPEND])
```

Create a file named <file> that can be included into a CMake listfile with the INCLUDE command. The file will contain a number of SET commands that will set all the variables needed for library dependency information. This should be the last command in the top level CMakeLists.txt file of the project. If the APPEND option is specified, the SET commands will be appended to the given file instead of replacing it.

install_files: Deprecated. Use the `install(FILES)` command instead.

This command has been superceded by the `install` command. It is provided for compatibility with older CMake code. The FILES form is directly replaced by the FILES form of the `install` command. The regexp form can be expressed more clearly using the GLOB form of the `file` command.

```
install_files(<dir> extension file file ...)
```

Create rules to install the listed files with the given extension into the given directory. Only files existing in the current source tree or its corresponding location in the binary tree may be listed. If a file specified already has an extension, that extension will be removed first. This is useful for providing lists of source files such as `foo.cxx` when you want the corresponding `foo.h` to be installed. A typical extension is '`.h`'.

```
install_files(<dir> regexp)
```

Any files in the current source directory that match the regular expression will be installed.

```
install_files(<dir> FILES file file ...)
```

Any files listed after the FILES keyword will be installed explicitly from the names given. Full paths are allowed in this form. The directory <dir> is relative to the installation prefix, which is stored in the variable `CMAKE_INSTALL_PREFIX`.

install_programs: Deprecated. Use the `install(PROGRAMS)` command instead.

This command has been superceded by the `install` command. It is provided for compatibility with older CMake code. The FILES form is directly replaced by the PROGRAMS form of the `INSTALL` command. The regexp form can be expressed more clearly using the GLOB form of the `FILE` command.

```
install_programs(<dir> file1 file2 [file3 ...])
install_programs(<dir> FILES file1 [file2 ...])
```

Create rules to install the listed programs into the given directory. Use the FILES argument to guarantee that the file list version of the command will be used even when there is only one argument.

```
install_programs(<dir> regexp)
```

In the second form any program in the current source directory that matches the regular expression will be installed. This command is intended to install programs that are not built by cmake, such as shell scripts. See the TARGETS form of the INSTALL command to create installation rules for targets built by cmake. The directory <dir> is relative to the installation prefix, which is stored in the variable CMAKE_INSTALL_PREFIX.

install_targets: Deprecated. Use the `install(TARGETS)` command instead.

This command has been superceded by the `install` command. It is provided for compatibility with older CMake code.

```
install_targets(<dir> [RUNTIME_DIRECTORY dir] target target)
```

Create rules to install the listed targets into the given directory. The directory <dir> is relative to the installation prefix, which is stored in the variable CMAKE_INSTALL_PREFIX. If RUNTIME_DIRECTORY is specified, then on systems with special runtime files (Windows DLL), the files will be copied to that directory.

link_libraries: Deprecated. Use the `target_link_libraries()` command instead.

Link libraries to all targets added later.

```
link_libraries(library1 <debug | optimized> library2 ...)
```

Specify a list of libraries to be linked into any following targets (typically added with the `add_executable` or `add_library` calls). This command is passed down to all subdirectories. The debug and optimized strings may be used to indicate that the next library listed is to be used only for that specific type of build.

make_directory: Deprecated. Use the `file(MAKE_DIRECTORY)` command instead.

```
make_directory(directory)
```

Creates the specified directory. Full paths should be given. Any parent directories that do not exist will also be created. Use with care.

remove: Deprecated. Use the list(REMOVE_ITEM) command instead.

```
remove (VAR VALUE VALUE ...)
```

Removes VALUE from the variable VAR. This is typically used to remove entries from a vector (e.g. semicolon separated list). VALUE is expanded.

subdir_depends: Deprecated. Does nothing.

```
subdir_depends (subdir dep1 dep2 ...)
```

Does not do anything. This command used to help projects order parallel builds correctly. This functionality is now automatic.

subdirs: Deprecated. Use the add_subdirectory() command instead.

Add a list of subdirectories to the build.

```
subdirs (dir1 dir2 ...[EXCLUDE_FROM_ALL exclude_dir1 ...]  
[PREORDER] )
```

Add a list of subdirectories to the build. The add_subdirectory command should be used instead of subdirs although subdirs will still work. This will cause any CMakeLists.txt files in the sub directories to be processed by CMake. Any directories after the PREORDER flag are traversed first by Makefile builds, the PREORDER flag has no effect on IDE projects. Any directories after the EXCLUDE_FROM_ALL marker will not be included in the top level Makefile or project file. This is useful for having CMake create Makefiles or projects for a set of examples in a project. You would want CMake to generate Makefiles or project files for all the examples at the same time, but you would not want them to show up in the top level project or be built each time make is run from the top.

use_mangled_mesa: Copy mesa headers for use in combination with system GL.

```
use_mangled_mesa (PATH_TO_MESA OUTPUT_DIRECTORY)
```

The path to mesa includes, should contain gl_mangle.h. The mesa headers are copied to the specified output directory. This allows mangled mesa headers to override other GL headers by being added to the include directory path earlier.

utility_source: Specify the source tree of a third-party utility.

```
utility_source(cache_entry executable_name  
             path_to_source [file1 file2 ...])
```

When a third-party utility's source is included in the distribution, this command specifies its location and name. The cache entry will not be set unless the path_to_source and all listed files exist. It is assumed that the source tree of the utility will have been built before it is needed.

When cross compiling CMake will print a warning if a utility_source() command is executed, because in many cases it is used to build an executable which is executed later on. This doesn't work when cross compiling, since the executable can run only on their target platform. So in this case the cache entry has to be adjusted manually so it points to an executable which is runnable on the build host.

variable_requires: Deprecated. Use the if() command instead.

Assert satisfaction of an option's required variables.

```
variable_requires(TEST_VARIABLE RESULT_VARIABLE  
                  REQUIRED_VARIABLE1  
                  REQUIRED_VARIABLE2 ...)
```

The first argument (TEST_VARIABLE) is the name of the variable to be tested, if that variable is false nothing else is done. If TEST_VARIABLE is true, then the next argument (RESULT_VARIABLE) is a variable that is set to true if all the required variables are set. The rest of the arguments are variables that must be true or not set to NOTFOUND to avoid an error. If any are not true, an error is reported.

write_file: Deprecated. Use the file(WRITE) command instead.

```
write_file(filename "message to write"... [APPEND])
```

The first argument is the file name, the rest of the arguments are messages to write. If the argument APPEND is specified, then the message will be appended.

NOTE 1: file(WRITE ... and file(APPEND ... do exactly the same as this one but add some more functionality.

NOTE 2: When using `write_file` the produced file cannot be used as an input to CMake (`CONFIGURE_FILE`, `source_file ...`) because it will lead to an infinite loop. Use `configure_file` if you want to generate input files to CMake.

Appendix D – Selected Modules

CMake Modules

This is the documentation for the modules and scripts that come with CMake. Using these modules you can check the computer system for installed software packages, features of the compiler and the existence of headers to name just a few.

AddFileDependencies: ADD_FILE_DEPENDENCIES(source_file depend_files...)

Adds the given files as dependencies to source_file

BundleUtilities:

A collection of CMake utility functions useful for dealing with .app bundles on the Mac and bundle-like directories on any OS. The following functions are provided by this module:

```
get_bundle_main_executable  
get_dotapp_dir  
get_bundle_and_executable  
get_bundle_all_executables  
get_item_key  
clear_bundle_keys  
set_bundle_key_values  
get_bundle_keys
```

```
copy_resolved_item_into_bundle
fixup_bundle_item
fixup_bundle
copy_and_fixup_bundle
verify_bundle_prerequisites
verify_bundle_symlinks
verify_app
```

Requires CMake 2.6 or greater because it uses function, break and PARENT_SCOPE. Also depends on GetPrerequisites.cmake.

CMakeBackwardCompatibilityCXX: define a bunch of backwards compatibility variables

```
CMAKE_ANSI_CXXFLAGS - flag for ansi c++
CMAKE_HAS_ANSI_STRING_STREAM - has <strstream>
INCLUDE(TestForANSIStreamHeaders)
INCLUDE(CheckIncludeFileCXX)
INCLUDE(TestForSTDNamespace)
INCLUDE(TestForANSIForScope)
```

CMakeDependentOption: Macro to provide an option dependent on other options.

This macro presents an option to the user only if a set of other conditions are true. When the option is not presented a default value is used, but any value set by the user is preserved for when the option is presented again. Example invocation:

```
CMAKE_DEPENDENT_OPTION(USE_FOO "Use Foo" ON
                      "USE_BAR;NOT USE_ZOT" OFF)
```

If USE_BAR is true and USE_ZOT is false, this provides an option called USE_FOO that defaults to ON. Otherwise, it sets USE_FOO to OFF. If the status of USE_BAR or USE_ZOT ever changes, any value for the USE_FOO option is saved so that when the option is re-enabled it retains its old value.

CMakeDetermineVSServicePack: Includes a public function for assisting users in trying to determine the

Visual Studio service pack in use. Sets the passed in variable to one of the following values or an empty string if unknown.

```
vc80
vc80sp1
```

```
vc90  
vc90sp1
```

Usage:

```
if (MSVC)
    include (CMakeDetermineVSServicePack)
    DetermineVSServicePack( my_service_pack )

    if( my_service_pack )
        message (STATUS "Detected: ${my_service_pack}")
    endif()
endif()
```

CMakeFindFrameworks: helper module to find OSX frameworks

CMakeForceCompiler:

This module defines macros intended for use by cross-compiling toolchain files when CMake is not able to automatically detect the compiler identification. The macro `CMAKE_FORCE_C_COMPILER` has the following signature:

```
CMAKE_FORCE_C_COMPILER(<compiler> <compiler-id>)
```

It sets `CMAKE_C_COMPILER` to the given compiler and the `cmake` internal variable `CMAKE_C_COMPILER_ID` to the given compiler-id. It also bypasses the check for working compiler and basic compiler information tests.

The macro `CMAKE_FORCE_CXX_COMPILER` has the following signature:

```
CMAKE_FORCE_CXX_COMPILER(<compiler> <compiler-id>)
```

It sets `CMAKE_CXX_COMPILER` to the given compiler and the `cmake` internal variable `CMAKE_CXX_COMPILER_ID` to the given compiler-id. It also bypasses the check for working compiler and basic compiler information tests.

So a simple toolchain file could look like this:

```
INCLUDE (CMakeForceCompiler)
SET (CMAKE_SYSTEM_NAME Generic)
CMAKE_FORCE_C_COMPILER  (chcl2 MetrowerksHicross)
CMAKE_FORCE_CXX_COMPILER (chcl2 MetrowerksHicross)
```

CMakePrintSystemInformation: print system information

This file can be used for diagnostic purposes just include it in a project to see various internal CMake variables.

CPack: Build binary and source package installers

The CPack module generates binary and source installers in a variety of formats using the cpack program. Inclusion of the CPack module adds two new targets to the resulting Makefiles, package and package_source, which build the binary and source installers, respectively. The generated binary installers contain everything installed via CMake's INSTALL command (and the deprecated INSTALL_FILES, INSTALL_PROGRAMS, and INSTALL_TARGETS commands).

For certain kinds of binary installers (including the graphical installers on Mac OS X and Windows), CPack generates installers that allow users to select individual application components to install. The contents of each of the components are identified by the COMPONENT argument of CMake's INSTALL command. These components can be annotated with user-friendly names and descriptions, inter-component dependencies, etc., and grouped in various ways to customize the resulting installer. See the cpack_add_* commands, described below, for more information about component-specific installations.

Before including the CPack module, there are a variety of variables that can be set to customize the resulting installers. The most commonly-used variables are:

CPACK_PACKAGE_NAME

The name of the package (or application). If not specified, defaults to the project name.

CPACK_PACKAGE_VENDOR

The name of the package vendor (e.g. "Kitware").

CPACK_PACKAGE_VERSION_MAJOR

Package major Version

CPACK_PACKAGE_VERSION_MINOR

Package minor Version

CPACK_PACKAGE_VERSION_PATCH

Package patch Version

CPACK_PACKAGE_DESCRIPTION_FILE

A text file used to describe the project. Used, for example, the introduction screen of a CPack-generated Windows installer to describe the project.

CPACK_PACKAGE_DESCRIPTION_SUMMARY

Short description of the project (only a few words).

CPACK_PACKAGE_FILE_NAME

The name of the package file to generate, not including the extension. For example, cmake-2.6.1-Linux-i686.

CPACK_PACKAGE_INSTALL_DIRECTORY

Installation directory on the target system, e.g., "CMake 2.5".

CPACK_RESOURCE_FILE_LICENSE

License file for the project, which will typically be displayed to the user (often with an explicit "Accept" button, for graphical installers) prior to installation.

CPACK_RESOURCE_FILE_README

ReadMe file for the project, which typically describes in some detail

CPACK_RESOURCE_FILE_WELCOME

Welcome file for the project, which welcomes users to this installer. Typically used in the graphical installers on Windows and Mac OS X.

CPACK_MONOLITHIC_INSTALL

Disables the component-based installation mechanism, so that all components are always installed.

CPACK_GENERATOR

List of CPack generators to use. If not specified, CPack will create a set of options (e.g., CPACK_BINARY_NSIS) allowing the user to enable/disable individual generators.

CPACK_OUTPUT_CONFIG_FILE

The name of the CPack configuration file for binary installers that will be generated by the CPack module. Defaults to CPackConfig.cmake.

CPACK_PACKAGE_EXECUTABLES

Lists each of the executables along with a text label, to be used to create Start Menu shortcuts on Windows. For example, setting this to the list `ccmake;CMake` will create a shortcut named "CMake" that will execute the installed executable `ccmake`.

CPACK_STRIP_FILES

List of files to be stripped. Starting with CMake 2.6.0 `CPACK_STRIP_FILES` will be a boolean variable which enables stripping of all files (a list of files evaluates to TRUE in CMake, so this change is compatible).

The following CPack variables are specific to source packages, and will not affect binary packages:

CPACK_SOURCE_PACKAGE_FILE_NAME

The name of the source package, e.g., `cmake-2.6.1`

CPACK_SOURCE_STRIP_FILES

List of files in the source tree that will be stripped. Starting with CMake 2.6.0 `CPACK_SOURCE_STRIP_FILES` will be a boolean variable which enables stripping of all files (a list of files evaluates to TRUE in CMake, so this change is compatible).

CPACK_SOURCE_GENERATOR

List of generators used for the source packages. As with `CPACK_GENERATOR`, if this is not specified then CPack will create a set of options (e.g., `CPACK_SOURCE_ZIP`) allowing users to select which packages will be generated.

CPACK_SOURCE_OUTPUT_CONFIG_FILE

The name of the CPack configuration file for source installers that will be generated by the CPack module. Defaults to `CPackSourceConfig.cmake`.

CPACK_SOURCE_IGNORE_FILES

Pattern of files in the source tree that won't be packaged when building a source package. This is a list of patterns, e.g., `/CVS/;\\.svn/;\\.swp$;\\.#/#.*~;cscope.*`

The following variables are specific to the graphical installers built on Windows using the Nullsoft Installation System.

CPACK_PACKAGE_INSTALL_REGISTRY_KEY

Registry key used when installing this project.

CPACK_NSIS_MUI_ICON

The icon file (.ico) for the generated install program.

CPACK_NSIS_MUI_UNIICON

The icon file (.ico) for the generated uninstall program.

CPACK_PACKAGE_ICON

A branding image that will be displayed inside the installer.

CPACK_NSIS_EXTRA_INSTALL_COMMANDS

Extra NSIS commands that will be added to the install Section.

CPACK_NSIS_EXTRA_UNINSTALL_COMMANDS

Extra NSIS commands that will be added to the uninstall Section.

CPACK_NSIS_COMPRESSOR

The arguments that will be passed to the NSIS SetCompressor command.

CPACK_NSIS MODIFY_PATH

If this is set to "ON", then an extra page will appear in the installer that will allow the user to choose whether the program directory should be added to the system PATH variable.

CPACK_NSIS_DISPLAY_NAME

The display name string that appears in the Windows Add/Remove Program control panel

CPACK_NSIS_PACKAGE_NAME

The title displayed at the top of the installer.

CPACK_NSIS_INSTALLED_ICON_NAME

A path to the executable that contains the installer icon.

CPACK_NSIS_HELP_LINK

URL to a web site providing assistance in installing your application.

CPACK_NSIS_URL_INFO_ABOUT

URL to a web site providing more information about your application.

CPACK_NSIS_CONTACT

Contact information for questions and comments about the installation process.

CPACK_NSIS_CREATE_ICONS_EXTRA

Additional NSIS commands for creating start menu shortcuts.

CPACK_NSIS_DELETE_ICONS_EXTRA

Additional NSIS commands to uninstall start menu shortcuts.

The following variable is specific to installers build on Mac OS X using PackageMaker:

CPACK OSX PACKAGE VERSION

The version of Mac OS X that the resulting PackageMaker archive should be compatible with. Different versions of Mac OS X support different features. For example, CPack can only build component-based installers for Mac OS X 10.4 or newer, and can only build installers that download component son-the-fly for Mac OS X 10.5 or newer. If left blank, this value will be set to the minimum version of Mac OS X that supports the requested features. Set this variable to some value (e.g., 10.4) only if you want to guarantee that your installer will work on that version of Mac OS X, and don't mind missing extra features available in the installer shipping with later versions of Mac OS X.

The following variables are for advanced uses of CPack:

CPACK_CMAKE_GENERATOR

What CMake generator should be used if the project is CMake project. Defaults to the value of CMAKE_GENERATOR; few users will want to change this setting.

CPACK_INSTALL_CMAKE_PROJECTS

List of four values that specify what project to install. The four values are: Build directory, Project Name, Project Component, Directory. If omitted, CPack will build an installer that installers everything.

CPACK_SYSTEM_NAME

System name, defaults to the value of \${CMAKE_SYSTEM_NAME}.

CPACK_PACKAGE_VERSION

Package full version, used internally. By default, this is built from CPACK_PACKAGE_VERSION_MAJOR, CPACK_PACKAGE_VERSION_MINOR, and CPACK_PACKAGE_VERSION_PATCH.

CPACK_TOPLEVEL_TAG

Directory for the installed files.

CPACK_INSTALL_COMMANDS

Extra commands to install components.

CPACK_INSTALL_DIRECTORIES

Extra directories to install.

Component-specific installation allows users to select specific sets of components to install during the install process. Installation components are identified by the COMPONENT argument of CMake's INSTALL commands, and should be further described by the following CPack commands:

```
cpack_add_component - Describes a CPack installation component  
named by the COMPONENT argument to a CMake INSTALL command.
```

```
cpack_add_component(compname  
[DISPLAY_NAME name]  
[DESCRIPTION description]  
[HIDDEN | REQUIRED | DISABLED ]  
[GROUP group]  
[DEPENDS comp1 comp2 ... ]  
[INSTALL_TYPES type1 type2 ... ]  
[DOWNLOADED]  
[ARCHIVE_FILE filename])
```

The cmake_add_component command describes an installation component, which the user can opt to install or remove as part of the graphical installation process. compname is the name of the component, as provided to the COMPONENT argument of one or more CMake INSTALL commands.

DISPLAY_NAME is the displayed name of the component, used in graphical installers to display the component name. This value can be any string.

DESCRIPTION is an extended description of the component, used in graphical installers to give the user additional information about the component. Descriptions can span multiple lines using "\n" as the line separator. Typically, these descriptions should be no more than a few lines long.

HIDDEN indicates that this component will be hidden in the graphical installer, so that the user cannot directly change whether it is installed or not.

REQUIRED indicates that this component is required, and therefore will always be installed. It will be visible in the graphical installer, but it cannot be unselected. (Typically, required components are shown greyed out).

DISABLED indicates that this component should be disabled (unselected) by default. The user is free to select this component for installation, unless it is also HIDDEN.

DEPENDS lists the components on which this component depends. If this component is selected, then each of the components listed must also be selected. The dependency information is encoded within the installer itself, so that users cannot install inconsistent sets of components.

GROUP names the component group of which this component is a part. If not provided, the component will be a standalone component, not part of any component group. Component groups are described with the `cpack_add_component_group` command, detailed below.

INSTALL_TYPES lists the installation types of which this component is a part. When one of these installations types is selected, this component will automatically be selected. Installation types are described with the `cpack_add_install_type` command, detailed below.

DOWNLOADED indicates that this component should be downloaded on-the-fly by the installer, rather than packaged in with the installer itself. For more information, see the `cpack_configure_downloads` command.

ARCHIVE_FILE provides a name for the archive file created by CPack to be used for downloaded components. If not supplied, CPack will create a file with some name based on `CPACK_PACKAGE_FILE_NAME` and the name of the component. See `cpack_configure_downloads` for more information.

```
cpack_add_component_group - Describes a group of related CPack
                             installation components.
```

```
cpack_add_component_group(groupname
                           [DISPLAY_NAME name]
                           [DESCRIPTION description]
                           [PARENT_GROUP parent]
                           [EXPANDED]
                           [BOLD_TITLE])
```

The `cpack_add_component_group` describes a group of installation components, which will be placed together within the listing of options. Typically, component groups allow the user to select/deselect all of the components within a single group via a single group-level option. Use component groups to reduce the complexity of installers with many options. `groupname` is an arbitrary name used to identify the group in the `GROUP` argument of the `cpack_add_component` command, which is used to place a component in a group. The name of the group must not conflict with the name of any component.

DISPLAY_NAME is the displayed name of the component group, used in graphical installers to display the component group name. This value can be any string.

DESCRIPTION is an extended description of the component group, used in graphical installers to give the user additional information about the components within that group. Descriptions can span multiple lines using "\n" as the line separator. Typically, these descriptions should be no more than a few lines long.

PARENT_GROUP, if supplied, names the parent group of this group. Parent groups are used to establish a hierarchy of groups, providing an arbitrary hierarchy of groups.

EXPANDED indicates that, by default, the group should show up as "expanded", so that the user immediately sees all of the components within the group. Otherwise, the group will initially show up as a single entry.

BOLD_TITLE indicates that the group title should appear in bold, to call the user's attention to the group.

```
cpack_add_install_type - Add a new installation type containing  
a set of predefined component selections to the graphical  
installer.
```

```
cpack_add_install_type(typename [DISPLAY_NAME name])
```

The cpack_add_install_type command identifies a set of preselected components that represents a common use case for an application. For example, a "Developer" install type might include an application along with its header and library files, while an "End user" install type might just include the application's executable. Each component identifies itself with one or more install types via the INSTALL_TYPES argument to cpack_add_component.

DISPLAY_NAME is the displayed name of the install type, which will typically show up in a drop-down box within a graphical installer. This value can be any string.

```
cpack_configure_downloads - Configure CPack to download selected  
components on-the-fly as part of the installation process.
```

```
cpack_configure_downloads(site  
[UPLOAD_DIRECTORY dirname]  
[ALL]  
[ADD REMOVE | NO_ADD_REMOVE])
```

The cpack_configure_downloads command configures installation-time downloads of selected components. For each downloadable component, CPack will create an archive

containing the contents of that component, which should be uploaded to the given site. When the user selects that component for installation, the installer will download and extract the component in place. This feature is useful for creating small installers that only download the requested components, saving bandwidth. Additionally, the installers are small enough that they will be installed as part of the normal installation process, and the "Change" button in Windows Add/Remove Programs control panel will allow one to add or remove parts of the application after the original installation. On Windows, the downloaded-components functionality requires the ZipDLL plug-in for NSIS, available at:

http://nsis.sourceforge.net/ZipDLL_plug-in

On Mac OS X, installers that download components on-the-fly can only be built and installed on system using Mac OS X 10.5 or later.

The site argument is a URL where the archives for downloadable components will reside, e.g., <http://www.cmake.org/files/2.6.1/installer/> All of the archives produced by CPack should be uploaded to that location.

UPLOAD_DIRECTORY is the local directory where CPack will create the various archives for each of the components. The contents of this directory should be uploaded to a location accessible by the URL given in the site argument. If omitted, CPack will use the directory CPackUploads inside the CMake binary directory to store the generated archives.

The ALL flag indicates that all components be downloaded. Otherwise, only those components explicitly marked as DOWNLOADED or that have a specified ARCHIVE_FILE will be downloaded. Additionally, the ALL option implies ADD_REMOVE (unless NO_ADD_REMOVE is specified).

ADD_REMOVE indicates that CPack should install a copy of the installer that can be called from Windows' Add/Remove Programs dialog (via the "Modify" button) to change the set of installed components. NO_ADD_REMOVE turns off this behavior. This option is ignored on Mac OS X.

CPackRPM: The builtin (binary) CPack RPM generator (UNIX only)

CPackRPM may be used to create RPM package using CPack. CPackRPM is a CPack generator thus it uses the CPACK_* variables used by CPack. However CPackRPM has specific features which are controlled by the specifics CPACK_RPM_* variables. You'll find a detailed usage on the wiki:

<http://www.cmake.org/Wiki/CMake:CPackPackageGenerators>

CheckCCompilerFlag: Check whether the C compiler supports a given flag.

```
CHECK_C_COMPILER_FLAG(<flag> <var>)
<flag> - the compiler flag
<var> - variable to store the result
```

This internally calls the `check_c_source_compiles` macro. See `help` for `CheckCSourceCompiles` for a listing of variables that can modify the build.

CheckCSourceCompiles: Check if the given C source code compiles.

```
CHECK_C_SOURCE_COMPILES(<code> <var> [FAIL_REGEX <fail-regex>])
<code> - source code to try to compile
<var> - variable to store whether the source code compiled
<fail-regex> - fail if test output matches this regex
```

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags
CMAKE_REQUIRED_DEFINITIONS = list of macros to define
                           e.g. (-DFOO=bar)
CMAKE_REQUIRED_INCLUDES = list of include directories
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

CheckCSourceRuns: Check if the given C source code compiles and runs.

```
CHECK_C_SOURCE_RUNS(<code> <var>)
<code> - source code to try to compile
<var> - variable to store the result
        (1 for success, empty for failure)
```

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags
CMAKE_REQUIRED_DEFINITIONS = list of macros to define
                           e.g. (-DFOO=bar)
CMAKE_REQUIRED_INCLUDES = list of include directories
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

CheckCXXCompilerFlag: Check whether the CXX compiler supports a given flag.

```
CHECK_CXX_COMPILER_FLAG(<flag> <var>)
<flag> - the compiler flag
<var> - variable to store the result
```

This internally calls the check_cxx_source_compiles macro. See help for CheckCXXSourceCompiles for a listing of variables that can modify the build.

CheckCXXSourceCompiles: Check if the given C++ source code compiles.

```
CHECK_CXX_SOURCE_COMPILES(<code> <var>
                           [FAIL_REGEX <fail-regex>])
<code> - source code to try to compile
<var> - variable to store whether the source code compiled
<fail-regex> - fail if test output matches this regex
```

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags
CMAKE_REQUIRED_DEFINITIONS = list of macros to define
                           e.g. (-DFOO=bar)
CMAKE_REQUIRED_INCLUDES = list of include directories
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

CheckCXXSourceRuns: Check if the given C++ source code compiles and runs.

```
CHECK_CXX_SOURCE_RUNS(<code> <var>)
<code> - source code to try to compile
<var> - variable to store the result
        (1 for success, empty for failure)
```

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags
CMAKE_REQUIRED_DEFINITIONS = list of macros to define
                            e.g. (-DFOO=bar)
CMAKE_REQUIRED_INCLUDES = list of include directories
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

CheckFortranFunctionExists: macro which checks if the Fortran function exists

```
CHECK_FORTRAN_FUNCTION_EXISTS(FUNCTION VARIABLE)
FUNCTION - the name of the Fortran function
VARIABLE - variable to store the result
```

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

CheckIncludeFile: macro which checks the include file exists.

```
CHECK_INCLUDE_FILE(INCLUDE VARIABLE)
INCLUDE - name of include file
VARIABLE - variable to return result
```

an optional third argument is the CFlags to add to the compile line or you can use CMAKE_REQUIRED_FLAGS

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags
CMAKE_REQUIRED_DEFINITIONS = list of macros to define
                            e.g. (-DFOO=bar)
CMAKE_REQUIRED_INCLUDES = list of include directories
```

CheckIncludeFileCXX: Check if the include file exists.

```
CHECK_INCLUDE_FILE_CXX (INCLUDE VARIABLE)
```

INCLUDE - name of include file
VARIABLE - variable to return result

An optional third argument is the CFlags to add to the compile line or you can use CMAKE_REQUIRED_FLAGS. The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags  
CMAKE_REQUIRED_DEFINITIONS = list of macros to define  
                           e.g. (-DFOO=bar)  
CMAKE_REQUIRED_INCLUDES = list of include directories
```

CheckIncludeFiles: Check if the files can be included

```
CHECK_INCLUDE_FILES (INCLUDE VARIABLE)  
INCLUDE - list of files to include  
VARIABLE - variable to return result
```

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags  
CMAKE_REQUIRED_DEFINITIONS = list of macros to define  
                           e.g. (-DFOO=bar)  
CMAKE_REQUIRED_INCLUDES = list of include directories
```

CheckLibraryExists: Check if the library exists with the specified function.

```
CHECK_LIBRARY_EXISTS (LIBRARY FUNCTION LOCATION VARIABLE)  
LIBRARY - the name of the library you are looking for  
FUNCTION - the name of the function  
LOCATION - location where the library should be found  
VARIABLE - variable to store the result
```

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags
CMAKE_REQUIRED_DEFINITIONS = list of macros to define
                            e.g. (-DFOO=bar)
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

CheckStructHasMember: Check if the given struct or class has the specified member variable

```
CHECK_STRUCT_HAS_MEMBER (STRUCT MEMBER HEADER VARIABLE)
STRUCT - the name of the struct or class you are interested in
MEMBER - the member which existence you want to check
HEADER - the header(s) where the prototype should be declared
VARIABLE - variable to store the result
```

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags
CMAKE_REQUIRED_DEFINITIONS = list of macros to define
                            e.g. (-DFOO=bar)
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

Example:

```
CHECK_STRUCT_HAS_MEMBER("struct timeval" tv_sec sys/select.h
                        HAVE_TIMEVAL_TV_SEC)
```

CheckSymbolExists: Check if the symbol exists in include files

```
CHECK_SYMBOL_EXISTS(SYMBOL FILES VARIABLE)
SYMBOL      - symbol
FILES       - include files to check
VARIABLE    - variable to return result
```

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags  
CMAKE_REQUIRED_DEFINITIONS = list of macros to define  
                                e.g. (-DFOO=bar)  
CMAKE_REQUIRED_INCLUDES = list of include directories  
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

CheckTypeSize: Check sizeof a type

```
CHECK_TYPE_SIZE(TYPE VARIABLE [BUILTIN_TYPES_ONLY])
```

Check if the type exists and determine size of type. If the type exists, the size will be stored to the variable. This also calls `check_include_file` for `sys/types.h` `stdint.h` and `stddef.h`, setting `HAVE_SYS_TYPES_H`, `HAVE_STDINT_H`, and `HAVE_STDEDEF_H`. This is because many types are stored in these include files.

VARIABLE	- variable to store size if the type exists.
HAVE_\${VARIABLE}	- does the variable exists or not
BUILTIN_TYPES_ONLY	- The third argument is optional and if it is set to the string <code>BUILTIN_TYPES_ONLY</code> this macro will not check for any header files.

The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags  
CMAKE_REQUIRED_DEFINITIONS = list of macros to define  
                                e.g. (-DFOO=bar)  
CMAKE_REQUIRED_INCLUDES = list of include directories  
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

CheckVariableExists: Check if the variable exists.

```
CHECK_VARIABLE_EXISTS(VAR VARIABLE)

VAR      - the name of the variable
VARIABLE - variable to store the result
```

This macro is only for C variables. The following variables may be set before calling this macro to modify the way the check is run:

```
CMAKE_REQUIRED_FLAGS = string of compile command line flags
CMAKE_REQUIRED_DEFINITIONS = list of macros to define
                           e.g. (-DFOO=bar)
CMAKE_REQUIRED_LIBRARIES = list of libraries to link
```

ExternalProject: Create custom targets to build projects in external trees

The 'ExternalProject_Add' function creates a custom target to drive download, update/patch, configure, build, install and test steps of an external project:

```
ExternalProject_Add(<name>      # Name for custom target
[DEPENDS projects...]    # Targets on which the project depends
[PREFIX dir]              # Root dir for entire project
[LIST_SEPARATOR sep]      # Sep to be replaced by ; in cmd lines
[TMP_DIR dir]             # Directory to store temporary files
[STAMP_DIR dir]           # Directory to store step timestamps

#--Download step-----
[DOWNLOAD_DIR dir]        # Directory to store downloaded files
[DOWNLOAD_COMMAND cmd...] # Command to download source tree
[CVS_REPOSITORY cvsroot]  # CVSROOT of CVS repository
[CVS_MODULE mod]          # Module to checkout from CVS repo
[CVS_TAG tag]              # Tag to checkout from CVS repo
[SVN_REPOSITORY url]      # URL of Subversion repo
[SVN_REVISION rev]        # Revision to checkout from SVN repo
[URL /.../src.tgz]         # Full path or URL of source

#--Update/Patch step-----
[UPDATE_COMMAND cmd...]   # Source work-tree update command
[PATCH_COMMAND cmd...]    # Command to patch downloaded source

#--Configure step-----
[SOURCE_DIR dir]          # Source dir to be used for build
```

```
[CONFIGURE_COMMAND cmd...] # Build tree configuration command
[CMAKE_COMMAND /....cmake] # Specify alternative cmake exec
[CMAKE_GENERATOR gen]      # Specify generator for native build
[CMAKE_ARGS args...]       # Arguments to CMake command line

---Build step-----
[BINARY_DIR dir]           # Specify build dir location
[BUILD_COMMAND cmd...]     # Command to drive the native build
[BUILD_IN_SOURCE 1]         # Use source dir for build dir

---Install step-----
[INSTALL_DIR dir]          # Installation prefix
[INSTALL_COMMAND cmd...]   # Command to drive install after
                           # build

---Test step-----
[TEST_BEFORE_INSTALL 1]    # Add test step before install step
[TEST_AFTER_INSTALL 1]     # Add test step after install step
[TEST_COMMAND cmd...]     # Command to drive test
)
```

The *_DIR options specify directories for the project, with default directories computed as follows. If the PREFIX option is given to ExternalProject_Add() or the EP_PREFIX directory property is set, then an external project is built and installed under the specified prefix:

TMP_DIR	= <prefix>/tmp
STAMP_DIR	= <prefix>/src/<name>-stamp
DOWNLOAD_DIR	= <prefix>/src
SOURCE_DIR	= <prefix>/src/<name>
BINARY_DIR	= <prefix>/src/<name>-build
INSTALL_DIR	= <prefix>

Otherwise, if the EP_BASE directory property is set then components of an external project are stored under the specified base:

TMP_DIR	= <base>/tmp/<name>
STAMP_DIR	= <base>/Stamp/<name>
DOWNLOAD_DIR	= <base>/Download/<name>
SOURCE_DIR	= <base>/Source/<name>
BINARY_DIR	= <base>/Build/<name>
INSTALL_DIR	= <base>/Install/<name>

If no PREFIX, EP_PREFIX, or EP_BASE is specified then the default is to set PREFIX to "<name>-prefix". Relative paths are interpreted with respect to the build directory corresponding to the source directory in which ExternalProject_Add is invoked.

If SOURCE_DIR is explicitly set to an existing directory the project will be built from it. Otherwise a download step must be specified using one of the DOWNLOAD_COMMAND, CVS_*, SVN_*, or URL options. The URL option may refer locally to a directory or source tarball, or refer to a remote tarball (e.g. <http://.../src.tgz>).

The 'ExternalProject_Add_Step' function adds a custom step to an external project:

```
ExternalProject_Add_Step(  
    <name> <step>          # Names of project and custom step  
    [COMMAND cmd...]        # Command line invoked by this step  
    [COMMENT "text..."]       # Text printed when step executes  
    [DEPENDS steps...]      # Steps on which this step depends  
    [DEPENDERS steps...]    # Steps that depend on this step  
    [DEPENDS files...]      # Files on which this step depends  
    [ALWAYS 1]               # No stamp file, step always runs  
    [WORKING_DIRECTORY dir] # Working directory for command  
)
```

The command line, comment, and working directory of every standard and custom step is processed to replace tokens <SOURCE_DIR>, <BINARY_DIR>, <INSTALL_DIR>, and <TMP_DIR> with corresponding property values.

The 'ExternalProject_Get_Property' function retrieves external project target properties:

```
ExternalProject_Get_Property(<name> [prop1 [prop2 [...]]])
```

It stores property values in variables of the same name. Property names correspond to the keyword argument names of 'ExternalProject_Add'.

FeatureSummary: Macros for generating a summary of enabled/disabled features

PRINT_ENABLED_FEATURES() - Print a summary of all enabled features. By default all successfull FIND_PACKAGE() calls will appear here, except the ones which used the QUIET keyword. Additional features can be added by appending an entry to the global ENABLED_FEATURES property. If SET_FEATURE_INFO() is used for that feature, the output will be much more informative.

PRINT_DISABLED_FEATURES() - Same as **PRINT_ENABLED_FEATURES()**, but for disabled features. It can be extended the same way by adding to the global property **DISABLED_FEATURES**.

SET_FEATURE_INFO(NAME DESCRIPTION [URL [COMMENT]]) - Use this macro to set up information about the named feature, which will then be displayed by **PRINT_ENABLED/DISABLED_FEATURES()**. For Example:

```
SET_FEATURE_INFO(LibXml2 "XML processing library."  
                  "http://xmlsoft.org/")
```

FindALSA: Find alsal

FindASPELL: Try to find ASPELL

FindAVIFILE: Locate AVIFILE library and include paths

FindBISON: Find bison executable and provides macros to generate custom build rules

FindBLAS: Find BLAS library

FindBZip2: Try to find BZip2

FindBoost: Try to find Boost include dirs and libraries

Usage of this module is as follows. Using Header-Only libraries from within Boost:

```
find_package( Boost 1.36.0 )  
if(Boost_FOUND)  
    include_directories(${Boost_INCLUDE_DIRS})  
    add_executable(foo foo.cc)  
endif()
```

Using actual libraries from within Boost:

```
set(Boost_USE_STATIC_LIBS    ON)  
set(Boost_USE_MULTITHREADED ON)  
find_package( Boost 1.36.0 COMPONENTS date_time filesystem  
             system ... )
```

```
if(Boost_FOUND)
    include_directories(${Boost_INCLUDE_DIRS})
    add_executable(foo foo.cc)
    target_link_libraries(foo ${Boost_LIBRARIES})
endif()
```

The components list needs to contain actual names of boost libraries only, such as "date_time" for "libboost_date_time". If you're using parts of Boost that contain header files only (e.g. foreach) you do not need to specify COMPONENTS.

You should provide a minimum version number that should be used. If you provide this version number and specify the REQUIRED attribute, this module will fail if it can't find the specified or a later version. If you specify a version number this is automatically put into the considered list of version numbers and thus doesn't need to be specified in the Boost_ADDITIONAL_VERSIONS variable (see below).

NOTE for Visual Studio Users: Automatic linking is used on MSVC & Borland compilers by default when including things in Boost. It's important to note that setting Boost_USE_STATIC_LIBS to OFF is NOT enough to get you dynamic linking, should you need this feature. Automatic linking typically uses static libraries with a few exceptions (Boost.Python is one).

Please see the section below near Boost_LIB_DIAGNOSTIC_DEFINITIONS for more details. Adding a TARGET_LINK_LIBRARIES() as shown in the example above appears to cause VS to link dynamically if Boost_USE_STATIC_LIBS gets set to OFF. It is suggested you avoid automatic linking since it will make your application less portable.

Boost_ADDITIONAL_VERSIONS

OK, so the Boost_ADDITIONAL_VERSIONS variable can be used to specify a list of boost version numbers that should be taken into account when searching for Boost. Unfortunately boost puts the version number into the actual filename for the libraries, so this variable will certainly be needed in the future when new Boost versions are released.

Currently this module searches for the following version numbers: 1.33, 1.33.0, 1.33.1, 1.34, 1.34.0, 1.34.1, 1.35, 1.35.0, 1.35.1, 1.36, 1.36.0, 1.36.1, 1.37, 1.37.0, 1.38, 1.38.0, 1.39, 1.39.0, 1.40, 1.40.0

NOTE: If you add a new major 1.x version in Boost_ADDITIONAL_VERSIONS you should add both 1.x and 1.x.0 as shown above. Official Boost include directories omit the 3rd version number from include paths if it is 0 although not all binary Boost releases do so.

```
set(Boost_ADDITIONAL_VERSIONS "0.99" "0.99.0" "1.78" "1.78.0")
```

Variables used by this module, they can change the default behaviour and need to be set before calling `find_package`:

Boost_USE_MULTITHREADED

Can be set to OFF to use the non-multithreaded boost libraries. If not specified, defaults to ON.

Boost_USE_STATIC_LIBS

Can be set to ON to force the use of the static boost libraries. Defaults to OFF.

Further documentation on other variables used by this module which you may want to set can be found in the module itself.

FindBullet: Try to find the Bullet physics engine

This module defines the following variables

BULLET_FOUND	- Was bullet found
BULLET_INCLUDE_DIRS	- the Bullet include directories
BULLET_LIBRARIES	- Link to this, by default it includes all bullet components (Dynamics, Collision, LinearMath, & SoftBody)

This module accepts the following variables:

BULLET_ROOT - Can be set to bullet install path or Windows build path

FindCABLE: Find CABLE

FindCUDA: Tools for building CUDA C files: libraries and build dependencies.

This script locates the NVIDIA CUDA C tools. It should work on linux, windows, and mac and should be reasonably up to date with CUDA C releases. This script makes use of the standard `find_package` arguments of `<VERSION>`, `REQUIRED` and `QUIET`. `CUDA_FOUND` will report if an acceptable version of CUDA was found.

FindCURL: Find curl

FindCVS:

FindCoin3D: Find Coin3D (Open Inventor)

FindCups: Try to find the Cups printing system

FindCurses: Find the curses include file and library

FindCxxTest: Find CxxTest

Find the CxxTest suite and declare a helper macro for creating unit tests and integrating them with CTest. For more details on CxxTest see <http://cxxtest.tigris.org>

INPUT Variables**CXXTEST_USE_PYTHON**

If true, the CXXTEST_ADD_TEST macro will use the Python test generator instead of Perl.

OUTPUT Variables**CXXTEST_FOUND**

True if the CxxTest framework was found

CXXTEST_INCLUDE_DIR

Where to find the CxxTest include directory

CXXTEST_PERL_TESTGEN_EXECUTABLE

The perl-based test generator.

CXXTEST_PYTHON_TESTGEN_EXECUTABLE

The python-based test generator.

MACROS for optional use by CMake users:**CXXTEST_ADD_TEST(<test_name> <gen_source_file>
<input_files_to_testgen...>)**

Creates a CxxTest runner and adds it to the CTest testing suite
Parameters:

test_name The name of the test**gen_source_file** The generated source filename to be generated by CxxTest**input_files_to_testgen** The list of header files containing the CxxTest::TestSuite's to be included in this runner

Sample usage:

```
find_package(CxxTest)
if(CXXTEST_FOUND)
    include_directories(${CXXTEST_INCLUDE_DIR})
    enable_testing()
    CXXTEST_ADD_TEST(unittest_foo foo_test.cc
                      ${CMAKE_CURRENT_SOURCE_DIR}/foo_test.h)
    target_link_libraries(unittest_foo foo) # as needed
endif()
```

This will (if CxxTest is found):

1. Invoke the testgen executable to autogenerated foo_test.cc in the binary tree from "foo_test.h" in the current source directory.
2. Create an executable and test called unittest_foo.

FindCygwin: this module looks for Cygwin

FindDCMTK: find DCMTK libraries

FindDart: Find DART

FindDevIL: This module locates the developer's image library.

FindDoxygen: This module looks for Doxygen and the path to Graphviz's dot

FindEXPAT: Find the native EXPAT headers and libraries.

FindFLEX: Find flex executable and provides a macro to generate custom build rules

FindFLTK: Find the native FLTK includes and library

FindFLTK2: Find the native FLTK2 includes and library

FindFreeType: Locate FreeType library

FindGCCXML: Find the GCC-XML front-end executable.

FindGDAL: Locate gdal

FindGIF:

FindGLUT: try to find glut library and include files

FindGTK: try to find GTK (and glib) and GTKGLArea

FindGTK2: FindGTK2.cmake

FindGTest: Locate the Google C++ Testing Framework.

FindGettext: Find GNU gettext tools

FindGnuTLS: Try to find the GNU Transport Layer Security library (gnutls)

FindGnuplot: this module looks for gnuplot

FindHDF5: Find HDF5, a library for reading and writing self describing array data.

This module invokes the HDF5 wrapper compiler that should be installed alongside HDF5. Depending upon the HDF5 Configuration, the wrapper compiler is called either h5cc or h5pcc. If this succeeds, the module will then call the compiler with the -show argument to see what flags are used when compiling an HDF5 client application.

FindHSPELL: Try to find HSPELL

FindHTMLHelp: This module looks for Microsoft HTML Help Compiler

FindITK: Find an ITK installation or build tree.

FindImageMagick: Find the ImageMagick binary suite.

FindJNI: Find JNI java libraries.

FindJPEG: Find the native JPEG includes and library.

FindJasper: Try to find the Jasper JPEG2000 library

FindJava: This module finds if Java is installed and determines where the include files and libraries are.

FindKDE3: Find the KDE3 include and library dirs, KDE preprocessors and define some macros.

FindKDE4: Find KDE4 and provide all necessary variables and macros to compile software for it.

FindLAPACK: Find the LAPACK library

FindLATEX: This module finds if Latex is installed and determines where the executables are.

FindLibXml2: Try to find LibXml2

FindLibXslt: Try to find LibXslt

FindLua50: Find Lua version 5.0

FindLua51: Find Lua version 5.1

FindMFC: Find MFC on Windows

FindMPEG: Find the native MPEG includes and library

FindMPEG2: Find the native MPEG2 includes and library

FindMPI: Message Passing Interface (MPI) module.

FindMatlab: this module looks for Matlab

FindMotif: Try to find Motif (or lesstif)

FindOpenAL:

FindOpenGL: Try to find OpenGL

FindOpenMP: Finds OpenMP support

FindOpenSSL: Try to find the OpenSSL encryption library

FindOpenSceneGraph: Comprehensive module to find OpenSceneGraph and all of its various parts.

FindOpenThreads: Find the OpenThreads, C++ based threading library.

FindPHP4: Find PHP4

FindPNG: Find the native PNG includes and library

FindPackageHandleStandardArgs:

```
FIND_PACKAGE_HANDLE_STANDARD_ARGS(NAME  
    (DEFAULT_MSG "Custom failure message") VAR1 ... )
```

This macro is intended to be used in FindXXX.cmake modules files. It handles the REQUIRED and QUIET argument to FIND_PACKAGE() and it also sets the <UPPERCASED_NAME>_FOUND variable. The package is found if all variables listed are TRUE. For example:

```
FIND_PACKAGE_HANDLE_STANDARD_ARGS(LibXml2 DEFAULT_MSG  
    LIBXML2_LIBRARIES  
    LIBXML2_INCLUDE_DIR)
```

LibXml2 is considered to be found, if both LIBXML2_LIBRARIES and LIBXML2_INCLUDE_DIR are valid. Then also LIBXML2_FOUND is set to TRUE. If it is not found and REQUIRED was used, it fails with FATAL_ERROR, independent whether QUIET was used or not. If it is found, the location is reported using the VAR1 argument, so here a message "Found LibXml2: /usr/lib/libxml2.so" will be printed out. If the second argument is DEFAULT_MSG, the message in the failure case will be "Could NOT find LibXml2", if you don't like this message you can specify your own custom failure message there.

FindPackageMessage:

```
FIND_PACKAGE_MESSAGE(<name> "message for user"  
    "find result details")
```

This macro is intended to be used in FindYYY.cmake modules files. It will print a message once for each unique find result. This is useful for telling the user where a package was found. The first argument specifies the name (YYY) of the package. The second argument specifies the message to display. The third argument lists details about the find result so that if they

change the message will be displayed again. The macro also obeys the QUIET argument to the `find_package` command. For example:

```
if(X11_FOUND)
    FIND_PACKAGE_MESSAGE(X11 "Found X11: ${X11_X11_LIB}"
                         "[${X11_X11_LIB}][${X11_INCLUDE_DIR}]")
else()
    ...
endif ()
```

FindPerl: Find perl

FindPerlLibs: Find Perl libraries

FindPhysFS: Locate the PhysFS library

FindPike: This module finds PIKE and determines where the include files and libraries are.

FindPkgConfig: a pkg-config module for CMake

Usage:

```
pkg_check_modules(<PREFIX> [REQUIRED] <MODULE> [<MODULE>]*)
    checks for all the given modules
```

```
pkg_search_module(<PREFIX> [REQUIRED] <MODULE> [<MODULE>]*)
    checks for given modules and uses the first working one
```

When the 'REQUIRED' argument was set, macros will fail with an error when module(s) could not be found

It sets the following variables:

```
PKG_CONFIG_FOUND      ... true if pkg-config works on the system
PKG_CONFIG_EXECUTABLE ... pathname of the pkg-config program
<PREFIX>_FOUND        ... set to 1 if module(s) exist
```

For the following variables two sets of values exist; first one is the common one and has the given PREFIX. The second set contains flags which are given out when `pkgconfig` was called with the '--static' option.

```

<XPREFIX>_LIBRARIES      ... only the libraries (w/o the '-l')
<XPREFIX>_LIBRARY_DIRS   ... the paths of the libraries
                           (w/o the '-L')
<XPREFIX>_LDFLAGS        ... all required linker flags
<XPREFIX>_LDFLAGS_OTHER  ... all other linker flags
<XPREFIX>_INCLUDE_DIRS   ... the '-I' preprocessor flags
                           (w/o the '-I')
<XPREFIX>_CFLAGS          ... all required cflags
<XPREFIX>_CFLAGS_OTHER   ... the other compiler flags

<XPREFIX> = <PREFIX>           for the common case
<XPREFIX> = <PREFIX>_STATIC for static linking

```

There are some special variables whose prefix depends on the count of given modules. When there is only one module, <PREFIX> stays unchanged. When there are multiple modules, the prefix will be changed to <PREFIX>_<MODNAME>:

```

<XPREFIX>_VERSION      ... version of the module
<XPREFIX>_PREFIX        ... prefix-directory of the module
<XPREFIX>_INCLUDEDIR    ... include-dir of the module
<XPREFIX>_LIBDIR        ... lib-dir of the module

<XPREFIX> = <PREFIX>  when |MODULES| == 1, else
<XPREFIX> = <PREFIX>_<MODNAME>

```

A <MODULE> parameter can have the following formats:

```

{MODNAME}            ... matches any version
{MODNAME}>={VERSION} ... at least version <VERSION> is required
{MODNAME}={VERSION} ... exactly version <VERSION> is required
{MODNAME}<={VERSION} ... modules must not be newer than
                       <VERSION>

```

Examples

```
pkg_check_modules (GLIB2    glib-2.0)

pkg_check_modules (GLIB2    glib-2.0>=2.10)
    requires at least version 2.10 of glib2 and defines e.g.
    GLIB2_VERSION=2.10.3

pkg_check_modules (FOO      glib-2.0>=2.10 gtk+-2.0)
    requires both glib2 and gtk2, and defines e.g.
    FOO_glib-2.0_VERSION=2.10.3
    FOO_gtk+-2.0_VERSION=2.8.20

pkg_check_modules (XRENDER REQUIRED xrender)
    defines e.g.:
    XRENDER_LIBRARIES=Xrender;X11
    XRENDER_STATIC_LIBRARIES=Xrender;X11;pthread;Xau;Xdmcp

pkg_search_module (BAR      libxml-2.0 libxml2 libxml>=2)
```

FindProducer: Find the producer library

FindProtobuf: Locate and configure the Google Protocol Buffers library.

FindPythonInterp: Find python interpreter

FindPythonLibs: Find python libraries

FindQt: Searches for all installed versions of QT.

FindQt3: Locate Qt 3 include paths and libraries

FindQt4: This module can be used to find Qt4.

FindQuickTime: Locate QuickTime library and includes.

FindRTI: Try to find M&S HLA RTI libraries and includes.

FindRuby: This module finds Ruby and determines where the include files and libraries are.

FindSDL: Locate the SDL Library and includes.

FindSDL_image: Locate SDL_image library.

FindSDL_mixer: Locate SDL_mixer library.

FindSDL_net: Locate SDL_net library.

FindSDL_sound: Locates the SDL_sound library

FindSDL_ttf: Locate SDL_ttf library.

FindSWIG: Find the SWIG wrapper generator.

FindSelfPackers: Find upx

FindSquish: This module can be used to find Squish.

FindSubversion: Extract information from a subversion client

The module defines the following variables:

```
Subversion_SVN_EXECUTABLE - path to svn command line client
Subversion_VERSION SVN - version of svn command line client
Subversion_FOUND - true if the command line client was found
```

If the command line client executable is found the macro

```
Subversion_WC_INFO(<dir> <var-prefix>)
```

is defined to extract information of a subversion working copy at a given location. The macro defines the following variables:

```
<var-prefix>_WC_URL - url of the repository (at <dir>)
<var-prefix>_WC_ROOT - root url of the repository
<var-prefix>_WC_REVISION - current revision
<var-prefix>_WC_LAST_CHANGED_AUTHOR - author of last commit
<var-prefix>_WC_LAST_CHANGED_DATE - date of last commit
<var-prefix>_WC_LAST_CHANGED_REV - revision of last commit
<var-prefix>_WC_LAST_CHANGED_LOG - last log of base revision
<var-prefix>_WC_INFO - output of command `svn info <dir>'
```

Example usage:

```
find_package(Subversion)
if(Subversion_FOUND)
    Subversion_WC_INFO(${PROJECT_SOURCE_DIR} Project)
    message ("Current revision is ${Project_WC_REVISION}")
    Subversion_WC_LOG(${PROJECT_SOURCE_DIR} Project)
    message ("Last changed log is ${Project_LAST_CHANGED_LOG}")
endif()
```

FindTCL: This module finds if Tcl is installed and locates the include files and libraries.

FindTIFF: Find the native TIFF includes and library.

FindTclStub: This module finds Tcl stub libraries.

FindTclsh: Find tclsh

FindThreads: This module determines the thread library of the system.

The following variables are set

```
CMAKE_THREAD_LIBS_INIT      - the thread library
CMAKE_USE_SPROC_INIT        - are we using sproc?
CMAKE_USE_WIN32_THREADS_INIT - using WIN32 threads?
CMAKE_USE_PTHREADS_INIT     - are we using pthreads
CMAKE_HP_PTHREADS_INIT      - are we using hp pthreads
```

FindUnixCommands: Find UNIX commands from cygwin

FindVTK: Find a VTK installation or build tree.

FindWget: Find wget

FindWish: Find wish installation

FindX11: Find X11 installation

FindXMLRPC: Find the native XMLRPC headers and libraries.

FindZLIB: Find the native ZLIB includes and library

Findosg*: Find a specific part of open scene graph. See the FindOpenSceneGraph module as well.

FindwxWidgets: Find a wxWidgets (a.k.a., wxWindows) installation.

FortranCInterface: Fortran/C Interface Detection

This module automatically detects the API by which C and Fortran languages interact. Variables indicate if the mangling is found:

```
FortranCInterface_GLOBAL_FOUND
  = Global subroutines and functions
FortranCInterface_MODULE_FOUND
  = Module subroutines and functions
    (declared by "MODULE PROCEDURE")
```

A function is provided to generate a C header file containing macros to mangle symbol names:

```
FortranCInterface_HEADER(<file>
  [MACRO_NAMESPACE <macro-ns>]
  [SYMBOL_NAMESPACE <ns>]
  [SYMBOLS [<module>:]<function> ...])
```

It generates in <file> definitions of the following macros:

```
#define FortranCInterface_GLOBAL (name,NAME) ...
#define FortranCInterface_GLOBAL_(name,NAME) ...
#define FortranCInterface_MODULE (mod,name, MOD,NAME) ...
#define FortranCInterface_MODULE_(mod,name, MOD,NAME) ...
```

These macros mangle four categories of Fortran symbols, respectively:

- Global symbols without '_': call mysub()
- Global symbols with '_': call my_sub()
- Module symbols without '_': use mymod; call mysub()
- Module symbols with '_': use mymod; call my_sub()

If mangling for a category is not known, its macro is left undefined. All macros require raw names in both lower case and upper case. The MACRO_NAMESPACE option replaces the default "FortranCInterface_" prefix with a given namespace "<macro-ns>". The SYMBOLS option lists symbols to mangle automatically with C preprocessor definitions:

```
<function>          ==> #define <ns><function> ...
<module>:<function> ==> #define <ns><module>_<function> ...
```

If the mangling for some symbol is not known then no preprocessor definition is created, and a warning is displayed. The SYMBOL_NAMESPACE option prefixes all preprocessor definitions generated by the SYMBOLS option with a given namespace "<ns>". Example usage:

```
include(FortranCInterface)
FortranCInterface_HEADER(FC.h MACRO_NAMESPACE "FC_")
```

This creates a "FC.h" header that defines mangling macros FC_GLOBAL(), FC_GLOBAL_(), FC_MODULE(), and FC_MODULE_(). Another example:

```
include(FortranCInterface)
FortranCInterface_HEADER(FCMangle.h
                        MACRO_NAMESPACE "FC_"
                        SYMBOL_NAMESPACE "FC_"
                        SYMBOLS mysub mymod:my_sub)
```

This creates a "FC.h" header that defines the same FC_*() mangling macros as the previous example plus preprocessor symbols FC_mysub and FC_mymod_my_sub. Another function is provided to verify that the Fortran and C/C++ compilers work together:

```
FortranCInterface_VERIFY([CXX] [QUIET])
```

It tests whether a simple test executable using Fortran and C (and C++ when the CXX option is given) compiles and links successfully. The result is stored in the cache entry FortranCInterface_VERIFIED_C (or FortranCInterface_VERIFIED_CXX if CXX is given) as a boolean. If the check fails and QUIET is not given the function terminates with a FATAL_ERROR message describing the problem. The purpose of this check is to stop a build early for incompatible compiler combinations.

FortranCInterface is aware of possible GLOBAL and MODULE manglings for many Fortran compilers, but it also provides an interface to specify new possible manglings. Set the variables

```
FortranCInterface_GLOBAL_SYMBOLS  
FortranCInterface_MODULE_SYMBOLS
```

before including FortranCInterface to specify manglings of the symbols "MySub", "My_Sub", "MyModule:MySub", and "My_Module:My_Sub". For example, the code:

```
set(FortranCInterface_GLOBAL_SYMBOLS mysub_ my_sub__ MYSUB_)  
set(FortranCInterface_MODULE_SYMBOLS  
    __mymodule_MOD_mysub __my_module_MOD_my_sub)  
include(FortranCInterface)
```

tells FortranCInterface to try given GLOBAL and MODULE manglings. (The carets point at raw symbol names for clarity in this example but are not needed.)

GetPrerequisites: See section 4.11

InstallRequiredSystemLibraries: See section 4.11

By including this file, all files in the CMAKE_INSTALL_DEBUG_LIBRARIES, will be installed with INSTALL_PROGRAMS into /bin for WIN32 and /lib for non-win32. If CMAKE_SKIP_INSTALL_RULES is set to TRUE before including this file, then the INSTALL command is not called. The user can use the variable CMAKE_INSTALL_SYSTEM_RUNTIME_LIBS to use a custom install command and install them into any directory they want. If it is the MSVC compiler, then the microsoft run time libraries will be found and automatically added to the CMAKE_INSTALL_DEBUG_LIBRARIES, and installed. If

CMAKE_INSTALL_DEBUG_LIBRARIES is set and it is the MSVC compiler, then the debug libraries are installed when available. If CMAKE_INSTALL_MFC_LIBRARIES is set then the MFC run time libraries are installed as well as the CRT run time libraries.

SelectLibraryConfigurations:

```
• select_library_configurations( basename )
```

This macro takes a library base name as an argument, and will choose good values for basename_LIBRARY, basename_LIBRARIES, basename_LIBRARY_DEBUG, and basename_LIBRARY_RELEASE depending on what has been found and set. If only basename_LIBRARY_RELEASE is defined, basename_LIBRARY, basename_LIBRARY_DEBUG, and basename_LIBRARY_RELEASE will be set to the release value. If only basename_LIBRARY_DEBUG is defined, then basename_LIBRARY, basename_LIBRARY_DEBUG and basename_LIBRARY_RELEASE will take the debug value.

If the generator supports configuration types, then basename_LIBRARY and basename_LIBRARIES will be set with debug and optimized flags specifying the library to be used for the given configuration. If no build type has been set or the generator in use does not support configuration types, then basename_LIBRARY and basename_LIBRARIES will take only the release values.

TestBigEndian: Define macro to determine endian type

```
TEST_BIG_ENDIAN(VARIABLE)
VARIABLE - variable to store the result to
```

TestCXXAcceptsFlag: Test CXX compiler for a flag

Check if the CXX compiler accepts a flag

```
Macro CHECK_CXX_ACCEPTS_FLAG(FLAGS VARIABLE) -
    checks if the flag is accepted
FLAGS - the flags to try
VARIABLE - variable to store the result
```

TestForANSIForScope: Check for ANSI for scope support

```
CMAKE_NO_ANSI_FOR_SCOPE - holds result
```

TestForANSIStreamHeaders: Test for compiler support of ANSI stream headers

Check if we they have the standard ansi stream files (without the .h)

```
CMAKE_NO_ANSI_STREAM_HEADERS - defined by the results
```

TestForSSTREAM:

```
CMAKE_NO_ANSI_STRING_STREAM - defined by the results
```

TestForSTDNamespace: Test for std:: namespace support

check if the compiler supports std:: on stl classes

```
CMAKE_NO_STD_NAMESPACE - defined by the results
```

UseEcos: This module defines variables and macros required to build eCos applications.

UseQt4: Use Module for QT4

Sets up C and C++ to use Qt 4. It is assumed that FindQt.cmake has already been loaded. See FindQt.cmake for information on how to load Qt 4 into your CMake project.

UseSWIG: SWIG module for CMake

Use_wxWindows:

UsewxWidgets: Convenience include for using wxWidgets library

Appendix E - Properties

This is the documentation for the properties supported by CMake. Properties can have different scopes. They can be assigned to a source file, a directory, a target, test, etc. By modifying the values of properties the behaviour of the build system can be customized.

Properties of Global Scope

ALLOW_DUPLICATE_CUSTOM_TARGETS: Allow duplicate custom targets to be created.

Normally CMake requires that all targets built in a project have globally unique logical names (see policy CMP0002). This is necessary to generate meaningful project file names in Xcode and VS IDE generators. It also allows the target names to be referenced unambiguously.

Makefile generators are capable of supporting duplicate custom target names. For projects that care only about Makefile generators and do not wish to support Xcode or VS IDE generators, one may set this property to true to allow duplicate custom targets. The property allows multiple add_custom_target command calls in different directories to specify the same target name. However, setting this property will cause non-Makefile generators to produce an error and refuse to generate the project.

DEBUG_CONFIGURATIONS: Specify which configurations are for debugging.

The value must be a semi-colon separated list of configuration names. Currently this property is used only by the `target_link_libraries` command (see its documentation for details). Additional uses may be defined in the future.

This property must be set at the top level of the project and before the first `target_link_libraries` command invocation. If any entry in the list does not match a valid configuration for the project the behavior is undefined.

DISABLED_FEATURES: List of features which are disabled during the CMake run.

List of features which are disabled during the CMake run. By default it contains the names of all packages which were not found. This is determined using the `<NAME>_FOUND` variables. Packages which are searched QUIET are not listed. A project can add its own features to this list. This property is used by the macros in `FeatureSummary.cmake`.

ENABLED_FEATURES: List of features which are enabled during the CMake run.

List of features which are enabled during the CMake run. By default it contains the names of all packages which were found. This is determined using the `<NAME>_FOUND` variables. Packages which are searched QUIET are not listed. A project can add its own features to this list. This property is used by the macros in `FeatureSummary.cmake`.

ENABLED_LANGUAGES: Read-only property that contains the list of currently enabled languages

Set to list of currently enabled languages.

FIND_LIBRARY_USE_LIB64_PATHS: Whether `FIND_LIBRARY` should automatically search lib64 directories.

`FIND_LIBRARY_USE_LIB64_PATHS` is a boolean specifying whether the `FIND_LIBRARY` command should automatically search the lib64 variant of directories called lib in the search path when building 64-bit binaries.

FIND_LIBRARY_USE_OPENBSD_VERSIONING: Whether `FIND_LIBRARY` should find OpenBSD-style shared libraries.

This property is a boolean specifying whether the `FIND_LIBRARY` command should find shared libraries with OpenBSD-style versioned extension: `".so.<major>.<minor>"`. The property is set to true on OpenBSD and false on other platforms.

GLOBAL_DEPENDS_DEBUG_MODE: Enable global target dependency graph debug mode.

CMake automatically analyzes the global inter-target dependency graph at the beginning of native build system generation. This property causes it to display details of its analysis to `stderr`.

GLOBAL_DEPENDS_NO_CYCLES: Disallow global target dependency graph cycles.

CMake automatically analyzes the global inter-target dependency graph at the beginning of native build system generation. It reports an error if the dependency graph contains a cycle that does not consist of all STATIC library targets. This property tells CMake to disallow all cycles completely, even among static libraries.

IN_TRY_COMPILE: Read-only property that is true during a try-compile configuration.

True when building a project inside a TRY_COMPILE or TRY_RUN command.

PACKAGES_FOUND: List of packages which were found during the CMake run.

List of packages which were found during the CMake run. Whether a package has been found is determined using the <NAME>_FOUND variables.

PACKAGES_NOT_FOUND: List of packages which were not found during the CMake run.

List of packages which were not found during the CMake run. Whether a package has been found is determined using the <NAME>_FOUND variables.

REPORT_UNDEFINED_PROPERTIES: If set, report any undefined properties to this file.

If this property is set to a filename then when CMake runs it will report any properties or variables that were accessed but not defined into the filename specified in this property.

RULE_LAUNCH_COMPILE: Specify a launcher for compile rules.

Makefile generators prefix compiler commands with the given launcher command line. This is intended to allow launchers to intercept build problems with high granularity. Non-Makefile generators currently ignore this property.

RULE_LAUNCH_CUSTOM: Specify a launcher for custom rules.

Makefile generators prefix custom commands with the given launcher command line. This is intended to allow launchers to intercept build problems with high granularity. Non-Makefile generators currently ignore this property.

RULE_LAUNCH_LINK: Specify a launcher for link rules.

Makefile generators prefix link and archive commands with the given launcher command line. This is intended to allow launchers to intercept build problems with high granularity. Non-Makefile generators currently ignore this property.

RULE_MESSAGES: Specify whether to report a message for each make rule.

This property specifies whether Makefile generators should add a progress message describing what each build rule does. If the property is not set the default is ON. Set the property to OFF to disable granular messages and report only as each target completes. This is intended to allow scripted builds to avoid the build time cost of detailed reports. If a

CMAKE_RULE_MESSAGES cache entry exists its value initializes the value of this property. Non-Makefile generators currently ignore this property.

TARGET_ARCHIVES_MAY_BE_SHARED_LIBS: Set if shared libraries may be named like archives.

On AIX shared libraries may be named "lib<name>.a". This property is set to true on such platforms.

TARGET_SUPPORTS_SHARED_LIBS: Does the target platform support shared libraries.

TARGET_SUPPORTS_SHARED_LIBS is a boolean specifying whether the target platform supports shared libraries. Basically all current general purpose OS do so, the exception are usually embedded systems with no or special OSs.

CMAKE_DELETE_CACHE_CHANGE_VARS: Internal property

Used to detect compiler changes, Do not set.

Properties on Directories

ADDITIONAL_MAKE_CLEAN_FILES: Additional files to clean during the make clean stage.

A list of files that will be cleaned as a part of the "make clean" stage.

CACHE_VARIABLES: List of cache variables available in the current directory.

This read-only property specifies the list of CMake cache variables currently defined. It is intended for debugging purposes.

CLEAN_NO_CUSTOM: Should the output of custom commands be left.

If this is true then the outputs of custom commands for this directory will not be removed during the "make clean" stage.

COMPILE_DEFINITIONS: Preprocessor definitions for compiling a directory's sources.

The **COMPILE_DEFINITIONS** property may be set to a semicolon-separated list of preprocessor definitions using the syntax VAR or VAR=value. Function-style definitions are not supported. CMake will automatically escape the value correctly for the native build system (note that CMake language syntax may require escapes to specify some values). This property may be set on a per-configuration basis using the name **COMPILE_DEFINITIONS_<CONFIG>** where <CONFIG> is an upper-case name (ex. "COMPILE_DEFINITIONS_DEBUG"). This property will be initialized in each directory by its value in the directory's parent.

CMake will automatically drop some definitions that are not supported by the native build tool. The VS6 IDE does not support definition values with spaces (but NMake does).

Disclaimer: Most native build tools have poor support for escaping certain values. CMake has work-arounds for many cases but some values may just not be possible to pass correctly. If a value does not seem to be escaped correctly, do not attempt to work-around the problem by adding escape sequences to the value. Your work-around may break in a future version of CMake that has improved escape support. Instead consider defining the macro in a (configured) header file. Then report the limitation.

COMPILE_DEFINITIONS_<CONFIG>: Per-configuration preprocessor definitions in a directory.

This is the configuration-specific version of COMPILE_DEFINITIONS. This property will be initialized in each directory by its value in the directory's parent.

DEFINITIONS: For CMake 2.4 compatibility only. Use COMPILE_DEFINITIONS instead.

This read-only property specifies the list of flags given so far to the add_definitions command. It is intended for debugging purposes. Use the COMPILE_DEFINITIONS instead.

EXCLUDE_FROM_ALL: Exclude the directory from the all target of its parent.

A property on a directory that indicates if its targets are excluded from the default build target. If it is not, then with a Makefile for example typing make will cause the targets to be built. The same concept applies to the default build of other generators.

IMPLICIT_DEPENDS_INCLUDE_TRANSFORM: Specify #include line transforms for dependencies in a directory.

This property specifies rules to transform macro-like #include lines during implicit dependency scanning of C and C++ source files. The list of rules must be semicolon-separated with each entry of the form "A_MACRO(%)=value-with-%" (the % must be literal). During dependency scanning occurrences of A_MACRO(...) on #include lines will be replaced by the value given with the macro argument substituted for '%'. For example, the entry

```
MYDIR(%)=<mydir/%>
```

will convert lines of the form

```
#include MYDIR(myheader.h)
```

```
#include <mydir/myheader.h>
```

allowing the dependency to be followed. This property applies to sources in all targets within a directory. The property value is initialized in each directory by its value in the directory's parent.

INCLUDE_DIRECTORIES: List of preprocessor include file search directories.

This read-only property specifies the list of directories given so far to the `include_directories` command. It is intended for debugging purposes.

INCLUDE_REGULAR_EXPRESSION: Include file scanning regular expression.

This read-only property specifies the regular expression used during dependency scanning to match include files that should be followed. See the `include_regular_expression` command.

LINK_DIRECTORIES: List of linker search directories.

This read-only property specifies the list of directories given so far to the `link_directories` command. It is intended for debugging purposes.

LISTFILE_STACK: The current stack of listfiles being processed.

This property is mainly useful when trying to debug errors in your CMake scripts. It returns a list of what list files are currently being processed, in order. So if one listfile does an `INCLUDE` command then that is effectively pushing the included listfile onto the stack.

MACROS: List of macro commands available in the current directory.

This read-only property specifies the list of CMake macros currently defined. It is intended for debugging purposes. See the `macro` command.

PARENT_DIRECTORY: Source directory that added current subdirectory.

This read-only property specifies the source directory that added the current source directory as a subdirectory of the build. In the top-level directory the value is the empty-string.

RULE_LAUNCH_COMPILE: Specify a launcher for compile rules.

See the global property of the same name for details. This overrides the global property for a directory.

RULE_LAUNCH_CUSTOM: Specify a launcher for custom rules.

See the global property of the same name for details. This overrides the global property for a directory.

RULE_LAUNCH_LINK: Specify a launcher for link rules.

See the global property of the same name for details. This overrides the global property for a directory.

TEST_INCLUDE_FILE: A cmake file that will be included when ctest is run.

If you specify TEST_INCLUDE_FILE, that file will be included and processed when ctest is run on the directory.

VARIABLES: List of variables defined in the current directory.

This read-only property specifies the list of CMake variables currently defined. It is intended for debugging purposes.

Properties on Targets

<CONFIG>_OUTPUT_NAME: Old per-configuration target file base name.

This is a configuration-specific version of OUTPUT_NAME. Use OUTPUT_NAME_<CONFIG> instead.

<CONFIG>_POSTFIX: Postfix to append to the target file name for configuration <CONFIG>.

When building with configuration <CONFIG> the value of this property is appended to the target file name built on disk. For non-executable targets, this property is initialized by the value of the variable CMAKE_<CONFIG>_POSTFIX if it is set when a target is created. This property is ignored on the Mac for Frameworks and App Bundles.

ARCHIVE_OUTPUT_DIRECTORY: Output directory in which to build ARCHIVE target files.

This property specifies the directory into which archive target files should be built. There are three kinds of target files that may be built: archive, library, and runtime. Executables are always treated as runtime targets. Static libraries are always treated as archive targets. Module libraries are always treated as library targets. For non-DLL platforms shared libraries are treated as library targets. For DLL platforms the DLL part of a shared library is treated as a runtime target and the corresponding import library is treated as an archive target. All Windows-based systems including Cygwin are DLL platforms. This property is initialized by the value of the variable CMAKE_ARCHIVE_OUTPUT_DIRECTORY if it is set when a target is created.

ARCHIVE_OUTPUT_NAME: Output name for ARCHIVE target files.

This property specifies the base name for archive target files. It overrides OUTPUT_NAME and OUTPUT_NAME_<CONFIG> properties. There are three kinds of target files that may be built: archive, library, and runtime. Executables are always treated as runtime targets. Static libraries are always treated as archive targets. Module libraries are always treated as library targets. For non-DLL platforms shared libraries are treated as library targets. For DLL platforms the DLL part of a shared library is treated as a runtime target and the corresponding

import library is treated as an archive target. All Windows-based systems including Cygwin are DLL platforms.

ARCHIVE_OUTPUT_NAME_<CONFIG>: Per-configuration output name for ARCHIVE target files.

This is the configuration-specific version of ARCHIVE_OUTPUT_NAME.

BUILD_WITH_INSTALL_RPATH: Should build tree targets have install tree rpaths.

BUILD_WITH_INSTALL_RPATH is a boolean specifying whether to link the target in the build tree with the INSTALL_RPATH. This takes precedence over SKIP_BUILD_RPATH and avoids the need for relinking before installation. This property is initialized by the value of the variable CMAKE_BUILD_WITH_INSTALL_RPATH if it is set when a target is created.

COMPILE_DEFINITIONS: Preprocessor definitions for compiling a target's sources.

The COMPILE_DEFINITIONS property may be set to a semicolon-separated list of preprocessor definitions using the syntax VAR or VAR=value. Function-style definitions are not supported. CMake will automatically escape the value correctly for the native build system (note that CMake language syntax may require escapes to specify some values). This property may be set on a per-configuration basis using the name COMPILE_DEFINITIONS_<CONFIG> where <CONFIG> is an upper-case name (ex. "COMPILE_DEFINITIONS_DEBUG").

CMake will automatically drop some definitions that are not supported by the native build tool. The VS6 IDE does not support definition values with spaces (but NMake does).

Disclaimer: Most native build tools have poor support for escaping certain values. CMake has work-arounds for many cases but some values may just not be possible to pass correctly. If a value does not seem to be escaped correctly, do not attempt to work-around the problem by adding escape sequences to the value. Your work-around may break in a future version of CMake that has improved escape support. Instead consider defining the macro in a (configured) header file. Then report the limitation.

COMPILE_DEFINITIONS_<CONFIG>: Per-configuration preprocessor definitions on a target.

This is the configuration-specific version of COMPILE_DEFINITIONS.

COMPILE_FLAGS: Additional flags to use when compiling this target's sources.

The COMPILE_FLAGS property sets additional compiler flags used to build sources within the target. Use COMPILE_DEFINITIONS to pass additional preprocessor definitions.

DEBUG_POSTFIX: See target property <CONFIG>_POSTFIX.

This property is a special case of the more-general <CONFIG>_POSTFIX property for the DEBUG configuration.

DEFINE_SYMBOL: Define a symbol when compiling this target's sources.

DEFINE_SYMBOL sets the name of the preprocessor symbol defined when compiling sources in a shared library. If not set here then it is set to target_EXPORTS by default (with some substitutions if the target is not a valid C identifier). This is useful for headers to know whether they are being included from inside their library or outside to properly setup dllexport/dllimport decorations.

ENABLE_EXPORTS: Specify whether an executable exports symbols for loadable modules.

Normally an executable does not export any symbols because it is the final program. It is possible for an executable to export symbols to be used by loadable modules. When this property is set to true CMake will allow other targets to "link" to the executable with the TARGET_LINK_LIBRARIES command. On all platforms a target-level dependency on the executable is created for targets that link to it. For non-DLL platforms the link rule is simply ignored since the dynamic loader will automatically bind symbols when the module is loaded. For DLL platforms an import library will be created for the exported symbols and then used for linking. All Windows-based systems including Cygwin are DLL platforms.

EXCLUDE_FROM_ALL: Exclude the target from the all target.

A property on a target that indicates if the target is excluded from the default build target. If it is not, then with a Makefile for example typing make will cause this target to be built. The same concept applies to the default build of other generators. Installing a target with EXCLUDE_FROM_ALL set to true has undefined behavior.

EchoString: A message to be displayed when the target is built.

A message to display on some generators (such as Makefiles) when the target is built.

FRAMEWORK: This target is a framework on the Mac.

If a shared library target has this property set to true it will be built as a framework when built on the mac. It will have the directory structure required for a framework and will be suitable to be used with the -framework option

Fortran_MODULE_DIRECTORY: Specify output directory for Fortran modules provided by the target.

If the target contains Fortran source files that provide modules and the compiler supports a module output directory this specifies the directory in which the modules will be placed. When this property is not set the modules will be placed in the build directory corresponding to the target's source directory. If the variable CMAKE_Fortran_MODULE_DIRECTORY is set when a target is created its value is used to initialize this property.

GENERATOR_FILE_NAME: Generator's file for this target.

An internal property used by some generators to record the name of project or dsp file associated with this target.

HAS_CXX: Link the target using the C++ linker tool (obsolete).

This is equivalent to setting the LINKER_LANGUAGE property to CXX. See that property's documentation for details.

IMPLICIT_DEPENDS_INCLUDE_TRANSFORM: Specify #include line transforms for dependencies in a target.

This property specifies rules to transform macro-like #include lines during implicit dependency scanning of C and C++ source files. The list of rules must be semicolon-separated with each entry of the form "A_MACRO(%)=value-with-%" (the % must be literal). During dependency scanning occurrences of A_MACRO(...) on #include lines will be replaced by the value given with the macro argument substituted for '%'. For example, the entry

```
MYDIR(%)=<mydir/%>
```

will convert lines of the form

```
#include MYDIR(myheader.h)
```

to

```
#include <mydir/myheader.h>
```

allowing the dependency to be followed. This property applies to sources in the target on which it is set.

IMPORTED: Read-only indication of whether a target is IMPORTED.

The boolean value of this property is true for targets created with the IMPORTED option to add_executable or add_library. It is false for targets built within the project.

IMPORTED_CONFIGURATIONS: Configurations provided for an IMPORTED target.

Lists configuration names available for an IMPORTED target. The names correspond to configurations defined in the project from which the target is imported. If the importing project uses a different set of configurations the names may be mapped using the MAP_IMPORTED_CONFIG_<CONFIG> property. Ignored for non-imported targets.

IMPORTED_IMPLIB: Full path to the import library for an IMPORTED target.

Specifies the location of the ".lib" part of a windows DLL. Ignored for non-imported targets.

IMPORTED_IMPLIB_<CONFIG>: Per-configuration version of IMPORTED_IMPLIB property.

This property is used when loading settings for the <CONFIG> configuration of an imported target. Configuration names correspond to those provided by the project from which the target is imported.

IMPORTED_LINK_DEPENDENT_LIBRARIES: Dependent shared libraries of an imported shared library.

Shared libraries may be linked to other shared libraries as part of their implementation. On some platforms the linker searches for the dependent libraries of shared libraries they are including in the link. This property lists the dependent shared libraries of an imported library. The list should be disjoint from the list of interface libraries in the IMPORTED_LINK_INTERFACE_LIBRARIES property. On platforms requiring dependent shared libraries to be found at link time CMake uses this list to add appropriate files or paths to the link command line. Ignored for non-imported targets.

IMPORTED_LINK_DEPENDENT_LIBRARIES_<CONFIG>: Per-configuration version of IMPORTED_LINK_DEPENDENT_LIBRARIES.

This property is used when loading settings for the <CONFIG> configuration of an imported target. Configuration names correspond to those provided by the project from which the target is imported. If set, this property completely overrides the generic property for the named configuration.

IMPORTED_LINK_INTERFACE_LANGUAGES: Languages compiled into an IMPORTED static library.

Lists languages of source files compiled to produce a STATIC IMPORTED library (such as "C" or "CXX"). CMake accounts for these languages when computing how to link a target to the imported library. For example, when a C executable links to an imported C++ static library CMake chooses the C++ linker to satisfy language runtime dependencies of the static library.

This property is ignored for targets that are not STATIC libraries. This property is ignored for non-imported targets.

IMPORTED_LINK_INTERFACE_LANGUAGES_<CONFIG>: Per-configuration version of IMPORTED_LINK_INTERFACE_LANGUAGES.

This property is used when loading settings for the <CONFIG> configuration of an imported target. Configuration names correspond to those provided by the project from which the target

is imported. If set, this property completely overrides the generic property for the named configuration.

IMPORTED_LINK_INTERFACE_LIBRARIES: Transitive link interface of an IMPORTED target.

Lists libraries whose interface is included when an IMPORTED library target is linked to another target. The libraries will be included on the link line for the target. Unlike the LINK_INTERFACE_LIBRARIES property, this property applies to all imported target types, including STATIC libraries. This property is ignored for non-imported targets.

IMPORTED_LINK_INTERFACE_LIBRARIES_<CONFIG>: Per-configuration version of IMPORTED_LINK_INTERFACE_LIBRARIES.

This property is used when loading settings for the <CONFIG> configuration of an imported target. Configuration names correspond to those provided by the project from which the target is imported. If set, this property completely overrides the generic property for the named configuration.

IMPORTED_LINK_INTERFACE_MULTIPLICITY: Repetition count for cycles of IMPORTED static libraries.

This is LINK_INTERFACE_MULTIPLICITY for IMPORTED targets.

IMPORTED_LINK_INTERFACE_MULTIPLICITY_<CONFIG>: Per-configuration repetition count for cycles of IMPORTED archives.

This is the configuration-specific version of IMPORTED_LINK_INTERFACE_MULTIPLICITY. If set, this property completely overrides the generic property for the named configuration.

IMPORTED_LOCATION: Full path to the main file on disk for an IMPORTED target.

Specifies the location of an IMPORTED target file on disk. For executables this is the location of the executable file. For bundles on OS X this is the location of the executable file inside Contents/MacOS under the application bundle folder. For static libraries and modules this is the location of the library or module. For shared libraries on non-DLL platforms this is the location of the shared library. For frameworks on OS X this is the location of the library file symlink just inside the framework folder. For DLLs this is the location of the ".dll" part of the library. For UNKNOWN libraries this is the location of the file to be linked. Ignored for non-imported targets.

IMPORTED_LOCATION_<CONFIG>: Per-configuration version of IMPORTED_LOCATION property.

This property is used when loading settings for the <CONFIG> configuration of an imported target. Configuration names correspond to those provided by the project from which the target is imported.

IMPORTED SONAME: The "soname" of an IMPORTED target of shared library type.

Specifies the "soname" embedded in an imported shared library. This is meaningful only on platforms supporting the feature. Ignored for non-imported targets.

IMPORTED SONAME <CONFIG>: Per-configuration version of IMPORTED SONAME property.

This property is used when loading settings for the <CONFIG> configuration of an imported target. Configuration names correspond to those provided by the project from which the target is imported.

IMPORT PREFIX: What comes before the import library name.

Similar to the target property PREFIX, but used for import libraries (typically corresponding to a DLL) instead of regular libraries. A target property that can be set to override the prefix (such as "lib") on an import library name.

IMPORT SUFFIX: What comes after the import library name.

Similar to the target property SUFFIX, but used for import libraries (typically corresponding to a DLL) instead of regular libraries. A target property that can be set to override the suffix (such as ".lib") on an import library name.

INSTALL NAME DIR: Mac OSX directory name for installed targets.

INSTALL_NAME_DIR is a string specifying the directory portion of the "install_name" field of shared libraries on Mac OSX to use in the installed targets.

INSTALL_RPATH: The rpath to use for installed targets.

A semicolon-separated list specifying the rpath to use in installed targets (for platforms that support it). This property is initialized by the value of the variable CMAKE_INSTALL_RPATH if it is set when a target is created.

INSTALL_RPATH_USE_LINK_PATH: Add paths to linker search and installed rpath.

INSTALL_RPATH_USE_LINK_PATH is a boolean that if set to true will append directories in the linker search path and outside the project to the INSTALL_RPATH. This property is initialized by the value of the variable CMAKE_INSTALL_RPATH_USE_LINK_PATH if it is set when a target is created.

LABELS: Specify a list of text labels associated with a target.

Target label semantics are currently unspecified.

LIBRARY_OUTPUT_DIRECTORY: Output directory in which to build LIBRARY target files.

This property specifies the directory into which library target files should be built. There are three kinds of target files that may be built: archive, library, and runtime. Executables are

always treated as runtime targets. Static libraries are always treated as archive targets. Module libraries are always treated as library targets. For non-DLL platforms shared libraries are treated as library targets. For DLL platforms the DLL part of a shared library is treated as a runtime target and the corresponding import library is treated as an archive target. All Windows-based systems including Cygwin are DLL platforms. This property is initialized by the value of the variable CMAKE_LIBRARY_OUTPUT_DIRECTORY if it is set when a target is created.

LIBRARY_OUTPUT_NAME: Output name for LIBRARY target files.

This property specifies the base name for library target files. It overrides OUTPUT_NAME and OUTPUT_NAME_<CONFIG> properties. There are three kinds of target files that may be built: archive, library, and runtime. Executables are always treated as runtime targets. Static libraries are always treated as archive targets. Module libraries are always treated as library targets. For non-DLL platforms shared libraries are treated as library targets. For DLL platforms the DLL part of a shared library is treated as a runtime target and the corresponding import library is treated as an archive target. All Windows-based systems including Cygwin are DLL platforms.

LIBRARY_OUTPUT_NAME_<CONFIG>: Per-configuration output name for LIBRARY target files.

This is the configuration-specific version of LIBRARY_OUTPUT_NAME.

LINKER_LANGUAGE: Specifies language whose compiler will invoke the linker.

For executables, shared libraries, and modules, this sets the language whose compiler is used to link the target (such as "C" or "CXX"). A typical value for an executable is the language of the source file providing the program entry point (main). If not set, the language with the highest linker preference value is the default. See documentation of CMAKE_<LANG>_LINKER_PREFERENCE variables.

LINK_FLAGS: Additional flags to use when linking this target.

The LINK_FLAGS property can be used to add extra flags to the link step of a target. LINK_FLAGS_<CONFIG> will add to the configuration <CONFIG>, for example, DEBUG, RELEASE, MINSIZEREL, RELWITHDEBINFO.

LINK_FLAGS_<CONFIG>: Per-configuration linker flags for a target.

This is the configuration-specific version of LINK_FLAGS.

LINK_INTERFACE_LIBRARIES: List public interface libraries for a shared library or executable.

By default linking to a shared library target transitively links to targets with which the library itself was linked. For an executable with exports (see the ENABLE_EXPORTS property) no default transitive link dependencies are used. This property replaces the default transitive link

dependencies with an explicit list. When the target is linked into another target the libraries listed (and recursively their link interface libraries) will be provided to the other target also. If the list is empty then no transitive link dependencies will be incorporated when this target is linked into another target even if the default set is non-empty. This property is ignored for STATIC libraries.

LINK_INTERFACE_LIBRARIES_<CONFIG>: Per-configuration list of public interface libraries for a target.

This is the configuration-specific version of LINK_INTERFACE_LIBRARIES. If set, this property completely overrides the generic property for the named configuration.

LINK_INTERFACE_MULTIPLICITY: Repetition count for STATIC libraries with cyclic dependencies.

When linking to a STATIC library target with cyclic dependencies the linker may need to scan more than once through the archives in the strongly connected component of the dependency graph. CMake by default constructs the link line so that the linker will scan through the component at least twice. This property specifies the minimum number of scans if it is larger than the default. CMake uses the largest value specified by any target in a component.

LINK_INTERFACE_MULTIPLICITY_<CONFIG>: Per-configuration repetition count for cycles of STATIC libraries.

This is the configuration-specific version of LINK_INTERFACE_MULTIPLICITY. If set, this property completely overrides the generic property for the named configuration.

LINK_SEARCH_END_STATIC: End a link line such that static system libraries are used.

Some linkers support switches such as -Bstatic and -Bdynamic to determine whether to use static or shared libraries for -LXXX options. CMake uses these options to set the link type for libraries whose full paths are not known or (in some cases) are in implicit link directories for the platform. By default the linker search type is left at -Bdynamic by the end of the library list. This property switches the final linker search type to -Bstatic.

LOCATION: Read-only location of a target on disk.

For an imported target, this read-only property returns the value of the LOCATION_<CONFIG> property for an unspecified configuration <CONFIG> provided by the target.

For a non-imported target, this property is provided for compatibility with CMake 2.4 and below. It was meant to get the location of an executable target's output file for use in add_custom_command. The path may contain a build-system-specific portion that is replaced at build time with the configuration getting built (such as "\$(ConfigurationName)" in VS). In CMake 2.6 and above add_custom_command automatically recognizes a target name in its

COMMAND and DEPENDS options and computes the target location. Therefore this property is not needed for creating custom commands.

LOCATION_<CONFIG>: Read-only property providing a target location on disk.

A read-only property that indicates where a target's main file is located on disk for the configuration <CONFIG>. The property is defined only for library and executable targets. An imported target may provide a set of configurations different from that of the importing project. By default CMake looks for an exact-match but otherwise uses an arbitrary available configuration. Use the MAP_IMPORTED_CONFIG_<CONFIG> property to map imported configurations explicitly.

MACOSX_BUNDLE: Build an executable as an application bundle on Mac OS X.

When this property is set to true the executable when built on Mac OS X will be created as an application bundle. This makes it a GUI executable that can be launched from the Finder. See the MACOSX_BUNDLE_INFO_PLIST target property for information about creation of the Info.plist file for the application bundle.

MACOSX_BUNDLE_INFO_PLIST: Specify a custom Info.plist template for a Mac OS X App Bundle.

An executable target with MACOSX_BUNDLE enabled will be built as an application bundle on Mac OS X. By default its Info.plist file is created by configuring a template called MacOSXBundleInfo.plist.in located in the CMAKE_MODULE_PATH. This property specifies an alternative template file name which may be a full path.

The following target properties may be set to specify content to be configured into the file:

```
MACOSX_BUNDLE_INFO_STRING  
MACOSX_BUNDLE_ICON_FILE  
MACOSX_BUNDLE_GUI_IDENTIFIER  
MACOSX_BUNDLE_LONG_VERSION_STRING  
MACOSX_BUNDLE_BUNDLE_NAME  
MACOSX_BUNDLE_SHORT_VERSION_STRING  
MACOSX_BUNDLE_BUNDLE_VERSION  
MACOSX_BUNDLE_COPYRIGHT
```

CMake variables of the same name may be set to affect all targets in a directory that do not have each specific property set. If a custom Info.plist is specified by this property it may of course hard-code all the settings instead of using the target properties.

MACOSX_FRAMEWORK_INFO_PLIST: Specify a custom Info.plist template for a Mac OS X Framework.

An library target with FRAMEWORK enabled will be built as a framework on Mac OS X. By default its Info.plist file is created by configuring a template called Mac OSX Framework Info.plist.in located in the CMAKE_MODULE_PATH. This property specifies an alternative template file name which may be a full path.

The following target properties may be set to specify content to be configured into the file:

```
MACOSX_FRAMEWORK_ICON_FILE  
MACOSX_FRAMEWORK_IDENTIFIER  
MACOSX_FRAMEWORK_SHORT_VERSION_STRING  
MACOSX_FRAMEWORK_BUNDLE_VERSION
```

CMake variables of the same name may be set to affect all targets in a directory that do not have each specific property set. If a custom Info.plist is specified by this property it may of course hard-code all the settings instead of using the target properties.

MAP_IMPORTED_CONFIG_<CONFIG>: Map from project configuration to IMPORTED target's configuration.

List configurations of an imported target that may be used for the current project's <CONFIG> configuration. Targets imported from another project may not provide the same set of configuration names available in the current project. Setting this property tells CMake what imported configurations are suitable for use when building the <CONFIG> configuration. The first configuration in the list found to be provided by the imported target is selected. If no matching configurations are available the imported target is considered to be not found. This property is ignored for non-imported targets.

OUTPUT_NAME: Output name for target files.

This sets the base name for output files created for an executable or library target. If not set, the logical target name is used by default.

OUTPUT_NAME_<CONFIG>: Per-configuration target file base name.

This is the configuration-specific version of OUTPUT_NAME.

POST_INSTALL_SCRIPT: Deprecated install support.

The PRE_INSTALL_SCRIPT and POST_INSTALL_SCRIPT properties are the old way to specify CMake scripts to run before and after installing a target. They are used only when the old INSTALL_TARGETS command is used to install the target. Use the INSTALL command instead.

PREFIX: What comes before the library name.

A target property that can be set to override the prefix (such as "lib") on a library name.

PRE_INSTALL_SCRIPT: Deprecated install support.

The PRE_INSTALL_SCRIPT and POST_INSTALL_SCRIPT properties are the old way to specify CMake scripts to run before and after installing a target. They are used only when the old INSTALL_TARGETS command is used to install the target. Use the INSTALL command instead.

PRIVATE_HEADER: Specify private header files in a FRAMEWORK shared library target.

Shared library targets marked with the FRAMEWORK property generate frameworks on OS X and normal shared libraries on other platforms. This property may be set to a list of header files to be placed in the PrivateHeaders directory inside the framework folder. On non-Apple platforms these headers may be installed using the PRIVATE_HEADER option to the install(TARGETS) command.

PROJECT_LABEL: Change the name of a target in an IDE.

Can be used to change the name of the target in an IDE like visual stuido.

PUBLIC_HEADER: Specify public header files in a FRAMEWORK shared library target.

Shared library targets marked with the FRAMEWORK property generate frameworks on OS X and normal shared libraries on other platforms. This property may be set to a list of header files to be placed in the Headers directory inside the framework folder. On non-Apple platforms these headers may be installed using the PUBLIC_HEADER option to the install(TARGETS) command.

RESOURCE: Specify resource files in a FRAMEWORK shared library target.

Shared library targets marked with the FRAMEWORK property generate frameworks on OS X and normal shared libraries on other platforms. This property may be set to a list of files to be placed in the Resources directory inside the framework folder. On non-Apple platforms these files may be installed using the RESOURCE option to the install(TARGETS) command.

RULE_LAUNCH_COMPILE: Specify a launcher for compile rules.

See the global property of the same name for details. This overrides the global and directory property for a target.

RULE_LAUNCH_CUSTOM: Specify a launcher for custom rules.

See the global property of the same name for details. This overrides the global and directory property for a target.

RULE_LAUNCH_LINK: Specify a launcher for link rules.

See the global property of the same name for details. This overrides the global and directory property for a target.

RUNTIME_OUTPUT_DIRECTORY: Output directory in which to build RUNTIME target files.

This property specifies the directory into which runtime target files should be built. There are three kinds of target files that may be built: archive, library, and runtime. Executables are always treated as runtime targets. Static libraries are always treated as archive targets. Module libraries are always treated as library targets. For non-DLL platforms shared libraries are treated as library targets. For DLL platforms the DLL part of a shared library is treated as a runtime target and the corresponding import library is treated as an archive target. All Windows-based systems including Cygwin are DLL platforms. This property is initialized by the value of the variable CMAKE_RUNTIME_OUTPUT_DIRECTORY if it is set when a target is created.

RUNTIME_OUTPUT_NAME: Output name for RUNTIME target files.

This property specifies the base name for runtime target files. It overrides OUTPUT_NAME and OUTPUT_NAME_<CONFIG> properties. There are three kinds of target files that may be built: archive, library, and runtime. Executables are always treated as runtime targets. Static libraries are always treated as archive targets. Module libraries are always treated as library targets. For non-DLL platforms shared libraries are treated as library targets. For DLL platforms the DLL part of a shared library is treated as a runtime target and the corresponding import library is treated as an archive target. All Windows-based systems including Cygwin are DLL platforms.

RUNTIME_OUTPUT_NAME_<CONFIG>: Per-configuration output name for RUNTIME target files.

This is the configuration-specific version of RUNTIME_OUTPUT_NAME.

SKIP_BUILD_RPATH: Should rpaths be used for the build tree.

SKIP_BUILD_RPATH is a boolean specifying whether to skip automatic generation of an rpath allowing the target to run from the build tree. This property is initialized by the value of the variable CMAKE_SKIP_BUILD_RPATH if it is set when a target is created.

SOURCES: Source names specified for a target.

Read-only list of sources specified for a target. The names returned are suitable for passing to the set_source_files_properties command.

SOVERSION: What version number is this target.

For shared libraries VERSION and SOVERSION can be used to specify the build version and api version respectively. When building or installing appropriate symlinks are created if the platform supports symlinks and the linker supports so-names. If only one of both is specified the missing is assumed to have the same version number. For shared libraries and executables on Windows the VERSION attribute is parsed to extract a "major.minor" version number. These numbers are used as the image version of the binary.

STATIC_LIBRARY_FLAGS: Extra flags to use when linking static libraries.

Extra flags to use when linking a static library.

SUFFIX: What comes after the library name.

A target property that can be set to override the suffix (such as ".so") on a library name.

TYPE: The type of the target.

This read-only property can be used to test the type of the given target. It will be one of STATIC_LIBRARY, MODULE_LIBRARY, SHARED_LIBRARY, EXECUTABLE or one of the internal target types.

VERSION: What version number is this target.

For shared libraries VERSION and SOVERSION can be used to specify the build version and api version respectively. When building or installing appropriate symlinks are created if the platform supports symlinks and the linker supports so-names. If only one of both is specified the missing is assumed to have the same version number. For executables VERSION can be used to specify the build version. When building or installing appropriate symlinks are created if the platform supports symlinks. For shared libraries and executables on Windows the VERSION attribute is parsed to extract a "major.minor" version number. These numbers are used as the image version of the binary.

VS_KEYWORD: Visual Studio project keyword.

Can be set to change the visual studio keyword, for example QT integration works better if this is set to Qt4VSv1.0.

VS_SCC_LOCALPATH: Visual Studio Source Code Control Provider.

Can be set to change the visual studio source code control local path property.

VS_SCC_PROJECTNAME: Visual Studio Source Code Control Project.

Can be set to change the visual studio source code control project name property.

VS_SCC_PROVIDER: Visual Studio Source Code Control Provider.

Can be set to change the visual studio source code control provider property.

WIN32_EXECUTABLE: Build an executable with a WinMain entry point on windows.

When this property is set to true the executable when linked on Windows will be created with a WinMain() entry point instead of just main(). This makes it a GUI executable instead of a console application. See the CMAKE_MFC_FLAG variable documentation to configure use of MFC for WinMain executables.

XCODE_ATTRIBUTE_<an-attribute>: Set Xcode target attributes directly.

Tell the Xcode generator to set '<an-attribute>' to a given value in the generated Xcode project. Ignored on other generators.

Properties on Tests

ENVIRONMENT: Specify environment variables that should be defined for running a test.

If set to a list of environment variables and values of the form MYVAR=value those environment variables will be defined while running the test. The environment is restored to its previous state after the test is done.

FAIL_REGULAR_EXPRESSION: If the output matches this regular expression the test will fail.

If set, if the output matches one of specified regular expressions, the test will fail. For example: PASS_REGULAR_EXPRESSION "[^a-z]Error;ERROR;Failed"

LABELS: Specify a list of text labels associated with a test.

The list is reported in dashboard submissions.

MEASUREMENT: Specify a CDASH measurement and value to be reported for a test.

If set to a name then that name will be reported to CDASH as a named measurement with a value of 1. You may also specify a value by setting MEASUREMENT to "measurement=value".

PASS_REGULAR_EXPRESSION: The output must match this regular expression for the test to pass.

If set, the test output will be checked against the specified regular expressions and at least one of the regular expressions has to match, otherwise the test will fail.

TIMEOUT: How many seconds to allow for this test.

This property if set will limit a test to not take more than the specified number of seconds to run. If it exceeds that the test process will be killed and ctest will move to the next test. This setting takes precedence over CTEST_TESTING_TIMEOUT.

WILL_FAIL: If set to true, this will invert the pass/fail flag of the test.

This property can be used for tests that are expected to fail and return a non zero return code.

Properties on Source Files

ABSTRACT: Is this source file an abstract class.

A property on a source file that indicates if the source file represents a class that is abstract. This only makes sense for languages that have a notion of an abstract class and it is only used by some tools that wrap classes into other languages.

COMPILE_DEFINITIONS: Preprocessor definitions for compiling a source file.

The COMPILE_DEFINITIONS property may be set to a semicolon-separated list of preprocessor definitions using the syntax VAR or VAR=value. Function-style definitions are not supported. CMake will automatically escape the value correctly for the native build system (note that CMake language syntax may require escapes to specify some values). This property may be set on a per-configuration basis using the name COMPILE_DEFINITIONS_<CONFIG> where <CONFIG> is an upper-case name (ex. "COMPILE_DEFINITIONS_DEBUG").

CMake will automatically drop some definitions that are not supported by the native build tool. The VS6 IDE does not support definition values with spaces (but NMake does). Xcode does not support per-configuration definitions on source files.

Disclaimer: Most native build tools have poor support for escaping certain values. CMake has work-arounds for many cases but some values may just not be possible to pass correctly. If a value does not seem to be escaped correctly, do not attempt to work-around the problem by adding escape sequences to the value. Your work-around may break in a future version of CMake that has improved escape support. Instead consider defining the macro in a (configured) header file. Then report the limitation.

COMPILE_DEFINITIONS_<CONFIG>: Per-configuration preprocessor definitions on a source file.

This is the configuration-specific version of COMPILE_DEFINITIONS. Note that Xcode does not support per-configuration source file flags so this property will be ignored by the Xcode generator.

COMPILE_FLAGS: Additional flags to be added when compiling this source file.

These flags will be added to the list of compile flags when this source file builds. Use COMPILE_DEFINITIONS to pass additional preprocessor definitions.

EXTERNAL_OBJECT: If set to true then this is an object file.

If this property is set to true then the source file is really an object file and should not be compiled. It will still be linked into the target though.

GENERATED: Is this source file generated as part of the build process.

If a source file is generated by the build process CMake will handle it differently in terms of dependency checking etc. Otherwise having a non-existent source file could create problems.

HEADER_FILE_ONLY: Is this source file only a header file.

A property on a source file that indicates if the source file is a header file with no associated implementation. This is set automatically based on the file extension and is used by CMake to determine certain dependency information should be computed.

KEEP_EXTENSION: Make the output file have the same extension as the source file.

If this property is set then the file extension of the output file will be the same as that of the source file. Normally the output file extension is computed based on the language of the source file, for example .cxx will go to a .o extension.

LABELS: Specify a list of text labels associated with a source file.

This property has meaning only when the source file is listed in a target whose LABELS property is also set. No other semantics are currently specified.

LANGUAGE: What programming language is the file.

A property that can be set to indicate what programming language the source file is. If it is not set the language is determined based on the file extension. Typical values are CXX C etc.

LOCATION: The full path to a source file.

A read only property on a SOURCE FILE that contains the full path to the source file.

MACOSX_PACKAGE_LOCATION: Place a source file inside a Mac OS X bundle or framework.

Executable targets with the MACOSX_BUNDLE property set are built as Mac OS X application bundles on Apple platforms. Shared library targets with the FRAMEWORK property set are built as Mac OS X frameworks on Apple platforms. Source files listed in the target with this property set will be copied to a directory inside the bundle or framework content folder specified by the property value. For bundles the content folder is "<name>.app/Contents". For frameworks the content folder is "<name>.framework/Versions/<version>". See the PUBLIC_HEADER, PRIVATE_HEADER, and RESOURCE target properties for specifying files meant for Headers, PrivateHeadres, or Resources directories.

OBJECT_DEPENDS: Additional files on which a compiled object file depends.

Specifies a semicolon-separated list of full-paths to files on which any object files compiled from this source file depend. An object file will be recompiled if any of the named files is newer than it.

This property need not be used to specify the dependency of a source file on a generated header file that it includes. Although the property was originally introduced for this purpose, it is no longer necessary. If the generated header file is created by a custom command in the same target as the source file, the automatic dependency scanning process will recognize the dependency. If the generated header file is created by another target, an inter-target

dependency should be created with the `add_dependencies` command (if one does not already exist due to linking relationships).

OBJECT_OUTPUTS: Additional outputs for a Makefile rule.

Additional outputs created by compilation of this source file. If any of these outputs is missing the object will be recompiled. This is supported only on Makefile generators and will be ignored on other generators.

SYMBOLIC: Is this just a name for a rule.

If `SYMBOLIC` (boolean) is set to true the build system will be informed that the source file is not actually created on disk but instead used as a symbolic name for a build rule.

WRAP_EXCLUDE: Exclude this source file from any code wrapping techniques.

Some packages can wrap source files into alternate languages to provide additional functionality. For example, C++ code can be wrapped into Java or Python etc using SWIG etc. If `WRAP_EXCLUDE` is set to true (1 etc) that indicates then this source file should not be wrapped.

Properties on Cache Entries

ADVANCED: True if entry should be hidden by default in GUIs.

This is a boolean value indicating whether the entry is considered interesting only for advanced configuration. The `mark_as_advanced()` command modifies this property.

HELPSTRING: Help associated with entry in GUIs.

This string summarizes the purpose of an entry to help users set it through a CMake GUI.

MODIFIED: Internal management property. Do not set or get.

This is an internal cache entry property managed by CMake to track interactive user modification of entries. Ignore it.

STRINGS: Enumerate possible STRING entry values for GUI selection.

For cache entries with type `STRING`, this enumerates a set of values. CMake GUIs may use this to provide a selection widget instead of a generic string entry field. This is for convenience only. CMake does not enforce that the value matches one of those listed.

TYPE: Widget type for entry in GUIs.

Cache entry values are always strings, but CMake GUIs present widgets to help users set values. The GUIs use this property as a hint to determine the widget type. Valid `TYPE` values are:

BOOL	= Boolean ON/OFF value.
PATH	= Path to a directory.
FILEPATH	= Path to a file.
STRING	= Generic string value.
INTERNAL	= Do not present in GUI at all.
STATIC	= Value managed by CMake, do not change.
UNINITIALIZED	= Type not yet specified.

Generally the TYPE of a cache entry should be set by the command which creates it (set, option, find_library, etc.).

VALUE: Value of a cache entry.

This property maps to the actual value of a cache entry. Setting this property always sets the value without checking, so use with care.

Appendix F – CMake Policies

CMP0000: A minimum required CMake version must be specified.

CMake requires that projects specify the version of CMake to which they have been written. This policy has been put in place so users trying to build the project may be told when they need to update their CMake. Specifying a version also helps the project build with CMake versions newer than that specified. Use the `cmake_minimum_required` command at the top of your main `CMakeLists.txt` file:

```
cmake_minimum_required(VERSION <major>.<minor>)
```

where "<major>.<minor>" is the version of CMake you want to support (such as "2.6"). The command will ensure that at least the given version of CMake is running and help newer versions be compatible with the project. See documentation of `cmake_minimum_required` for details.

Note that the command invocation must appear in the `CMakeLists.txt` file itself; a call in an included file is not sufficient. However, the `cmake_policy` command may be called to set policy `CMP0000` to OLD or NEW behavior explicitly. The OLD behavior is to silently ignore the missing invocation. The NEW behavior is to issue an error instead of a warning. An included file may set `CMP0000` explicitly to affect how this policy is enforced for the main `CMakeLists.txt` file.

This policy was introduced in CMake version 2.6.0.

CMP0001: CMAKE_BACKWARDS_COMPATIBILITY should no longer be used.

The OLD behavior is to check CMAKE_BACKWARDS_COMPATIBILITY and present it to the user. The NEW behavior is to ignore CMAKE_BACKWARDS_COMPATIBILITY completely.

In CMake 2.4 and below the variable CMAKE_BACKWARDS_COMPATIBILITY was used to request compatibility with earlier versions of CMake. In CMake 2.6 and above all compatibility issues are handled by policies and the cmake_policy command. However, CMake must still check CMAKE_BACKWARDS_COMPATIBILITY for projects written for CMake 2.4 and below.

This policy was introduced in CMake version 2.6.0. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the cmake_policy command to set it to OLD or NEW explicitly.

CMP0002: Logical target names must be globally unique.

Targets names created with add_executable, add_library, or add_custom_target are logical build target names. Logical target names must be globally unique because:

- Unique names may be referenced unambiguously both in CMake code and on make tool command lines.
- Logical names are used by Xcode and VS IDE generators to produce meaningful project names for the targets.

The logical name of executable and library targets does not have to correspond to the physical file names built. Consider using the OUTPUT_NAME target property to create two targets with the same physical name while keeping logical names distinct. Custom targets must simply have globally unique names (unless one uses the global property ALLOW_DUPLICATE_CUSTOM_TARGETS with a Makefiles generator).

This policy was introduced in CMake version 2.6.0. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the cmake_policy command to set it to OLD or NEW explicitly.

CMP0003: Libraries linked via full path no longer produce linker search paths.

This policy affects how libraries whose full paths are NOT known are found at link time, but was created due to a change in how CMake deals with libraries whose full paths are known. Consider the code

```
target_link_libraries(myexe /path/to/libA.so)
```

CMake 2.4 and below implemented linking to libraries whose full paths are known by splitting them on the link line into separate components consisting of the linker search path and the library name. The example code might have produced something like

```
... -L/path/to -lA ...
```

in order to link to library A. An analysis was performed to order multiple link directories such that the linker would find library A in the desired location, but there are cases in which this does not work. CMake versions 2.6 and above use the more reliable approach of passing the full path to libraries directly to the linker in most cases. The example code now produces something like

```
... /path/to/libA.so ....
```

Unfortunately this change can break code like

```
target_link_libraries(myexe /path/to/libA.so B)
```

where "B" is meant to find "/path/to/libB.so". This code is wrong because the user is asking the linker to find library B but has not provided a linker search path (which may be added with the `link_directories` command). However, with the old linking implementation the code would work accidentally because the linker search path added for library A allowed library B to be found.

In order to support projects depending on linker search paths added by linking to libraries with known full paths, the OLD behavior for this policy will add the linker search paths even though they are not needed for their own libraries. When this policy is set to OLD, CMake will produce a link line such as

```
... -L/path/to /path/to/libA.so -lB ...
```

which will allow library B to be found as it was previously. When this policy is set to NEW, CMake will produce a link line such as

```
... /path/to/libA.so -lB ...
```

which more accurately matches what the project specified.

The setting for this policy used when generating the link line is that in effect when the target is created by an `add_executable` or `add_library` command. For the example described above, the code

```
cmake_policy(SET CMP0003 OLD) # or cmake_policy(VERSION 2.4)
add_executable(myexe myexe.c)
target_link_libraries(myexe /path/to/libA.so B)
```

will work and suppress the warning for this policy. It may also be updated to work with the corrected linking approach:

```
cmake_policy(SET CMP0003 NEW) # or cmake_policy(VERSION 2.6)
link_directories(/path/to) # needed to find library B
add_executable(myexe myexe.c)
target_link_libraries(myexe /path/to/libA.so B)
```

Even better, library B may be specified with a full path:

```
add_executable(myexe myexe.c)
target_link_libraries(myexe /path/to/libA.so /path/to/libB.so)
```

When all items on the link line have known paths CMake does not check this policy so it has no effect.

Note that the warning for this policy will be issued for at most one target. This avoids flooding users with messages for every target when setting the policy once will probably fix all targets.

This policy was introduced in CMake version 2.6.0. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0004: Libraries linked may not have leading or trailing whitespace.

CMake versions 2.4 and below silently removed leading and trailing whitespace from libraries linked with code like

```
target_link_libraries(myexe " A ")
```

This could lead to subtle errors in user projects.

The OLD behavior for this policy is to silently remove leading and trailing whitespace. The NEW behavior for this policy is to diagnose the existence of such whitespace as an error. The setting for this policy used when checking the library names is that in effect when the target is created by an `add_executable` or `add_library` command.

This policy was introduced in CMake version 2.6.0. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0005: Preprocessor definition values are now escaped automatically.

This policy determines whether or not CMake should generate escaped preprocessor definition values added via `add_definitions`. CMake versions 2.4 and below assumed that only trivial values would be given for macros in `add_definitions` calls. It did not attempt to escape non-trivial values such as string literals in generated build rules. CMake versions 2.6 and above support escaping of most values, but cannot assume the user has not added escapes already in an attempt to work around limitations in earlier versions.

The OLD behavior for this policy is to place definition values given to `add_definitions` directly in the generated build rules without attempting to escape anything. The NEW behavior for this policy is to generate correct escapes for all native build tools automatically. See documentation of the `COMPILE_DEFINITIONS` target property for limitations of the escaping implementation.

This policy was introduced in CMake version 2.6.0. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0006: Installing `MACOSX_BUNDLE` targets requires a `BUNDLE DESTINATION`.

This policy determines whether the `install(TARGETS)` command must be given a `BUNDLE DESTINATION` when asked to install a target with the `MACOSX_BUNDLE` property set. CMake 2.4 and below did not distinguish application bundles from normal executables when installing targets. CMake 2.6 provides a `BUNDLE` option to the `install(TARGETS)` command that specifies rules specific to application bundles on the Mac. Projects should use this option when installing a target with the `MACOSX_BUNDLE` property set.

The OLD behavior for this policy is to fall back to the `RUNTIME DESTINATION` if a `BUNDLE DESTINATION` is not given. The NEW behavior for this policy is to produce an error if a bundle target is installed without a `BUNDLE DESTINATION`.

This policy was introduced in CMake version 2.6.0. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0007: `list` command no longer ignores empty elements.

This policy determines whether the list command will ignore empty elements in the list. CMake 2.4 and below list commands ignored all empty elements in the list. For example, a;b;;c would have length 3 and not 4. The OLD behavior for this policy is to ignore empty list elements. The NEW behavior for this policy is to correctly count empty elements in a list.

This policy was introduced in CMake version 2.6.0. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0008: Libraries linked by full-path must have a valid library file name.

In CMake 2.4 and below it is possible to write code like

```
target_link_libraries(myexe /full/path/to/somelib)
```

where "somelib" is supposed to be a valid library file name such as "libsomelib.a" or "somelib.lib". For Makefile generators this produces an error at build time because the dependency on the full path cannot be found. For VS IDE and Xcode generators this used to work by accident because CMake would always split off the library directory and ask the linker to search for the library by name (-lsomelib or somelib.lib). Despite the failure with Makefiles, some projects have code like this and build only with VS and/or Xcode. This version of CMake prefers to pass the full path directly to the native build tool, which will fail in this case because it does not name a valid library file.

This policy determines what to do with full paths that do not appear to name a valid library file. The OLD behavior for this policy is to split the library name from the path and ask the linker to search for it. The NEW behavior for this policy is to trust the given path and pass it directly to the native build tool unchanged.

This policy was introduced in CMake version 2.6.1. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0009: FILE_GLOB_RECURSE calls should not follow symlinks by default.

In CMake 2.6.1 and below, FILE_GLOB_RECURSE calls would follow through symlinks, sometimes coming up with unexpectedly large result sets because of symlinks to top level directories that contain hundreds of thousands of files.

This policy determines whether or not to follow symlinks encountered during a FILE_GLOB_RECURSE call. The OLD behavior for this policy is to follow the symlinks. The NEW behavior for this policy is not to follow the symlinks by default, but only if FOLLOW_SYMLINKS is given as an additional argument to the FILE command.

This policy was introduced in CMake version 2.6.2. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0010: Bad variable reference syntax is an error.

In CMake 2.6.2 and below, incorrect variable reference syntax such as a missing close-brace ("\${FOO}") was reported but did not stop processing of CMake code. This policy determines whether a bad variable reference is an error. The OLD behavior for this policy is to warn about the error, leave the string untouched, and continue. The NEW behavior for this policy is to report an error.

This policy was introduced in CMake version 2.6.3. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0011: Included scripts do automatic `cmake_policy` PUSH and POP.

In CMake 2.6.2 and below, CMake Policy settings in scripts loaded by the `include()` and `find_package()` commands would affect the includer. Explicit invocations of `cmake_policy(PUSH)` and `cmake_policy(POP)` were required to isolate policy changes and protect the includer. While some scripts intend to affect the policies of their includer, most do not. In CMake 2.6.3 and above, `include()` and `find_package()` by default PUSH and POP an entry on the policy stack around an included script, but provide a `NO_POLICY_SCOPE` option to disable it. This policy determines whether or not to imply `NO_POLICY_SCOPE` for compatibility. The OLD behavior for this policy is to imply `NO_POLICY_SCOPE` for `include()` and `find_package()` commands. The NEW behavior for this policy is to allow the commands to do their default `cmake_policy` PUSH and POP.

This policy was introduced in CMake version 2.6.3. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0012: The `if()` command can recognize named boolean constants.

In CMake versions prior to 2.6.5 the only boolean constants were 0 and 1. Other boolean constants such as true, false, yes, no, on, off, y, n, notfound, ignore (all case insensitive) were recognized in some cases but not all. In later versions of cmake these values are treated as boolean constants more consistently and should not be used as variable names. The OLD behavior for this policy is to allow variables to have names such as true and to dereference them. The NEW behavior for this policy is to treat strings like true as a boolean constant.

This policy was introduced in CMake version 2.6.5. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0013: Duplicate binary directories are not allowed.

CMake 2.6.3 and below silently permitted `add_subdirectory()` calls to create the same binary directory multiple times. During build system generation files would be written and then overwritten in the build tree and could lead to strange behavior. CMake 2.6.4 and above explicitly detect duplicate binary directories. CMake 2.6.4 always considers this case an error. In CMake 2.6.5 and above this policy determines whether or not the case is an error. The OLD behavior for this policy is to allow duplicate binary directories. The NEW behavior for this policy is to disallow duplicate binary directories with an error.

This policy was introduced in CMake version 2.6.5. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0014: Input directories must have CMakeLists.txt.

CMake versions before 2.8 silently ignored missing `CMakeLists.txt` files in directories referenced by `add_subdirectory()` or `subdirs()`, treating them as if present but empty. In CMake 2.8.0 and above this policy determines whether or not the case is an error. The OLD behavior for this policy is to silently ignore the problem. The NEW behavior for this policy is to report an error.

This policy was introduced in CMake version 2.8.0. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

CMP0015: The `set()` CACHE mode and `option()` command make the cache value visible.

In CMake 2.6 and below the CACHE mode of the `set()` command and the `option()` command did not expose the value from the named cache entry if it was already set both in the cache and as a local variable. This led to subtle differences between first and later configurations because a conflicting local variable would be overridden only when the cache value was first created. For example, the code

```
set(x 1)
set(before ${x})
set(x 2 CACHE STRING "X")
set(after ${x})
message(STATUS "${before}, ${after}")
```

would print "1,2" on the first run and "1,1" on future runs.

CMake 2.8.0 and above prefer to expose the cache value in all cases by removing the local variable definition, but this changes behavior in subtle cases when the local variable has a different value than that exposed from the cache. The example above will always print "1,2".

This policy determines whether the commands should always expose the cache value. The OLD behavior for this policy is to leave conflicting local variable values untouched and hide the true cache value. The NEW behavior for this policy is to always expose the cache value.

This policy was introduced in CMake version 2.8.0. CMake version 2.8.0-rc2 warns when the policy is not set and uses OLD behavior. Use the `cmake_policy` command to set it to OLD or NEW explicitly.

Index

A

ALL_BUILD target, 13, 109
ANT, 4
assert, 41
autoconf, 3
 converting to CMake, 121

B

batch commands, 103
binary packages, 150
Borland compiler, 5, 21, 23, 249
build
 configurations, 32, 67
 in-source. *See* in-source build
 out-of-source. *See* out-of-source build
build platform, 125
building your project, 19

C

cache
 advanced entries, 18, 30
 documentation, 30
 entries, 27, 29, 434
 find_* commands, 86
 FORCE option, 18, 32
 initializing, 17
 predefined options, 30
 purpose, 30
 set command, 29
 variable behavior, 31
case sensitivity, 9, 38
ccmake. *See* running CMake
CDash, 195
 adding notes, 205
 automatic submissions, 206
 backup, 230
 client setup, 199
 creating a new project, 196
 email, 225

expected builds, 224
filtering errors/warnings, 203
logging, 229
mobile support, 230
server setup, 218
sites, 226
specifying the server, 202
submitting results, 196
subprojects, 233
timing, 229
tutorial, 268
upgrading, 231

CMake
 benefits, 1
 command line, 287
 compiler selection, 15
 extending, 250
 history, 3
 plugin, 250
 porting, 241
 Structure, 21
 syntax, 8, 33
 tutorial, 255
 versions, 43

CMake Commands, 8, 9, 34, 301
 add_custom_command, 62, 105–8, 109, 110,
 111, 112, 123, 301
 add_custom_target, 24, 108, 109, 110, 111, 123,
 303
 add_definitions, 88, 95, 96, 304
 add_dependencies, 81, 304
 add_executable, 9, 10, 24, 25, 34, 121, 193, 305
 add_library, 24, 25, 34, 121, 306
 add_subdirectory, 28, 34, 115, 116, 122, 307
 add_test, 38, 39, 186, 187, 188, 192, 193, 307
 aux_source_directory, 308
 break, 42, 309
 build_command, 309
 cmake_minimum_required, 44, 51, 214, 309
 cmake_policy, 51, 310
 configure_file, 97, 98, 122, 311
 create_test_sourcelist, 192, 193, 311
 define_property, 312
 else, 10, 35
 elseif, 36
 enable_language, 313
 enable_testing, 313
 endif, 10, 35
 exec_program, 120, 217
 execute_process, 313
 export, 81, 314
 file, 89, 116, 119, 120, 315
 find_file, 29, 85, 318
 find_library, 10, 41, 46, 47, 85, 86, 320
 find_package, 48, 85, 92, 95, 323
 find_path, 46, 47, 85, 86, 329
 find_program, 85, 105, 332
 fltk_wrap_ui, 50, 335
 foreach, 29, 35, 38, 39, 335
 function, 38, 39, 336
 get_cmake_property, 336
 get_directory_property, 337
 get_filename_component, 193, 337
 get_property, 25, 26, 337
 get_source_file_property, 25, 338
 get_target_property, 24, 105, 106, 108, 339
 get_test_property, 339
 if, 10, 27, 35, 36, 37, 38, 40, 42, 339
 include, 28, 45, 46, 95, 343
 include_directories, 86, 99, 252, 344
 include_external_msproject, 344
 include_regular_expression, 82, 344
 install, 66, 68, 345
 link_directories, 350
 link_libraries, 25, 62
 list, 351
 load_cache, 352
 load_command, 252, 352
 macro, 35, 38, 39, 40, 41, 217, 352
 make_directory, 120
 mark_as_advanced, 18, 30, 119, 353
 math, 353
 message, 29, 122, 353
 obsolete/deprecated, 366
 option, 29, 95, 121, 354
 output_required_files, 82, 354
 project, 9, 10, 34, 82, 105, 119, 123, 242, 354
 qt_wrap_cpp, 355
 qt_wrap_ui, 355
 remove, 34, 193

- remove_definitions, 355
return, 42, 355
separate_arguments, 34, 356
set, 9, 10, 18, 24, 27, 28, 29, 31, 32, 34, 35, 46,
 356
set_directory_properties, 357
set_property, 25, 26, 357
set_source_file_properties, 358
set_source_files_properties, 25, 48, 62
set_target_properties, 24, 65, 359
set_tests_properties, 187, 188, 359
site_name, 359
source_group, 123, 359
string, 42, 62, 360
subdirs, 189
target_link_libraries, 25, 59, 60, 362
try_compile, 85, 88, 89, 92, 122, 131, 363
try_run, 85, 88, 89, 92, 364
unset, 365
variable_watch, 365
while, 35, 38, 39, 365
- CMake Modules. *See* modules
- CMake Policies. *See* policies
- CMake Properties, 26, 411
- ABSTRACT, 26
 - ADDITIONAL_MAKE_CLEAN_FILES, 26
 - COMPILE_DEFINITIONS, 96
 - COMPILE_DEFINITIONS_<CONFIG>, 97
 - COMPILE_FLAGS, 25
 - EXCLUDE_FROM_ALL, 26
 - GENERATED, 25
 - LABELS, 240
 - LINK_FLAGS, 25
 - list of, 27
 - LISTFILE_STACK, 26
 - OBJECT_DEPENDS, 26
 - WRAP_EXCLUDE, 26
- CMake Variables
- BUILD_SHARED_LIBS, 24
 - CMAKE_BUILD_TYPE, 32
 - CMAKE_BUILD_TYPE, 67
 - CMAKE_C_COMPILER, 16, 249
 - CMAKE_C_COMPILER_WORKS, 89
 - CMAKE_C_CREATE_SHARED_LIBRARY, 249
 - CMAKE_C_FLAGS, 16, 252
 - CMAKE_C_LINK_EXECUTABLE, 82
 - CMAKE_CFG_INDIR, 32, 62, 104
 - CMAKE_COMMAND, 104, 106, 107, 206
 - CMAKE_CONFIGURATION_TYPES, 32
 - CMAKE_CROSSCOMPILING, 127
 - CMAKE_CTEST_COMMAND, 192
 - CMAKE_CURRENT_BINARY_DIR, 105, 267
 - CMAKE_CURRENT_SOURCE_DIR, 105
 - CMAKE_CXX_COMPILE_OBJECT, 250
 - CMAKE_CXX_COMPILER, 16, 244, 245, 248, 249, 250
 - CMAKE_CXX_CREATE_SHARED_LIBRARIES, 244, 248, 249
 - CMAKE_CXX_CREATE_SHARED_MODULE, 249
 - CMAKE_CXX_CREATE_STATIC_LIBRARY, 245, 249, 250
 - CMAKE_CXX_FLAGS, 16, 32, 251
 - CMAKE_CXX_FLAGS_DEBUG, 245
 - CMAKE_CXX_FLAGS_DEBUG_INIT, 245
 - CMAKE_CXX_FLAGS_INIT, 245
 - CMAKE_CXX_FLAGS_MINSIZEREL, 245
 - CMAKE_CXX_FLAGS_MINSIZEREL_INIT, 245
 - CMAKE_CXX_FLAGS_RELEASE, 245
 - CMAKE_CXX_FLAGS_RELEASE_INIT, 245
 - CMAKE_CXX_FLAGS_RELWITHDEBINFO, 246
 - CMAKE_CXX_FLAGS_RELWITHDEBINFO_INIT, 246
 - CMAKE_CXX_LINK_EXECUTABLE, 82, 249
 - CMAKE_DL_LIBS, 245
 - CMAKE_END_TEMP_FILE, 249, 250
 - CMAKE_FIND_ROOT_PATH, 129
 - CMAKE_GENERATOR, 192, 207
 - CMAKE_HOST_*, 130
 - CMAKE_INSTALL_COMPONENT, 75
 - CMAKE_INSTALL_PREFIX, 15, 67, 75
 - CMAKE_INSTALL_RPATH, 148
 - CMAKE_LIBRARY_PREFIX, 105

CMAKE_MAJOR_VERSION, 44
CMAKE_MAKE_PROGRAM, 192
CMAKE_MODULE_PATH, 45
CMAKE_RANLIB, 248, 249
CMAKE_ROOT, 95, 252, 265
CMAKE_SHARED_LIBRARY_C_FLAGS, 249
.CMAKE_SHARED_LIBRARY_CREATE_C_F
LAGS, 245
CMAKE_SHARED_LIBRARY_CREATE_CX
X_FLAGS, 244, 245, 248
CMAKE_SHARED_LIBRARY_RUNTIME_C_
FLAG, 245
CMAKE_SHARED_LIBRARY_RUNTIME_C_
FLAG_SEP, 245
CMAKE_SHARED_MODULE_PREFIX, 105
CMAKE_SKIP_RPATH, 64
CMAKE_START_TEMP_FILE, 249, 250
CMAKE_SYSTEM_NAME, 127, 241, 242,
245
CMAKE_SYSTEM_PROCESSOR, 127, 142
CMAKE_SYSTEM_VERSION, 127, 241
CMAKE_TOOLCHAIN_FILE, 126
CMAKE_VERSION, 44
EXECUTABLE_OUTPUT_PATH, 105, 118,
119
LIBRARY_EXPORT, 61
LIBRARY_OUTPUT_PATH, 105, 118, 119
PROJECT_BINARY_DIR, 105, 119, 252
PROJECT_SOURCE_DIR, 37, 105, 109, 252,
253, 258
WIN32, 10, 24, 61, 123
cmake-gui, 11, 13, 15, 30
 pulldown option, 30
CMakeSetup. *See* cmake-gui
cmMakefile, 23, 24, 31
code coverage, 196, 200
comments, 10, 33
comparison, 37
compile flags, 25
compiler specification, 15
compilers, 5, 243
compiling CMake, 8
conditional statements, 35
configured header, 256, 260
converting autoconf projects to CMake, 121
converting Makefiles to CMake, 120
converting Visual Studio projects, 123
copying files. *See file*
CPack, 149
 command line, 298
 components, 163
 DESTDIR, 155
 escape characters, 152
 generators, 153
 tutorial, 267
 variables, 151
cross compiling, 125
 hello world, 136
 system introspection, 130
CTest, 189
 advanced, 212
 building tests, 191
 command line, 294
 options, 190
 properties, 187
 regular expressions, 188
 running the tests, 189
 tutorial, 260
 variables, 212
CTest Commands
 ctest_build, 214
 ctest_build_command, 213
 ctest_configure, 214
 ctest_empty_binary_directory, 213
 ctest_start, 214
 ctest_submit, 214
 ctest_test, 214
 ctest_update, 214
custom commands, 103, 105
custom targets, 108
CVS, 209
cygwin, 149, 173

D

DART. *See* CDash
dashboards, 17, 185, 194, 195
Debian, 150, 182

declspec, 61
dependencies, 82
dependency analysis, 16, 26
DESTDIR, 75, 149, 155
directory installation, 72
directory structures, 10, 115
 creating, 120
 properties, 26

E

editing CMakeLists files, 17
Emacs modes, 17
embedded devices, 125
environment variables, 2, 9, 16, 64, 217
escape character, 33
executing programs, 120

F

false, 38
file
 COMPONENT, 67
 copying, 104
 exclude, 74
 existance, 36
 globbing, 119
 installing, 66
 OPTIONAL, 68
 permissions, 67
 properties, 25
 regular expressions, 74
 remove, 104
 rename, 71
filtering errors and warnings, 203
final pass, 24
find modules. *See* modules
flags, 16, 25, 64
flow control, 34
FLTK, 50, 143
fortran, 242, 278
function, 35

functions, 39

G

generator, 15, 21, 23, 32, 207, 263
graphviz, 289

H

header files, configured, 97
hello world, 9, 136
history of CMake, 3

I

initial pass, 23
in-source build, 10, 117
installation scripts, 75
installer, 267
installing CMake
 Mac OS X, 7
 UNIX, 7
 Windows, 7
installing files. *See* files

J

JAM, 4

K

KDE, 3

L

languages, 34, 242, 247
LaTeX custom target, 109

libraries, 59
 API changes, 64
 exporting, 80
 importing, 78
 installing prerequisites, 76
 modules, 24
 pitfalls, 63
 rpath, 64
 shared, 24
 shared vs static, 59
 soname, 64
 static, 24, 25
 versions, 64

library configurations, 59

linking

- libraries, 56
- system libraries, 57
- Windows, 59

list. *See* variables

loaded commands, 26, 250, 254

looping constructs, 35

M

Mac OS X
 drag and drop installer, 178
 frameworks, 319, 322, 327
 Package Maker, 149, 176
 relocatable applications, 76
 Xcode. *See* Xcode
 macros, 39, 40
 Makefiles, 21, 23, 32
 build configurations, 32
 microcontroller, 140
 MinGW, 13, 143
 modules, 44, 373
 documentation, 47
 find, 45, 93
 find <Package>Config.cmake, 93
 find conventions, 93
 path, 45
 system introspection, 47, 87
 utility, 47

modules, loadable. *See* libraries

N

NMake, 23
 not found, 86
 NSIS, 149, 153, 156
 Nullsoft Installer, 149
 numeric comparison, 37

O

operator precedence, 38
 operators, 36
 out-of-source build, 10, 117

P

package configuration file, 99
 platforms, 4, 242
 PLEASE_FILL_OUT, 131
 plugins, 59, 250, 251
 policies, 50, 437
 scope, 52
 setting, 51
 support multiple versions, 56
 portability issues, 103
 portability of code, 87
 porting CMake, 241
 POSIX systems, 241
 project
 languages, 34
 name, 34
 properties. *See* CMake Properties
 proxies, 203
 purify, 82, 200, 203

Q

qmake, 4

Qt, 2, 49
 Qt CMake user interface. *See* cmake-gui
 quoting, 9

R

regression testing, 185
 regular expressions, 37, 42–43, 82, 190
 relocatable packages, 76, 150
 removing files. *See* file
 required package, 101
 RPATH, 64, 148
 RPM, 150, 183
 rule variables, 247
 running CMake, 10
 cmake, 13
 cmake-gui, 11
 command line, 15

S

SCons, 4
 shared library. *See* libraries
 shell commands, 103
 source files, 25, 431
 generated, 107
 source packages, 150, 154
 subprojects, 237
 SWIG, 2, 47–49
 syntax of CMake, 8, 33
 system introspection, 262, *See* modules

T

target platform, 126
 targets, 24
 custom, 108
 custom commands, 105
 exporting, 79
 importing, 78
 properties, 25, 417

testing, 185
 toolchain file, 126
 true, 38
 tutorial, 255

U

unit testing, 186
 UNIX
 libraries, 63
 Makefiles. *See* Makefiles
 symbols, 61
 utility modules. *See* modules

V

valgrind, 200, 203
 variable argument lists, 41
 variable scope, 28
 variables, 9, 27, 269
 cache, 29
 interaction with cache, 31
 list, 9, 29, 34
 PARENT_SCOPE, 28
 quoting, 9
 remove from list, 34
 rule, 247
 scope, 40
 types, 30
 versions, 43
 Vim mode, 17
 Visual Studio, 3, 5, 13, 15, 17, 21, 23, 32, 34, 66,
 104, 105, 107, 109, 112, 117
 converting to CMake, 123
 file groups, 123
 header files, 123

W

white space, 33
 Windows

- build directory, 119
- CDash, 208
- dllexport, 61
- executable, 24
- executables, 123
- installers. *See* NSIS
- *linking, 59
- manifest files, 163
- registry entries, 2, 86
- registry values, 9
- relocatable applications, 76
- run time libraries, 163
- symbols, 61
- WinMain, 24, 123
- wrapping C/C++. *See* SWIG

X

-
- Xcode, 3, 66, 274